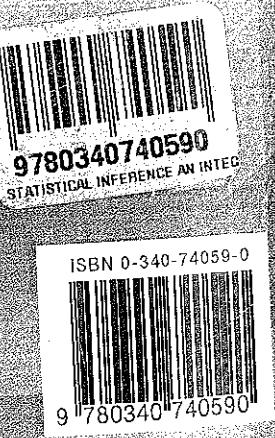


This book has been printed digitally and produced
in a standard design, so that Hodder Arnold can
ensure its continuing availability.

Statistical Inference: An Integrated Approach

HELIOS MIGON & DANI GAMERMAN

Hodder Arnold
www.hoddereducation.co.uk



519.5
M636S
c.1



544.5
M 63620
c. 1

Biblioteca Dep. Monograficas y Es. Vida
Ref. 060/08

**Statistical Inference:
an Integrated Approach**



H S Migon

and

D Gamerman

Federal University of Rio de Janeiro, Brazil

Proyecto Docente Cód. N° 6241



A member of the Hodder Headline Group
LONDON • NEW YORK • SYDNEY • AUCKLAND



av 102422
ch. 5294006

Para Mirna, Marcio, Marcelo e Renato (Helio)
Para meus pais Bernardo e Violeta (Dani)

First published in Great Britain in 1999 by
Arnold, a member of the Hodder Headline Group,
338 Euston Road, London NW1 3BH

<http://www.arnoldpublishers.com>

Co-published in the United States of America by
Oxford University Press Inc.,
198 Madison Avenue,
New York, NY 10016

© 1999 Dani Gamerman and Helio Migon

All rights reserved. No part of this publication may be reproduced or transmitted in any form or by any means, electronically or mechanically, including photocopying, recording or any information storage or retrieval system, without either prior permission in writing from the publisher or a licence permitting restricted copying. In the United Kingdom such licences are issued by the Copyright Licensing Agency: 90 Tottenham Court Road, London W1P 9HE.

Whilst the advice and information in this book are believed to be true and accurate at the date of going to press, neither the authors nor the publisher can accept any legal responsibility or liability for any errors or omissions that may be made.

British Library Cataloguing in Publication Data

A catalogue record for this book is available from the British Library

Library of Congress Cataloguing-in-Publication Data

A catalogue record for this book is available from the Library of Congress

ISBN 0 340 74059 0

1 2 3 4 5 6 7 8 9 10

Publisher: Nicki Dennis

Production Editor: Julie Delf

Production Controller: Sarah Kett

Cover Design: Terry Griffiths

Typeset in 10/12 pt Times by Focal Image Ltd, London

What do you think about this book? Or any other Arnold title?
Please send your comments to feedback.arnold@hodder.co.uk

Contents

Preface	vii
1 Introduction	1
1.1 Information	2
1.2 The concept of probability	3
1.3 An example	7
1.4 Basic results in linear algebra and probability	10
1.5 Notation	14
1.6 Outline of the book	15
Exercises	16
2 Elements of inference	19
2.1 Common statistical models	19
2.2 Likelihood-based functions	21
2.3 Bayes' theorem	26
2.4 Exchangeability	32
2.5 Sufficiency and the exponential family	35
2.6 Parameter elimination	41
Exercises	45
3 Prior distribution	53
3.1 Entirely subjective specification	53
3.2 Specification through functional forms	55
3.3 Conjugacy with the exponential family	59
3.4 The main conjugate families	60
3.5 Non-informative priors	65
3.6 Hierarchical priors	71
Exercises	73
4 Estimation	79
4.1 Introduction to decision theory	79
4.2 Classical point estimation	86
4.3 Comparison of estimators	92

vi Contents

4.4 Interval estimation	99
4.5 Estimation in the normal model	105
Exercises	116
5 Approximate and computationally intensive methods	125
5.1 The general problem of inference	125
5.2 Optimization techniques	126
5.3 Analytical approximations	133
5.4 Numerical integration methods	144
5.5 Simulation methods	147
Exercises	160
6 Hypothesis testing	167
6.1 Introduction	167
6.2 Classical hypothesis testing	168
6.3 Bayesian hypothesis testing	181
6.4 Hypothesis testing and confidence intervals	186
6.5 Asymptotic tests	188
Exercises	191
7 Prediction	197
7.1 Bayesian prediction	197
7.2 Classical prediction	201
7.3 Prediction in the normal model	203
7.4 Linear prediction	206
Exercises	208
8 Introduction to linear models	211
8.1 The linear model	211
8.2 Classical linear models	213
8.3 Bayesian linear models	218
8.4 Hierarchical linear models	224
8.5 Dynamic linear models	227
Exercises	232
Sketched solutions to selected exercises	235
List of distributions	247
References	253
Author index	256
Subject index	258

Preface

This book originated from the lecture notes of a course in statistical inference taught on the MSc programs in Statistics at UFRJ and IMPA (once). These notes have been used since 1987. During this period, various modifications were introduced until we arrived at this version, judged as minimally presentable.

The motivation to prepare this book came from two different sources. The first and more obvious one for us was the lack of texts in Portuguese, dealing with statistical inference to the desired depth. This motivation led us to prepare the first draft of this book in the Portuguese language in 1993. The second, and perhaps the most attractive as a personal challenge, was the perspective adopted in this text. Although there are various good books in the literature dealing with this subject, in none of them could we find an integrated presentation of the two main schools of statistical thought: the frequentist (or classical) and the Bayesian. This second motivation led to the preparation of this English version. This version has substantial changes with respect to the Portuguese version of 1993. The most notable one was the inclusion of a whole new chapter dealing with approximation and computationally intensive methods.

Generally, statistical books follow their author's point of view, presenting at most, and in separate sections, related results from the alternative approaches. In this book, our proposal was to show, wherever possible, the parallels existing between the results given by both methodologies. *Comparative Statistical Inference* by V. D. Barnett (1973) is the book that is closest to this proposal. It does not however present many of the basic inference results that should be included in a text proposing a wide study of the subject. Also we wanted to be as comprehensive as possible for our aim of writing a textbook in statistical inference.

This book is organized as follows. Chapter 1 is an introduction, describing the way we find most appropriate to think statistics: discussing the concept of information. Also, it briefly reviews basic results of probability and linear algebra. Chapter 2 presents some basic concepts of statistics such as sufficiency, exponential family, Fisher information, exchangeability and likelihood functions. Another basic concept specific to Bayesian inference is prior distribution, which is separately dealt with in Chapter 3.

Certain aspects of inference are individually presented in Chapters 4, 6, and 7. Chapter 4 deals with parameter estimation where, intentionally point and interval

estimation are presented as responses to the summarization question, and not as two unrelated procedures. The important results for the normal distribution are presented and also serve an illustrative purpose. Chapter 6 is about hypotheses testing problems under the frequentist approach and also under the various possible forms of the Bayesian paradigm.

In between them lies Chapter 5 where all approximation and computationally based results are gathered. The reader will find there at least a short description of the main tools used to approximately solve the relevant statistical problem for situations where an explicit analytic solution is not available. For this reason, asymptotic theory is also included in this chapter.

Chapter 7 covers prediction from both the frequentist and Bayesian points of view, and includes the linear Bayes method. Finally in Chapter 8, an introduction to normal linear models is made. Initially the frequentist approach is presented, followed by the Bayesian one. Based upon the latter approach, generalisations are presented leading to the hierarchical and dynamic models.

We tried to develop a critical analysis and to present the most important results of both approaches commenting on the positive and negative aspects of both. As has already been said, the level of this book is adequate for an MSc course in statistics, although we do not rule out the possibility of its use in an advanced undergraduate course aiming to compare the two approaches.

This book can also be useful for the more mathematically trained professionals from related areas of science such as economics, mathematics, engineering, operations research and epidemiology. The basic requirements are knowledge of calculus and probability, although basic notions of linear algebra are also used. As this book is intended as a basic text in statistical inference, various exercises are included at the end of each chapter. We have also included sketched solutions to some of the exercises and a list of distributions at the end of the book, for easy reference.

There are many possible uses of this book as a textbook. The first and most obvious one is to present all the material in the order it appears in the book and without skipping sections. This may be a heavy workload for a one semester course. In this case we suggest postponing Chapter 8 to a later course. A second option for exclusion in a first course is Chapter 5, although we strongly recommend it for anybody interested in the modern approach to statistics, geared towards applications. The book can also be used as a text for a course that is more strongly oriented towards one of the schools of thought. For a Bayesian route, follow Chapters 1, 2, 3, Sections 4.1, 4.4.1 and 4.5, Chapter 5, Sections 6.3, 6.4, 6.5, 7.1, 7.3.1, 7.4, 8.1, 8.3, 8.4 and 8.5. For a classical route, follow Chapter 1, Sections 2.1, 2.2, 2.5, 2.6, 4.2, 4.3, 4.4.2 and 4.5, Chapter 5, Sections 6.1, 6.2, 6.4, 6.5, 7.2, 7.3.2, 7.4, 8.1 and 8.2.

This book would not have been possible without the cooperation of various people. An initial and very important impulse was the typing of the original lecture notes in *TeX* by Ricardo Sandes Ehlers. Further help was provided by Ana Beatriz Soares Monteiro, Carolina Gomes, Eliane Amiune Camargo, Monica Magnanini

and Otávio Santos Figueiredo. Besides these, many of our former students helped with suggestions and criticism. Careful proofreading of this manuscript was made by our past MSc students and present colleagues Alexandra Mello Schmidt, Hédibert Freitas Lopes and Marco Antonio Rosa Ferreira. Many useful suggestions and comments were provided at this later stage by Steve Brooks, Eduardo Gutierrez-Peña and Gabriel Huerta. We also had the stimulus of several colleagues; in particular, we would like to mention Basílio de B. Pereira. We would also like to thank Nicki Dennis for her support and encouragement throughout all the stages of preparation of this book and for making us feel at home with Arnold, and Márcio N. Migon for the many comments on earlier versions of the book. The Brazilian research supporting agencies CNPq and FAPERJ and the Ministry of Science and Technology have helped by the continuing support of our research work. This support enabled the use in this project of the computational and bibliographic material they provided for our own research. Our families also played the important roles of support and understanding, especially in the weekends and late nights spent trying to meet deadlines! To all of them, our gratitude.

Finally, the subject of the book is not new and we are not claiming any originality here. We would like to think that we are presenting the subject in a way that is not favoured in many textbooks and that will help readers to have an integrated view of the subject. In our path to achieve this goal, we have been influenced by many researchers and books. We tried to acknowledge this influence by referring to these books whenever we felt it provided a description of a topic worth reading. Therefore, we tried to relate every major subject presented in our book to books that treated the subject in a more complete or more interesting way. In line with its textbook character, we opted to favour books rather than research papers as references. We would like to think of our book as a basis for discovery and will feel our task is accomplished whenever readers understand the subject through the book alone, its references or a combination of both.

H.S.M. & D.G.
Rio de Janeiro, December 1998

1

Introduction

Before beginning the study of statistics, it is relevant to characterize the scope of the area and the main issues involved in this study. We avoid defining the subject directly, which is a hard and polemical task. Some of the components involved in this area of science will be presented in the hope that in the end the reader will have a clear notion of the breadth of the subject under consideration. The fundamental problem towards which the study of statistics is addressed is that where randomness is present. The statistical methodology to deal with the resulting uncertainty is based on the elaboration of probabilistic models in order to summarize the relevant information available.

There are many concepts used in the last sentence that deserve a clarified explanation of their meaning in order to ensure a unified understanding. A model is a formal collection of coherent rules describing, in a simplified way, some real world problem. The language used to describe precisely a model in which uncertainty is present is probability. The meaning of summarization, in this context, refers to the ability to describe a problem in the most concise way (under the assumption that there are many possible ways to describe a problem). The art involved in model building is the desire to balance the need to include as many aspects of reality as possible while preventing it from being too complex. A related concept, which is useful to have in mind, is that of parsimony. This means that the model must have an optimal complexity. A very simple model can be misleading since it misses relevant aspects of reality. On the other hand, if it is highly complex it will be hard to understand and extract meaningful information from. From the previous discussion it is not difficult to guess that information is the main input for statistics. However, this is a hard concept to define.

Among the objectives of our study, we can identify the two main ones as being: to understand reality (estimation and hypothesis testing) and to make decisions (predictions). A strong limitation of many presentations of statistical inference is to be mainly concentrated in estimation and testing. In this context, it only deals with quantities that can never be observed in the present or in the future. Nevertheless, our view of statistics is that it must be concerned with observable quantities. In this way, one is able to verify the model adequacy in an irrefutable way. For the sake of completeness, however, we will present the main results available from all topics above.

1.1 Information

As we have already said, the notion of information is present in all the studies developed in statistics. Since uncertainty is one of the main ingredients in our models, we need to gather as much information as possible in order to reduce our initial uncertainty. A fundamental question which we are concerned with is about the type of information that is relevant and must be retained in the analysis. A possible reply to this question is that all available information is useful and must be taken into consideration. Another answer is to avoid arbitrariness and take into consideration only objective observation coming from a sampling process. In this way all the subjective information must be discarded.

These two points of view roughly form the bases for two different forms of statistical analysis: the Bayesian (or subjectivist) and the classical (or frequentist) approach, respectively. As we will see in the next section the divergence between these two approaches is much stronger, beginning with the interpretation of the concept of probability. This is always the starting point of a statistical model.

An example to illustrate and clarify these points is as follows.

Example. Consider the situation described by Berger (1985, p. 2) concerning the following experiments:

1. A fine musician, specializing in classical works, tell us that he is able to distinguish if Hayden or Mozart composed some classical song. Small excerpts of the compositions of both authors are selected at random and the experiment consists of playing them for identification by the musician. The musician makes 10 correct guesses in exactly 10 trials.
2. A drunken man says that he can correctly guess in a coin toss what face of the coin will fall down. Again, after 10 trials the man correctly guesses the outcomes of the 10 throws.
3. An old English lady is well known for her ability to distinguish whether a cup of tea is prepared by first pouring the milk or the tea. Ten cups filled with tea and milk, well mixed and in a random order, are presented to her. She correctly identifies all ten cups.

It is not difficult to see that the three experiments provide the same information and therefore any test to verify the authenticity of the persons' statements would result positive for all of them, with the same confidence.

This does not make any sense! We have more reasons to believe in the authenticity of the statement of the musician than of the old lady and, certainly, much more than of the drunken man. There is no doubt that the experimental outcome increases the veracity of the statements made. But we cannot reasonably say that we have the same confidence in the three assertions. By common sense, there is a long way to go before one accepts this conclusion.

1.2 The concept of probability

Although the definition of probability is well accepted by almost every statistician (ignoring some technical details), its interpretation or the sense attributed to it varies considerably. We mention here some of the more common interpretations: physical (or classical), frequentist and subjective.

- (1) Physical or classical. The probability of any event is the ratio between the number of favourable outcomes and the total number of possible experimental results. It is implicitly assumed that all the elementary events have the same chance. The concept of probability was first formulated based on these classical ideas, which are closely related to games of chance (cards, dice, coins, etc.), where the equal chance assumption is taken for granted.

The probability associated with more elaborate events is obtained just as a consequence of the probability of the elementary events. Obviously, this interpretation is too narrow to be used in general. Besides that, how can we recognize equal chance events? Finally, the notion of chance involves some probabilistic consideration and so the argument is in some way circular.

A similar interpretation is provided by the logical viewpoint that tries to ascertain relations between events based on logical reasoning. The main question is how to translate common scientific knowledge into undisputed, objective numbers representing probability of events.

- (2) Frequentist. The probability of an event A , denoted by $Pr(A)$ or $P(A)$, is given by

$$Pr(A) = \lim_{n \rightarrow \infty} \frac{m}{n}$$

where m is the number of times that A has occurred in n identical and independent experimental trials. This interpretation intends to be objective as far as it is based only on observable quantities. However, it is worth noting that:

- (i) The limit cannot be understood as a mathematical limit since given $\epsilon > 0$ and $N > 0$ there could well exist an $N_0 > N$ such that $|Pr(A) - (m/N_0)| > \epsilon$. This is improbable but not impossible.
- (ii) The concepts of identical and independent trial are not easy to define objectively and are in essence subjective.
- (iii) n does not go to infinity and so there is no way to ensure the existence of such a limit.

The scope of the two interpretations is limited to observable events and does not correspond to the concept used by common people. The human being evaluates (explicitly or implicitly) probabilities of observable and unobservable events.

Examples.

1. Consider the proposition $A = \text{'it will rain today'}$. A is typically non-observable at the moment I leave home but $Pr(A)$ is a legitimate and very

4 Introduction

- useful quantity to consider. If its numerical value is low then I will decide not to take an umbrella, I will prepare myself for a nice walk back home, etc.
2. Let A be ‘John has disease X ’. Although A can be an observable quantity after delicate and expensive surgery, John’s doctor can take a number of actions (including the surgery itself) based on the value he ascertains for $Pr(A)$.
 3. The proposition A is ‘John will get married to Mary’. Once again it makes sense to think about $Pr(A)$, especially if I have some sort of personal relationship with John and/or Mary.

In all the cases presented above the classical and frequentist interpretations do not make sense. A is always non-observable, unique and cannot be repeated under similar conditions.

- (3) Subjective. The probability of an event A is a measure of someone’s degree of belief in the occurrence of A . To emphasize its subjective character, it is better to denote this probability by $Pr(A | H)$ where H (for history) represents the available information set of the individual. For example, let A be the event ‘it is raining in Moscow’.

- (a) The easiest probability to associate with A for someone in Rio de Janeiro who does not know anything about the Moscow climate, is $Pr(A | H_1) = 0.5$, which is based on his or her body of knowledge H_1 .
- (b) On the other hand, someone in St. Petersburg could have stated

$$Pr(A | H_2) = \begin{cases} 0.75, & \text{if it is raining in St. Petersburg} \\ 0.25, & \text{if it is not.} \end{cases}$$

Note that, in contrast to someone in Rio, this person will typically have more information, contained in H_2 .

- (c) But for someone in Moscow

$$Pr(A | H_3) = \begin{cases} 1, & \text{if it is raining} \\ 0, & \text{otherwise} \end{cases}$$

because in this case, H_3 contains A !

It is worth pointing out that the values for $Pr(A | H)$ are not equal since they depend on the information H , which is different for each case. This interpretation of probability illustrated by the above example is called subjective and obeys the basic rules of probability. Note also that adopting the subjective interpretation we can associate probability for the cases unsolved by the other schools of thought. The remaining question is how to obtain its value for a given event or proposition based on a specified information set.

Probabilities can be evaluated directly or indirectly. One standard tool for direct probability evaluation can be an urn with 100 (or 1000, say) balls with two different colours: blue and red. For example, let us suppose that you want to assess

The concept of probability 5

the probability that the Canoas road (a very pleasant road along the mountains surrounding Rio) is closed because of a road accident. In the direct approach, you must compare this probability with the chance of drawing a red ball from the urn. If these probabilities were judged equal when the urn has 20 red balls, then the probability that the road is closed is 0.2.

For the case of indirect measurements, let us assume that we have two lottery systems: one involving the event you are interested in and the other is any direct evaluation instrument. Imagine the following lotteries:

1. Bet A : If the road is not blocked you win 5 monetary units, otherwise you do not win anything.
2. Bet B : If the ball drawn from the urn is red, you win 5 monetary units, otherwise nothing.

Considering the two lotteries offered to you, on which one do you prefer to bet? If you prefer A , it might be because there is a bigger chance to win the premium betting on A than on B . Now, let us make a small modification in the urn composition, to 10 red balls and 90 blue ones, so that the probability of a winning bet B is 0.1. If you still prefer A , redefine again the composition of the urn and continue until you become indifferent between bets A and B .

There are also difficulties associated with these forms of probability evaluations. In the first case, the difficulty is associated with the comparison of probabilities coming from different propositions. In the second case, the difficulty is caused by the introduction of monetary units and the evaluation of their respective utilities.

1.2.1 Assessing subjective probabilities

There are many alternative ways to determine subjective probabilities. De Finetti (1974) provides a very useful scheme based on the notion of squared error loss function. Let A be a proposition or an event identified with the value 1 when it is true and 0 otherwise. The probability p that we associate to A is obtained by minimizing the square error loss function.

$$(p - A)^2 = \begin{cases} (p - 1)^2, & \text{if } A = 1 \\ p^2, & \text{if } A = 0. \end{cases}$$

The basic properties of probability follow easily from this definition as will be shown below.

1. $p \in [0, 1]$

If $p > 1$ then $p^2 > 1$ and $(p - 1)^2 > 0$. Therefore the losses are always bigger than 1 and 0, the losses obtained by making $p = 1$. A similar argument is used to show that if $p < 0$ the losses will be bigger than those evaluated with $p = 0$. Then, minimization of the square error losses imposes $p \in [0, 1]$. Figure 1.1 arrives at the same conclusion graphically.

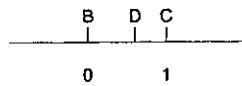


Fig. 1.1 The possible losses are given by \overline{BD}^2 and \overline{CD}^2 which are minimized if D is between B and C .

$$2. \Pr(\bar{A}) = 1 - \Pr(A)$$

The possible losses associated with the specification of $\Pr(A) = p$ and $\Pr(\bar{A}) = q$ are:

$$\begin{aligned} A &= 1 : (p - 1)^2 + q^2 \\ A &= 0 : p^2 + (q - 1)^2. \end{aligned}$$

As we have already seen in (1), the possible values of (p, q) are in the unit square. In Figure 1.2 line segments are drawn to describe the possible losses. The squared distance between two consecutive vertices represents the losses. It is clear from the figure that the losses are reduced by making $p + q = 1$.

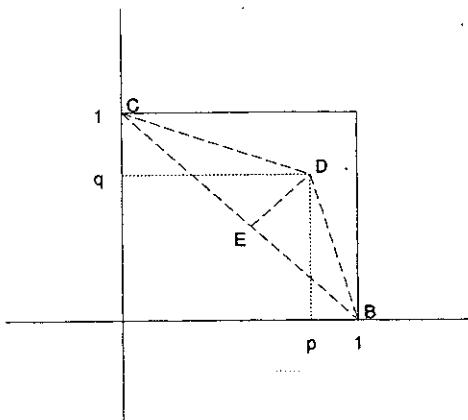


Fig. 1.2 The possible losses are given by \overline{BD}^2 when $A = 1$ and \overline{CD}^2 when $A = 0$, which are minimized if D is projected on E over the line $p + q = 1$.

$$3. \Pr(A \cap F) = \Pr(A | F)\Pr(F)$$

Define $\Pr(A | F)$ as the probability of A if $F = 1$. Denoting this probability by p , $\Pr(F)$ by q and $\Pr(A \cap F)$ by r , the total loss is given by $(p - A)^2F + (q - F)^2 + (r - AF)^2$. Its possible values are:

$$\begin{aligned} A = F = 1 &: (p - 1)^2 + (q - 1)^2 + (r - 1)^2 \\ A = 0, F = 1 &: p^2 + (q - 1)^2 + r^2 \end{aligned}$$

$$F = 0 : q^2 + r^2.$$

Note that (p, q, r) assume values in the unit cube. The same arguments used in (2) can be developed in the cube. Minimization of the three losses is attained when $p = r/q$.

1.3 An example

In this section a simple example will be presented with the main intention of anticipating many of the general questions to be discussed later on in the book. The problem to be described is extremely simple but is useful to illustrate some relevant ideas involved in statistical reasoning. Only very basic concepts on probability are enough for the reader to follow the classical and Bayesian inferences we will present. The interested reader is recommended to read the excellent paper by Lindley and Phillips (1976) for further discussion.

On his way to university one morning, one of the authors of this book (say, Helio) was stopped by a lady living in the neighbouring Maré slum. She was pregnant and anxious to know the chance of her seventh baby being male.

Initially, his reaction was to answer that the chance is $1/2$ and to continue his way to work. But the lady was so disappointed with the response that Helio decided to proceed as a professional statistician would. He asked her some questions about her private life and she told him that her big family was composed of five boys (M) and one girl (F) and the sequence in which the babies were born was MMMMMF. The lady was also emphatic in saying that all her pregnancies were consequence of her long relationship with the same husband. In fact, her disappointment with the naive answer was now understandable.

Our problem is to calculate the chance of the seventh baby being male, taking into consideration the specific experience of this young lady. How can we solve this problem?

Assuming that the order of the Ms and Fs in the outcome is not relevant to our analysis, it is enough or sufficient to take note that she had exactly five baby boys (5 Ms) and one baby girl (1 F). The question about the order of the births is brought up because people usually want to know if there is any sort of abnormality in the sequence. However, it seems reasonable to assume the births to be equally distributed in probabilistic terms.

Before proceeding with an analysis it is useful to define some quantities. Let us denote by X_i the indicator variable of a boy for the i th child, $i = 1, \dots, 7$ and let θ denote the common unknown probability of a boy, i.e., $\Pr(X_i = 1 | \theta) = \theta$ with $0 \leq \theta \leq 1$. Note that θ is a fixed but unknown quantity that does not exist in reality but becomes a useful summary of the situation under study.

A frequentist statistician would probably impose independence between the X_i 's. In doing that the only existing link between the X_i 's is provided by the value of θ . It seems reasonable at this stage to provide to the lady the value that one considers the most reasonable representation of θ . Given that he is only allowed to

8 Introduction

use the observed data in his analysis, this representative value of θ must be derived from the observations $X_i, i = 1, \dots, 6$.

There are many possible ways to do that. Let us start by the probabilistic description of the data. Given the assumptions above, it is not difficult to obtain

$$Pr(X_1 = 1, X_2 = 1, \dots, X_5 = 1, X_6 = 0 | \theta) = \theta^5(1 - \theta).$$

One can proceed on this choice of value for θ by finding the single value $\hat{\theta}$ that maximizes the above probability for the actual, observed data. It is a simple exercise to verify that, in the above case, this value is given by $\hat{\theta} = 5/6 = 0.83$. We would then say that she has 83% chance of giving birth to a boy.

There are other ways to proceed still based only on the observed data but more assumptions are now needed. One possible assumption is that the lady had previously decided that she only wanted to have six children, the last one being an unwanted pregnancy. In this case, the observed data can be summarized into the number Y of M's among the six births. It is clear that Y has a binomial distribution with size six and success probability θ , denoted by $Y \sim \text{bin}(6, \theta)$, and that $E(Y/6|\theta) = \theta$. Using frequentist arguments, one can reason that when we are able to observe many boys in a similar situation one would like to be correct on average. Therefore, one would estimate the value of θ as $Y/6$, the relative frequency of boys in the data. Given that the observed value of Y is 5, a reasonable estimate for θ is $\tilde{\theta} = 5/6$, coinciding with $\hat{\theta}$. There is no guarantee that the two approaches coincide in general but is reassuring that they did in this case.

When asking the lady if the assumption to stop at the sixth child was true, she said that her decision had to do with having had her first baby girl. In this case, the observed data should have been summarized by the number Z of M's she had until the first girl, and not by Y . (Even though the observed values of Z and Y are the same, their probability distributions are not.) It is not difficult to see that Z has a negative binomial distribution with size 1 and success probability $1 - \theta$, denoted $Z \sim NB(1, \theta)$, and that $E(Z|\theta) = \theta/(1 - \theta)$. Proceeding on Z with the reasoning used for Y leads to the estimation of θ by $5/6$ as in the previous cases.

The main message from this part of the example for choosing a representative value for θ is that there are many possible methods, two of which were applied above and while the first one did not depend on the way the data was observed, the second one did. These issues are readdressed at greater length in Chapters 2 and 4.

Another route that can be taken is to decide whether it is reasonable or not to discard $1/2$ as a possible value for θ . This can be done by evaluating how extreme (in discordance of the assumption $\theta = 1/2$) the observed value is. To see that, one can evaluate the probabilities that $Y \geq 5$ and $Z \geq 5$, depending on which stopping rule was used by the lady. The values of these probabilities are respectively given by 0.109 and 0.031. It is generally assumed that the cutoff point for measuring extremeness in the data is to have probabilities smaller than 0.05. It is interesting that in this case the stopping rule has a strong effect on the decision to discard the equal probabilities assumption. This will be readdressed in a more general form in Chapter 6.

An example 9

Intuition, however, leads to the belief that specification of the stopping rule is not relevant to solving our problem. This point can be more formally expressed in the following way: the unique relevant evidence is that in six births, the sex of the babies was honestly annotated as 1 F and 5 Ms. Furthermore, these outcomes occurred in the order specified before. This statement points to the conclusion that only the results that have effectively been observed are relevant for our analysis.

For a Bayesian statistician, the elements for the analysis are only the sequence of the observed results and a probability distribution describing the initial information about the chance of a baby being male. The experimental conditions were carefully described before and they guarantee that a result observed in any given birth is equivalent to that obtained in any other birth. The same is true for pairs, triplets, etc. of birth. This idea is formalized by the concept of exchangeability. The sequence of births is exchangeable if the order of the sex outcomes in the births is irrelevant. In the next chapter, we will define precisely the concept of exchangeability. For our present example this means that the probability of any sequence of r Ms and s Fs (subject to $r + s = n$) is the same as that of any other sequence with the same number of Ms and Fs.

Let us return to our original problem, that is, to calculate the chance of the seventh baby born being male based on the information gathered, namely that provided by the previous births. This probability is denoted by $Pr[X_7 = 1 | (5, 1)]$ where the pair $(5, 1)$ denote the number of births from each sex previously observed. Using basic notions of probability calculus we can obtain

$$\begin{aligned} Pr[X_7 = 1 | (5, 1)] &= \int_0^1 P[X_7 = 1, \theta | (5, 1)] d\theta \\ &= \int_0^1 P[X_7 = 1 | \theta, (5, 1)] p(\theta | (5, 1)) d\theta \\ &= \int_0^1 \theta p(\theta | (5, 1)) d\theta \\ &= E[\theta | (5, 1)] \end{aligned}$$

where the expected value is with respect to the distribution of θ given the past results. As we will see in the next chapter, this is the unique possible representation for our problem if the assumption of exchangeability of the sequences of births is acceptable. One of the elements involved in the above calculation is $p(\theta | (5, 1))$ which has not yet been defined. It has the interpretation, under the subjective approach, of the probability distribution of the possible values for θ after observing the data $(5, 1)$.

Let us suppose that before observing the values of $(5, 1)$, the subjective probability specification for θ can be represented by the density

$$p(\theta) = k\theta^{a-1}(1 - \theta)^{b-1}, \quad 0 \leq \theta \leq 1, \quad (a, b > 0)$$

which is a beta distribution with parameters a and b (see the list of distributions at the end of the book).

Note that

$$\begin{aligned}
 p(\theta | (5, 1)) &= \frac{p((5, 1), \theta)}{p((5, 1))} \\
 &= \frac{p((5, 1) | \theta)p(\theta)}{p((5, 1))} \\
 &\propto \theta^5(1 - \theta)^{a-1}(1 - \theta)^{b-1}, \\
 &\quad \text{since } p((5, 1)) \text{ does not depend on } \theta \\
 &\propto \theta^{5+a-1}(1 - \theta)^{1+b-1}
 \end{aligned}$$

and the stopping rule or the sample space, in classical language, is irrelevant because it gets incorporated into the proportionality constant. Furthermore, for any experimental result the final distribution will be a beta. So we can complete the calculations to obtain

$$Pr[X_7 = 1 | (5, 1)] = E[\theta | (5, 1)] = \frac{a+5}{a+b+6}.$$

We still have a problem to be solved. What are the values of a and b ? Suppose, in the case of births of babies, that our initial opinion about the chances associated with M and F are symmetric and concentrated around 0.5. This means that the distribution of θ is symmetrically distributed with mean 0.5 and with high probability around the mean. We can choose in the family of beta distributions that one with $a = b = 2$, for instance. With this specification, $E(\theta) = 0.5$, $P(0.4 < \theta < 0.6) = 0.3$ and the probability of the seventh birth being a boy is $7/10 = 0.70$.

If we have been thinking of an experiment with honest coins instead of births of babies, we could have chosen a beta(50,50) which is symmetrically distributed but much more concentrated around 0.5 than the beta(2,2). This is a clear representation of the fact that we know much more about coins than about the sex of new birth, even before observing the data. Under this specification of the beta, $Pr(0.4 < \theta < 0.6) = 0.8$ and the chance of the 7th outcome being a head would be $55/106 = 0.52$. Evidently this result is closer to 0.5, which seems reasonable since the sex of babies and coins are quite different things. In Chapter 3 we will come back to the discussion of how to specify and determine this prior distribution in any given statistical problem.

1.4 Basic results in linear algebra and probability

This is a book concerned with the study of statistics. In order to do that, a few simple results from linear algebra and probability theory will be extensively used. We thought it might be useful to have the main ones cited here to provide the reader with the basic results he will be using throughout the book. We will start with the basic

results concerning densities and probability functions of collections of random variables and will then define the multivariate normal distribution to motivate the importance of linear algebra results about matrices and their connection with distributions.

1.4.1 Probability theory

Let $\mathbf{X} = (X_1, \dots, X_p)$, $\mathbf{Y} = (Y_1, \dots, Y_q)$ and $\mathbf{Z} = (Z_1, \dots, Z_r)$ be three random vectors defined over sample spaces \mathcal{X} , \mathcal{Y} and \mathcal{Z} , respectively ($p, q, r \geq 1$). Assume for simplicity that they are all continuous with joint probability density function $p(\mathbf{x}, \mathbf{y}, \mathbf{z})$. The marginal and conditional densities will be denoted by their relevant arguments. So, for example, $p(\mathbf{x})$ denotes the marginal density of \mathbf{X} and $p(\mathbf{z}|\mathbf{x}, \mathbf{y})$ denotes the conditional density of $\mathbf{Z}|\mathbf{X}, \mathbf{Y}$. Then, the following equations hold:

$$\begin{aligned}
 p(\mathbf{x}) &= \int_{\mathcal{Y}} p(\mathbf{x}, \mathbf{y}) d\mathbf{y} \\
 p(\mathbf{x}|\mathbf{y}) &= \frac{p(\mathbf{x}, \mathbf{y})}{p(\mathbf{y})} \\
 p(\mathbf{x}|\mathbf{z}) &= \int_{\mathcal{Y}} p(\mathbf{x}, \mathbf{y}|\mathbf{z}) d\mathbf{y} \\
 p(\mathbf{x}|\mathbf{y}, \mathbf{z}) &= \frac{p(\mathbf{x}, \mathbf{y}|\mathbf{z})}{p(\mathbf{y}|\mathbf{z})}
 \end{aligned}$$

There are many possible combinations of these results but most can be derived easily from one of the above results. In fact, all results below are valid under more general settings with some components of the vectors being discrete and others continuous. The only change in the notation is the replacement of integrals by summations over the relevant parameter space.

These relations define a structure of dependence between random variables. An interesting concept is conditional independence. \mathbf{X} and \mathbf{Y} are said to be conditionally independent given \mathbf{Z} if $p(\mathbf{x}, \mathbf{y}|\mathbf{z}) = p(\mathbf{x}|\mathbf{z})p(\mathbf{y}|\mathbf{z})$. Conditional dependence structures can be graphically displayed with the use of influence diagrams. Figure 1.3 shows some possible representations involving three random variables. These diagrams can be very useful in establishing probability structures for real problems that tend to have many more than just three variables.

Important elements to describe a distribution are its mean vector and variance-covariance matrix. The mean vector μ of a collection of random variables $\mathbf{X} = (X_1, \dots, X_p)$ has components $\mu_i = E(X_i)$, $i = 1, \dots, p$. The variance-covariance matrix Σ of \mathbf{X} has elements $\sigma_{ij} = Cov(X_i, X_j)$. It is clearly a symmetric matrix since $Cov(X_i, X_j) = Cov(X_j, X_i)$, for every possible pair (i, j) .

These moments are well-defined for any well defined distribution although they may not exist for some distributions. Therefore, one can evaluate the conditional mean of $X_3|X_1, X_2$ and X_5 by calculating the mean of X_3 under the con-

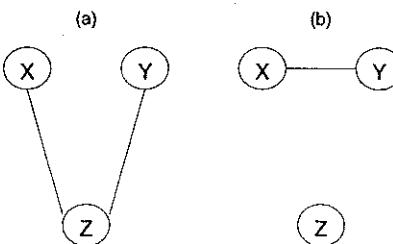


Fig. 1.3 Possible influence diagrams for three random variables: (a) X and Y are conditionally independent given Z ; (b) (X, Y) is independent of Z .

ditional distribution of $X_3|X_1, X_2$ and X_5 . Likewise, one can evaluate the joint marginal variance–covariance matrix of X_1, X_2 and X_5 by evaluating the variance–covariance matrix of the joint marginal distribution of (X_1, X_2, X_5) .

Another useful result concerns transformation of random vectors. Let $\mathbf{X} = (X_1, \dots, X_p)$ and $\mathbf{Y} = (Y_1, \dots, Y_p)$ be p -dimensional random vectors defined over continuous spaces \mathcal{X} and \mathcal{Y} , respectively. Assume further that \mathbf{X} and \mathbf{Y} are uniquely related by the 1-to-1 transformation $\mathbf{Y} = \mathbf{g}(\mathbf{X})$ with inverse function $\mathbf{X} = \mathbf{h}(\mathbf{Y})$ and these functions are at least differentiable. As a consequence, $\mathcal{Y} = \mathbf{g}(\mathcal{X})$ and $\mathcal{X} = \mathbf{h}(\mathcal{Y})$. Then the densities $p_X(\mathbf{y})$ and $p_Y(\mathbf{y})$ are related via

$$p_Y(\mathbf{y}) = p_X(\mathbf{h}(\mathbf{y})) \left| \frac{\partial \mathbf{h}(\mathbf{y})}{\partial \mathbf{y}} \right|, \quad \mathbf{y} \in \mathcal{Y},$$

where

$$\left| \frac{\partial \mathbf{h}(\mathbf{y})}{\partial \mathbf{y}} \right|$$

is the absolute value of the Jacobian of \mathbf{h} , the determinant of the matrix of derivatives of \mathbf{h} with respect to \mathbf{y} . This matrix has element (i, j) given by $\partial h_i(\mathbf{y})/\partial y_j$, $\forall(i, j)$. In the case of scalar X and Y , the relation becomes

$$p_Y(y) = p_X(h(y)) \left| \frac{\partial h(y)}{\partial y} \right|, \quad y \in \mathcal{Y}.$$

An important distribution where some of these results can be used is the multivariate normal distribution with mean vector μ and variance–covariance matrix Σ , denoted by $N(\mu, \Sigma)$ with density

$$(2\pi)^{-p/2} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu)' \Sigma^{-1} (\mathbf{x} - \mu) \right\}, \quad \mathbf{x} \in R^p,$$

where $|\mathbf{A}|$ denotes the determinant of \mathbf{A} . The scalar version ($p = 1$) of this density will simply be referred to as a normal distribution with density

$$(2\pi\sigma^2)^{-1/2} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}, \quad x \in R.$$

When $\mu = 0$ and $\sigma^2 = 1$, the distribution is referred to as standard normal.

It can be shown that the quadratic form in the exponent of the density has a χ^2 distribution. Also, the normal distribution is preserved under linear transformations. The multivariate normal distribution is completely characterized by its parameters μ and especially Σ . If the components are uncorrelated with $\sigma_{ij} = 0$, $\forall i \neq j$, then they are also independent. In any case, the correct understanding of the variance–covariance structure of a distribution is vital, for example, to obtain some properties such as marginal and conditional distributions (see exercises). In order to do that, a few basic results about matrices will be reviewed below.

There are many other distributions that will be considered in some detail in the book. They will be introduced in the text as they become needed.

1.4.2 Linear algebra

Let \mathbf{A} be a real matrix of order $r \times p$, $p, r \geq 1$. Denote the matrix element in row i and column j by a_{ij} , $i = 1, \dots, r$ and $j = 1, \dots, p$. If $p = r$, the matrix is said to be square of order p . Such a matrix is said to be symmetric if $a_{ij} = a_{ji}$, for every possible pair (i, j) . In this case, the transpose of \mathbf{A} , denoted by \mathbf{A}' , is $\mathbf{A}' = \mathbf{A}$.

Let $\mathbf{Y} = (Y_1, \dots, Y_r)'$ be a random vector defined by the linear transformation $\mathbf{Y} = \mathbf{c} + \mathbf{C} \mathbf{X}$ of another random vector $\mathbf{X} = (X_1, \dots, X_p)'$ where \mathbf{c} and \mathbf{C} are an r -dimensional vector and an $r \times p$ matrix of constants. Then the expectation and variance–covariance matrix of \mathbf{Y} are respectively given by $E(\mathbf{Y}) = \mathbf{c} + \mathbf{C}E(\mathbf{X})$ and $V(\mathbf{Y}) = \mathbf{C}\Sigma\mathbf{C}'$. As an example, let $\mathbf{Y} = \mathbf{1}_p' \mathbf{X} = \sum_i X_i$ where $\mathbf{1}_p$ is the p -dimensional vector of 1's. Then, \mathbf{Y} will have mean $\mathbf{1}_p'E(\mathbf{X}) = \sum_i E(X_i)$ and variance $\mathbf{1}_p'V(\mathbf{X})\mathbf{1}_p = \sum_{i,j} Cov(X_i, X_j)$.

A matrix \mathbf{A} is said to be positive (non-negative) definite if $\mathbf{b}'\mathbf{A}\mathbf{b} > (\geq) 0$, for every non-null vector $\mathbf{b} = (b_1, \dots, b_p)'$. Similarly, negative (non-positive) definite matrices can be defined by replacement of the $>$ (\geq) sign by the $<$ (\leq) sign. A trivial example of a positive definite matrix is the p -dimensional identity matrix, denoted by \mathbf{I}_p with diagonal elements equal to 1 and off-diagonal elements equal to 0. It is very easy to check that $\mathbf{b}'\mathbf{I}_p\mathbf{b} = b_1^2 + \dots + b_p^2$ which must be larger than zero because \mathbf{b} is non-null.

Variance–covariance matrices are always non-negative definite and usually are positive definite matrices. To see that, assume that Σ is the variance–covariance matrix of $\mathbf{X} = (X_1, \dots, X_p)'$. Then, a non-null random variable $Z = \mathbf{b}'\mathbf{X}$ can be formed. This variable will have variance $V(Z) = \mathbf{b}'\Sigma\mathbf{b}$ which is necessarily non-negative. Varying \mathbf{b} over all possible values of \mathbf{b} in R^p (excluding the origin) proves the result.

Positive definiteness defines an ordering over matrices. Using the notation $\mathbf{A} > (\geq) \mathbf{0}$ for a positive (non-negative) matrix \mathbf{A} allows one to denote by $\mathbf{A} > (\geq) \mathbf{B}$ the fact that the matrix $\mathbf{A} - \mathbf{B} > (\geq) \mathbf{0}$. This ordering makes sense in the context of probability. Let \mathbf{A} and \mathbf{B} be the variance–covariance matrices of independent p -dimensional random vectors \mathbf{X} and \mathbf{Y} . Then $\mathbf{A} > \mathbf{B}$ implies that

14 Introduction

$V(\mathbf{b}'\mathbf{X}) - V(\mathbf{b}'\mathbf{Y}) = \mathbf{b}'\mathbf{A}\mathbf{b} - \mathbf{b}'\mathbf{B}\mathbf{b} = \mathbf{b}'(\mathbf{A} - \mathbf{B})\mathbf{b} > 0$, for every non-null vector \mathbf{b} . So, there is a sense in which matrices can be compared in magnitude and one can say that \mathbf{A} is *larger* than \mathbf{B} .

Symmetric positive-definite matrices are said to be non-singular, because they have a non-null determinant. This implies that they have full rank and all their rows are linearly independent. Likewise, singular matrices have null determinant, which means that they do not have full rank and some of their rows can be represented as linear combinations of the other rows. Non-singular matrices also have a well-defined inverse matrix.

A square matrix \mathbf{A} of order p with p -dimensional rows $\mathbf{a}_i = (a_{i1}, \dots, a_{ip})'$, for $i = 1, \dots, p$ is said to be orthogonal if $\mathbf{a}_i'\mathbf{a}_j = 1$, if $i = j$ and 0, if $i \neq j$. Note that if \mathbf{A} is orthogonal then $\mathbf{A}'\mathbf{A} = \mathbf{A}\mathbf{A}' = \mathbf{I}_p$. Therefore, orthogonal matrices can be shown to be full rank with inverse \mathbf{A}^{-1} and $\mathbf{A}' = \mathbf{A}^{-1}$. There are many methods available in linear algebra for iteratively constructing an orthogonal matrix from given starting rows.

In most statistical applications, matrices are non-negative definite and symmetric. The square root matrix of \mathbf{A} denoted by $\mathbf{A}^{1/2}$, can then be defined and satisfies $\mathbf{A}^{1/2}\mathbf{A}^{1/2} = \mathbf{A}$. One of the most common methods of finding the square root matrix is called the Choleski decomposition.

1.5 Notation

Before effectively beginning the study of statistical inference it will be helpful to make some general comments. Firstly, it is worth noting that here we will deal with regular models only, that is, the random quantities involved in our models are of the discrete or continuous type. A unifying notation will be used and the distinction will be clear from the context. Then, if \mathbf{X} is a random vector with distribution function denoted by $F(\mathbf{x})$, its probability (or density) function will be denoted by $p(\mathbf{x})$ if \mathbf{X} is discrete (or continuous) and we will assume that

$$\int dF(\mathbf{x}) = \int p(\mathbf{x})d\mathbf{x}$$

independently of \mathbf{X} being continuous or discrete, with the integral symbol representing a sum in the discrete case. In addition, as far as the probability (or density) function is defined from the distribution of \mathbf{X} , we will use the notation $\mathbf{X} \sim p$ meaning that \mathbf{X} has distribution p or, being more precise, a distribution whose probability (or density) function is p . Similarly, $\mathbf{X} \xrightarrow{\mathcal{D}} p$ will be used to denote that \mathbf{X} converges in distribution to a random variable with density p .

In general, the observables are denoted by the capital letters of the alphabet (X, Y, \dots), as usual, and their observed values by lower case letters (x, y, \dots). Known quantities are denoted by the first letters of the alphabet (A, B, \dots), and the greek letters (θ, λ, \dots) are used to describe unobservable quantities. Matrices, observed or not, will be denoted by capitals. Additionally, vectors and matrices will be

Outline of the book 15

distinguished from scalars by denoting the first ones in bold face. Results will generally be presented for the vector case and whenever the specialization to the scalar case is not immediate, they will be presented again in a univariate version.

The distribution of X is denoted by $P(X)$. The expected value of $\mathbf{X}|Y$ is denoted by $E(\mathbf{X}|Y)$, $E_{\mathbf{X}|Y}(\mathbf{X})$ or even $E_{\mathbf{X}|Y}(\mathbf{X}|Y)$ and the variance of $\mathbf{X}|Y$ is denoted by $V(\mathbf{X}|Y)$, $V_{\mathbf{X}|Y}(\mathbf{X})$ or even $V_{\mathbf{X}|Y}(\mathbf{X}|Y)$. The indicator function, denoted by $I_x(A)$, assumes the values

$$I_x(A) = \begin{cases} 1, & \text{if } x \in A \\ 0, & \text{otherwise.} \end{cases}$$

1.6 Outline of the book

The purpose of this book is to present an integrated approach to statistical inference at an intermediate level by discussion and comparison of the most important results of the two main schools of statistical thought: frequentist and Bayesian. With that in mind, the results are presented for multiparameter models. Whenever needed, the special case of a single parameter is presented and derivations are sometimes made at this level to help understanding. Also, most of the examples are presented at this level. It is hoped that they will provide motivation for the usefulness of the results in the more general setting.

Presentation of results for the two main schools of thought are made in parallel as much as possible. Estimation and prediction are introduced with the Bayesian approach followed by the classical approach. The presentation of hypothesis testing and linear models goes the other way round. All chapters, including this introduction, contain a set of exercises at the end. These are included to help the student practice his/her knowledge of the material covered in the book. The exercises are divided according to the section of the chapter they refer to, even though many exercises contain a few items which cover a few different sections. At the end of the book, a selection of exercises from all chapters have their solution presented. We tried to spread these exercises evenly across the material contained in the chapter.

The material of the book will be briefly presented below. The book is composed of eight chapters that can broadly be divided into three parts. The first part contains the first three chapters introducing basic concepts needed for statistics. The second part is composed of Chapters 4, 5 and 6 which discuss in an integrated way the standard topics of estimation and hypothesis testing. The final two chapters deal with other important topics of inference: prediction and linear models.

Chapter 1 consisted of an introduction with the aim of providing the flavour of and intuition for the task ahead. In doing so, it anticipated at an elementary level many of the points to be addressed later in the book.

Chapter 2 presents the main ingredients used in statistical inference. The key concepts of likelihood, sufficiency, posterior distribution, exponential family, Fisher information and exchangeability are introduced here. The issue of

16 Introduction

parameter elimination leading to marginal, conditional and profile likelihoods is presented here too.

A key element of statistical inference for the Bayesian approach is the use of prior distributions. These are separately presented and discussed in Chapter 3. Starting from an entirely subjective specification, we move on to functional form specifications where conjugate priors play a very important role. Then, non-informative priors are presented and illustrated. Finally, the structuring of a prior distribution in stages with the so-called hierarchical form is presented.

Chapter 4 deals with parameter estimation where, intentionally, point and interval estimation are presented as different responses to the same summarization question, and not as two unrelated procedures. Different methods of estimation (maximum likelihood, method of moments, least squares) are presented. The classical results in estimation are shown to numerically coincide with Bayesian results obtained using vague prior distributions in many of the problems considered for the normal observational model.

Chapter 5 deals with approximate and computationally intensive methods of inference. These results are useful when an explicit analytic treatment is not available. Maximization techniques including Newton–Raphson, Fisher scoring and EM algorithms are presented. Asymptotic theory is presented and includes the delta method and Laplace approximations. Quadrature integration rules are also presented here. Finally, methods based on simulation are presented. They include bootstrap and its Bayesian versions, Monte Carlo integration and MCMC methods.

Chapter 6 is about hypothesis testing problems under the frequentist approach and also under the various forms of the Bayesian paradigm. Various test procedures are presented and illustrated for the models with normal observations. Tests based on the asymptotic results of the previous chapter are also presented.

Chapter 7 deals with prediction of unknown quantities to be observed. The prediction analysis is covered from the classical and Bayesian point of view. Linear models are briefly introduced here and provide an interesting example of prediction. This chapter also includes linear Bayes methods by relating them to prediction in linear models.

Chapter 8 deals with linear models. Initially, the frequentist inference for linear models is presented, followed by the Bayesian one. Generalizations based on the Bayesian approach are presented leading to hierarchical and dynamic models. Also, a brief introduction to generalized linear models is presented.

Exercises

§ 1.2

1. Consider the equation $P(A \cap F) = P(A | F)P(F)$ in the light of the de Finetti loss function setup with the three losses associated with events $A = 0, F = 1, A = 1, F = 1$ and $F = 0$ and respective probabilities p, q and r . Show that losses are all minimized when $p = r/q$.

Exercises 17

§ 1.3

2. Consider the example of the pregnant lady.
 - Show that by proceeding on Z with the reasoning used for Y leads to the estimation of θ by $5/6$.
 - Evaluate the probabilities that $Y \geq 5$ and $Z \geq 5$ and show that the values of these probabilities are respectively given by 0.109 and 0.031.
3. Consider again the example of the pregnant lady. Repeat the evaluation of the $P(X_{r+s+1} = 1 | (r, s))$ assuming now that
 - the observed values of (r, s) are $(15, 3)$ using beta priors with parameters $(a, b) = (1, 1)$ and $(a, b) = (5, 5)$. Compare the results obtained.
 - it is known that her 7th pregnancy will produce twins and the observed value of (r, s) is $(5, 1)$.

§ 1.4

4. Let $X|Y \sim \text{bin}(Y, \pi)$ and let $Y \sim \text{Pois}(\lambda)$.
 - Show that $E(X) = E[E(X|Y)] = \lambda\pi$ and that
$$V(X) = E[V(X|Y)] + V[E(X|Y)] = \lambda\pi.$$
 - Show that $X \sim \text{Pois}(\lambda\pi)$ and that $Y - X|X \sim \text{Pois}[\lambda(1 - \pi)]$.
5. Let $X|Y \sim N(0, Y^{-1})$ and $Y \sim G(a/2, b/2)$. Obtain the marginal distribution of X and the conditional distribution of $Y|X$.
6. Show that normality is preserved under linear transformations, i.e., if $\mathbf{X} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ and $\mathbf{Y} = \mathbf{c} + \mathbf{C}\mathbf{X}$ then $\mathbf{Y} \sim N(\mathbf{c} + \mathbf{C}\boldsymbol{\mu}, \mathbf{C}\boldsymbol{\Sigma}\mathbf{C}')$. Apply the result to obtain the marginal distribution of any subvector of \mathbf{X} .
7. Show that if

$$\begin{pmatrix} \mathbf{X} \\ \mathbf{Y} \end{pmatrix} \sim N \left[\begin{pmatrix} \boldsymbol{\mu}_X \\ \boldsymbol{\mu}_Y \end{pmatrix}, \begin{pmatrix} \boldsymbol{\Sigma}_X & \boldsymbol{\Sigma}_{XY} \\ \boldsymbol{\Sigma}'_{XY} & \boldsymbol{\Sigma}_Y \end{pmatrix} \right]$$

then $\mathbf{X}|Y \sim N(\boldsymbol{\mu}_{X|Y}, \boldsymbol{\Sigma}_{X|Y})$ where $\boldsymbol{\mu}_{X|Y} = \boldsymbol{\mu}_X + \boldsymbol{\Sigma}_{XY}\boldsymbol{\Sigma}_Y^{-1}(\mathbf{Y} - \boldsymbol{\mu}_Y)$ and $\boldsymbol{\Sigma}_{X|Y} = \boldsymbol{\Sigma}_X - \boldsymbol{\Sigma}_{XY}\boldsymbol{\Sigma}_Y^{-1}\boldsymbol{\Sigma}_{XY}$.

8. Show that if X_1, \dots, X_p are independent standard normal variables then

$$\sum_{i=1}^p X_i^2 \sim \chi_p^2.$$

9. Show that if $\mathbf{X} = (X_1, \dots, X_p)' \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ and $\mathbf{Y} = (\mathbf{X} - \boldsymbol{\mu})'\boldsymbol{\Sigma}^{-1}(\mathbf{X} - \boldsymbol{\mu})$ then $\mathbf{Y} \sim \chi_p^2$.
Hint: Define $\mathbf{Z} = \mathbf{A}(\mathbf{X} - \boldsymbol{\mu})$ where the matrix \mathbf{A} satisfies $\mathbf{A}'\mathbf{A} = \boldsymbol{\Sigma}^{-1}$ and use the result from the previous exercise.
10. Let \mathbf{A} , and \mathbf{B} be non-singular symmetric matrices of orders p and q and \mathbf{C} a $p \times q$ matrix. Show that

$$(a) (A + CBC')^{-1} = A^{-1} - A^{-1}C(C'A^{-1}C + B^{-1})^{-1}C'A^{-1}$$

$$(b) \begin{pmatrix} A & C \\ C' & B \end{pmatrix}^{-1} = \begin{pmatrix} A^{-1} + ED^{-1}E' & -ED^{-1} \\ -D^{-1}E' & D^{-1} \end{pmatrix}$$

where $D = B^{-1} - C'A^{-1}C$ and $E = A^{-1}C$.

2

Elements of inference

In this chapter, the basic concepts needed for the study of Bayesian and classical statistics will be described. In the first section, the most commonly used statistical models are presented. They will provide the basis for the presentation of most of the material of this book. Section 2.2 introduces the fundamental concept of likelihood function. A theoretically sound and operationally useful definition of measures of information is also given in this Section. The Bayesian point of view is introduced in Section 2.3. The Bayes' theorem acts as the basic rule in this inferential procedure. The next section deals with the concept of exchangeability. This is a strong and useful concept as will be seen in the following chapter. Other basic concepts, like sufficiency and the exponential family, are presented in Section 2.5. Finally, in Section 2.6, the multiparametric case is presented and the main concepts are revised and extended from both the Bayesian and the classical points of view. Particular attention is given to the problem of parameter elimination in order to make inference with respect to the remaining parameters.

2.1 Common statistical models

Although the nature of statistical applications is only limited by our ability to formulate probabilistic models, there are a few models that are more frequently used in statistics. There is a number of reasons for it: first, because they are more commonly found in applications; second, because they are the simplest models that can be entertained in non-trivial applications; finally, because they provide a useful starting point in the process of building up models.

The first class of models considered is a random sample from a given distribution. They are followed by the location model, the scale model and the location-scale model. Excellent complementary reading for this topic is the book of Bickel and Doksum (1977).

The most basic situation of observations is the case of a homogeneous population from a distribution F_θ , depending on the unknown quantity θ . Knowledge of the value of θ is vital for the understanding and description of this population and we would need to extract information from it to accomplish this task. Typically in this case, a random sample X_1, \dots, X_n is drawn from this population and we hope to

build strategies to ascertain the value of θ from the values of the sample.

In this case, the observations X_1, \dots, X_n are independent and identically distributed (iid, in short) with common distribution F_θ . Assuming that F_θ has density or probability function f , they are probabilistically described through

$$p(x_1, \dots, x_n | \theta) = \prod_{i=1}^n f(x_i | \theta).$$

Example. Consider a series of measurements made about an unknown quantity θ . Unfortunately, measurements are made with imprecise devices which means that there are errors that should be taken into account. These errors are a result of many (small) contributions and are more effectively described in terms of a probability distribution. This leads to the construction of a model in the form $X_i = \theta + e_i$, $i = 1, \dots, n$. The e_i 's represent the measurement errors involved. If the experiment is performed with care with measurements being collected independently and using the same procedures, the e_i 's will form a random sample from the distribution of errors F_e . For the same reason, the X_i 's will also be iid with joint density

$$p(x_1, \dots, x_n | \theta) = \prod_{i=1}^n f(x_i | \theta) = \prod_{i=1}^n f_e(x_i - \theta)$$

where f_e is the density of the error distribution.

Definition (Location model). X has a location model if a function f and a quantity θ exist such that the distribution of X given θ satisfies $p(x | \theta) = f(x - \theta)$. In this case, θ is called a location parameter.

Examples.

1. Normal with known variance

In this case the density is $p(x | \theta) = (2\pi\sigma^2)^{-0.5} \exp\{-0.5\sigma^{-2}(x - \theta)^2\}$, which is a function of $x - \theta$.

2. Cauchy with known scale parameter

The Cauchy distribution is the Student t distribution with 1 degree of freedom. In this case, the density is $p(x | \theta) = \{\pi\sigma[1 + (x - \theta)\sigma^2]\}^{-1}$, which is a function of $x - \theta$, too.

3. Multivariate normal with known variance–covariance matrix

In this case the density is

$$p(\mathbf{x} | \theta) = (2\pi)^{-p/2} |\Sigma|^{-1/2} \exp\{-(\mathbf{x} - \theta)' \Sigma^{-1} (\mathbf{x} - \theta)/2\},$$

which is a function of $\mathbf{x} - \theta$. Note that an iid sample from the $N(\theta, \sigma^2)$ distribution is a special case with $\theta = \theta\mathbf{1}$ and $\Sigma = \sigma^2\mathbf{I}$.

Definition (Scale model). X has a scale model if a function f and a quantity σ exist such that the distribution of X is given by

$$p(x | \sigma) = \frac{1}{\sigma} f\left(\frac{x}{\sigma}\right)$$

In this case, σ is called a scale parameter.

Examples.

1. Exponential with parameter θ

The density $p(x | \theta) = \theta \exp(-\theta x)$ is in the form $\sigma^{-1} f(x/\sigma)$ with $\theta = \sigma^{-1}$ and $f(u) = e^{-u}$.

2. Normal with known mean θ

The density $p(x | \sigma^2) = (2\pi)^{-1/2} \sigma^{-1} \exp\{-(x - \theta)^2 / 2\}$ is in the form $\sigma^{-1} f(x/\sigma)$.

Definition (Location and scale model). X has location and scale model if there are a function f and quantities θ and σ such that the distribution of X given (θ, σ) satisfies

$$p(x | \theta, \sigma) = \frac{1}{\sigma} f\left(\frac{x - \theta}{\sigma}\right).$$

In this case, θ is called the location–parameter and σ the scale parameter.

Some examples in the location–scale family are the normal and the Cauchy distributions. Once again, the location part of the model can also be multivariate.

2.2 Likelihood-based functions

Most statistical work is based on functions that are constructed from the probabilistic description of the observations. In this section, these functions are defined and their relevance to statistical inference is briefly introduced. These functions will be heavily used in later chapters, where their importance will be fully appreciated. We start with the likelihood function and then present Fisher measures of information and the score function.

2.2.1 Likelihood function

The likelihood function of θ is the function that associates the value $p(\mathbf{x} | \theta)$ to each θ . This function will be denoted by $l(\theta; \mathbf{x})$. Other common notations are $l_{\mathbf{x}}(\theta)$, $l(\theta | \mathbf{x})$ and $l(\theta)$. It is defined as follows

$$\begin{aligned} l(\cdot; \mathbf{x}) : \Theta &\rightarrow R^+ \\ \theta &\rightarrow l(\theta; \mathbf{x}) = p(\mathbf{x} | \theta). \end{aligned}$$

The likelihood function associates to each value of θ , the probability of an observed value \mathbf{x} for X . Then, the larger the value of l the greater are the chances associated to the event under consideration, using a particular value of θ . Therefore, by fixing the value of \mathbf{x} and varying θ we observe the plausibility (or likelihood) of each value of θ . The concept of likelihood function was discussed by Fisher,

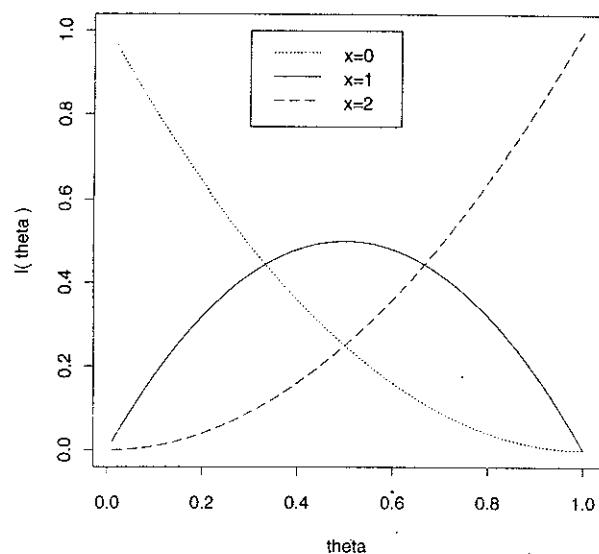


Fig. 2.1 Likelihood function of the example for different values of x .

Barnard and Kalbeifesh among many others. The likelihood function is also of fundamental importance in many theories of statistical inference. Note that even though $\int_R p(x | \theta) dx = 1$, $\int_{\Theta} l(\theta; x) d\theta = k \neq 1$, in general.

Example. $X \sim \text{bin}(2, \theta)$

$$p(x | \theta) = l(\theta; x) = \binom{2}{x} \theta^x (1 - \theta)^{2-x} \quad x = 0, 1, 2; \theta \in \Theta = (0, 1)$$

but

$$\int_{\Theta} l(\theta; x) d\theta = \binom{2}{x} \int_0^1 \theta^x (1 - \theta)^{2-x} d\theta = \binom{2}{x} B(x+1, 3-x) = \frac{1}{3} \neq 1.$$

Note that:

1. If $x = 1$ then $l(\theta; x = 1) = 2\theta(1 - \theta)$. The value of θ with highest likelihood or, in other words, the most likely (or probable) value of θ is $1/2$.
2. If $x = 2$ then $l(\theta; x = 2) = \theta^2$, the most likely value of θ is 1 .
3. If $x = 0$ then $l(\theta; x = 0) = (1 - \theta)^2$, the most likely value is again 0 .

These likelihood functions are plotted in Figure 2.1.

The notion of likelihood can also be introduced from a slightly more general perspective. This broader view will be useful in more general observation contexts

than those considered so far. These include cases where observations are only obtained in an incomplete way.

Denoting by E an observed event and assuming the probabilistic description of E depends on an unknown quantity θ , one can define the likelihood function of θ based on the observation E as $l(\theta; E) \propto \Pr(E|\theta)$ where the symbol \propto is to be read as *is proportional to*. If E is of a discrete nature, there is no difference with respect to the previous definition.

Example. Let X_1, \dots, X_n be a collection of iid random variables with a common Bernoulli distribution with success probability θ . Let E be any observation of the X_1, \dots, X_n consisting of x successes. Then, $l(\theta; E) \propto \Pr(E|\theta) = \theta^x (1 - \theta)^{n-x}$.

For continuous observations, assume that any observed value x is an approximation of the real value due to rounding errors. Therefore, the observation x in fact corresponds to the observed event $E = \{x : a \leq x \leq a + \Delta\}$ for given values of a and $\Delta > 0$, which do not depend on θ . In this case, $\Pr(E|\theta) = F(a + \Delta) - F(a)$ where F is the distribution function of the observation. For typical applications, the value of Δ is very small and one can approximate $F(a + \Delta) - F(a) = p(x|\theta)\Delta$. Therefore, $l(\theta; E) \propto p(x|\theta)$. This definition can be extended for a vector of observations with minor, technical adaptations to the reasoning above to lead to $l(\theta; E) \propto p(x|\theta)$.

The likelihood function leads to the likelihood principle, which states that all the information provided by the experiment X is summarized by the likelihood function. This principle draws a clear line that separates inference schools. On the same side lie the Bayesian and likelihood approaches, that do not violate this principle, and on the other side the frequentist approach which is based on the probabilities implied by the sampling distribution of X . In this way, it takes into consideration all the possible values of X .

2.2.2 Fisher information

We have already mentioned that the understanding and measuring of information is one of the key aspects of the statistical activity. In this section, the most commonly accepted measures of information are introduced. They have important connections with the notion of sufficiency, to be defined later in this chapter, and will prove to be very useful in the sequel. In fact, they consist of a series of related measures generally known as Fisher information measures.

Definition. Let X be a random vector with probability (density) function $p(x | \theta)$. The expected Fisher information measure of θ through X is defined by

$$I(\theta) = E_{X|\theta} \left[-\frac{\partial^2 \log p(X | \theta)}{\partial \theta^2} \right].$$

If $\theta = (\theta_1, \dots, \theta_p)$ is a parametric vector then the Fisher expected information matrix of θ through \mathbf{X} can be defined by

$$\mathbf{I}(\theta) = E_{\mathbf{X}|\theta} \left[-\frac{\partial^2 \log p(\mathbf{X} | \theta)}{\partial \theta \partial \theta'} \right]$$

with elements $I_{ij}(\theta)$ given by

$$I_{ij}(\theta) = E_{\mathbf{X}|\theta} \left[-\frac{\partial^2 \log p(\mathbf{X} | \theta)}{\partial \theta_i \partial \theta_j} \right], \quad i, j = 1, \dots, p.$$

The information measure defined this way is related to the mean value of the curvature of the likelihood. The larger this curvature is, the larger is the information content summarized in the likelihood function and so the larger will $I(\theta)$ be. Since the curvature is expected to be negative, the information value is taken as minus the curvature. The expectation is taken with respect to the sample distribution. The **observed** Fisher information corresponds to minus the second derivative of the log likelihood:

$$J_{\mathbf{X}}(\theta) = -\frac{\partial^2 \log p(\mathbf{x} | \theta)}{\partial \theta \partial \theta'}$$

and is interpreted as a local measure of the information content while its expected value, the expected Fisher information, is a global measure. Both $J_{\mathbf{X}}(\theta)$ and $I(\theta)$ will be used later on in connection with Bayesian and frequentist estimation. The motivation for these definitions will be clear from the following example.

Example. Let $X \sim N(\theta, \sigma^2)$, with σ^2 unknown. It is easy to get $I(\theta) = J_x(\theta) = \sigma^{-2}$, the normal precision. Then, the observed and expected Fisher information measures with respect to θ obtained from the observation X coincide with the precision, which we had previously tried to identify with information.

There are many properties that can be derived from the Fisher information. One of the most useful ones concerns the additivity of the information with respect to independent observations, or more generally, sources of information.

Lemma. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a collection of independent random variables with distribution $p_i(x_i | \theta)$, $i = 1, \dots, n$. Let $J_{\mathbf{X}}$ and J_{x_i} be the observed information measures obtained through \mathbf{X} and X_i , $i = 1, \dots, n$, respectively. Let I and I_i be the expected information measures obtained through \mathbf{X} and X_i , $i = 1, \dots, n$, respectively. Then,

$$J_{\mathbf{X}}(\theta) = \sum_{i=1}^n J_{x_i}(\theta) \quad \text{and} \quad I(\theta) = \sum_{i=1}^n I_i(\theta).$$

The lemma states that the total information obtained from independent observations is the sum of the information of the individual observations. This provides

further intuition about the appropriateness of the Fisher measures as actual summaries of information.

Proof. $p(\mathbf{X} | \theta) = \prod_{i=1}^n p_i(X_i | \theta)$ and therefore $\log p(\mathbf{X} | \theta) = \sum_{i=1}^n \log p_i(X_i | \theta)$. Then,

$$-\frac{\partial^2 \log p(\mathbf{X} | \theta)}{\partial \theta \partial \theta'} = -\sum_{i=1}^n \frac{\partial^2 \log p_i(X_i | \theta)}{\partial \theta \partial \theta'}$$

which proves the result about observed information. Taking expectation with respect to $\mathbf{X} | \theta$ on both sides, gives

$$\begin{aligned} I(\theta) &= E \left[-\sum_{i=1}^n \frac{\partial^2 \log p_i(X_i | \theta)}{\partial \theta \partial \theta'} \mid \theta \right] \\ &= \sum_{i=1}^n E \left[-\frac{\partial^2 \log p_i(X_i | \theta)}{\partial \theta \partial \theta'} \mid \theta \right] \\ &= \sum_{i=1}^n I_i(\theta). \end{aligned}$$

□

Another very important statistic involved in the study of the likelihood function is the score function.

Definition. The score function of X , denoted as $U(X; \theta)$, is defined as

$$U(X; \theta) = \frac{\partial \log p(X | \theta)}{\partial \theta}.$$

In the case of a parametric vector $\theta = (\theta_1, \dots, \theta_p)^T$, the score function is also a vector $\mathbf{U}(X; \theta)$ with components $U_i(X; \theta) = \partial \log p(X | \theta) / \partial \theta_i$, $i = 1, \dots, p$.

The score function is very relevant for statistical inference as will be shown in the next chapters. The following lemma shows an alternative way to obtain the Fisher information based on the score function.

Lemma. Under certain regularity conditions,

$$I(\theta) = E_{\mathbf{X}|\theta}[U^2(\mathbf{X}; \theta)].$$

In the case of a vector parameter θ , the result becomes

$$\mathbf{I}(\theta) = E_{\mathbf{X}|\theta}[\mathbf{U}(\mathbf{X}; \theta)\mathbf{U}'(\mathbf{X}; \theta)].$$

Although we shall not go into the technical details of the regularity conditions, the main reasons for their presence is to ensure that differentiation of the likelihood can be performed over the entire parameter space and integration and differentiation can be interchanged.

Proof. Using the equality $\int p(\mathbf{X} | \theta) d\mathbf{X} = 1$ and differentiating both sides with respect to θ , it follows, after interchanging the integration and differentiation operators, that

$$\begin{aligned} \mathbf{0} &= \int \frac{\partial p(\mathbf{X} | \theta)}{\partial \theta} d\mathbf{X} = \int \frac{1}{p(\mathbf{X} | \theta)} \frac{\partial p(\mathbf{X} | \theta)}{\partial \theta} p(\mathbf{X} | \theta) d\mathbf{X} \\ &= \int \frac{\partial \log p(\mathbf{X} | \theta)}{\partial \theta} p(\mathbf{X} | \theta) d\mathbf{X}. \end{aligned}$$

Therefore the score function has expected value equal to a zero vector. Differentiating again with respect to θ and interchanging integration and differentiation we have

$$\begin{aligned} \mathbf{0} &= \int \frac{\partial \log p(\mathbf{X} | \theta)}{\partial \theta} \left[\frac{\partial p(\mathbf{X} | \theta)}{\partial \theta} \right]' d\mathbf{X} + \int \frac{\partial^2 \log p(\mathbf{X} | \theta)}{\partial \theta \partial \theta'} p(\mathbf{X} | \theta) d\mathbf{X} \\ &= \int \left[\frac{\partial \log p(\mathbf{X} | \theta)}{\partial \theta} \right] \left[\frac{\partial \log p(\mathbf{X} | \theta)}{\partial \theta} \right]' p(\mathbf{X} | \theta) d\mathbf{X} - I(\theta). \end{aligned}$$

The result follows straightforwardly. \square

2.3 Bayes' theorem

We have seen that the statistical inference problem can be stated as having an unknown, unobserved quantity of interest θ assuming values in a set denoted by Θ . θ can be a scalar, a vector or a matrix. Until now, the only relevant source of inference was provided by the probabilistic description of the observations. In this section, we will formalize the use of other sources of information in statistical inference. This will define the Bayesian approach to inference.

Let H (for history) denote the initial available information about some parameter of interest. Assume further that this initial information is expressed in probabilistic terms. It can then be summarized through $p(\theta | H)$ and, if the information content of H is enough for our inferential purpose, this is all that is needed. In this case the description of our uncertainty about θ is complete.

Depending upon the relevance of the question we are involved with, H may not be sufficient and, in this case, it must be augmented. The main tool used in this case is experimentation. Assume a vector of random quantities \mathbf{X} related to θ can be observed thus providing further information about θ . (If \mathbf{X} is not random then a functional relationship relating it to θ should exist. We can then evaluate the value of θ and the problem is trivially solved.) Before observing \mathbf{X} , we should know the sampling distribution of \mathbf{X} given by $p(\mathbf{X} | \theta, H)$, where the dependence on θ , central to our argument, is clearly stated. After observing the value of \mathbf{X} , the amount of information we have about θ has changed from H to $H^* = H \cap \{\mathbf{X} = \mathbf{x}\}$. In fact, H^* is a subset of H (a refinement on H was performed).

Now the information about θ is summarized by $p(\theta | \mathbf{x}, H)$ and the only remaining question left is how to pass from $p(\theta | H)$ to $p(\theta | \mathbf{x}, H)$. From Section 1.4, one can write

$$p(\theta | \mathbf{x}, H) = \frac{p(\theta, \mathbf{x} | H)}{p(\mathbf{x} | H)} = \frac{p(\mathbf{x} | \theta, H)p(\theta | H)}{p(\mathbf{x} | H)}$$

where

$$p(\mathbf{x} | H) = \int_{\Theta} p(\mathbf{x}, \theta | H) d\theta.$$

The result presented above is known as Bayes' theorem. This theorem was introduced by the Rev. Thomas Bayes in two papers in 1763 and 1764, published after his death, as mentioned in Barnett (1973). As we can see the function in the denominator does not depend upon θ and so, as far as the quantity of interest θ is concerned, it is just a constant. Therefore, Bayes' theorem can be rewritten in its more usual form

$$p(\theta | \mathbf{x}) \propto p(\mathbf{x} | \theta)p(\theta).$$

The dependence on H is dropped, for simplicity of notation, since it is a common factor to all the terms. Nevertheless, it should not be forgotten. The above formula is valid for discrete and continuous, scalar, vector and matrix quantities. The theorem provides a rule for updating probabilities about θ , starting from $p(\theta)$ and leading to $p(\theta | \mathbf{x})$. This is the reason why the above distributions are called prior and posterior distributions, respectively.

To recover the removed constant in the former equation, it is enough to notice that densities must integrate to 1 and to rewrite it as

$$p(\theta | \mathbf{x}) = kp(\mathbf{x} | \theta)p(\theta)$$

where

$$1 = \int_{\Theta} p(\theta | \mathbf{x}) d\theta = k \int_{\Theta} p(\mathbf{x} | \theta)p(\theta) d\theta$$

and hence

$$\begin{aligned} k^{-1} &= p(\mathbf{x} | H) = \int_{\Theta} p(\mathbf{x} | \theta)p(\theta) d\theta \\ &= E_{\theta}[p(\mathbf{x} | \theta)]. \end{aligned}$$

This is the predictive (or marginal) distribution of \mathbf{X} . As before, after removing the dependence on H , it can be denoted by $p(\mathbf{x})$. This is the expected distribution of \mathbf{X} (under the prior) and it behaves like a prediction, for a given H . So, before observing \mathbf{X} it is useful to verify the prior adequacy through the predictions it provides for \mathbf{X} . After observing \mathbf{X} , it serves to test the model as a whole. An observed value of \mathbf{X} with a low predictive probability is an indication that the stated model is not providing good forecasts. This is evidence that something unexpected happened. Either the model must be revised or an aberrant observation occurred.

2.3.1 Prediction

Another relevant aspect that follows from the calculations presented above is that we obtain an automatic way to make predictions for future observations. If we want to predict \mathbf{Y} , whose probabilistic description is $P(\mathbf{Y} | \theta)$, we have

$$\begin{aligned} p(\mathbf{y} | \mathbf{x}) &= \int_{\Theta} p(\mathbf{y}, \theta | \mathbf{x}) d\theta \\ &= \int_{\Theta} p(\mathbf{y} | \theta, \mathbf{x}) p(\theta | \mathbf{x}) d\theta \\ &= \int_{\Theta} p(\mathbf{y} | \theta) p(\theta | \mathbf{x}) d\theta, \end{aligned}$$

where the last equality follows from the independence between \mathbf{X} and \mathbf{Y} , once θ is given. This conditional independence assumption is typically present in many statistical problems. Also, it follows from the last equation that

$$p(\mathbf{y} | \mathbf{x}) = E_{\theta | \mathbf{x}}[p(\mathbf{y} | \theta)].$$

It is always useful to concentrate on prediction rather than on estimation because the former is verifiable. The reason for the difference is that \mathbf{Y} is observable and θ is not. This concept can be further explored by reading the books of Aitchison and Dunsmore (1975) and Geisser (1993).

Example. John goes to the doctor claiming some discomfort. The doctor is led to believe that he may have the disease A. He then takes some standard procedures for this case: he examines John, carefully observes the symptoms and prescribes routine laboratory examinations.

Let θ be the unknown quantity of interest indicating whether John has disease A or not. The doctor assumes that $P(\theta = 1 | H) = 0.7$. H in this case contains the information John gave him and all other relevant knowledge he has obtained from former patients. To improve the evidence about the illness, the doctor asks John to undertake an examination. Examination X is related to θ and provides an uncertain result, of the positive/negative type, with the following probability distribution

$$\begin{cases} P(X = 1 | \theta = 0) = 0.40, & \text{positive test without disease} \\ P(X = 1 | \theta = 1) = 0.95, & \text{positive test with disease.} \end{cases}$$

Suppose that John goes through the examination and that its result is $X = 1$. So, for the doctor, the probability that John has the disease is

$$\begin{aligned} P(\theta = 1 | X = 1) &\propto l(\theta = 1; X = 1) P(\theta = 1) \\ &\propto (0.95)(0.7) = 0.665 \end{aligned}$$

and the probability that he does not have the disease is

$$\begin{aligned} P(\theta = 0 | X = 1) &\propto l(\theta = 0; X = 1) P(\theta = 0) \\ &\propto (0.40)(0.30) = 0.120. \end{aligned}$$

The normalizing constant, such that the total probability adds to 1, is calculated so that $k(0.665) + k(0.120) = 1$ and $k = 1/0.785$. Consequently

$$P(\theta = 1 | X = 1) = 0.665/0.785 = 0.847$$

and

$$P(\theta = 0 | X = 1) = 0.120/0.785 = 0.153.$$

The information $X = 1$ increases, for the doctor, the probability that John has the disease A from 0.7 to 0.847. This is not too much since the probability that the test would give a positive result even if John were not ill was reasonably high. So the doctor decides to ask John to undertake another test Y , again of the positive/negative type, where

$$\begin{cases} P(Y = 1 | \theta = 0) = 0.04 \\ P(Y = 1 | \theta = 1) = 0.99. \end{cases}$$

Note that the probability of this new test yielding a positive result given that John doesn't have the illness is very small. Although this test might be more expensive, its results are more efficient.

The posterior distribution of θ given X , $P(\theta | X)$, will be the prior distribution for the Y test. Before observing the result of test Y , it is useful to ask ourselves what will the predictive distribution be, that is, what are the values of $P(Y = y | X = 1)$, for $y = 0, 1$. As we have already seen, in the discrete case,

$$p(Y = y | X = x) = \sum_{j=0}^1 p(Y = y | \theta = j) p(\theta = j | X = x)$$

and so

$$P(Y = 1 | X = 1) = (0.04)(0.153) + (0.99)(0.847) = 0.845$$

and

$$P(Y = 0 | X = 1) = 1 - P(Y = 1 | X = 1) = 0.155.$$

Let us suppose now that John undertook test Y and the observed result was $Y = 0$. This is a reasonably unexpected result as the doctor only gave it a 15.5% chance. He should rethink the model based on this result. In particular, he might want to ask himself

1. Did 0.7 adequately reflect his $P(\theta = 1 | H)$?
2. Is test X really so unreliable? Is the sample distribution of X correct?
3. Is the test Y so powerful?
4. Have the tests been carried out properly?

30 Elements of inference

In any case, the doctor has to re-evaluate the chances of $\theta = 1$ considering that as well as the knowledge that $X = 1$ he now also knows that $Y = 0$. By Bayes' theorem

$$\begin{aligned} P(\theta = 1 | X = 1, Y = 0) &\propto l(\theta = 1; Y = 0)P(\theta = 1 | X = 1) \\ &\propto (0.01)(0.847) = 0.008 \end{aligned}$$

and

$$\begin{aligned} P(\theta = 0 | X = 1, Y = 0) &\propto l(\theta = 0; Y = 0)P(\theta = 0 | X = 1) \\ &\propto (0.96)(0.155) = 0.147. \end{aligned}$$

Note that now $P(\theta | X = 1)$ is working as the prior distribution for the experiment Y . The normalizing constant is $0.008 + 0.147 = 0.155$ and

$$P(\theta = 1 | Y = 0, X = 1) = 0.008/0.155 = 0.052$$

and

$$P(\theta = 0 | Y = 0, X = 1) = 0.147/0.155 = 0.948.$$

Summarizing the doctor's findings chronologically,

$$P(\theta = 1) = \begin{cases} 0.7, & \text{before the tests } X \text{ and } Y \\ 0.847, & \text{after } X \text{ and before } Y \\ 0.052, & \text{after } X \text{ and } Y. \end{cases}$$

The doctor can then decide his course of action as these updating operations took place.

2.3.2 Sequential nature of Bayes' theorem

Bayes' theorem is nothing more than a rule for updating probabilities. From an experimental result X_1 with probability distribution $p_1(x_1 | \theta)$ (and consequently, likelihood $l_1(\theta; x_1)$), it follows that

$$p(\theta | x_1) \propto l_1(\theta; x_1)p(\theta).$$

After observing another experimental result X_2 with probability $p_2(x_2 | \theta)$ not depending on X_1 , we have

$$\begin{aligned} p(\theta | x_2, x_1) &\propto l_2(\theta; x_2)p(\theta | x_1) \\ &\propto l_2(\theta; x_2)l_1(\theta; x_1)p(\theta). \end{aligned}$$

Repeating this procedure n times, after observing X_3, \dots, X_n related to θ through the observational distributions $p_i(x_i | \theta)$, for $i = 3, \dots, n$, we get

$$p(\theta | x_n, \dots, x_1) \propto l_n(\theta; x_n)p(\theta | x_{n-1}, \dots, x_1)$$

Bayes' theorem 31

or alternatively

$$p(\theta | x_n, x_{n-1}, \dots, x_1) \propto \left[\prod_{i=1}^n l_i(\theta; x_i) \right] p(\theta)$$

and it is not difficult to see that the order in which the observations are processed by the theorem is irrelevant. In fact, the observations can be processed in one batch, on an individual basis or through disjoint subgroups. The final result is always the same as long as the observations are conditionally independent (on θ). The sequential nature of Bayesian inference is deeply explored by Lindley (1965).

Another basic result corresponds to the case of normal observations with unknown mean, which is used in many practical situations. If the mean is described by a normal prior distribution, the posterior distribution will also be a normal distribution.

Theorem 2.1 (Normal prior and observation). Let $\theta \sim N(\mu, \tau^2)$ and $X | \theta \sim N(\theta, \sigma^2)$, with σ^2 known. Therefore, the posterior distribution of θ is $(\theta | X = x) \sim N(\mu_1, \tau_1^2)$ where

$$\mu_1 = \frac{\tau^{-2}\mu + \sigma^{-2}x}{\tau^{-2} + \sigma^{-2}} \quad \text{and} \quad \tau_1^{-2} = \tau^{-2} + \sigma^{-2}.$$

Defining the precision as the reciprocal of the variance, it follows from the theorem that the posterior precision is the sum of the prior and likelihood precisions and does not depend on x . Interpreting the precision as a measure of information and defining $w = \tau^{-2}/(\tau^{-2} + \sigma^{-2}) \in (0, 1)$, w measures the relative information contained in the prior distribution with respect to the total information (prior plus likelihood). Then one can write

$$\mu_1 = w\mu + (1 - w)x$$

which is the weighted mean of prior and likelihood means.

Proof. From Bayes' theorem, it follows that

$$\begin{aligned} p(\theta | x) &\propto l(\theta; x)p(\theta) \\ &\propto \exp \left\{ -\frac{1}{2\sigma^2}(x - \theta)^2 - \frac{1}{2\tau^2}(\theta - \mu)^2 \right\} \\ &\propto \exp \left\{ -\frac{\theta^2}{2\sigma^2} - \frac{\theta^2}{2\tau^2} + \frac{x\theta}{\sigma^2} + \frac{\mu\theta}{\tau^2} \right\} \\ &= \exp \left\{ -\frac{\theta^2}{2} \left(\frac{1}{\sigma^2} + \frac{1}{\tau^2} \right) + \theta \left(\frac{x}{\sigma^2} + \frac{\mu}{\tau^2} \right) \right\}. \end{aligned}$$

where in the first step all constants involved were included in the proportionality factor. Now let $\tau_1^2 = (\tau^{-2} + \sigma^{-2})^{-1}$ and $\mu_1 = (\sigma^{-2}x + \mu\tau^{-2})\tau_1^2$. Substituting

into the above expression gives

$$\begin{aligned} p(\theta | x) &\propto \exp \left\{ -\frac{\theta^2}{2\tau_1^2} + \frac{\theta\mu_1}{\tau_1^2} \right\} \\ &\propto \exp \left\{ -\frac{1}{2\tau_1^2}(\theta - \mu_1)^2 \right\} \\ &\propto \frac{1}{\sqrt{2\pi\tau_1^2}} \exp \left\{ -\frac{1}{2\tau_1^2}(\theta - \mu_1)^2 \right\}. \end{aligned}$$

It is easy to recognize that the last term in the above expression corresponds to a normal density. Therefore, the last proportionality constant is equal to 1 and $(\theta | x) \sim N(\mu_1, \tau_1^2)$. \square

Example (Box and Tiao, 1992). Two physicists, A and B, want to determine the value of a physical constant θ . The physicist A with large experience in the area specifies his prior as $\theta \sim N(900, (20)^2)$. On the other side, the physicist B, not so experienced on the subject, stated a more uncertain prior $\theta \sim N(800, (80)^2)$. It is easy to obtain that for the physicist A, $Pr(\theta \in (860, 940)) = 0.95$ and for the physicist B, $Pr(\theta \in (640, 960)) = 0.95$. Both physicists agree to make an evaluation of θ , denoted by X , using a calibrated device with sampling distribution $(X | \theta) \sim N(\theta, (40)^2)$. After observing the value $X = 850$, the posterior distributions of θ can be obtained using Theorem 2.1 and the values stated above as

1. physicist A: $(\theta | X = 850) \sim N(890, (17.9)^2)$
2. physicist B: $(\theta | X = 850) \sim N(840, (35.7)^2)$.

The inferential procedure of the two physicists can be summarized in Figure 2.2. It is worth noting that due to the different initial uncertainties, the same experiment provides very little additional information for A, although the uncertainty of B has been substantially reduced. Identifying precision (the inverse of the variance) with information we have that the information about θ for physicist A increases from 0.0025 to 0.00312 since the likelihood precision was 0.000625 (an increase of 25%). For physicist B, it increases from 0.000156 to 0.000781 (an increase of 400%).

2.4 Exchangeability

Exchangeability is a very important concept introduced by de Finetti (1937). It is weaker than independence but it is just as useful, as will be shown below.

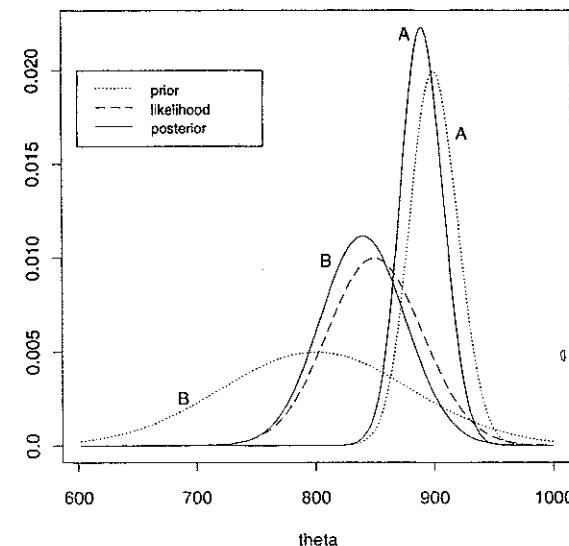


Fig. 2.2 Prior and posterior densities and likelihood for θ for the physicist's example.

Definition. Let $\mathcal{K} = \{k_1, \dots, k_n\}$ be a permutation of $\{1, \dots, n\}$. Random quantities X_1, \dots, X_n are exchangeable if the $n!$ permutations $(X_{k_1}, \dots, X_{k_n})$ have the same n -dimensional probability distribution. An infinite sequence of random quantities is exchangeable if any finite subsequence is exchangeable.

One immediate consequence of exchangeability is that all marginal distributions must be the same. To see this, consider any two distinct permutations \mathcal{K} and \mathcal{K}' of an exchangeable sequence of random variables, that therefore must have the same probability. Let k_1 and k'_1 be the first index of the two permutations, respectively. If both sides of this equality are integrated with respect to all the components but the first, one gets that the marginal distribution of X_{k_1} and $X_{k'_1}$ must be the same. Since one is free to choose the values of k_1 and k'_1 , this means that all marginal distributions are equal.

A sequence (finite or not) of iid random quantities is trivially exchangeable, although the reciprocal is not true, in general.

Examples.

1. Consider an urn with m balls, r with number 1 and $m - r$ with number 0. Balls are drawn from the urn, one at time, without replacement. Let X_k denote the number associated with the k th ball selected. Then, X_1, \dots, X_n , $n \leq m$ is an exchangeable sequence, but the X_i 's are not independent.

2. Let X_1, X_2, \dots be a sequence of Bernoulli trials with unknown success probability θ . The classical assumption is that the X_k 's are iid. For a Bayesian, if θ is unknown, the information content of the j th observation can modify your belief about the distribution of X_k . The example of the previous chapter illustrated this point. If the experiments are judged similar in some sense, the hypothesis of exchangeability is acceptable, while marginal independence is not.

The relevance of the concept of exchangeability is due to the fact that, although it is based on weak assumptions, it allows one to state a very powerful result, known as the representation theorem of de Finetti. The theorem is stated here without proof, but it can be found in de Finetti (1937).

Theorem 2.2. To all infinite sequences of exchangeable random quantities $\{X_n, n = 1, 2, \dots\}$ assuming values $\{0, 1\}$ there corresponds a distribution F in $(0, 1)$ such that for all n and $k \leq n$,

$$P[(k, n - k)] = \int_0^1 \theta^k (1 - \theta)^{n-k} dF(\theta)$$

where $(k, n - k)$ denotes the event that k of the X_i 's are 1 and the other $n - k$ of the X_i 's are 0. Note that, due to the exchangeability assumption, any sequence with k 1's and $n - k$ 0's also admits a representation according to the above theorem.

A very simple outline of the proof of this representation theorem, worth reading, can be found in Heath and Sudderth (1976). The importance of the theorem is that it provides further backing for the Bayesian approach. If one is willing to consider a collection of 0-1 observations as exchangeable then one is prepared to rephrase their model into a sampling Bernoulli model with success probability θ that itself is random with probability distribution F . The theorem however does not tell us anything about the distribution F . For example, we could have:

1. a degenerate distribution: $P(\theta = \theta_0) = 1$, for some θ_0 , implying that

$$P[(k, n - k)] = \theta_0^k (1 - \theta_0)^{n-k}$$

2. a discrete distribution: $P(\theta = \theta_i) = p_i$, if $\theta = \theta_i, i = 1, \dots, s$ with $p_i \geq 0$ and $\sum p_i = 1$, implying that

$$P[(k, n - k)] = \sum_{i=1}^s p_i \theta_i^k (1 - \theta_i)^{n-k}$$

3. a continuous distribution: $\theta \sim \text{beta}(a, b)$, implying that

$$\begin{aligned} P[(k, n - k)] &= \int_0^1 \theta^k (1 - \theta)^{n-k} \frac{\theta^{a-1} (1 - \theta)^{b-1}}{B(a, b)} d\theta \\ &= \frac{1}{B(a, b)} \int_0^1 \theta^{a+k-1} (1 - \theta)^{b+n-k-1} d\theta \\ &= \frac{B(a+k, b+n-k)}{B(a, b)}. \end{aligned}$$

The exchangeability concept has already been extended to many other distributions with the inclusion of some additional hypotheses; see Bernardo and Smith (1994) for a review of the main results. The definition is the same as presented before with removal of the constraint imposed on the sample space. In particular, if we introduce the hypothesis of symmetry of the distributions and invariance under linear transformations, it is not difficult to show that the joint density of any finite subsequence is given by

$$p(x_1, \dots, x_n) = \int_0^\infty \int_{-\infty}^\infty \prod_{i=1}^n p_N(x_i; \theta, \sigma^2) dF(\theta, \sigma^2)$$

where θ is a quantity varying in R , σ a quantity assuming values in R^+ , $p_N(\cdot; a, b)$ is the density of a $N(a, b)$ distribution and F is a distribution function in $R \times R^+$. Exchangeability, along with invariance now, leads to a representation where a normal sampling distribution is obtained with parameters having some prior distribution F . It is worth noting that F once again is not specified.

Another useful extension, well explored in recent years, is the concept of partial exchangeability. In this case the exchangeability holds only under certain conditions. For example, we can define some groups of variables where exchangeability is valid only within each group. This concept can be extended to more general cases than group classification and is formalized as follows.

Definition. Let $\{X_i, i = 1, 2, \dots, n\}$ be any sequence of random quantities and \mathcal{K} be any permutation of $\{1, 2, \dots, n\}$. We say that \mathbf{X} is partially exchangeable if there are quantities $\{Z_i, i = 1, 2, \dots, n\}$ such that the distribution of $(\mathbf{X} | \mathbf{Z})$ is the same as that of $(\mathbf{X}_{\mathcal{K}} | \mathbf{Z}_{\mathcal{K}})$, for any permutation \mathcal{K} , where for any vector $\mathbf{c} = (c_1, \dots, c_n)$, $\mathbf{c}_{\mathcal{K}} = (c_{\mathcal{K}_1}, \dots, c_{\mathcal{K}_n})$.

The main idea behind this definition is that when the indexes of the X_i 's are permuted, the resulting vector will have the same conditional distribution, as far as the same permutation is applied to the Z_i 's indexes. This is clearly a weaker concept than exchangeability. The case $Z_i = 1, \forall i$, corresponds to the exchangeability definition. Another interesting and less trivial case is when there are s groups and each Z_i takes on a value in $\{1, \dots, s\}$ to identify the group the observation X_i belongs to. In this case, one has exchangeability within each group but not globally.

The extension of this concept to countable sequences is immediate. The notion of partial exchangeability will be returned to in Chapter 3 when the related concept of hierarchical prior is introduced and in Chapter 6 when inference for hierarchical models is discussed.

2.5 Sufficiency and the exponential family

As we have said before, one of the main goals of statistics is to summarize information. An important question is to know if, given a set of observations \mathbf{X} sampled

to get information about a parameter of interest θ , there are statistics, i.e. functions of the observations \mathbf{X} , that summarize all the information contained in \mathbf{X} .

Definition (classical). Let \mathbf{X} be a random quantity with probability (density) function $p(\mathbf{x} | \theta)$. Then, the statistic $T = T(\mathbf{X})$ is sufficient for the parameter θ if

$$p(\mathbf{x} | t, \theta) = p(\mathbf{x} | t).$$

The definition states that given T, \mathbf{X} does not bring any additional information about θ . From the Bayesian point of view this means that \mathbf{X} and θ are independent conditionally on T . The main point of the definition is that, after observing T , one can forget the rest of the data if one is only interested in gathering information about θ . The concept of sufficient statistics was introduced by Fisher and was studied by Lehmann, Scheffe and Bahadur in the 1950s as pointed out by DeGroot (1970).

The strength of the definition lies in the possibility of finding sufficient statistics of a smaller dimension than data \mathbf{X} thus implying savings in information storage. In some cases, it is possible to find sufficient statistics with fixed dimension independent of the sample size. In these cases the dimension reduction and consequent storage saving can be huge if large sample sizes are considered.

Theorem 2.3. If $T = T(\mathbf{X})$ is a sufficient statistic for θ , then

$$p(\theta | \mathbf{x}) = p(\theta | t), \text{ for all priors } p(\theta)$$

Proof. $p(\mathbf{x} | \theta) = p(\mathbf{x}, t | \theta)$, if $t = T(\mathbf{x})$ and 0, if $t \neq T(\mathbf{x})$. So,

$$\begin{aligned} p(\mathbf{x} | \theta) &= p(\mathbf{x} | t, \theta) p(t | \theta) \\ &= p(\mathbf{x} | t) p(t | \theta), \quad \text{by the definition of sufficiency.} \end{aligned}$$

But, by Bayes' theorem,

$$\begin{aligned} p(\theta | \mathbf{x}) &\propto p(\mathbf{x} | \theta) p(\theta) \\ &= p(\mathbf{x} | t) p(t | \theta) p(\theta) \\ &\propto p(t | \theta) p(\theta), \quad \text{since } p(\mathbf{x} | t) \text{ does not depend on } \theta \\ &\propto p(\theta | t). \end{aligned}$$

Then $p(\theta | \mathbf{x}) = k p(\theta | t)$, for some $k > 0$.

Additionally,

$$1 = \int_{\Theta} p(\theta | \mathbf{x}) d\theta = k \int_{\Theta} p(\theta | t) d\theta = k$$

and so $p(\theta | \mathbf{x}) = p(\theta | t)$. \square

This theorem leads to a possible Bayesian definition of sufficiency below.

Definition (Bayesian). The statistic $T(\mathbf{X})$ is sufficient for θ if there is a function f such that

$$p(\theta | \mathbf{x}) \propto f(\theta, t).$$

A useful insight into sufficiency is gained through the notion of partitions of the sample space. Let $T = T(\mathbf{X})$ be a p -dimensional statistic and $A_t = \{\mathbf{x} : T(\mathbf{x}) = t\}$. The collection of sets $\{A_t : t \in R^p\} = \{A_t\}$ is a partition if $A_t \cap A_{t'} = \emptyset$, $\forall t, t' \in R^p$ and $\cup_t A_t = S$, the sample space. A partition induced by a sufficient statistic is called a sufficient partition.

The equivalence between the classical and Bayesian definitions follows easily from the theorem presented below for the classical definition of sufficiency.

Theorem 2.4 (Neyman's factorization criterion). The statistic T is sufficient for θ if and only if

$$p(\mathbf{x} | \theta) = f(t, \theta)g(\mathbf{x})$$

where f and g are non-negative functions.

Proof. (\Rightarrow) We have already seen that $p(\mathbf{x} | \theta) = p(\mathbf{x} | t)p(t | \theta)$. Then it is enough to define $g(\mathbf{x}) = p(\mathbf{x} | t) = p(\mathbf{x} | T(\mathbf{x}))$ and $f(t, \theta) = p(t | \theta)$ completing this part of the proof.

(\Leftarrow) Conversely, we have that $p(\mathbf{x} | \theta) = f(t, \theta)g(\mathbf{x})$. Defining $A_t = \{\mathbf{x} : T(\mathbf{x}) = t\}$, the probability (density) function of $T | \theta$ is

$$\begin{aligned} p(t | \theta) &= \int_{A_t} p(\mathbf{x} | \theta) d\mathbf{x} \\ &= f(t, \theta) \int_{A_t} g(\mathbf{x}) d\mathbf{x} \\ &= f(t, \theta)G(x), \quad \text{for some function } G \end{aligned}$$

and so, $f(t, \theta) = p(t | \theta)/G(x)$. On the other hand, from the hypothesis of the theorem, $f(t, \theta) = p(\mathbf{x} | \theta)/g(\mathbf{x})$. Equating the two forms for $f(t, \theta)$ leads to

$$\frac{p(\mathbf{x} | \theta)}{p(t | \theta)} = \frac{g(\mathbf{x})}{G(x)}.$$

Since $p(\mathbf{x} | t, \theta) = p(\mathbf{x} | \theta)/p(t | \theta)$, then

$$p(\mathbf{x} | t, \theta) = \frac{g(\mathbf{x})}{G(x)} = p(\mathbf{x} | t)$$

since it does not depend on θ . Thus, T is sufficient for θ . \square

Neyman's factorization criterion is the tool usually used for identifying sufficient statistics.

We can now show that the two concepts of sufficiency are equivalent, since

1. (classical \implies Bayesian) follows trivially from Theorem 2.4;
2. (Bayesian \implies classical)

$$p(\theta | \mathbf{x}) = \frac{p(\mathbf{x} | \theta)p(\theta)}{p(\mathbf{x})} = f(\mathbf{t}, \theta)p(\theta), \quad \text{by hypothesis.}$$

So, $p(\mathbf{x} | \theta) = f(\mathbf{t}, \theta)p(\mathbf{x})$ which, by the factorization criterion, is equivalent to saying that \mathbf{T} is a sufficient statistic.

Definition. Suppose that \mathbf{X} has density $p(\mathbf{x} | \theta)$. Then $\mathbf{T}(\mathbf{X})$ is an ancillary statistic for θ if $p(\mathbf{t} | \theta) = p(\mathbf{t})$.

In this case, \mathbf{T} does not provide any information for θ although it is a function of \mathbf{X} , which is related to θ . Ancillarity can be understood as an antonym for sufficiency.

Sufficiency is a basic concept in classical statistic although it is not so relevant for the Bayesian approach. On the other hand, from the applied point of view this is also not a very useful concept since even small perturbations in the model can imply the loss of sufficiency.

Examples.

1. Let $\mathbf{X} = (X_1, \dots, X_n)$ be observations with values 0 or 1, where $P(X_i = 1 | \theta) = \theta$.

$$p(\mathbf{x} | \theta) = \theta^t(1 - \theta)^{n-t}, \quad \text{with } t = \sum_{i=1}^n x_i.$$

From the factorization criterion it follows that $T(\mathbf{X}) = \sum_{i=1}^n X_i$ is a sufficient statistic. In this case, it is also possible to conclude straightforwardly from the definition of sufficiency and using some combinatorial arguments, that $T(\mathbf{X})$ is a sufficient statistic since $p(\mathbf{x} | T(\mathbf{x}) = t) = \left[\binom{n}{t} \right]^{-1}$, which does not depend on θ .

2. Let X_1, X_2, \dots, X_n be iid conditional on θ with common density $p(x_i | \theta)$. Then:

$$p(x_1, \dots, x_n | \theta) = \prod_{i=1}^n p(x_i | \theta).$$

The order statistics are defined as $Y_1 = X_{(1)} = \min_i X_i$, $Y_2 = X_{(2)} = \text{second smallest sample value}, \dots, Y_n = X_{(n)} = \max_i X_i$. Since the order of the terms does not alter the product and to each x_i there corresponds a unique y_i (assuming continuity),

$$\prod_{i=1}^n p(x_i | \theta) \propto \prod_{i=1}^n p(y_i | \theta).$$

Then, with $g(\mathbf{x}) = 1$, $\mathbf{t} = (y_1, \dots, y_n)$ and $f(\mathbf{t}, \theta) = \prod_{i=1}^n p(y_i | \theta)$, the factorization criterion holds and $\mathbf{T} = (Y_1, \dots, Y_n)$ is a sufficient statistic for θ . Note that the dimension of \mathbf{T} depends upon the sample size. In this case no dimensionality reduction was achieved and the definition becomes deprived of its strength. It is also clear that the sample \mathbf{X} itself is trivially a sufficient statistic for θ .

The application of the sufficiency concept developed above is not necessarily useful. It is only relevant when the dimension of the sufficient statistic is significantly smaller than the sample size. An interesting question is how to obtain a sufficient statistic with maximal reduction of the sample data. Such a statistic is known in the literature as a minimal sufficient statistic.

Definition. Let $\mathbf{X} \sim p(\mathbf{x} | \theta)$. The statistic $\mathbf{T}(\mathbf{X})$ is a minimal sufficient statistic for θ if it is a sufficient statistic for θ and a function of every other sufficient statistic for θ .

If $\mathbf{S}(\mathbf{X})$ is a statistic obtained as a bijective function of a sufficient statistic $\mathbf{T}(\mathbf{X})$ then \mathbf{S} is also a sufficient statistic. On the other hand, the minimal sufficient statistic is unique, apart from bijective transformation of itself.

Definition. Two elements \mathbf{x} and \mathbf{y} of the sample space are information equivalent if and only if the ratio $p(\mathbf{x} | \theta)/p(\mathbf{y} | \theta)$ or equivalently $I(\theta; \mathbf{x})/I(\theta; \mathbf{y})$ does not depend on θ .

Information equivalence defines an equivalence relation of the elements of the sample space. Therefore, it defines a partition of the sample space. This partition is called a minimal sufficient partition. It can be shown that a sufficient statistic is minimal if and only if the partition it defines on the sample space is minimal sufficient.

Example. Let X_1, \dots, X_n be iid Poisson variables with mean λ and define $T(\mathbf{X}) = \sum_{i=1}^n X_i$. Then,

$$p(\mathbf{x} | \lambda) = \prod_{i=1}^n p(x_i | \lambda) = \prod_{i=1}^n e^{-\lambda} \frac{\lambda^{x_i}}{x_i!} = e^{-n\lambda} \frac{\lambda^{T(\mathbf{X})}}{\prod_i x_i!}.$$

Therefore, $p(\mathbf{x} | \lambda)/p(\mathbf{y} | \lambda) = \lambda^{T(\mathbf{X})-T(\mathbf{Y})} \prod_i (y_i!)/x_i!$, which does not depend on λ if and only if $T(\mathbf{x}) = T(\mathbf{y})$. Hence, $T(\mathbf{X})$ is a minimal sufficient statistic for λ .

Another interesting question is whether there are families of distributions admitting fixed dimension sufficient statistics for θ . Fortunately, for a large class of distributions, the dimension of the sufficient statistic \mathbf{T} is equal to the number of parameters. Maximum summarization is obtained when we have one sufficient statistic for each parameter. Subject to some weak regularity conditions, all distributions with the dimension of \mathbf{T} equal to the number of the parameters belong to the exponential family.

Definition. The family of distributions with probability (density) function $p(x | \theta)$ belongs to the exponential family with r parameters if $p(x | \theta)$ can be written as

$$p(x | \theta) = a(x) \exp \left\{ \sum_{j=1}^r U_j(x) \phi_j(\theta) + b(\theta) \right\}, \quad x \in \mathcal{X} \subset R$$

and \mathcal{X} does not depend on θ .

By the factorization criterion, $U_1(X), \dots, U_r(X)$ are sufficient statistics for θ (when a single X is observed). For a size n sample of \mathbf{X} we have

$$p(\mathbf{x} | \theta) = \left[\prod_{i=1}^n a(x_i) \right] \exp \left\{ \sum_{j=1}^r \left[\sum_{i=1}^n U_j(x_i) \right] \phi_j(\theta) + nb(\theta) \right\}$$

which belongs to the exponential family too, with $a(\mathbf{x}) = \prod_{i=1}^n a(x_i)$ and $U_j(\mathbf{X}) = \sum_{i=1}^n U_j(X_i)$, $j = 1, \dots, r$. So, $\mathbf{T} = (T_1, \dots, T_r)$ with $T_j = U_j(\mathbf{X})$, $j = 1, 2, \dots, r$ is a sufficient statistic for θ .

The exponential family is very rich and includes most of the distributions more commonly used in statistics. Among the most important distributions not included in this family we have the uniform distribution (which has sufficient statistics with dimension not depending on the sample size) and the Student t distribution (which have none). Darmois, Koopman and Pitman have independently shown that among families satisfying some regularity conditions, a sufficient statistic of fixed dimension will only exist for the exponential family.

Examples.

1. Bernoulli(θ)

$$\begin{aligned} p(x | \theta) &= \theta^x (1-\theta)^{1-x} I_x(\{0, 1\}), \\ &= \exp \left\{ x \log \left(\frac{\theta}{1-\theta} \right) + \log(1-\theta) \right\} I_x(\{0, 1\}). \end{aligned}$$

For a sample \mathbf{x} we have,

$$\begin{aligned} p(\mathbf{x} | \theta) &= \prod_{i=1}^n \theta^{x_i} (1-\theta)^{1-x_i} I_{x_i}(\{0, 1\}) \\ &= \exp \left\{ \sum_{i=1}^n x_i \log \left(\frac{\theta}{1-\theta} \right) + n \log(1-\theta) \right\} I_{\mathbf{x}}(\{0, 1\}^n). \end{aligned}$$

Then, a Bernoulli belongs to the one parameter exponential family with $a(\mathbf{x}) = I_{\mathbf{x}}(\{0, 1\}^n)$, $b(\theta) = n \log(1-\theta)$, $\phi(\theta) = \log[\theta/(1-\theta)]$ and $U(\mathbf{X}) = \sum_{i=1}^n X_i$. So, U is a sufficient statistic for θ as we have seen before.

2. Poisson (λ)

$$P(x | \lambda) = \frac{e^{-\lambda} \lambda^x}{x!} I_x(\{0, 1, \dots\}),$$

which in the exponential form is

$$p(x | \lambda) = \frac{1}{x!} \exp\{-\lambda + x \log \lambda\} I_x(\{0, 1, \dots\}).$$

For a sample \mathbf{x} we have,

$$p(\mathbf{x} | \lambda) = \frac{1}{\prod_{i=1}^n x_i!} \exp \left\{ \sum_{i=1}^n x_i \log \lambda - n\lambda \right\} I_{\mathbf{x}}(\{0, 1, \dots\}^n).$$

So, the Poisson belongs to the one parameter exponential family with $a(\mathbf{x}) = I_{\mathbf{x}}(\{0, 1, \dots\}^n) / \prod_{i=1}^n x_i!$, $b(\lambda) = -n\lambda$, $\phi(\lambda) = \log \lambda$ and $U(\mathbf{X}) = \sum_{i=1}^n X_i$. Then U is sufficient for λ .

3. Normal(μ, σ^2)

$$\begin{aligned} p(x | \mu, \sigma^2) &= \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{(x - \mu)^2}{2\sigma^2} \right\} \\ &= \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ \frac{\mu}{\sigma^2}x - \frac{1}{2\sigma^2}x^2 - \frac{\mu^2}{2\sigma^2} \right\} \\ &= \frac{1}{\sqrt{2\pi}} \exp \left\{ \frac{\mu}{\sigma^2}x - \frac{1}{2\sigma^2}x^2 - \frac{\mu^2}{2\sigma^2} - \frac{1}{2} \log \sigma^2 \right\}. \end{aligned}$$

For a sample \mathbf{x} it follows,

$$\begin{aligned} p(\mathbf{x} | \mu, \sigma^2) &= \frac{1}{(2\pi)^{n/2}} \\ &\times \exp \left\{ \frac{\mu}{\sigma^2} \sum_{i=1}^n x_i - \frac{1}{2\sigma^2} \sum_{i=1}^n x_i^2 - \frac{n}{2} \left(\frac{\mu^2}{\sigma^2} + \log \sigma^2 \right) \right\}. \end{aligned}$$

Then, the normal distribution is a member of the exponential family with a bidimensional parameter $\theta = (\mu, \sigma^2)$, $a(\mathbf{x}) = (2\pi)^{-n/2}$, $b(\theta) = -(n/2)[(\mu^2/\sigma^2) + \log \sigma^2]$, $\phi_1(\theta) = \mu^2/\sigma^2$, $\phi_2(\theta) = -1/2\sigma^2$, $U_1(\mathbf{X}) = \sum_{i=1}^n X_i$ and $U_2(\mathbf{X}) = \sum_{i=1}^n X_i^2$. So, $U = (U_1, U_2)$ is sufficient for (μ, σ^2) .

2.6 Parameter elimination

Sometimes models need to be developed with the inclusion of various parameters, many of which are included regardless of our particular interest in them. These parameters are often included to describe relevant aspects of the reality we are modelling, even though they are not related to our main concerns about the problem. To simplify the discussion, the parametric vector can be broadly split into two

subvectors: θ containing the parameters of interest and ϕ , the components that despite being present in the model are not of our immediate concern. The first subvector is called the parameter of interest and the second one is known as the nuisance parameter. Usually, it is our desire to eliminate ϕ from the analysis as soon as possible, in order to concentrate efforts on θ .

From the Bayesian point of view, this is done using probability rules. Once having observed $\mathbf{X} = \mathbf{x}$, we get $p(\theta, \phi | \mathbf{x})$, and we can easily calculate

1. Marginal posterior distributions:

$$p(\theta | \mathbf{x}) = \int_{\Phi} p(\theta, \phi | \mathbf{x}) d\phi$$

and

$$p(\phi | \mathbf{x}) = \int_{\Theta} p(\theta, \phi | \mathbf{x}) d\theta$$

where Θ and Φ are the respective parameter spaces of θ and ϕ .

2. Conditional posterior distributions:

$$p(\theta | \phi, \mathbf{x}) = \frac{p(\theta, \phi | \mathbf{x})}{p(\phi | \mathbf{x})} \propto p(\theta, \phi | \mathbf{x})$$

and

$$p(\phi | \theta, \mathbf{x}) = \frac{p(\theta, \phi | \mathbf{x})}{p(\theta | \mathbf{x})} \propto p(\theta, \phi | \mathbf{x})$$

where these conditional distributions are well defined if the corresponding denominators are non-zero, that is, for values of θ and ϕ with strictly positive marginal posterior density. The above calculations use the fact that the terms in the denominator do not depend on the quantity of interest and can be included in the proportionality constant.

3. Marginal likelihood functions:

The likelihood function is defined as $l(\theta, \phi; \mathbf{x}) = p(\mathbf{x} | \theta, \phi)$. In a similar way, one can define the marginal likelihood functions as:

$$\begin{aligned} l(\theta; \mathbf{x}) &= p(\mathbf{x} | \theta) = \int_{\Phi} p(\mathbf{x}, \phi | \theta) d\phi \\ &= \int_{\Phi} p(\mathbf{x} | \phi, \theta) p(\phi | \theta) d\phi \end{aligned}$$

and

$$l(\phi; \mathbf{x}) = \int_{\Theta} p(\mathbf{x} | \theta, \phi) p(\theta | \phi) d\theta$$

Conceptually, there is no difficulty in defining and obtaining any of the above quantities although sometimes there are difficulties in analytically solving these integrals. The same is not true for classical or frequentist inference. Some special rules must be stated leading to *ad hoc* procedures to solve the stated problem.

Many efforts are geared in classical inference towards a proper definition of marginal likelihood. This is a relevant problem to the frequentist statistician and has received much research attention, beginning with the work of Bartlett in the 1930s and further developed by Kalbfleisch and Sprott in the 1970s, as discussed in Cox and Hinkley (1974). Many proposals to express the marginal likelihood $l(\theta; \mathbf{x})$ are based in substituting ϕ in the (total) likelihood by some particular value. Often, this is taken as the value that maximizes the likelihood. Denoting this value by $\hat{\phi}(\theta)$, because it can depend on θ , we get the relative or profile likelihood for θ as

$$l_P(\theta; \mathbf{x}) = l(\theta, \hat{\phi}(\theta); \mathbf{x}).$$

Some authors suggest corrections in this expression taking into consideration measures of information associated with ϕ .

There are other frequentist definitions for the marginal and conditional likelihood. Suppose that the vector \mathbf{X} , or some one-to-one transformation of it, is partitioned into (\mathbf{T}, \mathbf{U}) , with joint density given by

$$p(\mathbf{t}, \mathbf{u} | \theta, \phi) = p(\mathbf{t} | \theta, \phi) p(\mathbf{u} | \mathbf{t}, \theta, \phi).$$

The likelihood function of θ, ϕ is given by the left-hand side of the equation above and the right-hand side terms can also be written in likelihood terms as

$$l(\theta, \phi; \mathbf{t}, \mathbf{u}) = l(\theta, \phi; \mathbf{t}) l(\theta, \phi; \mathbf{u} | \mathbf{t}).$$

The first term on the right-hand side is called the marginal likelihood and the second, the conditional likelihood. These forms of likelihood are useful when some of the parametric components can be eliminated. For example, if \mathbf{T} is such that $l(\theta, \phi; \mathbf{t}) = l(\theta; \mathbf{t})$ only this term is used to make inferences about θ . For conditional likelihood, this form is related to sufficient statistics because if \mathbf{T} is sufficient for ϕ , with θ fixed, then $l(\theta, \phi; \mathbf{u} | \mathbf{t}) = l(\theta; \mathbf{u} | \mathbf{t})$. Again only this term is used to make inferences about θ . The question is how much information is lost when ignoring the other part of the likelihood.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from a $N(\theta, \sigma^2)$ and define $\phi = \sigma^{-2}$. The parameter vector (θ, ϕ) is unknown and we assume that the main interest lies in the mean of the population. The precision ϕ is essentially a nuisance parameter that we would like to eliminate from the analysis.

Suppose that the prior distribution is such that $n_0 \sigma_0^2 \phi \sim \chi_{n_0}^2$ or equivalently $\phi \sim G(n_0/2, n_0 \sigma_0^2/2)$ and ϕ is independent of θ a priori. In these conditions we have

$$p(\phi | \theta) = p(\phi) \propto \phi^{n_0/2-1} \exp \left\{ -\frac{n_0 \sigma_0^2}{2} \phi \right\}.$$

On the other hand,

$$p(\mathbf{x} | \theta, \phi) \propto \phi^{n/2} \exp \left\{ -\frac{\phi}{2} \sum_{i=1}^n (x_i - \theta)^2 \right\}$$

but

$$\begin{aligned}\sum_{i=1}^n (x_i - \theta)^2 &= \sum_{i=1}^n [(x_i - \bar{x}) + (\bar{x} - \theta)]^2 \\ &= \sum_{i=1}^n (x_i - \bar{x})^2 + n(\bar{x} - \theta)^2 \\ &= n[s^2 + (\bar{x} - \theta)^2]\end{aligned}$$

where

$$s^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2 / n.$$

Therefore, the marginal likelihood of θ is

$$\begin{aligned}l(\theta; \mathbf{x}) &= \int_{\Phi} p(\mathbf{x} | \phi, \theta) p(\phi | \theta) d\phi \\ &\propto \int_0^{\infty} \phi^{\frac{n_0}{2}-1} \exp\left\{-\frac{\phi}{2} n_0 \sigma_0^2\right\} \phi^{\frac{n}{2}} \exp\left\{-\frac{\phi}{2}[ns^2 + n(\bar{x} - \theta)^2]\right\} d\phi \\ &= \int_0^{\infty} \phi^{\frac{n+n_0}{2}-1} \exp\left\{-\frac{\phi}{2}[ns^2 + n(\bar{x} - \theta)^2 + n_0 \sigma_0^2]\right\} d\phi \\ &= \frac{\Gamma[(n+n_0)/2]}{[(ns^2 + n(\bar{x} - \theta)^2 + n_0 \sigma_0^2)/2]^{(n+n_0)/2}} \\ &= \frac{\Gamma[(n+n_0)/2]}{[(ns^2 + n_0 \sigma_0^2)/2]^{(n+n_0)/2}} \left[1 + \frac{n(\bar{x} - \theta)^2}{ns^2 + n_0 \sigma_0^2}\right]^{-\frac{n+n_0}{2}}.\end{aligned}$$

Letting $n_0 \rightarrow 0$, which corresponds to vague prior information (as will be seen in Chapter 3), gives

$$\begin{aligned}l(\theta; \mathbf{x}) &\rightarrow k \left[1 + \frac{(\bar{x} - \theta)^2}{s^2}\right]^{-n/2} \\ &= k \left[1 + \frac{T^2(\mathbf{x}, \theta)}{n-1}\right]^{-[(n-1)+1]/2} \quad \text{where } T(\mathbf{X}, \theta) = \frac{\bar{X} - \theta}{\sqrt{\frac{s^2}{n-1}}}.\end{aligned}$$

Interpreting the likelihood as proportional to the probability density function of T , it follows that $T \sim t_{n-1}(0, 1)$ (see list of distributions).

The sampling distribution of \bar{X} , the minimal sufficient statistic for θ , is used in the frequentist inference. This distribution depends on σ^2 . Substituting it by the estimator S^2 , the minimal sufficient statistic for ϕ leads to $T(\mathbf{X}, \theta)$ with a $t_{n-1}(0, 1)$ sampling distribution as will be seen in Chapter 4. Therefore, if the adopted prior distribution for θ is proportional to a constant then the Bayesian and classical results will agree numerically at least, although theoretically they were obtained using different arguments.

Exercises

§ 2.1

1. For each of the following distributions verify if the model is location, scale or location-scale.
 - (a) $t_\alpha(\mu, \sigma^2)$, α known;
 - (b) Pareto (x_0, α) , with α fixed, with density $p(x|x_0) = \alpha x_0^\alpha / x^{1+\alpha}$, $x > x_0$, ($a, x_0 > 0$);
 - (c) uniform distribution on $(\theta - 1, \theta + 1)$;
 - (d) uniform distribution on $(-\theta, \theta)$.

§ 2.2

2. Let X have sampling density $p(x|\theta)$ with $\theta \in \Theta \subset R$. Prove or give a counterexample to the assertion that if the sampling distribution $p(x|\theta)$ has a unique maximum then the likelihood $l(\theta; x)$ also has a unique maximum. Generalize the result to the case of a vector of observations \mathbf{X} and also for a parameter vector θ .
3. A situation that usually occurs in lifetime or survival analysis is to have observations that are censored because of time and/or cost restrictions on the experiment. One common situation occurs when the experiment is run only until time $T > 0$. If the observational unit is observed until time T , it is uncensored but if it is censored then all one can extract from the experiment is that the lifetime of this unit is larger than T . Assuming that a random sample X_1, \dots, X_n from a density $f(x|\theta)$ is observed, show that
 - (a) the likelihood is

$$l(\theta) = \prod_{i=1}^n [f(x_i|\theta)]^{1-\chi_i} [1 - F(T|\theta)]^{\chi_i}$$

where F is the distribution function of the observations and χ_i is the censoring indicator, $i = 1, \dots, n$, taking values 0, when failure is observed, and 1, when censoring takes place;

- (b) in the case that $f(x|\theta) = \theta e^{-\theta x}$, $l(\theta) = \theta^m e^{-\theta U}$ where $m \leq n$ is the number of uncensored observations and $U = (n-m)T + \sum_i (1-\chi_i)x_i$ is the total time on test.

4. Let X_1, \dots, X_n be iid random quantities from the Weibull distribution, denoted by $\text{Wei}(\alpha, \beta)$ ($\alpha, \beta > 0$), with

$$p(x|\alpha, \beta) = \beta \alpha x^{\alpha-1} \exp(-\beta x^\alpha), \quad \alpha > 0, \beta > 0$$

- (a) Obtain the likelihood function, the score function and the observed and expected Fisher information matrix for the pair of parameters (α, β) .
- (b) The Weibull distribution is sometimes parametrized in terms of α and $\theta = 1/\beta^\alpha$. Repeat item (a) for the pair of parameters (α, θ) .

§ 2.3

5. Return to the example about John's illness and consider the same distributions.
- Which test result makes the doctor more certain about John's illness? Why?
 - The test X is applied and provides the result $X = 1$. Suppose the doctor is not satisfied with the available evidence and decides to ask for another replication of test X and again the result is $X = 1$. What is now the probability that John has disease A ?
 - What is the minimum number of repetitions of the test X which allows the doctor to ensure that John has the disease with 99.9% probability. What are the results of these replications that guarantee this?
6. Suppose that $X | \theta \sim N(\theta, 1)$ (for example, X is a measurement of a physical constant θ made with an instrument with variance 1). The prior distribution for θ elicited by the scientist A corresponds to a $N(5, 1)$ distribution and the scientist B elicits a $N(15, 1)$ distribution. The value $X = 6$ was observed.
- What prior fits the data better?
 - What kind of comparison can be done between the two scientists?
7. Classify the following assertions as TRUE or FALSE, briefly justifying your answer.
- The posterior distribution is always more precise than the prior because it is based on more information.
 - When X_2 is observed after X_1 , the prior distribution before observing X_2 has to be necessarily the posterior distribution after observing X_1 .
 - The predictive density is the prior expected value of the sampling distribution.
 - The smaller the prior information, the bigger the influence of the sample in the posterior distribution.
8. A test to verify if a driver is driving in a drunken state has 0.8 chance of being correct, that is, to provide a positive result when in fact the driver has a high level of alcohol in his/her blood or negative result when it is below the acceptable limit. A second test is only applied to the suspected cases. This never fails if the driver is not drunk, but has only a 10% chance of error with drunk drivers. If 25% of all the drivers stopped by the police are *above the limit*, calculate:
- the proportion of drivers stopped that have to be submitted to a second test;
 - the posterior probability that this driver really has the high level of alcohol in his blood informed by the two tests;
 - the proportion of drivers that will be submitted only to the first test.

9. (DeGroot, 1970, p. 152) The random variables X_1, \dots, X_k are such that $k - 1$ of them have probability function h and one has probability function g . X_j will have the probability function g with probability α_j , $j = 1, \dots, k$, where $\alpha_j > 0$, $\forall j$ and $\sum_{j=1}^k \alpha_j = 1$. What is the probability that X_1 has probability function g given that:

- $X_1 = x$ was observed?
- $X_i = x, i \neq 1$ was observed?

10. Let $X | \theta, \mu \sim N(\theta, \sigma^2)$, σ^2 known and $\theta | \mu \sim N(\mu, \tau^2)$, τ^2 known and $\mu \sim N(0, 1)$. Obtain the following distributions:

- $(\theta | x, \mu)$;
- $(\mu | x)$;
- $(\theta | x)$.

11. Let $(X|\theta) \sim N(\theta, 1)$ be observed. Suppose that your prior is such that θ is $N(\mu, 1)$ or $N(-\mu, 1)$ with equal probabilities. Write the prior distribution and find the posterior after observing $X = x$. Show that

$$\mu' = E(\theta | x) = \frac{x}{2} + \frac{\mu}{2} \frac{1 - \exp(-\mu x)}{1 + \exp(-\mu x)}$$

and draw a graph of μ' as a function of x .

12. The standard Cauchy density function is $p(x|\theta) = (1/\pi)\{1/[1+(x-\theta)^2]\}$ and is similar to $N(\theta, 1)$ and can be used in its place in many applications. Find the modal equation (the first order condition to the maximum) of the posterior supposing that the prior is $p(\theta) = 1/\pi(1+\theta^2)$.

- Solve it for $x = 0$ and $x = 3$.
- Compare with the results obtained assuming that $(x|\theta) \sim N(\theta, 1)$ and $\theta \sim N(0, 1)$.

13. Assume that an observation vector \mathbf{X} has multivariate normal distribution, introduced in Chapter 1, with mean vector $\boldsymbol{\mu}$ and variance-covariance matrix $\boldsymbol{\Sigma}$. Assuming that $\boldsymbol{\Sigma}$ is known and the prior distribution is $\boldsymbol{\mu} \sim N(\boldsymbol{\mu}_0, \mathbf{B}_0)$ obtain the posterior distribution for $\boldsymbol{\mu}$.

§ 2.4

14. Let $\mathbf{X} = (X_1, \dots, X_n)$ be an exchangeable sample of 0-1 observations. Show that
- $E[X_i] = E[X_j], \forall i, j = 1, \dots, n$;
 - $V[X_i] = V[X_j], \forall i, j = 1, \dots, n$;
 - $Cov(X_i, X_j) = Cov(X_k, X_l), \forall i, j, k, l = 1, \dots, n$.
15. Let $\mathbf{X} = (X_1, \dots, X_n)$ be an exchangeable sample of 0-1 observations and $T = \sum_{i=1}^n X_i$. Show that
- $P(T = t) = \int_0^1 \binom{n}{t} \theta^t (1-\theta)^{n-t} p(\theta) d\theta, \quad t = 1, \dots, n$.

- (b) $E(T) = nE(\theta)$.

Hint: in (b), use the definition of $E(T)$ and exchange the summation and integration signs.

16. Let $\theta_1, \dots, \theta_k$ be the probability that patients I_1, \dots, I_k have the disease B . After summarizing all the available information the doctor concludes that

- (a) The patients can be divided in two groups, the first containing the patients I_1, \dots, I_j , $j < k$ and the second with patients I_{j+1}, \dots, I_k .
- (b) The patients in the same group are similar, that is to say they are indistinguishable with respect to B .
- (c) There is no relevant information relating these two groups.

Use the idea of partial exchangeability to construct a prior distribution for $\theta = (\theta_1, \dots, \theta_k)$. If instead of (c), there was information relating the two groups, what modifications would this imply in the prior for θ ?

§ 2.5

17. Let $X = (X_1, \dots, X_n)$ be a random sample from $U(\theta_1, \theta_2)$, that is,

$$p(x | \theta_1, \theta_2) = \frac{1}{\theta_2 - \theta_1}, \quad \theta_1 \leq x \leq \theta_2.$$

Let $T(X) = (X_{(1)}, X_{(n)})$, obtain its joint distribution and show that it is a sufficient statistic for $\theta = (\theta_1, \theta_2)$.

18. Let X be a random sample from $P(X | \theta)$. Show that if $T = T(X)$ is a sufficient statistic for θ and $S(X)$ is a 1-to-1 function of T then $S(X)$ is also a sufficient statistic for θ .

19. Let X_1, \dots, X_n be a random sample from $P(X | \theta_1, \theta_2)$. Show that if T_1 is sufficient for θ_1 when θ_2 is known and T_2 is sufficient for θ_2 when θ_1 is known, then $T = (T_1, T_2)$ is sufficient for $\theta = (\theta_1, \theta_2)$.

20. Verify whether the following distributions belong to the exponential family. If so, determine the functions a , b , u and ϕ .

- (a) $\text{bin}(n, \theta)$, n known;
 - (b) $\exp(\theta)$;
 - (c) $G(\alpha, \beta)$;
 - (d) beta (α, β) ;
 - (e) $N(\mu, \Sigma)$, Σ known.
21. Which of the following distribution families are members of the exponential family? Obtain the minimal sufficient statistic for those belonging to the exponential family.
- (a) $p(x | \theta) = 1/9$, $x \in \{0.1 + \theta, \dots, 0.9 + \theta\}$;
 - (b) the family of $N(\theta, \theta^2)$ distributions;
 - (c) the family of $N(\theta, \theta)$ distributions, with $\theta > 0$;
 - (d) $p(x | \theta) = 2(x + \theta)/(1 + 2\theta)$, $x \in (0, 1)$, $\theta > 0$;

- (e) the distribution family of $X | X \neq 0$ where $X \sim \text{bin}(n, \theta)$;

- (f) $f(x | \theta) = \theta/(1+x)^{1+\theta}$, $x \in R^+$;

- (g) $f(x | \theta) = \theta^x \log \theta / (\theta - 1)$, $x \in (0, 1)$;

- (h) $f(x | \theta) = (1/2) \exp(-|x - \theta|)$, $x \in R$.

22. Let X_1, \dots, X_n be a random sample from $N(\mu, \sigma^2)$, with σ^2 unknown. Show, using the classical definition, that the sample mean \bar{X} is a sufficient statistic for μ .

Hint: It is enough to show that $E(X | \bar{X})$ and $V(X | \bar{X})$ is not a function of μ . Why?

23. Let (X_1, X_2, X_3) be a vector with probability function

$$\frac{n!}{\prod_{i=1}^3 x_i!} \prod_{i=1}^3 p_i^{x_i}, \quad x_i \geq 0, \quad x_1 + x_2 + x_3 = n$$

where $p_1 = \theta^2$, $p_2 = 2\theta(1-\theta)$, $p_3 = (1-\theta)^2$ and $0 \leq \theta \leq 1$.

- (a) Verify whether this distribution belongs to the exponential family with k parameters. If this is true, what is the value of k ?

- (b) Obtain the minimal sufficient statistic for θ .

24. Using the same notation adopted for the one parameter exponential family,

- (a) show that

$$E[U(X)] = -\frac{b'(\theta)}{\phi'(\theta)} \quad \text{and} \quad V[U(X)] = \frac{b'(\theta)\phi''(\theta) - \phi'(\theta)b''(\theta)}{[\phi'(\theta)]^3}.$$

Hint: From the relationship $\int p(x | \theta) dx = 1$, differentiate both sides with respect to θ .

- (b) Verify that the result in (a) is correct for the case where $X \sim \exp(\theta)$ by direct evaluation of $E[U(X)]$ and $V[U(X)]$.

25. Show that information equivalence defines an equivalence relation of the elements of the sample space.

26. Let X_1, X_2, X_3 be iid Bernoulli variables with success probability θ and define $T = T(X) = \sum_{i=1}^n X_i$, $T_1 = X_1$ and $T_2 = (T, T_1)$. Note that the sample space is $\mathcal{S} = \{0, 1\}^3$.

- (a) Obtain the partitions induced by T , T_1 and T_2 .

- (b) Show that T_2 is a sufficient statistic.

- (c) Prove that T is a minimal sufficient statistic for θ but T_2 isn't by showing that T induces a minimal sufficient partition on \mathcal{S} but T_2 does not.

27. Consider a sample $\mathbf{X} = (X_1, \dots, X_n)$ from a common density $p(x | \theta)$ and let \mathbf{T} be the vector of order statistics from the sample.

50 Elements of inference

- (a) Prove that the sample \mathbf{X} is always sufficient for θ
 (b) Obtain the factor $g(\mathbf{x})$ in the factorization criterion for the sufficient statistic \mathbf{T} .
28. Consider a sample $\mathbf{X} = (X_1, \dots, X_n)$ from a uniform distribution on the interval $[\theta_1, \theta_2]$, $\theta_1 < \theta_2$, so that $\theta = (\theta_1, \theta_2)$.
- Show that this distribution does not belong to the exponential family.
 - Obtain a sufficient statistics of fixed size.
 - Specialize the results above for the cases that θ_1 is known and θ_2 is known.
29. Consider a sample $\mathbf{X} = (X_1, \dots, X_n)$ from a $t_v(\mu, \sigma^2)$, and $\theta = (\nu, \mu, \sigma^2)$.
- Show that this distribution does not belong to the exponential family.
 - Show that it is not possible to obtain a sufficient statistics for θ of fixed size.
 - Show that the results above are retained even for the cases when some of the components of θ are known.

§ 2.6

30. Let X and Y be independent random variables Poisson distributed with parameters θ and ϕ , respectively, and suppose that the prior distribution is $p(\theta, \phi) = p(\theta)p(\phi) \propto k$. Let $\psi = \theta/(\theta + \phi)$ and $\xi = \theta + \phi$ be a parametric transformation of (θ, ϕ) .
- Obtain the prior for (ψ, ξ) .
 - Show that $\psi | x, y \sim \text{beta}(x+1, y+1)$ and $\xi | x, y \sim G(x+y+2, 1)$ are independent.
 - Show that the conditional distribution X given $X + Y$ depends only on ψ , that is $p(x | x+y, \psi, \xi) = p(x | x+y, \psi)$ and that the distribution of $X + Y$ depends only on ξ .
 - Show that $X + Y$ is a sufficient statistic for ξ , X is a sufficient statistic for ψ , given $X + Y$, and that $(X, X + Y)$ is a sufficient statistic for (ψ, ξ) .
 - Obtain the marginal likelihoods of ψ and ξ .
 - To make an inference about ξ a statistician decides to use the fact presented in item (d). Show that the posterior is identical to that obtained in (b). Does it mean that $X + Y$ does not contain information about ψ ?
31. Suppose that X has density $f(x | \theta)$ where $\theta = (\theta_1, \theta_2, \theta_3)$ and the prior for θ is built up as $p(\theta) = g(\theta_1, \theta_2 | \theta_3)h(\theta_3)$ where g and h are densities. Obtain the marginal likelihood $f(x | \theta_2, \theta_3)$ as a function of f , g and h .
32. Let $\mathbf{X} = (X_1, X_2)$ where $X_1 = (X_{11}, \dots, X_{1m})$ and $X_2 = (X_{21}, \dots, X_{2n})$ are samples from the $\exp(\theta_1)$ and $\exp(\theta_2)$ distributions respectively. Suppose that independent $G(a_i, b_i)$ priors are assumed, $i = 1, 2$ and define $\psi = \theta_1/\theta_2$.

Exercises 51

- Obtain the distribution of (θ_1, θ_2) given $X = \mathbf{x}$.
- Repeat item (a), assuming now that $a_i, b_i \rightarrow 0$, $i = 1, 2$.
- Using the posterior obtained in item (b), show that

$$\frac{\bar{x}_1}{\bar{x}_2} \psi | X = \mathbf{x} \sim F(2m, 2n)$$

where

$$\bar{x}_1 = \frac{\sum x_{1j}}{m} \quad \text{and} \quad \bar{x}_2 = \frac{\sum x_{2j}}{n}.$$

Hint: complete the transformation with $\psi_1 = \theta_2$.

33. Let (X_1, X_2, X_3) be a random vector with trinomial distribution with parameter $\theta = (\theta_1, \theta_2, \theta_3)$ where $\theta_3 = 1 - \theta_1 - \theta_2$ and assume that the prior for θ is constant.
- Define $\lambda = \theta_1/(\theta_1 + \theta_2)$ and $\psi = \theta_1 + \theta_2$ and obtain their priors.
 - Obtain the marginal likelihood of ψ .
 - Show that $X_1 + X_2$ is a sufficient statistic for ψ .
34. A machine emits particles following a Poisson process with mean intensity of λ particles per unit time. Each particle generates a $N(\theta, 1)$ signal. A signal detector registers the particles. Unfortunately, the detector only registers the positive signals (making it impossible to observe the number n of emitted particles).
- Obtain the distribution of $k|n$ where k is the number of particles registered.
 - Show that the likelihood $l(\theta, \lambda)$ based on the observation of just one registered signal ($k = 1$) assuming the value $x_1 > 0$ during a unit interval is given by

$$\phi(x_1 - \theta)\lambda\Phi(\theta)\exp\{-\lambda\Phi(\theta)\}$$

where ϕ and Φ are the density and the cumulative distribution function of the standard normal.

Hint: Obtain the joint distribution of (x_1, k, n) and eliminate n by integration.

- Obtain the profile likelihood of θ , that is, $l(\theta, \hat{\lambda}(\theta))$ where $\hat{\lambda}(\theta)$ maximizes the likelihood of λ supposing θ known.
- Supposing that the prior is $p(\theta, \lambda) \propto k$, obtain the marginal likelihood of θ .

3

Prior distribution

In this chapter, different specification forms of the prior distribution will be discussed. Apart from the interpretation of probability, this is the only novelty introduced by the Bayesian analysis, relative to the frequentist approach. It can be seen as an element implied from exchangeability by de Finetti's representation theorem. It is determined in a subjective way, although it is not forbidden to use past experimental data to set it. The only requirement is that this distribution should represent the knowledge about θ before observing the results of the new experiment. In this chapter, alternative forms of assessing the prior distribution will be discussed. In Section 3.1 entirely subjective methods for direct assessment of the prior will be presented. An indirect approach, via functional forms, is discussed in Section 3.2. The parameters of those functional forms, known as *hyperparameters*, must be specified in correspondence with the subjective information available. The conjugate distribution will be introduced in Section 3.3 and the most common families will be presented in Section 3.4. The concept of reference prior and different forms of building up non-informative priors will be presented in Section 3.5. Finally, hierarchical prior specification will be discussed in Section 3.6.

3.1 Entirely subjective specification

Let θ be an unknown quantity and consider its possible values. If it is discrete, a prior probability for each possible value of θ can be evaluated directly. Also one may use some auxiliary tools, like lotteries or roulettes, as described in Chapter 1. De Finetti (1974) characterizes subjective probability through the consideration of betting and scoring rules.

The continuous case is slightly more complicated. Some suggestions are:

1. *The histogram approach:* first, the range of values of θ is divided into intervals, and prior probabilities for θ belonging to each interval are specified, as in the discrete case. Hence, a histogram for θ is built up and a smooth curve can be fitted to obtain the prior density of θ . Note that the number of intervals involved is arbitrarily chosen. Although the probability in the tails of the prior distribution are often very small, they can influence the subsequent inference. This is a relevant aspect in prior elicitation that deserves

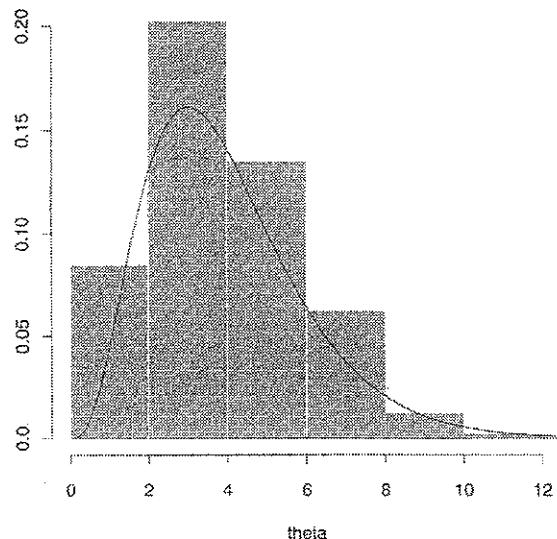


Fig. 3.1 Histogram representing (subjective) probabilities of the intervals I_1, I_2, I_3, I_4, I_5 and I_6 with a fitted density.

some caution. Figure 3.1 shows one such elicitation exercise for a positive quantity θ .

2. *The distribution function approach:* First, let us define percentiles. z_α is the $100\alpha\%$ percentile (α quantile) of X if $P(X \leq z_\alpha) = \alpha$, $\alpha \in [0, 1]$.

The median of X , denoted by m , is the 50% percentile, that is $P(X \leq m) = 0.5$.

The collection of all percentiles of X describes the distribution function of X . In this approach, some percentiles are subjectively assessed, as in the discrete case and later a smooth curve is fitted to the distribution function of θ as in Figure 3.2. This approach is less used since it is easier to identify a distribution through its density than through its distribution function.

3. *Relative likelihood approach:* The procedure is similar to the histogram approach but, instead of intervals, it evaluates the relative chances of isolated points. Even though $Pr(\theta = \theta_0) = 0$, $\forall \theta_0 \in \Theta$, this can be done because

$$Pr(\theta = \theta_0 \mid \theta = \theta_0 \text{ or } \theta = \theta_1) = \frac{p(\theta_0)}{p(\theta_0) + p(\theta_1)}$$

where $p(\theta)$ is the prior density of θ . Then a set of values proportional to the prior density of θ can be evaluated. For example, if $\theta = 2$ is three times more probable than $\theta = 1$ and $\theta = 3$ is twice more likely than $\theta = 1$, then we have $p(2) = 3p(1) = 1.5p(3)$. Again a smooth curve can be fitted to these points. One problem still outstanding is the evaluation of the

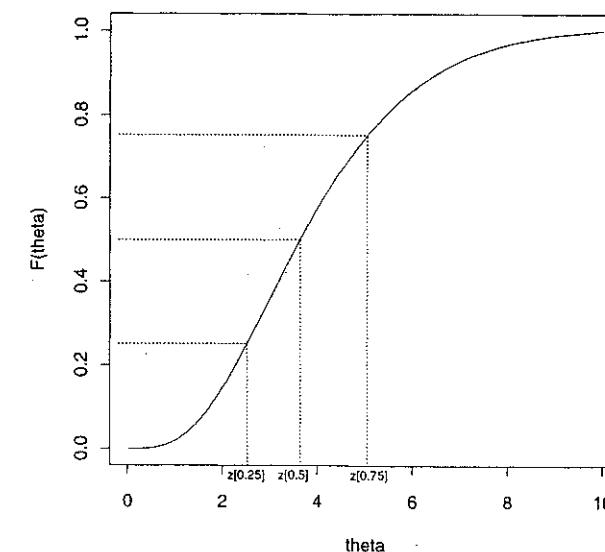


Fig. 3.2 Distribution function fitted to the quartiles $z_{0.25}$, $z_{0.5}$ and $z_{0.75}$.

normalization constant. Note that every density must integrate to 1 and, by construction, this curve does not necessarily satisfy this requirement.

These concepts are well described in the book by Berger (1985).

3.2 Specification through functional forms

The prior knowledge about θ can be used to specify a prior density with a particular functional form. A parametric family of densities can, for instance, be defined. Although very often this family can make the analysis easier, one must be careful and make sure that the chosen density really represents the available information. For example, we could make the following assumptions about θ :

- θ is symmetrically distributed with respect to the mode;
- its density decays fast (say, exponentially) when far away from the mode;
- intervals far from the mode have irrelevant probabilities.

These considerations can characterize, at least approximately, a normal distribution with parameters, generically called hyperparameters, determined in correspondence with the information expressed in H . These ideas may be put in a more general framework and have led to a systematic approach of determination of prior distributions. The most relevant case corresponds to a conjugate family of distributions.

We have seen in Theorem 2.1 that if the observational distribution is $(X | \theta) \sim N(\theta, \sigma^2)$ and the prior is $\theta \sim N(\mu, \tau^2)$, then the posterior distribution is also normal, with mean μ_1 and variance τ_1^2 . So, if we start with a normal prior we end up with a normal posterior. The main advantage of this approach is the ease of the resulting analysis. Among other things, this allows for the possibility of exploring the sequential aspect of the Bayesian paradigm. Every new normal observation that is obtained only leads to changes in the parameters of the (new) posterior distribution. No new analytic calculations are required.

Definition. Let $\mathcal{F} = \{p(x|\theta), \theta \in \Theta\}$ be a family of a sampling or observational distributions. A class \mathcal{P} of distributions is said to be a conjugate family with respect to \mathcal{F} if for all $p \in \mathcal{F}$ and $p(\theta) \in \mathcal{P}$ then $p(\theta | x) \in \mathcal{P}$.

Thus, we can say that the class of normal distributions is a conjugate family with respect to the class of normal (sampling) distributions. Some caution is necessary when using the notion of conjugacy:

- The class \mathcal{P} can be very broad

For example, take $\mathcal{P} = \{\text{all distributions}\}$ and \mathcal{F} to be any family of sampling distributions. It is easy to see that \mathcal{P} is conjugate with respect to \mathcal{F} since any posterior will be a member of \mathcal{P} . In this context, the definition of conjugacy does not have any practical appeal and is useless.

- The class \mathcal{P} could be very narrow

Suppose, for example, that

$$\mathcal{P} = \{p : p(\theta = \theta_0) = 1, \theta_0 \in A\}$$

for some non-null set A . This means that \mathcal{P} consists only of distributions concentrated on a single point. Whatever the information provided by the sample the posterior distribution would be the same as the prior because if we know, a priori, that $\theta = \theta_0$ with certainty, nothing will remove this certainty. That is

$$p(\theta|x) \propto l(\theta) p(\theta) = \begin{cases} l(\theta) \times 1, & \text{if } \theta = \theta_0 \\ l(\theta) \times 0, & \text{if } \theta \neq \theta_0. \end{cases}$$

Then it follows that

$$p(\theta|x) = \begin{cases} k \times l(\theta), & \text{if } \theta = \theta_0 \\ 0, & \text{if } \theta \neq \theta_0. \end{cases}$$

As $\int p(\theta|x) d\theta = 1$, we must have that $p(\theta|x) = 1$ if and only if (iff, in short) $\theta = \theta_0$. Hence, \mathcal{P} is conjugate for any distribution family and again the definitions are not helpful.

The last consideration illustrates, in an extreme situation, another very important aspect of prior specification. When a null probability is given to a particular subset

of the possible values of θ , no observed information will change this specification, even if it is proved to be obviously inadequate. In order to avoid this incoherent statement, it is strongly recommended that the statistician always associates a non-null prior probability to every possible value of θ , even if some of them are judged very unlikely. Dennis Lindley refers to this recommendation as *Cromwell's rule*.

Therefore the class \mathcal{P} must be broad enough to ensure elicitation of the convenient prior and, at the same time, restricted enough in order that the definition be useful. A general procedure for obtaining conjugate prior families is illustrated in the following example.

Example (Bernoulli trials). Let $(X_i|\theta) \sim \text{Ber}(\theta)$, $i = 1, \dots, n$. The joint sampling density is

$$p(\mathbf{x} | \theta) = \theta^t (1 - \theta)^{n-t} \quad \text{where } t = \sum_{i=1}^n x_i, \quad x_i = 0, 1, i = 1, \dots, n$$

defining a class of distribution parameterized by $\theta \in (0, 1)$. From Bayes' theorem, we know that the posterior density of θ given \mathbf{x} is given by

$$\begin{aligned} p(\theta | \mathbf{x}) &\propto p(\mathbf{x} | \theta) p(\theta) \\ &\propto \theta^t (1 - \theta)^{n-t} p(\theta). \end{aligned}$$

It is worth pointing out that $p(\theta)$ and $p(\theta | \mathbf{x})$ are related through the likelihood function. The conjugate prior can then be obtained by mimicking the kernel of the likelihood function. In this example the likelihood kernel is of the form $\theta^a (1 - \theta)^b$. This is also the kernel of the beta family of distributions, introduced in Chapter 1. Taking the prior distribution as a beta(α, β) and combining with the likelihood, the posterior distribution is

$$\begin{aligned} p(\theta | \mathbf{x}) &\propto \theta^t (1 - \theta)^{n-t} \theta^{\alpha-1} (1 - \theta)^{\beta-1} \\ &\propto \theta^{\alpha+t-1} (1 - \theta)^{\beta+n-t-1}. \end{aligned}$$

Therefore $(\theta | \mathbf{x}) \sim \text{beta}(\alpha + t, \beta + n - t)$ which belongs to the same family of distributions used for the prior. So, the beta family is conjugate with respect to the Bernoulli sampling model. It is not difficult to show that the same is true for binomial, geometric and negative binomial sampling distributions. The proportionality constant for the posterior density is given by $1/B(\alpha + t, \beta + n - t)$.

We can now discuss the setting of the conjugate prior family from a practical point of view for the general case of any given random sample. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a sample from $p(\mathbf{x} | \theta)$, $\theta \in \Theta$, and consider the density function of \mathbf{X} , $p(\mathbf{x} | \theta)$. The family \mathcal{P} is said to be closed under multiplication or sampling if for all $p_1, p_2 \in \mathcal{P}$, there is a k such that $kp_1 p_2 \in \mathcal{P}$.

Example. Let \mathcal{P} be the class of gamma distributions. If $p_i, i = 1, 2$, denote the gamma densities with parameters $(a_i, b_i), i = 1, 2$, then

$$\begin{aligned} p_1 \times p_2 &= \frac{\Gamma(a_1)}{b_1^{a_1}} x^{a_1-1} e^{-b_1 x} \frac{\Gamma(a_2)}{b_2^{a_2}} x^{a_2-1} e^{-b_2 x} \\ &\propto x^{a_1+a_2-1} e^{-(b_1+b_2)x} \end{aligned}$$

which is proportional to another gamma density with parameters a_1+a_2 and b_1+b_2 . It can be shown that the same result is true for the class of beta distributions.

Closure under multiplication is very important in the search for conjugate families. If the kernel of the likelihood can be identified with the kernel (of a member) of a given family of distributions and this family is closed under multiplication then prior and posterior will necessarily belong to the same family and conjugacy is obtained. The concept of conjugacy was formalized by Raiffa and Schlaifer (1961). They also studied many of the families that are presented in the next section.

With the definition of closure under sampling in hand it becomes easy to specify a procedure to determine the conjugate class. It consists of

1. identifying the class \mathcal{P} of distribution for θ whose members are proportional to $l(\theta; \mathbf{x})$;
2. verifying if \mathcal{P} is closed under sampling.

If, in addition, for any given likelihood $l(\theta; \mathbf{x})$ obtained from a family \mathcal{F} , there exists a constant k defining a density p as $p(\theta) = k l(\theta; \mathbf{x})$, then the family \mathcal{P} of all such densities p is said to be a natural conjugate family with respect to the sampling model with likelihood function l .

Example (continued). Setting $k^{-1} = B(t+1, n-t+1)$ implies that

$$k l(\theta; \mathbf{x}) = \frac{1}{B(t+1, n-t+1)} \theta^t (1-\theta)^{n-t}$$

which has the form of a beta($t+1, n-t+1$) density. Therefore, the class of beta distributions with integer parameters is a natural conjugate family to the Bernoulli sampling model. Nothing substantial is lost however if this class is enlarged to the class of all beta distributions, including all positive values for the parameters. This new class, strictly speaking, is no longer a natural conjugate family. Nevertheless it keeps the essence of the definition and is used in practice.

Natural conjugate families are especially useful because an objective meaning can be attributed to the hyperparameters involved. Revisiting the above example, suppose that n_0 hypothetical (or not) trials were previously made, with t_0 successes. Then, the likelihood l^* of this hypothetical experiment would be $l^*(\theta) \propto \theta^{t_0} (1-\theta)^{n_0-t_0}$. If our (subjective) prior information is equivalent to that provided by the experiment described above, the prior for θ will be a beta with hyperparameters t_0+1 and n_0-t_0+1 .

3.3 Conjugacy with the exponential family

The one-parameter exponential family includes many of the most common probability distributions. An essential characteristic of this family is that there exists a sufficient statistic with fixed dimension. The conjugate class \mathcal{P} to the one-parameter exponential family is easy to characterize. Following the reasoning behind natural conjugacy, it is not difficult to see that members of this class have density

$$p(\theta) \propto \exp\{\alpha\phi(\theta) + \beta b(\theta)\}$$

and so

$$p(\theta | x) \propto \exp\{[\alpha + u(x)]\phi(\theta) + [\beta + 1]b(\theta)\}.$$

Denoting the constant involved in the definition of $p(\theta)$ by $k(\alpha, \beta)$, the constant associated to $p(\theta | x)$ will be $k(\alpha + u(x), \beta + 1)$. Using k as defined above it is easy to obtain $p(x)$ without explicitly calculating $\int p(x | \theta) p(\theta) d\theta$. From the equation $p(x) p(\theta | x) = p(x | \theta) p(\theta)$, it follows that

$$p(x) = \frac{p(x | \theta) p(\theta)}{p(\theta | x)}.$$

Substituting the densities previously obtained we get

$$p(x) = \frac{a(x) \exp\{u(x)\phi(\theta) + b(\theta)\} k(\alpha, \beta) \exp\{\alpha\phi(\theta) + \beta b(\theta)\}}{k(\alpha + u(x), \beta + 1) \exp\{[\alpha + u(x)]\phi(\theta) + [\beta + 1]b(\theta)\}}$$

and after some simplification we arrive at

$$p(x) = \frac{a(x) k(\alpha, \beta)}{k(\alpha + u(x), \beta + 1)}.$$

A straightforward extension of the Bernoulli example is the binomial model. In that case, it follows that

$$\begin{aligned} p(x) &= \frac{\binom{n}{x} \theta^x (1-\theta)^{n-x} B^{-1}(\alpha, \beta) \theta^{\alpha-1} (1-\theta)^{\beta-1}}{B^{-1}(\alpha+x, \beta+n-x) \theta^{\alpha+x-1} (1-\theta)^{\beta+n-x-1}} \\ &= \binom{n}{x} \frac{B(\alpha+x, \beta+n-x)}{B(\alpha, \beta)}, \text{ for } x = 0, 1, \dots, n, n \geq 1. \end{aligned}$$

This is the beta-binomial distribution.

In general, from a sample of size n of the exponential family we obtain the joint density

$$p(\mathbf{x} | \theta) = \left[\prod_{i=1}^n a(x_i) \right] \exp \left\{ \left[\sum_{i=1}^n u(x_i) \right] \phi(\theta) + nb(\theta) \right\}.$$

The use of a conjugate prior $p(\theta) = k(\alpha, \beta) \exp\{\alpha\phi(\theta) + \beta b(\theta)\}$ leads to the posterior density

$$\begin{aligned} p(\theta | \mathbf{x}) &= k \left(\alpha + \sum_{i=1}^n u(x_i), \beta + n \right) \\ &\times \exp \left\{ \left[\alpha + \sum_{i=1}^n u(x_i) \right] \phi(\theta) + [\beta + n]b(\theta) \right\} \end{aligned}$$

and the marginal or predictive distribution is

$$p(\mathbf{x}) = \frac{[\prod u(x_i)] k(\alpha, \beta)}{k(\alpha + \sum u(x_i), \beta + n)}.$$

3.4 The main conjugate families

The main members of the exponential family will be presented in this section. The results obtained previously will be applied to these particular cases and the resulting conjugate families obtained. Some of these families were presented before.

3.4.1 Binomial distribution

The family of beta distributions is conjugate to the binomial (or Bernoulli) model as we have shown above.

3.4.2 Normal distribution with known variance

Theorem 2.1 stated that the normal distribution family is conjugate to the normal model, based on a single observation. For the case of a sample of size n , we have seen in Section 2.6 that

$$l(\theta; \mathbf{x}) \propto \exp \left\{ -\frac{n}{2\sigma^2} (\bar{x} - \theta)^2 \right\}$$

where the terms involving σ^2 were incorporated into the proportionality constant. So, the likelihood has the same form as that based on a single observation x , substituting x by \bar{x} and σ^2 by σ^2/n . Another way to say this is to note that \bar{X} is a sufficient statistic for θ and so the likelihood based on the observed value of \bar{X} , which is distributed as $N(\theta, \sigma^2/n)$, is proportional to the likelihood obtained with the individual observations \mathbf{X} . Therefore, the result presented in Theorem 2.1 is still true, with the substitutions mentioned above, i.e., the posterior distribution of θ given \mathbf{x} is $N(\mu_1, \tau_1^2)$, with

$$\mu_1 = \frac{n\sigma^{-2}\bar{x} + \tau^{-2}\mu}{n\sigma^{-2} + \tau^{-2}} \quad \text{and} \quad \tau_1^{-2} = n\sigma^{-2} + \tau^{-2}.$$

3.4.3 Poisson distribution

Suppose that $\mathbf{X} = (X_1, \dots, X_n)$ is a random sample from the Poisson distribution with parameter θ , denoted $\text{Pois}(\theta)$. Its joint probability function is

$$p(\mathbf{x} | \theta) = \prod_{i=1}^n p(x_i | \theta) = \prod_{i=1}^n \frac{e^{-\theta} \theta^{x_i}}{x_i!}$$

and the likelihood function assumes the form

$$l(\theta | \mathbf{x}) \propto e^{-n\theta} \theta^{\sum x_i}.$$

Its kernel has the form $\theta^a e^{-b\theta}$ characterizing a gamma family of distributions. We have already seen in the previous section that the gamma family is closed under sampling. Then the conjugate prior distribution of θ will be $\theta \sim G(\alpha, \beta)$. The posterior density will be

$$p(\theta | \mathbf{x}) \propto \theta^{\alpha + \sum x_i - 1} \exp\{-(\beta + n)\theta\}$$

corresponding to the $G(\alpha + \sum x_i, \beta + n)$ density. The calculation of the predictive distribution, using the method described before, is left as an exercise.

3.4.4 Exponential distribution

Suppose that $\mathbf{X} = (X_1, \dots, X_n)$ is a random sample from the exponential distribution with parameter θ , denoted by $\text{Exp}(\theta)$. Its joint density function is

$$p(\mathbf{x} | \theta) = \theta^n \exp \left\{ -\theta \sum_{i=1}^n x_i \right\}.$$

The form of the likelihood allows recognition of the kernel of the gamma family as a conjugate distribution for θ . Assuming a $G(\alpha, \beta)$ prior, the posterior will have the form

$$\begin{aligned} p(\theta | \mathbf{x}) &\propto \theta^n \exp \left\{ -\theta \sum_{i=1}^n x_i \right\} \theta^{\alpha-1} \exp\{-\beta\theta\} \\ &\propto \theta^{\alpha+n-1} \exp \left\{ -(\beta + \sum x_i)\theta \right\} \end{aligned}$$

which is the density of a $G(\alpha + n, \beta + \sum x_i)$ distribution.

3.4.5 Multinomial distribution

Denote by $\mathbf{X} = (X_1, \dots, X_p)$ and $\boldsymbol{\theta} = (\theta_1, \dots, \theta_p)$, respectively, the number of observed cases and the probabilities associated with each of p categories in a sample of size n . Assume that the following constraints are true: $\sum_{i=1}^p X_i = n$

and $\sum_{i=1}^p \theta_i = 1$. \mathbf{X} is said to have a multinomial distribution with parameters n and $(\theta_1, \dots, \theta_p)$. The joint probability function of the p counts \mathbf{X} is

$$p(\mathbf{x} | \boldsymbol{\theta}) = \frac{n!}{\prod_{i=1}^p x_i!} \prod_{i=1}^p \theta_i^{x_i}.$$

It is not difficult to show that this distribution also belongs to the exponential family. The likelihood function for $\boldsymbol{\theta}$ is $l(\boldsymbol{\theta}) \propto \prod \theta_i^{x_i}$. Its kernel is the same as the kernel of the density of a Dirichlet distribution. The Dirichlet family with integer parameters a_1, \dots, a_p is natural conjugate with respect to the multinomial sampling distribution. Again, little is lost by extending natural conjugacy over all Dirichlet distributions.

The posterior distribution will then be

$$p(\boldsymbol{\theta} | \mathbf{x}) \propto \left[\prod_{i=1}^p \theta_i^{x_i} \right] \left[\prod_{i=1}^p \theta_i^{a_i-1} \right] = \prod_{i=1}^p \theta_i^{x_i+a_i-1}.$$

and, as anticipated, this posterior is also a Dirichlet distribution with parameters a_1+x_1, \dots, a_p+x_p which is denoted by $(\boldsymbol{\theta} | \mathbf{x}) \sim D(a_1+x_1, \dots, a_p+x_p)$. From the above results about the Dirichlet, it is not difficult to obtain the proportionality constant as

$$\frac{\Gamma(a+n)}{\prod_{i=1}^p \Gamma(a_i+x_i)}.$$

This conjugate analysis generalizes the analysis for binomial samples with beta priors.

3.4.6 Normal distribution with known mean and unknown variance

Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample of size n from the $N(\theta, \sigma^2)$, θ known, and $\phi = \sigma^{-2}$. In this case the joint density function will be:

$$l(\phi; \mathbf{x}) = p(\mathbf{x} | \theta, \phi) \propto \phi^{n/2} \exp \left\{ -\frac{\phi}{2} n s_0^2 \right\} \quad \text{where} \quad s_0^2 = \frac{1}{n} \sum (x_i - \theta)^2.$$

The conjugate prior may have the kernel of $l(\phi; \mathbf{x})$, which is in the gamma distribution form. As the gamma family is closed under sampling, we can consider a $G(n_0/2, n_0\sigma_0^2/2)$ prior distribution or, equivalently, that $n_0\sigma_0^2\phi$ has a χ^2 distribution with n_0 degrees of freedom. The posterior distribution of ϕ is obtained using Bayes' theorem,

$$\begin{aligned} p(\phi | \mathbf{x}) &\propto l(\phi; \mathbf{x}) p(\phi) \\ &\propto \phi^{n/2} \exp\{-ns_0^2\phi/2\} \phi^{(n_0/2)-1} \exp\{-n_0\sigma_0^2\phi/2\} \\ &= \phi^{[(n_0+n)/2]-1} \exp\{-(n_0\sigma_0^2 + ns_0^2)\phi/2\}. \end{aligned}$$

The above expression corresponds to the kernel of the gamma distribution as expected. Therefore:

$$\phi | \mathbf{x} \sim G \left(\frac{n_0 + n}{2}, \frac{n_0\sigma_0^2 + ns_0^2}{2} \right)$$

or equivalently

$$(n_0\sigma_0^2 + ns_0^2)\phi | \mathbf{x} \sim \chi_{n_0+n}^2.$$

Then it follows that the gamma or the χ^2 family of distributions is conjugate with respect to the normal sampling model with θ known and σ^2 unknown.

3.4.7 Normal distribution with unknown mean and variance

The conjugate prior distribution for (θ, ϕ) will be presented in two stages. First, the following conditional distribution of θ given ϕ will be considered:

$$(\theta | \phi) \sim N[\mu_0, (c_0\phi)^{-1}]$$

and the marginal prior for ϕ is as stated before, that is,

$$n_0\sigma_0^2\phi \sim \chi_{n_0}^2 \quad \text{or} \quad \phi \sim G \left(\frac{n_0}{2}, \frac{n_0\sigma_0^2}{2} \right)$$

where (n_0, σ_0^2) and (μ_0, c_0) are obtained from the initial information H . This distribution is usually called normal-gamma or normal- χ^2 with parameters $(\mu_0, c_0, n_0, \sigma_0^2)$ and joint density given by:

$$\begin{aligned} p(\theta, \phi) &= p(\theta | \phi)p(\phi) \\ &\propto \phi^{1/2} \exp \left\{ -\frac{c_0\phi}{2}(\theta - \mu_0)^2 \right\} \phi^{n_0/2-1} \exp \left(-\frac{n_0\sigma_0^2\phi}{2} \right) \\ &= \phi^{(n_0+1)/2-1} \exp \left\{ -\frac{\phi}{2} [n_0\sigma_0^2 + c_0(\theta - \mu_0)^2] \right\}. \end{aligned}$$

The marginal prior distribution of θ can be obtained by integration, using the following result:

$$\int_0^\infty \phi^{a-1} e^{-b\phi} d\phi = \frac{\Gamma(a)}{b^a}.$$

Application of the result gives

$$\begin{aligned} p(\theta) &\propto \int_0^\infty \phi^{(n_0+1)/2-1} \exp \left\{ -\frac{\phi}{2} [n_0\sigma_0^2 + c_0(\theta - \mu_0)^2] \right\} d\phi \\ &= \frac{\Gamma((n_0+1)/2)}{\{[n_0\sigma_0^2 + c_0(\theta - \mu_0)^2]/2\}(n_0+1)/2} \\ &\propto [n_0\sigma_0^2 + c_0(\theta - \mu_0)^2]^{-(n_0+1)/2} \end{aligned}$$

since the $\Gamma(\cdot)$ term does not depend on θ . Rearranging terms gives

$$p(\theta) \propto \left[1 + \frac{(\theta - \mu_0)^2}{n_0(\sigma_0^2/c_0)} \right]^{-(n_0+1)/2}$$

which is the kernel of the Student t distribution with n_0 degrees of freedom, location parameter μ_0 and scale parameter σ_0^2/c_0 , denoted by $t_{n_0}(\mu_0, \sigma_0^2/c_0)$.

The conditional distribution of $\phi \mid \theta$ can be obtained from the joint distribution of (θ, ϕ) and is a $G\{(n_0+1)/2, [n_0\sigma_0^2 + c_0(\theta - \mu_0)^2]/2\}$, or equivalently, $[n_0\sigma_0^2 + c_0(\theta - \mu_0)^2]\phi \mid \theta \sim \chi_{2(n_0+1)}^2$.

The joint distribution of a random sample $\mathbf{X} = (X_1, \dots, X_n)$ is

$$\begin{aligned} p(\mathbf{x} \mid \theta, \phi) &= \prod_{i=1}^n \phi^{1/2} \exp \left\{ -\frac{\phi}{2}(x_i - \theta)^2 \right\} \\ &\propto \phi^{n/2} \exp \left\{ -\frac{\phi}{2} [ns^2 + n(\bar{x} - \theta)^2] \right\} \end{aligned}$$

where

$$s^2 = \frac{1}{n} \sum (x_i - \bar{x})^2$$

as we have seen in Section 2.4. The above expression has the same kernel as the normal-gamma density for (θ, ϕ) . Next, it is necessary to check if the normal-gamma family is closed under sampling. It is not difficult to verify that it is. The posterior distribution will then be

$$\begin{aligned} p(\theta, \phi \mid \mathbf{x}) &\propto p(\mathbf{x} \mid \theta, \phi)p(\theta, \phi) \\ &\propto \phi^{[(n+n_0+1)/2]-1} \\ &\quad \times \exp \left\{ -\frac{\phi}{2} [n_0\sigma_0^2 + ns^2 + c_0(\theta - \mu_0)^2 + n(\bar{x} - \theta)^2] \right\}. \end{aligned}$$

It is not difficult to show that

$$c_0(\theta - \mu_0)^2 + n(\theta - \bar{x})^2 = (c_0 + n)(\theta - \mu_1)^2 + \frac{c_0n}{c_0 + n} (\mu_0 - \bar{x})^2$$

where $\mu_1 = (c_0\mu_0 + n\bar{x})/(c_0 + n)$. Thus it follows that the posterior density for (θ, ϕ) is proportional to

$$\begin{aligned} &\phi^{[(n+n_0+1)/2]-1} \\ &\times \exp \left\{ -\frac{\phi}{2} [n_0\sigma_0^2 + ns^2 + \frac{c_0n}{c_0 + n} (\mu_0 - \bar{x})^2 + (c_0 + n)(\theta - \mu_1)^2] \right\}. \end{aligned}$$

Therefore, the joint posterior for $(\theta, \phi \mid \mathbf{x})$ is normal-gamma with parameters $(\mu_1, c_1, n_1, \sigma_1^2)$ given by

$$\mu_1 = \frac{c_0\mu_0 + n\bar{x}}{c_0 + n}$$

$$c_1 = c_0 + n$$

$$n_1 = n_0 + n$$

$$n_1\sigma_1^2 = n_0\sigma_0^2 + ns^2 + \frac{c_0n}{c_0 + n} (\mu_0 - \bar{x})^2.$$

Prior and posterior distributions are members of the same family. So, the normal-gamma family is conjugate with respect to the normal sampling model when θ and σ^2 are both unknown. Table 3.1 summarizes the distributions involved in the Bayesian analysis of the normal models with unknown mean and variance.

Table 3.1 Summary of the distributions

	Prior	Posterior
$\theta \mid \phi$	$N(\mu_0, (c_0\phi)^{-1})$	$N(\mu_1, (c_1\phi)^{-1})$
ϕ	$n_0\sigma_0^2\phi \sim \chi_{n_0}^2$	$n_1\sigma_1^2\phi \sim \chi_{n_1}^2$
θ	$t_{n_0}(\mu_0, \sigma_0^2/c_0)$	$t_{n_1}(\mu_1, \sigma_1^2/c_1)$
$\phi \mid \theta$	$[n_0\sigma_0^2 + c_0(\theta - \mu_0)^2]\phi \sim \chi_{n_0+1}^2$	$[n_1\sigma_1^2 + c_1(\theta - \mu_1)^2]\phi \sim \chi_{n_1+1}^2$

3.5 Non-informative priors

Many statisticians show concern about the nature of the prior distribution. This is mainly due to an influence from the frequentist point of view. They typically maintain that the prior distribution is arbitrary and alters the conclusions about the statistical problem at hand. Therefore, they argue that prior information is not acceptable for use in a scientific context. In this section, the concept of non-informative or reference prior will be presented in an effort to reconcile these arguments with the Bayesian point of view. The idea behind these priors comes from the desire to make statistical inference based on a minimum of subjective prior information. This minimum is clearly a relative concept and should take into consideration, for example, the sample information content.

Another context where the concept of a reference prior may be useful was outlined in the example of the two physicists in Chapter 2. Let us suppose that two scientists have strong and divergent prior opinions about an unknown quantity and that is not possible to reconcile these initial opinions. This is a situation where it is necessary to produce a ‘neutral’ analysis, introducing a referential. Another plausible justification to support a reference analysis is the usual expectation that the evidence from the experiment is stronger than the prior.

Initially, uniform priors were proposed to represent situations where little or no initial information is available or, even, if it is available and we do not wish to use it. So, $p(\theta) \propto k$ for θ varying in a given subset of the real line means that none of the particular values of θ is preferred (Bayes, 1763). This choice brings some difficulties with it. The first one is that $p(\theta)$ is not a proper distribution if the range

of values of θ is unbounded. This means that $\int p(\theta) d\theta \rightarrow \infty$ which goes against the basic rules of probability. Also, if $\phi = \phi(\theta)$ is a one-to-one transformation of θ and if θ is uniformly distributed, then by the theorem of variable transformation, the density of ϕ is

$$p(\phi) = p(\theta(\phi)) \left| \frac{d\theta}{d\phi} \right| \propto \left| \frac{d\theta}{d\phi} \right|$$

which is only constant if ϕ is defined by a linear transformation. However, the same assumptions leading to the specification of $p(\theta) \propto k$ should also lead to $p(\phi) \propto k$, which contradicts the above deduction. Ideally, we would like to state an invariant rule that would not violate results about variable transformation.

In practice, we are concerned with the posterior distribution, which is often proper, even when the prior distribution is not. In this case, one doesn't need to give much relevance to the impropriety of the prior distribution. Careful examination must be carried out to make sure the posterior is actually proper to proceed confidently with the analysis.

The class of non-informative prior proposed by Jeffreys (1961) is invariant but in many cases leads to improper distributions. This class of priors is extensively used by Box and Tiao (1992). Intuitively, it tries to provide as little prior information as possible, relative to the sample information. It is not surprising therefore that it should depend on Fisher information measures.

Definition. Consider an observation X with probability (density) function $p(x | \theta)$. The Jeffreys non-informative prior has density given by

$$p(\theta) \propto [I(\theta)]^{1/2}, \quad \theta \in \Theta.$$

In the multivariate case, the density is given by

$$p(\theta) \propto |\mathbf{I}(\theta)|^{1/2}.$$

Lemma. The Jeffreys prior $p(\theta) \propto [I(\theta)]^{1/2}$ is invariant under one-to-one transformations, that is, if $\phi = \phi(\theta)$ is a one-to-one transformation of θ , then the Jeffreys prior for ϕ is $p(\phi) \propto [I(\phi)]^{1/2}$.

Proof. Let $\phi = \phi(\theta)$ be a one-to-one transformation of θ . Taking the derivative of $\log p(X | \phi)$ with respect to ϕ , it follows that

$$\frac{\partial \log p(X | \phi)}{\partial \phi} = \frac{\partial \log p(X | \phi(\theta))}{\partial \theta} \frac{\partial \theta}{\partial \phi}$$

where $\theta = \theta(\phi)$ is the inverse transformation of ϕ . To obtain the Fisher information of a parameter, the log likelihood of this parameter needs to be differentiated twice. This gives

$$\frac{\partial^2 \log p(X | \phi)}{\partial \phi^2} = \frac{\partial \log p(X | \phi(\theta))}{\partial \theta} \frac{\partial^2 \theta}{\partial \phi^2} + \frac{\partial^2 \log p(X | \phi(\theta))}{\partial \theta^2} \left(\frac{\partial \theta}{\partial \phi} \right)^2.$$

Multiplying both sides by (-1) and calculating the expected value with respect to $p(x | \theta)$, gives

$$\begin{aligned} I(\phi) &= E_{X|\theta} \left[\frac{\partial \log p(X | \theta)}{\partial \theta} \right] \frac{\partial^2 \theta}{\partial \phi^2} + I(\theta) \left(\frac{\partial \theta}{\partial \phi} \right)^2 \\ &= I(\theta) \left(\frac{\partial \theta}{\partial \phi} \right)^2 \end{aligned}$$

since $E_{X|\theta} [\partial \log p(X | \theta) / \partial \theta] = 0$, as seen in Chapter 2. Therefore, $I^{1/2}(\phi) = I^{1/2}(\theta) |\partial \theta / \partial \phi|$. By the rules of probability, if θ has density proportional to $I^{1/2}(\theta)$ then ϕ has density

$$p(\phi) \propto I^{1/2}(\theta(\phi)) |\partial \theta / \partial \phi| = I^{1/2}(\phi)$$

and the specification is invariant to one-to-one transformation. \square

Corollary. The same result is true for the multiparameter case, that is, the Jeffreys prior is invariant under one-to-one transformations.

There is only one transformation ψ of θ which satisfies the invariance rule and has constant density. This transformation is easily obtained by making

$$p(\psi) \propto I^{1/2}(\theta) |\partial \theta / \partial \psi| \propto k$$

or

$$|\partial \theta / \partial \psi| \propto I^{-1/2}(\theta) \implies |\partial \psi / \partial \theta| \propto I^{1/2}(\theta) \implies \psi \propto \int^\theta I^{1/2}(u) du.$$

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a sample of $\text{Pois}(\theta)$ variables. The joint density of the observations is

$$p(\mathbf{x} | \theta) = \frac{e^{-n\theta} \theta^{\sum x_i}}{\prod x_i!}.$$

Taking logarithm, it follows that

$$\log p(\mathbf{x} | \theta) = -n\theta + \sum_{i=1}^n x_i \log \theta - \log \prod_{i=1}^n x_i!.$$

The first and second order derivatives of the log likelihood are

$$\frac{\partial \log p(\mathbf{x} | \theta)}{\partial \theta} = -n + \frac{\sum_{i=1}^n x_i}{\theta} \quad \text{and} \quad \frac{\partial^2 \log p(\mathbf{x} | \theta)}{\partial \theta^2} = -\frac{\sum x_i}{\theta^2}.$$

Then, the Fisher information will be:

$$I(\theta) = E_{\mathbf{X}|\theta} \left[\frac{\sum X_i}{\theta^2} \right] = \frac{1}{\theta^2} \sum E(X_i) = \frac{n\theta}{\theta^2} = \frac{n}{\theta}$$

and the non-informative prior is $p(\theta) \propto \theta^{-1/2}$. So the posterior density will be

$$\begin{aligned} p(\theta | \mathbf{x}) &\propto p(\mathbf{x} | \theta)p(\theta) \propto e^{-n\theta}\theta^{\sum_i x_i}\theta^{-1/2} \\ &= e^{-n\theta}\theta^{\sum_i x_i - 1/2}, \end{aligned}$$

that is, $\theta | \mathbf{x} \sim G(\sum_i x_i + 1/2, n)$, or alternatively, $2n\theta | \mathbf{x} \sim \chi^2_{2\sum_i x_i + 1}$. The transformation leading to the uniform prior is

$$\phi \propto \int_0^\theta u^{-1/2} du = 2u^{1/2}|_0^\theta \propto \theta^{1/2}.$$

The non-informative prior is frequently obtained from the conjugate prior by letting the scale parameters go to zero and keeping the other ones constant. In the above example, it can be noted that the reference prior is the limit of the gamma distribution (the natural conjugate prior for the Poisson model) with $\theta \sim G(1/2, \epsilon)$, $\epsilon \rightarrow 0$.

The Jeffreys prior specification was alternatively obtained in the univariate case by Bernardo (1979). He called them reference priors and they are defined as the distributions that maximize the amount of unknown information about θ in an infinite number of replications of the experiment. The amount of unknown information about θ in n replications of the experiment is defined as

$$\begin{aligned} I(\mathbf{X}_n, \theta) &= E_{(\mathbf{X}_n, \theta)} \left[\log \frac{p(\theta | \mathbf{X}_n)}{p(\theta)} \right] \\ &= E_{\mathbf{X}_n} \left[E_{\theta | \mathbf{X}_n} \left[\log \frac{p(\theta | \mathbf{X}_n)}{p(\theta)} \right] \right] \end{aligned}$$

where $\mathbf{X}_n = (X_1, \dots, X_n)$. The amount of unknown information about θ in an infinite number of replications of the experiment is obtained as the limit of the information based on n replications when $n \rightarrow \infty$.

In the multivariate case, Bernardo proposed a modification to Jeffreys rule. He suggested a two-stage procedure. The parametric vector is divided into two components: θ denoting the parameters of interest and ϕ , the nuisance parameters. First, a (conditional) reference prior distribution $p(\phi | \theta)$ is obtained. This prior is used to eliminate the parameters ϕ from the likelihood and gives a marginal likelihood $p(\mathbf{x} | \theta)$ as we have seen in Section 2.4. Then, this likelihood is used to obtain the (marginal) reference prior $p(\theta)$. Finally, the complete reference prior is obtained by the multiplication rule $p(\phi, \theta) = p(\phi | \theta)p(\theta)$. This procedure seems to provide better results than that proposed by Jeffreys, although it depends on an arbitrary partition of the parametric vector. Therefore, reference priors are not invariant to the choice of the parameter partition. Another similar procedure, based on information maximization, was proposed by Zellner (1971).

There are other difficulties associated with the reference prior distribution besides its specification not being unique and often leading to an improper density. It can lead to incoherent inferences in the sense that if analysis conditional on the

sample is replaced by analysis conditional on a sufficient statistic, which should not affect inferences, it could lead to different posterior distributions for some parameter transformations. This is obviously an unpleasant situation that fortunately does not occur very often.

Another drawback of Jeffreys priors is that they do not satisfy the likelihood principle. Since it is based on the experiment, a Jeffreys prior produces different results for equal likelihoods. A famous example illustrating that the non-informative prior depends on the sample models is provided by Bernoulli trials with success probability θ . If the sample design is such that n fixed Bernoulli trials are made and the number of success observed, then $X \sim \text{bin}(n, \theta)$ and

$$p(x | \theta) = \binom{n}{x} \theta^x (1-\theta)^{n-x}$$

which implies that

$$\log p(x | \theta) = \log \binom{n}{x} + x \log \theta + (n-x) \log(1-\theta).$$

Therefore, the first and second derivatives of the log likelihood are

$$\frac{\partial \log p(x | \theta)}{\partial \theta} = \frac{x}{\theta} - \frac{n-x}{1-\theta} \quad \text{and} \quad \frac{\partial^2 \log p(x | \theta)}{\partial \theta^2} = -\frac{x}{\theta^2} - \frac{n-x}{(1-\theta)^2}.$$

Then the expected information measure is

$$I_B(\theta) = \frac{E(X | \theta)}{\theta^2} + \frac{E(n-X | \theta)}{(1-\theta)^2} = \frac{n}{\theta(1-\theta)}$$

and the non-informative prior is $p_B(\theta) \propto \theta^{-1/2}(1-\theta)^{-1/2}$. Note that this prior is a beta(1/2, 1/2) distribution and it is not improper.

Now, suppose that the sample scheme consists in observing the number of replications until s successes are obtained. The observation now is Y with negative binomial distribution denoted by $Y \sim NB(s, \theta)$ and

$$p(y | \theta) = \binom{n-1}{s-1} \theta^s (1-\theta)^{y-s}$$

which implies that the first and second derivatives with respect to θ are

$$\frac{\partial \log p(y | \theta)}{\partial \theta} = \frac{s}{\theta} - \frac{y-s}{1-\theta} \quad \text{and} \quad \frac{\partial^2 \log p(y | \theta)}{\partial \theta^2} = -\frac{s}{\theta^2} - \frac{y-s}{(1-\theta)^2}.$$

The expected information is

$$\begin{aligned} I_{BN}(\theta) &= \frac{s}{\theta^2} + \frac{E(Y-s | \theta)}{(1-\theta)^2} \\ &= \frac{s}{\theta(1-\theta)}, \quad \text{since } E(Y | \theta) = \frac{s}{\theta} \end{aligned}$$

and the non-informative prior is $p_{BN}(\theta) \propto \theta^{-1}(1-\theta)^{-1/2}$. This prior is the limit of a beta ($\epsilon, 1/2$) when $\epsilon \rightarrow 0$ and then it is an improper distribution.

Suppose now that in 10 experiments, three successes were observed. Therefore, $I(\theta; x=3) \propto \theta^3(1-\theta)^7$ and $I(\theta; y=10) \propto \theta^3(1-\theta)^7$. The sample information, namely the likelihood, about θ is the same for both models. Although the information is the same, the prior and consequently the posterior distributions are different, with

$$p(\theta | x=3) \propto \theta^{2.5}(1-\theta)^{6.5} \quad \text{and} \quad p(\theta | y=10) \propto \theta^2(1-\theta)^{6.5}$$

showing, then, that the non-informative prior violates the likelihood principle.

These criticisms only reinforce the point that non-informative priors must be used with extreme care. Nevertheless, if proper care is taken, they provide a useful benchmark in a number of problems and may be a useful input to the analysis. Applications of non-informative priors in a few frequently used models described in Chapter 2 are presented below, starting with the location model.

If X has the location model, then

$$\frac{\partial \log p(x | \theta)}{\partial \theta} = \frac{\partial \log f(x - \theta)}{\partial \theta} = -\frac{f'(x - \theta)}{f(x - \theta)} \quad \text{where } f' = \frac{\partial f}{\partial \theta}.$$

Then by the last lemma,

$$I(\theta) = E_{X|\theta} \left[\left(-\frac{f'(X - \theta)}{f(X - \theta)} \right)^2 \right].$$

Making the transformation $U = X - \theta$, then

$$I(\theta) = E_U \left(-\frac{f'(U)}{f(U)} \right)^2$$

which does not depend on θ . So, $I(\theta) = k$ and then $p(\theta) \propto k$. It is easy to see that this result is also true for a parameter vector θ .

One way to justify this prior is through model invariance directly. Working with an observation vector \mathbf{X} and location parameter θ is equivalent to working with observation vector $\mathbf{Y} = \mathbf{X} + \mathbf{c}$ and location parameter $\eta = \theta + \mathbf{c}$ for any given constant \mathbf{c} . We can then insist on the same non-informative prior specification for θ as for η . It is not difficult to show that the only distribution satisfying this requirement has density $p(\theta) \propto k$.

If X has the scale model, then

$$\begin{aligned} \frac{\partial \log p(x | \sigma)}{\partial \sigma} &= \frac{\partial \log[\sigma^{-1}f(x/\sigma)]}{\partial \sigma} \\ &= \frac{\partial}{\partial \sigma} \left[-\log \sigma + \log f\left(\frac{x}{\sigma}\right) \right] \\ &= -\frac{1}{\sigma} + \frac{\partial \log f(x/\sigma)}{\partial \sigma} \left(-\frac{x}{\sigma^2} \right) \\ &= -\frac{1}{\sigma} \left[1 + \frac{x}{\sigma} \frac{f'(x/\sigma)}{f(x/\sigma)} \right] \quad \text{where } f' = \frac{\partial f}{\partial \sigma}. \end{aligned}$$

Therefore the information measure about σ will be

$$\begin{aligned} I(\sigma) &= \frac{1}{\sigma^2} E_{X|\sigma} \left[\left(1 + \frac{X}{\sigma} \frac{f'(X/\sigma)}{f(X/\sigma)} \right)^2 | \sigma \right] \\ &= \frac{1}{\sigma^2} E \left(1 + U \frac{f'(U)}{f(U)} \right)^2 \end{aligned}$$

after the transformation $U = X/\sigma$. Since the distribution of U does not depend on σ then $I(\sigma) = k \times \sigma^{-2}$ and $p(\sigma) \propto \sigma^{-1}$ is the non-informative prior distribution.

Once again, model invariance can be invoked directly by assuming equivalence between a model with observation X and scale parameter σ and a model with observation $Y = cX$ and scale parameter $\eta = c\sigma$. Insisting on the same non-informative prior for σ and η leads to $p(\sigma) \propto \sigma^{-1}$.

If X has a location-scale model, a reference prior (θ, σ) can be obtained following the procedure proposed by Bernardo with θ being the parameter of interest and σ the nuisance parameter. This partition corresponds to the majority of cases of interest. Now, if θ is supposed known, we are restricted to a scale model and its prior is $p(\sigma | \theta) \propto \sigma^{-1}$. The distribution of $X | \theta$ is now obtained as

$$p(x | \theta) = \int p(x | \sigma, \theta) p(\sigma | \theta) d\sigma$$

where we can observe that the dependence on X and θ will continue to be of the form $f(x - \theta)$. Therefore, in a location-scale model, the reference prior is

$$p(\theta, \sigma) = p(\theta)p(\sigma | \theta) \propto \frac{1}{\sigma}.$$

This prior is also recommended by Jeffreys (1961) although is not the one implied by direct application of his rule.

3.6 Hierarchical priors

A good strategy to specify the prior distribution or for a better description of an experimental situation, is often to divide it into stages or into a hierarchy. The idea of using a hierarchical structure with multistage was formalized by Lindley and Smith (1972). This way, the prior specification is made in two phases:

1. structural, for the division into stages;
2. subjective, for quantitative specification at each stage.

Example. Suppose that Y_1, \dots, Y_n are such that $Y_i \sim N(\theta_i, \sigma^2)$, with σ^2 known. Among many possibilities depending on the situation under study, many choices are available for specification of the prior for $\theta = (\theta_1, \dots, \theta_n)$. The following options can be used:

- θ_i 's are independent, that is, $p(\theta) = \prod_i p(\theta_i)$.
- θ_i 's are a sample from a population with $p(\theta | \lambda)$ where λ contains the parameters describing the population.

So, for the last option,

$$p(\theta | \lambda) = \prod_{i=1}^n p(\theta_i | \lambda).$$

This specification corresponds to the first stage. To complete the prior setting, it is necessary to specify the second stage: the distribution of λ , $p(\lambda)$. Note that $p(\lambda)$ corresponds to the second stage and does not depend on the first stage.

One can then obtain the marginal prior distribution of θ by

$$p(\theta) = \int p(\theta, \lambda) d\lambda = \int p(\theta | \lambda)p(\lambda) d\lambda = \int \prod_{i=1}^n p(\theta_i | \lambda)p(\lambda) d\lambda.$$

Note that the θ_i 's are supposed exchangeable as their subscripts are irrelevant in terms of this prior. Since the distribution of λ is independent of the first stage, it can be stated as:

1. Concentrated: $p(\lambda = \lambda_0) = 1$.
2. Discrete: $p(\lambda = \lambda_j) = p_j$, $j = 1, \dots, k$, with $\sum_j p_j = 1$. In this case the distribution of θ will be a finite mixture of the densities $p(\theta | \lambda_j)$ with weights p_j , $j = 1, \dots, k$.
3. Continuous: as before, the distribution of θ will be a continuous mixture of $p(\theta | \lambda)$ with weights given by $p(\lambda)$.

If the first stage prior assumes that $\theta_i \sim N(\mu, \tau^2)$, $i = 1, \dots, n$, then $\lambda = (\mu, \tau^2)$. Assuming that $p(\tau^2 = \tau_0^2) = 1$ and μ is normally distributed then θ has a multivariate normal distribution. On the other hand, assuming that $p(\mu = \mu_0) = 1$ and τ^{-2} has a gamma prior distribution implies that θ has a multivariate Student t distribution.

This subdivision into stages is a probabilistic strategy that allows easy identification and specification of coherent priors. Nothing prevents these ideas from going further into the hierarchy. For example, the distribution of λ can depend on ϕ . In this case,

$$p(\theta) = \int_{\Phi} \int_{\Lambda} p(\theta | \lambda)p(\lambda | \phi)p(\phi) d\lambda d\phi.$$

The parameters λ and ϕ are called hyperparameters and are introduced to ease the prior specification. Theoretically one can state as many states as one thinks are necessary to improve prior specification. In practice, it is very hard to interpret the parameters of third or higher stages, so it is common practice to use a non-informative prior for these levels.

The concept of hierarchical modelling will be returned to in Chapter 8, at least for the normal case. That chapter provides an introduction to more elaborate models where the full strength of prior specification will be better appreciated.

Exercises

§ 3.1

1. Let θ represent the maximum temperature at your house door in September.
 - Determine, subjectively, the 0.25 and 0.5 quantiles of your prior distribution for θ .
 - Obtain the normal density that best fits these quantiles.
 - Find subjectively (without using the normal density obtained in (b)) the 0.1 quantile of your prior distribution for θ . Is it consistent with the normal obtained in (b)? What can you conclude from this fact?
2. Let θ be the probability that a football team from Rio de Janeiro will be the winner of the next Brazilian championship. Supposing that θ does not vary in time, build a prior distribution for θ based on past information.

§ 3.2

3. Show that the classes of beta and normal-gamma distributions are closed under sampling.

§§ 3.3/3.4

4. Show that the beta family is conjugate with respect to the binomial, geometric and negative binomial sampling distributions.
5. For four pairs of a rare specimen of bird that nested last season, the number of eggs per nest n and the number of eggs hatched Y were observed, providing the data $n = 2, 3, 3$ and 4 and $y = 1, 2, 3$ and 3 . For a fifth pair nesting this season in similar conditions, $n_5 = 3$ eggs were observed and are about to hatch. Let θ be the probability that an egg is hatched.
 - Obtain the likelihood function for θ , based on the observations $y = (y_1, \dots, y_4)$.
 - Assess a conjugate prior that in your opinion is adequate and calculate the posterior of $\theta | y$.
 - State a probabilistic model for Y_5 , the number of eggs hatched in the fifth nest and obtain its predictive distribution.
6. Suppose that a random sample $\mathbf{X} = (X_1, \dots, X_n)$ from the $N(\theta, \sigma^2)$ is observed with θ known and that a χ^2 prior for σ^{-2} is used. If the prior coefficient of variation (CV) of σ^{-2} is equal to 0.5, what must the value of n be to ensure that the posterior CV reduces to 0.1?

Note: The coefficient of variation of X (CV) is defined by $\sigma/|\mu|$, where $\mu = E(X)$ and $\sigma^2 = \text{var}(X)$ represent the mean and the standard deviation of X , respectively.

7. Let X_1, \dots, X_n be a random sample of the $N(\theta, \phi^{-1})$ distribution and consider the conjugate prior distribution for θ and ϕ .
- Determine the parameters $(\mu_0, c_0, n_0, \sigma_0)$ of the prior distribution, knowing that $E(\theta) = 0$, $Pr(|\theta| < 1.412) = 0.5$, $E(\phi) = 2$ and $E(\phi^2) = 5$.
 - In a sample of size $n = 10$, $\bar{X} = 1$ and $\sum_{i=1}^n (X_i - \bar{X})^2 = 8$ were observed. Determine the posterior distribution of θ and sketch a graph of the prior, posterior and likelihood functions with ϕ fixed.
 - Obtain $Pr(|Y| > 1 | \mathbf{x})$, where Y is a new observation taken from the same population.
8. A random sample X_1, \dots, X_n is selected from the $N(\theta, \sigma^2)$ distribution, with σ^2 known. The prior distribution for θ is a $N(\mu_0, \sigma_0^2)$. What must the sample size be to reduce the variance
- of the posterior distribution of θ to σ_0^2/k ($k > 1$)?
 - of the predictive distribution of Y , a future observation drawn from the same population to σ_0^2/k ($k > 1$)?
9. Consider the sampling model $\mathbf{X} \sim N(\theta, \sigma^2 \mathbf{I}_p)$, where the p -dimensional mean vector θ and the scalar σ^2 are known. Show that the distribution family given by $\theta | \sigma^2 \sim N(\mu, \sigma^2 \mathbf{C}_0)$ and $n_0 \sigma_0^2 / \sigma^2 \sim \chi_{n_0}^2$ is conjugate to the model presented above, generalizing the results obtained in Section 3.3 for univariate normal distributions.
10. Let X_1, \dots, X_n be a random sample from the $\text{Pois}(\theta)$ distribution.
- Determine the conjugate prior parameters for θ assuming that $E(\theta) = 4$ and $CV(\theta) = 0.5$ and determine n such that $V(\theta | \mathbf{x}) < 0.01$.
 - Show that the posterior mean is of the form
- $$\gamma_n \bar{X}_n + (1 - \gamma_n) \mu_0,$$
- where $\mu_0 = E(\theta)$ and that $\gamma_n \rightarrow 1$ when $n \rightarrow \infty$.
11. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample of the $U(0, \theta)$ distribution.
- Show that the Pareto family of distributions, with parameters a and b , and density $p(\theta) = ab^a/\theta^{1+a}$, $\theta > b$, $(a, b > 0)$, is a conjugate family to the uniform.
 - Obtain the mode, mean and median of the posterior distribution of θ .
12. Consider the conjugate model Poisson–gamma with $n = 1$. Obtain the predictive distribution using
- the usual integration procedures;
 - the approach described at the end of Section 3.3;

(c) calculate also the mean and variance of this distribution.

13. Show that

$$c_0(\theta - \mu_0)^2 + n(\theta - \bar{x})^2 = (c_0 + n)(\theta - \mu_1)^2 + \frac{c_0 n}{c_0 + n} (\mu_0 - \bar{x})^2$$

where $\mu_1 = (c_0 \mu + n \bar{x})/(c_0 + n)$.

§ 3.5

14. Consider the observation of a sample $\mathbf{X} = (X_1, \dots, X_n)$ with probability (or density) function $p(x | \theta)$.

- (a) Show that

$$E_{\theta|\mathbf{x}} \left[\log \frac{p(\theta | \mathbf{x})}{p(\theta)} \right] \geq 0 \quad \forall \mathbf{x}$$

with equality obtained only when $p(\theta | \mathbf{x}) = p(\theta)$.

- (b) Interpret the above result.

15. Let $X_i \sim p(x_i | \theta_i)$ and $p_i(\theta_i)$ the non-informative prior for θ_i , for $i = 1, \dots, p$. Assuming that the X_i 's are independent, show that the non-informative Jeffreys prior for $\theta = (\theta_1, \dots, \theta_p)$ is given by $\prod_{i=1}^p p_i(\theta_i)$.

16. Consider a random sample of size n from the Pareto distribution with parameters θ and b , respectively.

- (a) Show that there is a sufficient statistic of fixed dimension for θ .

- (b) Obtain the non-informative prior for θ . Is it improper?

17. Suppose that $X | \theta \sim \text{Exp}(\theta)$ and that the prior for θ is non-informative.

- (a) Obtain the predictive distribution of X and show that $p(x)$ and $p(x | \theta)$ are monotonically decreasing in x .

- (b) Calculate the mode and the median of the sampling distribution and of the predictive distribution.

18. Suppose that $\mathbf{X} = (X_1, X_2, X_3)$ has a trinomial distribution with parameters n and $\boldsymbol{\pi} = (\pi_1, \pi_2, \pi_3)$, where $\pi_1 + \pi_2 + \pi_3 = 1$ and n is known, with density given by

$$f(x_1, x_2 | \theta_1, \theta_2) = \frac{n!}{x_1! x_2! (1 - x_1 - x_2)!} \theta_1^{x_1} \theta_2^{x_2} (1 - \theta_1 - \theta_2)^{n-x_1-x_2}$$

where $x_i = 0, 1, \dots, n$, $i = 1, 2$, $0 \leq x_1 + x_2 \leq n$. Show that the non-informative Jeffreys prior for $\boldsymbol{\pi}$ is $p(\boldsymbol{\pi}) \propto [\pi_1 \pi_2 (1 - \pi_1 - \pi_2)]^{-1/2}$.

19. Suppose that the lifetimes of n bulbs are exponentially distributed with mean θ .

- (a) Obtain a non-informative prior for θ and show that it is improper.

- (b) Suppose that the times are observed up until r failures have occurred. Obtain the likelihood expression.

- (c) Show that, if no bulbs fail before a pre-specified time limit $c > 0$ when observation stops, then the posterior distribution is also improper.
20. Suppose that θ has non-informative prior $p(\theta) \propto k$. Show that $\phi = a\theta + b$, $a \neq 0$, also has prior $p(\phi) \propto k$. Suppose now that θ has non-informative prior $p(\theta) \propto \theta^{-1}$, $\theta > 0$. Show that $\phi = \theta^a$, $a \neq 0$, also has prior $p(\phi) \propto \phi^{-1}$ and that $\psi = \log \theta$ has prior $p(\psi) \propto k$.
21. Let $\mathbf{X} = (X_1, X_2, X_3)$ be a random vector with trinomial distribution with parameters n and $(\theta_1, \theta_2, \theta_3)$. The statistician decides to reparametrize the problem defining $\lambda = \theta_1/(\theta_1 + \theta_2)$ and $\psi = \theta_1 + \theta_2$. (This is a valid procedure since the transformation from (θ_1, θ_2) to (λ, ψ) is 1-to-1.)
- Write the density of \mathbf{X} as a function of λ and ψ .
 - Show that $T = X_1 + X_2$ is a sufficient statistic for ψ . Interpret these results in terms of the inference about ψ .
 - Obtain the non-informative prior for ψ based on the result proved in (b).
 - Interpreting ψ as the success probability in an experiment and supposing that in n repetitions of the experiment t successes were observed, what is the probability of a future experiment being a success?
22. Show that the Jeffreys prior $p(\theta) \propto |I(\theta)|^{1/2}$ is invariant under one-to-one transformations, that is, if $\phi = \phi(\theta)$ is a one-to-one transformation of θ , then the Jeffreys prior for ϕ is $p(\phi) \propto |\det I(\phi)|^{1/2}$.
23. Assuming that to work with an observation vector \mathbf{X} and location parameter θ is equivalent to working with an observation vector $\mathbf{Y} = \mathbf{X} + \mathbf{c}$ and location parameter $\eta = \theta + \mathbf{c}$ for any given constant \mathbf{c} and insisting on the same non-informative prior specification for θ as for η , show that the only possible distribution for θ has density $p(\theta) \propto k$.
24. Repeat the above exercise under the conditions of the scale model to show that the non-informative prior must have density in the form $p(\sigma) \propto \sigma^{-1}$ by
- assuming equivalence between the model with observation X and scale parameter σ and the model with observation $Y = cX$ and scale parameter $\eta = c\sigma$;
 - transforming the problem into a location model with observation $Z = \log X$ and location $\xi = \log \sigma$.

§ 3.6

25. Assume that the first stage prior specifies that $\theta_i \sim N(\mu, \tau^2)$, $i = 1, \dots, n$, and define $\lambda = (\mu, \tau^2)$.
- Assuming that $p(\tau^2 = \tau_0^2) = 1$ and μ is normally distributed, prove that θ has a multivariate normal distribution.
 - Assuming that $p(\mu = \mu_0) = 1$ and τ^{-2} has a gamma prior distribution, prove that θ has a multivariate Student t distribution.

26. (DeGroot, 1970, p. 154) Suppose that the prior distribution $p(\theta)$ for θ is built up hierarchically as follows

- If $\xi = i$, then the prior density for θ is $p_i(\theta)$, $i = 1, \dots, k$.
- The distribution of ξ is $p(\xi = i) = c_i$, $i = 1, \dots, k$.

Suppose also that X , with density $p(x | \theta)$, is observed and define Ψ_i as the class, containing p_i , of conjugate distributions to the sampling distribution of X , $i = 1, \dots, k$, and Ψ as the class of distributions given by

$$\{p : p(\theta) = \sum_{i=1}^k \beta_i p_i(\theta) \text{ and } p_i \in \Psi_i\}.$$

- What is the mathematical expression for the prior of $p(\theta)$?
- Show that Ψ is conjugate to the distribution of X , that is, there exist constants b_i , $i = 1, \dots, k$, such that

$$p(\theta | x) = \sum_{i=1}^k b_i p_i(\theta | x).$$

- Obtain the relationship between b_i and the prior and posterior probabilities of $\xi = i$, $i = 1, \dots, k$.

Hint: Define $h_i(x) = \int p(x | \theta) p_i(\theta) d\theta$ and $p_i(\theta | x) = p(x | \theta) p_i(\theta) / h_i(x)$, $i = 1, \dots, k$.

27. The IQ's of a sample of n senior undergraduate students of statistics at UFRJ are represented respectively by θ_i , $i = 1, \dots, n$, and the common unknown mean for all the final year students at UFRJ is denoted by μ . Suppose that the θ_i 's constitute a random sample of the population of IQ's, with unknown mean but with known variance b . A useful test to assess IQ's is applied, providing the independent observations Y_1, \dots, Y_n , where $Y_i | \theta_i \sim N(\theta_i, a)$, with a known.

- Build a hierarchical prior for the parameters $\theta_1, \dots, \theta_n$.
- Calculate $p(\mu | y_1, \dots, y_n)$ and obtain $E(\mu | y_1, \dots, y_n)$.
- Obtain $p(\theta_i | \mu, y_1, \dots, y_n)$ and $E(\theta_i | \mu, y_1, \dots, y_n)$, for $i = 1, \dots, n$.
- Obtain $E(\theta_i | y_1, \dots, y_n)$, for $i = 1, \dots, n$.

28. Suppose that the prior distribution for $(\theta_1, \dots, \theta_n)$ is such that the θ_i 's constitute a random sample from a $N(\mu, \tau^2)$ distribution and $\mu | \tau^2 \sim N(\mu_0, \tau^2/c)$. Obtain the prior distribution for $(\theta_1, \dots, \theta_n)$ and, in particular, calculate the covariance between θ_i and θ_j , $1 \leq i, j \leq n$, with $i \neq j$ supposing that

- τ^2 is known;
- τ^2 is unknown with prior distribution $\tau^{-2} \sim G(\alpha, \beta)$.

4

Estimation

One of the central problems of statistical inference is discussed in this chapter. The general problem of decision making is briefly described to motivate estimation as a special case. Estimation is then treated from both Bayesian and frequentist points of view.

In Section 4.1, the decision problem is defined and the concepts of loss function and Bayes risk are discussed. Different loss functions are considered in the definition of different decision problems. The Bayes estimators ensuing from these losses are presented and their advantages and disadvantages discussed. The more important classical methods of estimation, namely maximum likelihood, minimal least squares and moments, are presented in Section 4.2. Properties of these methods are extensively discussed. In Section 4.3, methods of comparison of these estimators are defined. The concepts of bias, (classical) risk and consistency of an estimator are introduced. Following a discussion on point estimation, interval estimation is presented in Section 4.4. Finally, the results are applied in Section 4.5 to the estimation of mean and variance in the normal model. Results concerning approximate methods of estimation, including asymptotics, are deferred to the next chapter.

4.1 Introduction to decision theory

Consider the posterior density exhibited in Figure 4.1. This density is not completely uncommon and illustrates the difficulties that may be associated in the learning process involved in a statistical procedure. Nevertheless, this density contains all that is available in terms of a probabilistic description of our information about a quantity of interest. Any attempt to summarize the information contained in this density must be made with caution. The graph of the posterior density is the best description of the inferential process. Sometimes, however, it is useful to summarize further the information into a few numerical figures for communication purposes. The simplest possible case is point estimation where one seeks to determine a single value of the unknown quantity of interest θ that summarizes the entire information of the distribution. Denote this value by $\hat{\theta}$, the point estimator of θ . As we will see below, it is easier to understand the choice of

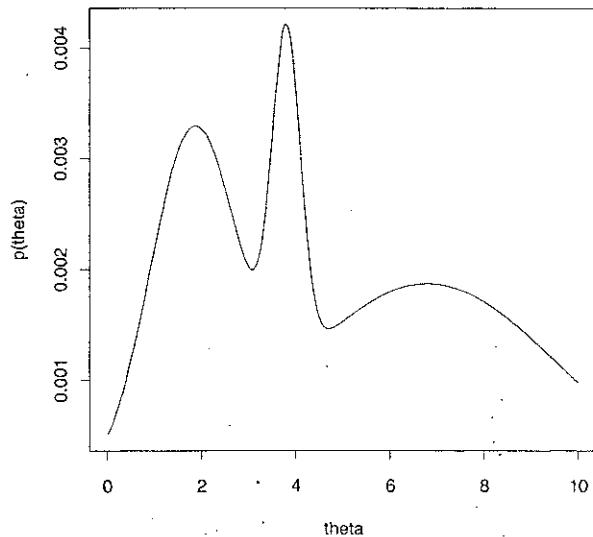


Fig. 4.1 Posterior density of θ with three distinct regions: the first containing around 30% of the total probability, the second with 10% and the third with around 60%. The mode of this density is 3.78, the mean is 4.54 and the median is 4.00.

the value of $\hat{\theta}$ in the context of decision theory.

A decision problem is completely specified by the description of three spaces:

1. parameter (or states of the nature) space Θ ;
2. space of possible results of an experiment Ω ;
3. space of possible actions \mathcal{A} .

A decision rule δ is a function defined in Ω with values in \mathcal{A} , that is $\delta : \Omega \rightarrow \mathcal{A}$. A loss function may be associated to each decision $\delta(x)$ and each possible value of $\theta \in \Theta$. It can be interpreted as the punishment that one suffers for taking decision δ when the value of the parameter is θ . This function from $\Theta \times \mathcal{A}$ with values in \mathbb{R}^+ will be denoted by $L(\delta, \theta)$.

Definition. The risk of a decision rule, denoted by $R(\delta)$, is the expected posterior loss given by $R(\delta) = E_{\theta|x}[L(\delta, \theta)]$.

The importance of the risk is the introduction of a measure that enables one to rank different decision rules.

Definition. A decision rule δ^* is optimal if it has minimum risk, namely $R(\delta^*) < R(\delta)$, $\forall \delta$. This rule is called the Bayes rule and its risk is called the Bayes risk.

Example. Suppose that a doctor must decide if a patient (for example, John from previous chapters) with a given disease must undergo surgery or not. The states of nature are: John is sick ($\theta = 1$) or not ($\theta = 0$). Let us simplify the problem by assuming that the doctor will only prescribe surgery ($\delta = 1$) if he thinks John is sick. This way, a decision rule δ directly related to the value of θ gets established. Unfortunately, the value of the parameter θ is unknown for the decision maker, the doctor. Table 4.1 is a possible representation of the losses (measured in a hypothetical monetary unit) associated with all combination of values of the action and state of nature.

Table 4.1 Losses associated with the doctor problem

θ	δ	
	no surgery - 0	surgery - 1
healthy	0	500
sick	1000	100

These losses represent the subjective evaluation of the decision maker with respect to the combination of actions and states of nature. The smallest loss is null which occurs when the patient is not sick and does not undergo surgery. The largest loss occurs when the patient is sick but is not prescribed surgery. This implies a loss in the doctor's reputation and may even lead to legal problems. He evaluates this loss at 1000 monetary units. Note that all losses are non-negative as defined and do not take into account the doctor's fee which is constant or at least should be immaterial to the problem considered.

As must be clear by now, the decision must be guided by taking into consideration the uncertainty about the unknowns involved in the problem. In this case, the unknown is θ and let us assume that its uncertainty is described by its updated distribution, $Pr(\theta = 1) = \pi$ and $Pr(\theta = 0) = 1 - \pi$, for $0 \leq \pi \leq 1$. This can be a prior distribution or a posterior distribution, obtained after a few tests have been carried out on the patient. Evaluation of the risk of an action δ is straightforward and

$$R(\delta = 0) = E_{\theta}[L(\delta = 0, \theta)] = 0(1 - \pi) + 1000\pi = 1000\pi$$

$$R(\delta = 1) = E_{\theta}[L(\delta = 1, \theta)] = 500(1 - \pi) + 100\pi = 500 - 400\pi.$$

As can be seen from Figure 4.2, the two actions have equal risk if $R(\delta = 0) = R(\delta = 1)$, which happens iff $1000\pi = 500 - 400\pi$, or $\pi = 5/14$. For $\pi < 5/14$, the risk associated with $\delta = 0$ is smaller than the risk associated with $\delta = 1$. In this case, $\delta = 0$ is the Bayes rule and the Bayes risk is 1000π . For $\pi > 5/14$, the problem is reversed; the Bayes rule is $\delta = 1$ and the Bayes risk is $500 - 400\pi$.

In summary, the doctor's strategy must be to prescribe John surgery iff $\pi > 5/14$. This example shows how sensitive the decision is to the choice of priors. It is

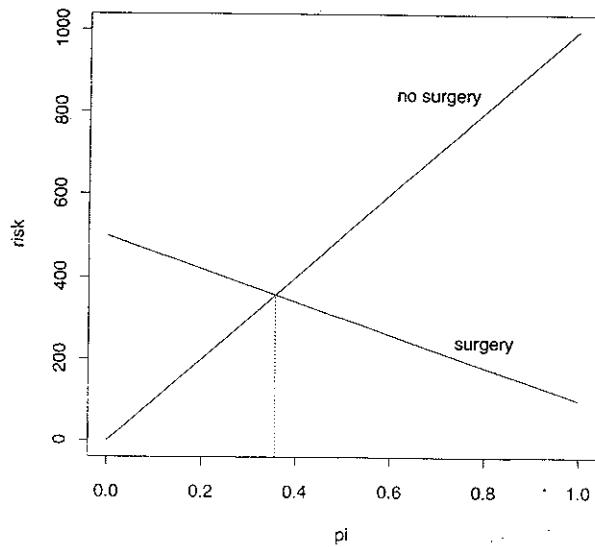


Fig. 4.2 Risks associated with the two actions: surgery or not, as a functions of the probability of disease π .

important to study also the sensitivity of the decision to the choice of loss function. Estimation is clearly dependent on the specified losses and variation to their values may lead to different decisions.

Definition. An estimator is an optimal decision rule with respect to a given loss function. Its observed value is called an estimate.

This definition is broad enough to be useful in a classical perspective, where other optimality criteria will be introduced.

In what follows, most of the presentation will be concentrated on symmetric loss functions of the form $L(\delta, \theta) = h(\delta - \theta)$, for some function h . These are the most commonly used loss functions. In general, $\Theta \subset R$ and the loss functions are continuous.

Lemma. Let $L_1(\delta, \theta) = (\delta - \theta)^2$ be the loss associated with the estimation of θ by δ . (This loss is usually known as quadratic loss.) The estimator of θ is $\delta_1 = E(\theta)$, the mean of the updated distribution of θ .

Proof. We have to calculate the risk and show that δ_1 minimizes it. So, $R(\delta) = E[(\delta - \theta)^2] = E\{[(\delta - \delta_1) + (\delta_1 - \theta)]^2\}$, where $\delta_1 = E(\theta)$. Therefore,

$$\begin{aligned} R(\delta) &= E_\theta[(\delta - \delta_1)^2] + E_\theta[(\delta_1 - \theta)^2] + 2E_\theta[(\delta - \delta_1)(\delta_1 - \theta)] \\ &\approx (\delta - \delta_1)^2 + E_\theta[(\delta_1 - \theta)^2] + 2(\delta - \delta_1)E_\theta[\delta_1 - \theta] \end{aligned}$$

$$\begin{aligned} &= (\delta - \delta_1)^2 + E_\theta[(\delta_1 - \theta)^2], \quad \text{since } \delta_1 = E(\theta) \\ &= (\delta - \delta_1)^2 + V(\theta) \end{aligned}$$

and the risk is minimized for $\delta = \delta_1$. In this case, the Bayes risk is $R(\delta_1) = V(\theta)$ and $R(\delta_1) \leq R(\delta)$, $\forall \delta$, with equality iff $\delta_1 = \delta$. \square

The quadratic loss is sometimes criticized for introducing a penalty that increases strongly with the estimation error $\delta - \theta$. In many cases, it is desirable to have a loss function that does not overly emphasize large estimation errors. The next lemma presents the estimator associated with the absolute loss function, which considers punishments increasing linearly with the estimation error.

Lemma. Let $L_2(\delta, \theta) = |\delta - \theta|$ be the loss associated with the estimation of θ . The estimator of θ is $\delta_2 = \text{med}(\theta)$, the median of the updated distribution of θ .

The proof of this lemma is more cumbersome and will be left as an exercise.

Another form to reduce the effect of large estimation errors is to consider loss functions that remain constant whenever $|\delta - \theta| > k$ for some k arbitrary. There is some freedom for options of suitable values of k . The most common choice is the limiting value as $k \rightarrow 0$. This loss function associates a fixed loss when an error is committed, irrespective of its magnitude. This loss is usually known as the 0-1 loss.

Lemma. Let $L_3(\delta, \theta) = \lim_{\varepsilon \rightarrow 0} I_{|\delta-\theta|}([\varepsilon, \infty))$. The estimator of θ is $\delta_3 = \text{mode}(\theta)$, the mode of the updated distribution of θ .

Proof (for the θ continuous case).

$$\begin{aligned} E[L_3(\delta, \theta)] &= \lim_{\varepsilon \rightarrow 0} \left[\int_{-\infty}^{\delta-\varepsilon} 1 \cdot p(\theta) d\theta + \int_{\delta-\varepsilon}^{\delta+\varepsilon} 0 \cdot p(\theta) d\theta + \int_{\delta+\varepsilon}^{\infty} 1 \cdot p(\theta) d\theta \right] \\ &= \lim_{\varepsilon \rightarrow 0} \left[1 - \int_{\delta-\varepsilon}^{\delta+\varepsilon} p(\theta) d\theta \right] \\ &= 1 - \lim_{\varepsilon \rightarrow 0} P(\delta - \varepsilon < \theta < \delta + \varepsilon). \end{aligned}$$

But $\lim_{\varepsilon \rightarrow 0} P(\delta - \varepsilon < \theta < \delta + \varepsilon) d\theta = p(\delta)$. Note that $E[L_3]$ is minimized when $p(\delta)$ is maximized. Hence, $\delta_3 = \text{mode}(\theta)$. \square

When the updated distribution is the posterior, the estimator associated with the 0-1 loss is the posterior mode. This is also referred to as the generalized maximum likelihood estimator (GMLE). In the next section, we will see the reason for the name. The GMLE is the easiest estimator to be obtained among the estimators presented so far. In the continuous case, it typically involves finding the solution

84 Estimation

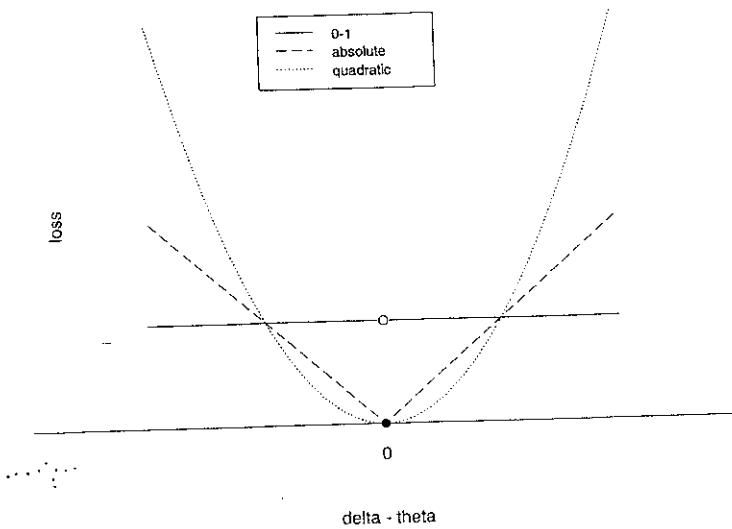


Fig. 4.3 Loss functions: quadratic, — — —; absolute, ······; 0-1, — — —.

of the equation $d p(\theta | \mathbf{x}) / d\theta = 0$. Figure 4.3 illustrates the variation of the loss functions considered here as a function of the estimation error.

Many of these results can be generalized to the multivariate case. Apart from the absolute value that has no clear extension to the multivariate case, the quadratic and 0-1 loss can be respectively extended by,

$$L_1(\delta, \theta) = (\delta - \theta)'(\delta - \theta)$$

and

$$L_3(\delta, \theta) = \lim_{\text{vol}(\mathbf{A}) \rightarrow 0} I_{|\delta - \theta| < \text{vol}(\mathbf{A})}$$

where \mathbf{A} is a region containing the origin and $\text{vol}(\mathbf{A})$ is the volume of the region \mathbf{A} . It is not difficult to show that the Bayes estimators of θ under loss functions L_1 and L_3 are respectively given by the joint mean and joint mode of the updated distribution of θ . These concepts are treated in greater depth by Berger (1985), DeGroot (1970) and Ferguson (1967).

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a sample from a $N(\theta, \sigma^2)$ distribution and $\phi = \sigma^{-2}$. We have previously seen that in a conjugate analysis, the posterior distribution is $\theta | \phi \sim N(\mu_1, (c_1\phi)^{-1})$ and $n_1\sigma_1^2\phi \sim \chi_{n_1}^2$.

To ease the derivation of the mode of the joint distribution, it is usual to work with the logarithm of $p(\theta, \phi | \mathbf{x})$ given by

$$\log p(\theta, \phi | \mathbf{x}) = k - \frac{\phi}{2} \left[c_1(\theta - \mu_1)^2 + n_1\sigma_1^2 \right] + \left(\frac{n_1 + 1}{2} - 1 \right) \log \phi$$

and differentiate it with respect to θ and ϕ leading to

$$\frac{\partial \log p(\theta, \phi | \mathbf{x})}{\partial \theta} = -\frac{\phi}{2} [2c_1(\theta - \mu_1)]$$

and

$$\frac{\partial \log p(\theta, \phi | \mathbf{x})}{\partial \phi} = -\frac{c_1(\theta - \mu_1)^2 + n_1\sigma_1^2}{2} + \left(\frac{n_1 + 1}{2} - 1 \right) \frac{1}{\phi}.$$

Making $\partial \log p(\theta, \phi | \mathbf{x}) / \partial \theta = 0$ gives $\theta = \mu_1$ as a critical point and $\partial \log p(\mu_1, \hat{\phi} | \mathbf{x}) / \partial \phi = 0$ gives $\hat{\phi} = \sigma_1^{-2}(n_1 - 1)/n_1$. The second order conditions are satisfied as

$$\begin{aligned} \frac{\partial^2 \log p(\theta, \phi | \mathbf{x})}{\partial^2 \theta} \Big|_{\theta=\mu_1, \phi=\hat{\phi}} &= -c_1\hat{\phi} < 0 \\ \frac{\partial^2 \log p(\theta, \phi | \mathbf{x})}{\partial^2 \phi} \Big|_{\theta=\mu_1, \phi=\hat{\phi}} &= -\left(\frac{n_1 + 1}{2} - 1 \right) \frac{1}{\hat{\phi}^2} < 0 \\ \frac{\partial^2 \log p(\theta, \phi | \mathbf{x})}{\partial \theta \partial \phi} \Big|_{\theta=\mu_1, \phi=\hat{\phi}} &= 0. \end{aligned}$$

Therefore, $(\mu_1, \hat{\phi})$ is the mode of the joint posterior distribution of (θ, ϕ) .

The above calculations do not guarantee that μ_1 is the maximum of the marginal distribution of θ and $\hat{\phi}$ is the maximum of the marginal distribution of ϕ . In this example, it is easy to see that μ_1 is also the marginal mode since the marginal distribution of θ is a Student-t centred at μ_1 . This automatically implies that μ_1 is the mean and the median of the marginal posterior distribution of θ .

However, the same is not true for $\hat{\phi}$. It was shown in Section 3.4.7 that $\phi | \mathbf{x} \sim G(n_1/2, n_1\sigma_1^2/2)$. This distribution has mode

$$\tilde{\phi} = \frac{n_1 - 2}{n_1\sigma_1^2} \neq \hat{\phi}.$$

Note also that the posterior mean of ϕ is σ_1^{-2} and the median cannot be explicitly evaluated.

Another important consequence from probability theory is that the mode and the mean are not invariant under transformations. Let $\sigma^2 = \phi^{-1}$ and denote the mode of σ^2 by $\tilde{\sigma}^2$. $\tilde{\sigma}^{-1}$ is not the joint nor the marginal mode of σ^2 . To evaluate the mode of σ^2 , the posterior distribution of σ^2 must be obtained. For the case of the marginal distribution,

$$p(\sigma^2 | \mathbf{x}) = p(\phi(\sigma^2) | \mathbf{x}) \left| \frac{d\phi}{d\sigma^2} \right| \quad \text{where} \quad \left| \frac{d\phi}{d\sigma^2} \right| = \left| -\frac{1}{\sigma^4} \right| = \frac{1}{\sigma^4} = (\sigma^{-2})^2$$

$$\propto (\sigma^{-2})^{(n_1/2)-1+2} \exp \left(-\frac{n_1\sigma_1^2}{2\sigma^2} \right)$$

and its logarithm is $\log p(\sigma^2 | \mathbf{x}) = k - \left(\frac{n_1}{2} + 1\right) \log \sigma^2 - \frac{n_1 \sigma_1^2}{2\sigma^2}$. Differentiating with respect to σ^2 :

$$\frac{d \log p(\sigma^2 | \mathbf{x})}{d\sigma^2} \Big|_{\sigma^2=\tilde{\sigma}^2} = -\left(\frac{n_1}{2} + 1\right) \frac{1}{\tilde{\sigma}^2} + \frac{n_1 \sigma_1^2}{2\tilde{\sigma}^4} = 0.$$

The solution of the equation is $\tilde{\sigma}^2 = \frac{n_1 \sigma_1^2}{n_1 + 2} \neq \frac{n_1 \sigma_1^2}{n_1 - 2} = \tilde{\phi}^{-1}$.

The second order condition guarantees the maximum as

$$\frac{d^2 \log p(\tilde{\sigma}^2 | \mathbf{x})}{d(\sigma^2)^2} = \left(\frac{n_1}{2} + 1\right) \frac{1}{\tilde{\sigma}^4} - 2 \frac{n_1 \sigma_1^2}{2\tilde{\sigma}^6} = -\frac{1}{2} \frac{(n_1 + 2)^3}{(n_1 \sigma_1^2)^2} < 0.$$

4.2 Classical point estimation

In the Bayesian methodology, point estimation is always dealt with by minimization of the expected loss. In the classical perspective, many methods have been proposed in an effort to make them adequate to a variety of problems. In this section, the three most important methods will be cited: the method of maximum likelihood, the method of minimum least squares and the method of moments. We will also briefly present non-parametric estimation. Classical point estimation is covered in a clear and concise way in the books by Cox and Hinkley (1974) and Silvey (1970). We recommend both books to the interested reader.

4.2.1 Maximum likelihood

This is currently the most used method of estimation in classical inference. Its use is justified in many instances and it is useful to note that it is entirely based on the likelihood function. Therefore it does not violate the likelihood principle. In addition, there is a good body of theory developed in a variety of situations. Its intuitive appeal can be grasped in the very simple example below.

Example. Consider a situation where all that is known about an unknown quantity of interest θ is that its value is either $1/4$ or $3/4$. Assume now that a $0-1$ random variable X is observed and the success ($X = 1$) probability of X is either $1/6$, when $\theta = 1/4$ or $4/5$, when $\theta = 3/4$. This probabilistic setup is summarized in Table 4.2.

Observe that the sum of each column is 1 but the sum of each line is not. Given the value of X , the likelihood function of θ can be constructed as $l(\theta; x) = p(x | \theta)$.

- If $X = 1$ is observed, $l(1/4; x = 1) = 1/6 < 4/5 = l(3/4; x = 1)$. This means that the model with $\theta = 3/4$ attached a larger probability to the observed event than the model with $\theta = 1/4$ and therefore seems more plausible or likely. If we had to choose an estimate for θ after observing $X = 1$, we would probably opt for the value $3/4$.

Table 4.2 Table of probabilities

x	θ	
	$1/4$	$3/4$
0	$5/6$	$1/5$
1	$1/6$	$4/5$

- If $X = 0$ is observed, the same reasoning would lead to the choice of the value $1/4$ for θ since $l(1/4; x = 0) = 5/6 > 1/5 = l(3/4; x = 0)$.

So, in the above example, we have opted to estimate θ by the value that maximizes the likelihood function for every value of x . This simple and powerful idea is the basis of the method.

Definition. Consider the observation of \mathbf{X} with joint density $p(\mathbf{x}|\theta)$. The maximum likelihood estimator (MLE, in short) of θ is the value of $\theta \in \Theta$ that maximizes $l(\theta; \mathbf{X})$. The usual notation for the MLE of θ is $\hat{\theta}$. Its observed value is called the maximum likelihood estimate.

In most cases, θ varies continuously over an interval or, more generally, over a region of R^p , for some p . In these cases, irrespective of whether \mathbf{X} contains discrete variables or not, the MLE can typically be found by solving the equation $\partial l(\theta; \mathbf{X})/\partial \theta = 0$, or equivalently $\partial \log l(\theta; \mathbf{X})/\partial \theta = 0$, where θ is a vector of 0's.

We are now in position to compare the GMLE, presented in the previous section, with the MLE. The GLME generalizes the MLE just as the posterior density generalizes the likelihood function. Recall that $p(\theta | \mathbf{X}) \propto l(\theta; \mathbf{X}) p(\theta)$. In the special case $p(\theta) \propto k$ it follows that $p(\theta | \mathbf{X}) \propto l(\theta; \mathbf{X})$. Therefore, the value of θ that maximizes the posterior (the GMLE) also maximizes the likelihood. So, the MLE is the GMLE in the case $p(\theta) \propto k$.

Despite their similarity, it is important to stress the distinction between classical and Bayesian estimators. The first ones are statistics and therefore have a sampling distribution based on which their properties will be established. Bayesian estimators are based on the posterior distribution which is always conditional on the value of the observed sample and therefore their properties are based on the posterior distribution, an entirely different object. Nevertheless, they can be seen as functions of the observed sample and in this way compared numerically with classical estimators.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a sample from the $N(\theta, \sigma^2)$ distribution. The likelihood function is

$$l(\theta, \sigma^2; \mathbf{X}) = \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right)^n \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \theta)^2 \right\}$$

with logarithm

$$\log l(\theta, \sigma^2; \mathbf{X}) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \left[\sum_{i=1}^n (X_i - \bar{X})^2 + n(\bar{X} - \theta)^2 \right].$$

Differentiating with respect to θ and σ^2 gives

$$\frac{\partial \log l(\theta, \sigma^2; \mathbf{X})}{\partial \theta} = \frac{1}{2\sigma^2} 2n(\bar{X} - \theta)$$

and

$$\frac{\partial \log l(\theta, \sigma^2; \mathbf{X})}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2(\sigma^2)^2} [(n-1)S^2 + n(\bar{X} - \theta)^2]$$

where

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2.$$

Observe that the denominator used in the expression of S^2 , $(n-1)$, was modified with respect to the value used in previous chapters (n).

Equating derivatives to 0 leads to $\hat{\theta} = \bar{X}$ and $\hat{\sigma}^2 = (n-1)S^2/n$ as critical points.

Differentiating again gives

$$\begin{aligned} \frac{\partial^2 \log l(\theta, \sigma^2; \mathbf{X})}{\partial \theta^2} &= -\frac{n}{\sigma^2} < 0, \quad \forall (\theta, \sigma^2) \in R \times R^+, \\ \frac{\partial^2 \log l(\theta, \sigma^2; \mathbf{X})}{\partial \theta \partial \sigma^2} \Big|_{\theta=\hat{\theta}, \sigma^2=\hat{\sigma}^2} &= 0 \end{aligned}$$

and

$$\begin{aligned} \frac{\partial^2 \log l(\theta, \sigma^2; \mathbf{X})}{\partial (\sigma^2)^2} \Big|_{\theta=\hat{\theta}, \sigma^2=\hat{\sigma}^2} &= \frac{n}{2(\hat{\sigma}^2)^2} - \frac{1}{(\hat{\sigma}^2)^3} [(n-1)S^2 + n(\bar{X} - \hat{\theta})^2] \\ &= -\frac{n}{2} \frac{1}{(\hat{\sigma}^2)^2} < 0. \end{aligned}$$

Therefore, $(\hat{\theta}, \hat{\sigma}^2)$ is the MLE of (θ, σ^2) .

This result can also be established from the similarities between GMLE and MLE and calculations in the previous section. Note that by taking $c_0 \rightarrow 0$, $\sigma_0^2 \rightarrow 0$ and $n_0 = 2$, the prior becomes constant and the above result becomes a special case of the derivations of the previous section.

There are many interesting properties of the MLE's that can be easily shown:

1. The MLE is invariant to 1-to-1 transformations (unlike the GMLE). If $\hat{\theta}$ is the MLE of θ and $\phi = \phi(\theta)$ is a 1-to-1 function of θ then the MLE of ϕ is

given by $\hat{\phi} = \phi(\hat{\theta})$. To see that, remember that $\hat{\theta}$ maximizes $l(\theta)$ and that the likelihood of ϕ is $l^*(\phi) = l^*(\phi(\theta)) = l(\theta)$. Therefore, if $\hat{\theta}$ maximizes l , it will also maximize $l^*(\phi(\theta))$ by uniqueness of the transformation and $\phi(\hat{\theta})$ will maximize l^* . Consequently, $\hat{\phi}$ maximizes l^* and therefore $\hat{\phi}$ is the MLE of ϕ .

Example. In the case of a $N(\theta, \sigma^2)$ distribution, the MLE of σ^2 is $\hat{\sigma}^2$. Therefore the MLE of the standard deviation σ is $\hat{\sigma}$ and the MLE of precision $\phi = \sigma^{-2}$ is $\hat{\phi} = \hat{\sigma}^{-2}$.

2. As briefly mentioned before, the MLE does not depend on the sampling plan. If different experiments \mathcal{E}_1 and \mathcal{E}_2 lead to respective likelihood functions $l_1(\theta)$ and $l_2(\theta)$ and $l_1 = k l_2$ for some $k > 0$ that does not depend on θ , their MLE will be the same. Therefore, the MLE does not violate the likelihood principle.
3. The MLE may not exist. Assume that X_1, \dots, X_n is a sample from the $U(0, \theta)$ distribution, $\theta > 0$. The likelihood function is

$$l(\theta; \mathbf{X}) = p(\mathbf{X} | \theta) = \frac{1}{\theta^n} \prod_{i=1}^n I_\theta(X_i, \infty) = \frac{1}{\theta^n} I_\theta(T, \infty)$$

where $T = \max_i X_i$. As the likelihood function is a strictly decreasing function of θ , its maximum is attained at the lower value of its domain, the interval (T, ∞) . As the interval is open, the function does not have a maximum and θ does not have an MLE. This technical difficulty is easily remedied by considering without loss of generality closed intervals. In this case, the MLE of θ is $T = \max_i X_i$. Note, however, that $Pr(T < \theta) = 1$ and therefore the MLE will underestimate θ with certainty.

4. The MLE may not be unique. Assume that X_1, \dots, X_n is a sample from the $U(\theta, \theta+1)$ distribution, $\theta \in R$. The likelihood function is

$$l(\theta; \mathbf{X}) = \prod_{i=1}^n I_\theta(X_i - 1, X_i) = I_\theta(T_2 - 1, T_1)$$

where $T_1 = \min_i X_i$ and $T_2 = \max_i X_i$. Therefore, the MLE of θ will be any value in the interval $(T_2 - 1, T_1)$, if it exists, because the likelihood function is constant over that region.

5. MLE and Bayes estimators depend on the sample only through minimal sufficient statistics. By the factorization criterion, $l(\theta; \mathbf{X}) = g(\mathbf{X})f(\mathbf{T}, \theta)$, where \mathbf{T} is a sufficient statistic. Maximization of $l(\theta; \mathbf{X})$ is therefore equivalent to maximization of $f(\mathbf{T}, \theta)$. Since f only depends on the sample through \mathbf{T} , the MLE will have to be a function of \mathbf{T} . The same reason applies for the Bayes estimators. As this is valid for every sufficient statistic, it must be valid for the minimal one.

To obtain the (G)MLE, one must obtain the maximum of the likelihood function (posterior density, respectively). This task can be seen as an optimization problem. Assuming further that θ varies continuously over a region, simplifies the task considerably. The problem can typically be solved by solving the equation $\partial l(\theta; \mathbf{X})/\partial\theta = \mathbf{0}$ (or $\partial p(\theta|\mathbf{x})/\partial\theta = \mathbf{0}$). In many cases, it is easier to work with the logarithm. Concentrating on the likelihood from now on, let $L(\theta; \mathbf{X}) = \log l(\theta; \mathbf{X})$. The problem can then be rephrased in terms of finding the solution of

$$U(\mathbf{X}; \theta) = \frac{\partial L(\theta; \mathbf{X})}{\partial\theta} = \mathbf{0},$$

where U is the score function introduced in Chapter 2. Note that

$$\frac{\partial \log p(\theta|\mathbf{X})}{\partial\theta} = \frac{\partial L(\theta; \mathbf{X})}{\partial\theta} + \frac{\partial \log p(\theta)}{\partial\theta} = U(\mathbf{X}; \theta) + \frac{\partial \log p(\theta)}{\partial\theta}$$

and can thus be referred to as the generalized score function. In some cases, the above equation can be analytically solved and the roots of the (generalized) score function found explicitly. In other applications, typically involving a highly dimensional θ , no analytical solution can be found. Algorithms for solving this problem will be presented in the next chapter.

4.2.2 Method of least squares

Assume now that $\mathbf{Y} = (Y_1, \dots, Y_n)$ for a random sample such that $E(Y_i | \theta) = f_i(\theta)$ and $V(Y_i | \theta) = \sigma^2$. One can rewrite each Y_i as

$$Y_i = f_i(\theta) + e_i \quad \text{where } E(e_i) = 0 \text{ and } V(e_i) = \sigma^2, i = 1, \dots, n.$$

One possible criterion for the estimation of θ is to minimize the observation errors e_i 's incurred. There are many ways to account globally for the errors. Given the assumption of homoscedasticity (equal error variances), it seems fair to account for all the errors in the same way and with the same weight. Also, it seems reasonable to account for the errors symmetrically to avoid penalizing more positive or negative errors. One possibility is to attempt minimization of the sum of the absolute errors. In fact, this choice is as plausible as the choice of the absolute loss function in the context of Bayesian estimation. However, for historical and mathematical reasons, the criterion preferred in many cases is to account for the squared errors.

Therefore, the estimation criterion can be stated as the minimization of

$$S(\theta) = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (Y_i - f_i(\theta))^2.$$

Note that by forming the vector $\mathbf{f}(\theta) = (f_1(\theta), \dots, f_n(\theta))'$, the quadratic form can be rewritten as $S(\theta) = (\mathbf{Y} - \mathbf{f}(\theta))'(\mathbf{Y} - \mathbf{f}(\theta))$. The value of θ that minimizes $S(\theta)$ is called the least squares estimator (LSE, in short) of θ . Once again, minimization is achieved by solving the equation $\partial S(\theta)/\partial\theta = \mathbf{0}$.

Example. Simple linear regression. Assume one knows that his/her variable of interest Y is affected by the values of another known quantity X and that this dependence is linear on the mean. One model for this setup is to take $E(Y_i | \theta) = \theta_0 + \theta_1 X_i$ where $\theta = (\theta_0, \theta_1)$. The quadratic form is given by

$$S(\theta_0, \theta_1) = \sum_{i=1}^n (Y_i - \theta_0 - \theta_1 X_i)^2.$$

Differentiation gives

$$\begin{aligned}\frac{\partial S(\theta)}{\partial\theta_0} &= -2 \sum_{i=1}^n (Y_i - \theta_0 - \theta_1 X_i) = -n(\bar{Y} - \theta_0 - \theta_1 \bar{X}) \\ \frac{\partial S(\theta)}{\partial\theta_1} &= -2 \sum_{i=1}^n X_i(Y_i - \theta_0 - \theta_1 X_i) = -n(\bar{XY} - \theta_0 \bar{X} - \theta_1 \bar{X^2})\end{aligned}$$

where $\bar{g}(X, Y)$ generically denotes $(1/n) \sum_{i=1}^n g(X_i, Y_i)$. It is not difficult to show that the least squares estimator of (θ_0, θ_1) is

$$(\hat{\theta}_0, \hat{\theta}_1) = \left(\bar{Y} - \hat{\theta}_1 \bar{X}, \frac{\bar{XY} - \bar{X}\bar{Y}}{\bar{X^2} - \bar{X}^2} \right).$$

One important support for the method is the fact that it coincides with the MLE if the error distribution is normal. This method can be extended to the case of error variances that are unequal due to constants, namely $V(e_i) = w_i^{-1} \sigma^2$. It seems reasonable in this case to take into account the different variabilities and weigh more heavily in the sum the more precise observations, those with larger values of w_i . This modification of the criterion leads to the weighted LSE obtained by minimization of

$$S(\theta) = \sum_{i=1}^n w_i (Y_i - f_i(\theta))^2.$$

The sum can again be written in matrix notation as $(\mathbf{Y} - \mathbf{f}(\theta))' \mathbf{W} (\mathbf{Y} - \mathbf{f}(\theta))$ where \mathbf{W} is the $n \times n$ diagonal matrix with elements w_1, \dots, w_n . Note that in its full generality, the matrix of weights \mathbf{W} does not even need to be diagonal by allowing correlated observation errors.

The previous method (without weights) is also called ordinary least squares. It can be obtained as a special case where $\mathbf{W} = \sigma^2 \mathbf{I}_n$, the $n \times n$ identity matrix.

One useful application of this method is in the case of spurious observations. We would not want our analysis to be influenced by observations that are known to be discordant from the rest of the observations. Reduction of the effect of these variables is achieved by setting smaller values of w_i for them. In the limit, $w_i \rightarrow 0$ and the observation has no influence in the estimation. This line of reasoning forms the basis of robustness studies.

In all the cases above, σ^2 is estimated by σ_{LS}^2 given by

$$(\mathbf{Y} - \mathbf{f}(\theta_{LS}))' \mathbf{W} (\mathbf{Y} - \mathbf{f}(\theta_{LS}))$$

where θ_{LS} is the weighted LSE of θ . Sometimes, n is replaced by $n - p$, where p is the dimension of θ .

4.2.3 Method of moments

Assume again a random sample X_1, \dots, X_n from a distribution $p(x | \theta)$ with moments of order k given by $\mu_k = E(X^k | \theta)$, $k = 1, 2, \dots$. The method of moments recommends estimation of μ_k by

$$\hat{\mu}_k = \frac{1}{n} \sum_{i=1}^n X_i^k.$$

In words, the populational moments are estimated by the sample moments. Any other function of θ is estimated in the same way by taking into account its relation with the population moments. For example, $\mu_1 = E(X | \theta)$ is estimated by $\hat{\mu}_1 = \bar{X}$, $\mu_2 = E(X^2 | \theta)$ is estimated by \bar{X}^2 and the populational variance $V(X | \theta)$, once it is written as $\mu_2 - \mu_1^2$, is estimated by $\hat{\mu}_2 - \hat{\mu}_1^2 = \hat{\sigma}^2$. Mean and variance estimators coincide with MLE in the normal case.

For any distribution with finite μ_k , the laws of large numbers ensure that $\hat{\mu}_k \rightarrow \mu_k$, with probability 1 as $n \rightarrow \infty$. So, estimators obtained by the method of moments enjoy good asymptotic properties.

4.2.4 Empirical distribution function

This is a non-parametric method which involves no knowledge of the distribution function to be estimated. This estimator is useful at least as an initial estimator, thus providing some insight into the form of the distribution function.

Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from an unknown distribution function F . Recall from the definition that $F(x) = P(X \leq x)$. The empirical distribution function is denoted by \hat{F} and is given by

$$\hat{F}(x) = \frac{\#\{X_i \leq x\}}{n}.$$

Observe that just like F , \hat{F} is non-decreasing and contained in the interval $[0, 1]$. It is interesting to note that $\hat{F}(x)$ can be written as \bar{Z} where $Z_i = I_{X_i}(-\infty, x]$, $i = 1, \dots, n$. The populational quantity equivalent to \bar{Z} is the proportion of population elements that are $\leq x$. This is given by $P(X \leq x) = F(x)$. So, the empirical distribution function is a form of method of moments estimator of the distribution function. Other properties of \hat{F} will be described in the sequel.

4.3 Comparison of estimators

Given that there is no unifying criterion for the choice of frequentist estimators in any given problem, it is important that a set of criteria is established to compare

them. The main criteria for comparison are: bias, (frequentist) risk or mean squared error and consistency. Much effort was concentrated on this area during the 1950s and 1960s. These studies lead to the characterization of uniformly minimum variance unbiased (UMVU, for short) estimators.

4.3.1 Bias

Definition. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from $p(x | \theta)$ and $\delta = \delta(\mathbf{X})$ an estimator of $\mathbf{h}(\theta)$, for any given function \mathbf{h} . δ is an unbiased estimator of $\mathbf{h}(\theta)$ if $E[\delta | \theta] = \mathbf{h}(\theta)$, $\forall \theta$. The estimator δ is said to be biased otherwise. In this case, the bias is denoted by $b(\theta)$ and defined as $b(\theta) = E[\delta | \theta] - \mathbf{h}(\theta)$.

The frequentist interpretation of the definition is that after repeating sampling of \mathbf{X} from $p(x | \theta)$ many times, averaging the corresponding values of δ will produce $\mathbf{h}(\theta)$ as a result. This is a desirable property because one formulates an estimator δ in an effort to obtain the value of $\mathbf{h}(\theta)$. The difficulty is that in most cases only a single sample \mathbf{X} is observed for time and/or financial restrictions. Note also that unbiased estimation is always related to a given parametric function; an estimator can be biased with respect to a given function but unbiased with respect to another one.

Example. None of the three parametric methods of estimation proposed in the previous section can guarantee unbiased estimators. The empirical distribution function however is an unbiased estimator of the distribution function.

4.3.2 Risk

Definition. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from $p(x | \theta)$ and denote now by $\delta = \delta(\mathbf{X})$ an estimator of $\mathbf{h}(\theta)$. The frequentist risk of the estimator δ is defined as $R_\delta(\theta) = E_{\mathbf{X}|\theta}[L(\delta(\mathbf{X}), \theta)]$. In the case of a quadratic loss function L , the risk is given by $R_\delta(\theta) = E_{\mathbf{X}|\theta}[(\delta - \mathbf{h}(\theta))'(\delta - \mathbf{h}(\theta))]$ and is also called the mean squared error (MSE, in short). In the scalar case, the MSE reduces to $E_{\mathbf{X}|\theta}[\delta - h(\theta)]^2$.

Comparing with the Bayes risk, once again we see the presence of an expected loss. The change with respect to the approach in the evaluation of the expectation must be stressed. The Bayesian risk considers expectation with respect to the posterior distribution of $\theta | \mathbf{x}$ whereas here expectations are taken with respect to the sampling distribution of $\mathbf{X} | \theta$. Although the dependence on \mathbf{x} causes no harm to the Bayesian estimators, the dependence on θ will cause additional problems, to be described below.

In terms of risk, the estimation task resumes in finding the estimator of smallest risk. Let $\delta_1 = \delta_1(\mathbf{X})$ and $\delta_2 = \delta_2(\mathbf{X})$ be estimators of $\mathbf{h}(\theta)$. Their respective risks are $R_{\delta_1}(\theta)$ and $R_{\delta_2}(\theta)$ and δ_1 is better than δ_2 if $R_{\delta_1}(\theta) \leq R_{\delta_2}(\theta)$, for all θ with strict inequality for at least one value of θ . An estimator is admissible if there is no better estimator than it. When an estimator is unbiased and its variance

is uniformly smaller for all possible values of θ over all possible estimators, it is referred to as a uniformly minimal unbiased estimator (UMVU, in short).

If the estimator is unbiased, its quadratic risk is given by $\text{tr}[\mathbf{V}(\delta|\theta)]$, the trace of its sampling covariance matrix. If it is biased, the quadratic risk is given by $\text{tr}[\mathbf{V}(\delta)] + [\mathbf{b}(\theta)']\mathbf{b}(\theta)$. In the case of a scalar θ , the quadratic risk of an unbiased estimator is given by its sampling variance $V(\delta|\theta)$ and if δ is biased, its quadratic risk is given by $V(\delta|\theta) + b^2(\theta)$.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution with σ^2 known and $h(\theta) = \theta$. Taking $\delta_1(\mathbf{X}) = \bar{X}$ and $\delta_2(\mathbf{X}) = X_1$ gives

$$E[\delta_1(\mathbf{X}) | \theta] = E(\bar{X} | \theta) = \frac{1}{n} \sum_{i=1}^n E(X_i | \theta) = \frac{n\theta}{n} = \theta$$

$$E[\delta_2(\mathbf{X}) | \theta] = E(X_1 | \theta) = \theta$$

and therefore δ_1 and δ_2 are unbiased estimators of θ . Therefore, their quadratic risks will coincide with their sampling variances and will be respectively given by

$$\begin{aligned} R_{\delta_1}(\theta) &= V(\bar{X} | \theta) = \frac{\sigma^2}{n} \\ R_{\delta_2}(\theta) &= V(X_1) = \sigma^2. \end{aligned}$$

Of course, $R(\delta_1) < R(\delta_2)$, if $n > 1$, for all values of θ and therefore δ_1 is better than δ_2 .

It is not always possible to find an estimator that completely dominates the other ones in terms of risk. In the ideal situation one would eliminate θ by suitably weighting the risks over their different values. This is performed naturally in the Bayesian context. Here however θ is fixed and no such natural weighting scheme exists. An alternative is to consider the worst possible risk for each estimator and choose the estimator with smallest worst possible risk. This is the definition of the minimax estimator.

Returning to the example, it is not entirely surprising that δ_1 is better than δ_2 . After all, δ_1 seems to be using the sample information better than δ_2 . Once again, the key concept here is sufficiency and the result is formalized in the Rao–Blackwell theorem below.

Theorem 4.1. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from $p(x|\theta)$, $\delta = \delta(\mathbf{X})$ an unbiased estimator of $h(\theta)$, for some function h and $\mathbf{T} = \mathbf{T}(\mathbf{X})$ a sufficient statistic for θ . Then, $\delta^* = \delta^*(\mathbf{X}) = E(\delta | \mathbf{T})$ is an unbiased estimator of $h(\theta)$ with $\mathbf{V}(\delta^* | \theta) \leq \mathbf{V}(\delta | \theta)$,¹ $\forall \theta$. In the case of a scalar $h(\theta)$, the result states that $V(\delta^* | \theta) \leq V(\delta | \theta)$.

¹Recall from Chapter 1 that if \mathbf{A} and \mathbf{B} are squared matrices of the same dimension, $\mathbf{A} \leq \mathbf{B}$ means that the matrix $\mathbf{A} - \mathbf{B}$ is a non-positive definite matrix.

Before proving the result there are a few important comments to make. The first one is that the theorem states that whenever one finds an unbiased estimator, it can always be improved in terms of risk by conditioning on a sufficient statistic. Also, the conditional expectation used in the definition of δ^* does not introduce θ because of the definition of a sufficient statistic.

Proof. Initially note that δ^* is unbiased because

$$E[\delta^*(\mathbf{X}) | \theta] = E\{E[\delta(\mathbf{X}) | \mathbf{T}(\mathbf{X})] | \theta\} = E[\delta(\mathbf{X}) | \theta] = h(\theta).$$

Finally note that

$$\mathbf{V}(\delta^* | \theta) = \mathbf{V}(\delta | \theta) - E[\mathbf{V}(\delta | \mathbf{T}) | \theta].$$

As $\mathbf{V}(\delta | \mathbf{T}) \geq 0$, its expectation is also non-negative positive and therefore $\mathbf{V}(\delta | \theta) \geq \mathbf{V}(\delta^* | \theta)$. \square

The important message of the theorem is that estimators have their risks reduced if they are functions of sufficient statistics. Risks are indeed reduced by properties of non-negative definite matrices (see Chapter 1). Yet again, maximal improvement in terms of risk is achieved if minimal sufficient statistics are used. A related interesting question is to know if the reduction in risk is the smallest possible. The search for maximal reduction is helped in a sense by the concept of complete families of distributions.

Definition. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from $p(x | \theta)$, $\mathbf{T} = \mathbf{T}(\mathbf{X})$ any statistic and g any function of \mathbf{T} . The family of distributions of \mathbf{T} is complete if for all θ ,

$$E(g(\mathbf{T}) | \theta) = \mathbf{0} \Rightarrow g(\mathbf{T}) = \mathbf{0}, \quad \text{with probability 1.}$$

Verification of completeness of families directly from the definition is cumbersome. Fortunately, for exponential families with k parameters, it can be shown that the family of distributions of the k -dimensional statistic $(U_1(\mathbf{X}), \dots, U_k(\mathbf{X}))$ is complete if the variation space of $(\phi_1(\theta), \dots, \phi_k(\theta))$ is k -dimensional. The definitions of the U_i 's and ϕ_i 's were given in Section 2.3. The proof of this result requires elements that are beyond the scope of the study of this book and will therefore be omitted. The interested reader is referred to Lehmann (1986).

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution. Then, $U_1 = \sum X_i$, $U_2 = \sum X_i^2$, $\phi_1 = \theta/\sigma^2$ and $\phi_2 = -1/2\sigma^2$ and (ϕ_1, ϕ_2) vary over a bidimensional space. Therefore, the family of distributions of (U_1, U_2) is complete. If one assumes that $\theta = \sigma^2$, the space of variation of (ϕ_1, ϕ_2) is reduced to a single dimension and the family of distributions of (U_1, U_2) is no longer complete.

The concept of completeness is useful to ensure uniqueness of the UMVU estimator. It can be shown that the UMVU estimator is unique in the presence of complete families because if δ_1 and δ_2 are unbiased estimators and functions of the minimal sufficient statistic T then $E[\delta_1^* - \delta_2^* | \theta] = \mathbf{0}$. From completeness, this means that $\delta_1^* - \delta_2^* = \mathbf{0}$, with probability 1. Therefore, δ_1^* and δ_2^* must be equal.

Another interesting aspect of risk calculation is the existence of a lower bound for the quadratic risk of unbiased estimators. This result is known as the Cramer–Rao inequality. This inequality, central in the theory of unbaised estimation, is due to Fisher although it was independently stated in its present form by Cramer and Rao in the 1940s, as cited in Cox and Hinkley (1974).

Theorem 4.2. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from $p(x | \theta)$ and δ an unbiased estimator of $\mathbf{h}(\theta)$, for some function \mathbf{h} . Assume further that $\{x : p(x | \theta) > 0\}$ does not depend on θ , the differentials $\partial p(x | \theta)/\partial\theta$ and $\partial\mathbf{h}(\theta)/\partial\theta$ exist, $E(\delta | \theta)$ is differentiable under the integral sign and that the Fisher information $I(\theta)$ is finite. Then

$$V(\delta | \theta) \geq \frac{\partial\mathbf{h}(\theta)}{\partial\theta} [I(\theta)]^{-1} \left(\frac{\partial\mathbf{h}(\theta)}{\partial\theta} \right)'.$$

In the case of a scalar θ , the inequality reduces to

$$V[\delta | \theta] \geq \frac{[dh(\theta)/d\theta]^2}{I(\theta)}.$$

Proof (of the scalar case). $E[\delta | \theta] = \int \delta(x)p(x | \theta)dx = h(\theta)$ because δ is unbiased. Differentiating both sides with respect to θ gives

$$\begin{aligned} \frac{dh(\theta)}{d\theta} &= \frac{\partial}{\partial\theta} \int \delta(x)p(x | \theta)dx \\ &= \int \delta(x) \frac{\partial p(x | \theta)}{\partial\theta} dx, \text{ where interchange of signs is valid by hypothesis} \\ &= \int \delta(x) \frac{1}{p(x | \theta)} \frac{\partial p(x | \theta)}{\partial\theta} p(x | \theta) dx \\ &= E \left[\left(\delta(\mathbf{X}) \frac{\partial \log p(\mathbf{X} | \theta)}{\partial\theta} \right) | \theta \right] \\ &= E[\delta(\mathbf{X}) U(\mathbf{X}; \theta) | \theta]. \end{aligned}$$

As previously seen in Section 2.4, $E[U(\mathbf{X}; \theta)] = 0$ and

$$\begin{aligned} \frac{dh(\theta)}{d\theta} &= E[(\delta(\mathbf{X}) - h(\theta))U(\mathbf{X}; \theta) | \theta] \\ &= Cov[(\delta(\mathbf{X}), U(\mathbf{X}; \theta)) | \theta]. \end{aligned}$$

Since the absolute value of the correlation between two random variables is never larger than 1, the squared covariance will never be larger than the product of the

two variances. Hence,

$$[dh(\theta)/d\theta]^2 \leq V[\delta(\mathbf{X}) | \theta]V[U(\mathbf{X}; \theta) | \theta].$$

But

$$V[U(\mathbf{X}; \theta) | \theta] = E[U^2(\mathbf{X}; \theta) | \theta] = I(\theta),$$

completing the proof. \square

The proof of the theorem in the multiparameter case is left as an exercise.

Observe that the unbiased estimator attains the lower bound when it has maximal correlation with the score function. In other words, when there are functions c and d of θ such that

$$\delta(\mathbf{X}) = c(\theta)U(\mathbf{X}; \theta) + d(\theta),$$

with probability 1. Taking expectation of both sides with respect to $\mathbf{X} | \theta$ gives that $\delta(\mathbf{X})$ is an unbiased estimator of $d(\theta)$ and therefore $d = \mathbf{h}$. In this case, $c(\theta) = I^{-1}(\theta)$.

Also, when the MLE is unbiased, it attains the Cramer–Rao lower bound. This can be seen by solving the above equation for θ . This leads to

$$U(\mathbf{X}; \theta) = \frac{\partial \log p(\mathbf{X} | \theta)}{\partial\theta} = \frac{\delta(\mathbf{X}) - d(\theta)}{c(\theta)}.$$

Equating to $\mathbf{0}$ implies that $\delta(\mathbf{X})$ is the MLE of $d(\theta)$. But we have already seen that $d = \mathbf{h}$ and, by hypothesis, δ is unbiased for $\mathbf{h}(\theta)$. Hence, it attains the Cramer–Rao lower bound.

Definition. The estimator δ of $\mathbf{h}(\theta)$ is said to be efficient if it is unbiased and its variance attains the Cramer–Rao lower bound, for all θ . The efficiency of an unbiased scalar estimator is given by the ratio between the Cramer–Rao bound and its variance.

Note that there is no guarantee that UMVU estimators will attain the Cramer–Rao bound and they may have their efficiency smaller than 1. However, the converse is true with efficient estimators being necessarily UMVU.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $\text{Pois}(\theta)$ distribution. Then

$$\log p(\mathbf{X} | \theta) = -n\theta + \sum_{i=1}^n X_i \log \theta - \sum_{i=1}^n \log X_i!$$

and therefore $U(\mathbf{X}; \theta) = -n + \sum X_i/\theta$. As estimators cannot possibly depend on the parameter they are supposed to estimate, define $c(\theta) = \theta/n$ and $d(\theta) = \theta$. This way, it is immediate that \bar{X} is an efficient estimator of its mean θ . In fact, any linear function of \bar{X} is an efficient estimator of the respective linear function of θ . More than that, these are the unique efficient estimators that can be found in the presence of a random sample from the $\text{Pois}(\theta)$ distribution.

4.3.3 Consistency

It is to be expected that the information contained in the sample increases with an increase in the sample size. This is certainly true at least for the Fisher measures of information. One would then expect that reasonable estimators will tend to get closer and closer to their estimands. This subsection discusses theoretical properties of the estimators as the sample size gets larger and larger. The relevant question is how close are the estimator and its estimand. Related questions of interest are: Is the bias getting smaller when sample size increases? Is the variance getting smaller as well? These questions will be deferred to the next chapter.

Definition. Let $\mathbf{X}_n = (X_1, \dots, X_n)$ be a random sample of size n from $p(x|\theta)$ and $\delta_n(\mathbf{X})$ an estimator of $\mathbf{h}(\theta)$ based on a sample of size n . As the sample size n varies, a sequence of estimators for $\mathbf{h}(\theta)$ is obtained. This sequence is said to be (weakly) consistent for $\mathbf{h}(\theta)$ if $\delta_n(\mathbf{X}) \rightarrow \mathbf{h}(\theta)$, in probability, when $n \rightarrow \infty$.

In practice, the definition is shortened by saying that the estimator is or is not consistent instead of a sequence. The definition means that $\forall \epsilon > 0$, $P(|\delta_n(\mathbf{X}) - \mathbf{h}(\theta)| > \epsilon) \rightarrow 0$, when $n \rightarrow \infty$. This result is usually denoted by $\text{plim } \delta_n(\mathbf{X}) = \mathbf{h}(\theta)$. As an example, \hat{F}_n (the empirical distribution function) is consistent for F . The three important questions to ask about an estimator are: is it unbiased, how large is its risk and is it consistent?

When Bayes estimators are considered as functions of the sample \mathbf{X}_n instead of its observed value \mathbf{x}_n , they can be studied for their sampling properties just like any other estimator. In particular, it can be shown that Bayes estimators are invariably biased. Also, it may be reasoned that as the sample size increases, the influence of any non-degenerate prior becomes smaller and Bayes estimators will become closer to the MLE. So, intuitively, one can expect them to inherit all the properties of the MLE, irrespective of the loss function used.

Example. Let $\mathbf{X}_n = (X_1, \dots, X_n)$ be a random sample from the $\text{Ber}(\theta)$ distribution, with $\theta > 0$. We know that the MLE of θ is $\hat{\theta}_n = \bar{X}_n$ and that, by the laws of large numbers, $\bar{X}_n \rightarrow \theta$, in probability and almost surely. Therefore, \bar{X}_n is a consistent estimator of θ .

In the case of a conjugate prior beta(α, β) and quadratic loss function, the Bayes estimator is $\delta_n^*(\mathbf{x}_n) = (\alpha + n\bar{x}_n)/(\alpha + \beta + n)$ which converges to \bar{x}_n when $n \rightarrow \infty$. Therefore, $|\delta_n^*(\mathbf{X}_n) - \bar{X}_n| \rightarrow 0$ in probability and as $\bar{X}_n \rightarrow \theta$, almost surely, δ_n^* is also consistent.

It follows readily from Tchebychev's inequality that

$$\Pr(|\delta_n(\mathbf{X}) - \mathbf{h}(\theta)| > \epsilon) < \frac{E\{[\delta_n(\mathbf{X}) - \mathbf{h}(\theta)]'[\delta_n(\mathbf{X}) - \mathbf{h}(\theta)] \mid \theta\}}{\epsilon^2}, \quad \forall \epsilon > 0$$

but $E\{[\delta_n(\mathbf{X}) - \mathbf{h}(\theta)]'[\delta_n(\mathbf{X}) - \mathbf{h}(\theta)] \mid \theta\} = R_{T_n}(\theta)$. So, if a sequence of estimators has quadratic risk tending to 0, the estimator is consistent.

One can also define strong consistency of a sequence of estimators if the convergence for the parameter is almost sure instead of in probability. The theory of probability assures us that almost sure convergence implies convergence in probability and therefore strong consistency implies weak consistency. It can be shown that the MLE is strongly consistent under the same regularity conditions of the Cramer–Rao inequality. Nevertheless, the concept of weak consistency retains the essence of what is needed and will be maintained in the sequel.

4.4 Interval estimation

4.4.1 Bayesian approach

Returning to the Bayesian point of view, the most adequate form to express available information about unknown parameters is through the posterior distribution. Despite its coherent specification through expected loss functions, point estimation presents some inconvenient features. The main restriction is that it simplifies the multitude of information from a distribution into a single figure. It is important at least to have some information about how precise the specification of this figure is. One possibility is to associate point estimators with a measure of the uncertainty about them. So, for the mean, one can use the variance or the coefficient of variation. For the mode, the observed information given by the curvature at the mode is usually adequate. Finally, for the median, the interquartile distance could be used.

In this section, another line of work is sought. The aim is to provide a compromise between the complete posterior distribution and a single figure extracted from it. This compromise is reached by providing a range of values extracted from the posterior distribution. Typically, one attaches a probability to this region and when the probability is large one gets a good idea of the probable or likely values of the unknown of interest. Ideally, one would like to report a region of values of θ that is as small as possible but that contains as much probability as possible. The size of the interval informs us about the dispersion of the values of θ .

Definition. Let θ be an unknown quantity defined in Θ . A region $C \subset \Theta$ is a $100(1 - \alpha)\%$ credibility or Bayesian confidence region for θ if $\Pr(\theta \in C | \mathbf{x}) \geq 1 - \alpha$. In this case, $1 - \alpha$ is called the credibility or confidence level. In the scalar case, the region C is usually given by an interval, $[c_1, c_2]$ say, hence the name.

It should be clear from the above definition that the intervals are defined by simple probability evaluation over the posterior distribution of θ . Many Bayesian authors reject the use of the word confidence for Bayesian intervals. As will be seen shortly, confidence has a very precise meaning in the definition of frequentist intervals that differs substantially from the meaning given here. These authors consider it important to dissociate the concepts. In the sequel, we will refer to confidence intervals whenever we refer to the method in general or in the classical context and will use the word credibility in reference to Bayesian intervals.

Note that C is never an interval in the multidimensional case. Even in the uniparameter case, there is nothing in the definition enforcing the region C to be an interval. Therefore, there is a slightly misleading use of the word interval.

The above probability is evaluated over the updated distribution of θ which will be taken from now on as the posterior. In general, one would want both α and C to be as small as possible. This in turn implies that the posterior distribution is as concentrated as possible. The requirement of a larger posterior probability than the confidence level is essentially technical. It is mainly due to the use in discrete distributions where it is not always possible to find a region that exactly satisfies the probability required by a given level. In many cases, the inequality can be taken as an equality thus implying that the region C will be as small as possible.

Note also that credibility intervals are invariant under 1-to-1 transformations of the parameter. So, if C is a $100(1 - \alpha)\%$ credibility interval for θ and $\phi = \phi(\theta)$ is a 1-to-1 transformation of θ then $\phi(C)$, the image of C under ϕ , is a $100(1 - \alpha)\%$ credibility interval for ϕ . This useful property is also shared by frequentist confidence intervals.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution with σ^2 known. The non-informative prior for θ is $p(\theta) \propto k$ and the likelihood is

$$l(\theta; \mathbf{x}) \propto \exp\left\{-\frac{n}{2\sigma^2}(\theta - \bar{x})^2\right\}$$

providing the posterior $p(\theta | \mathbf{x}) \propto l(\theta; \mathbf{x})p(\theta) \propto l(\theta; \mathbf{x})$ and therefore $\theta | \mathbf{x} \sim N(\bar{x}, \sigma^2/n)$ or equivalently $\sqrt{n}(\theta - \bar{x})/\sigma | \mathbf{x} \sim N(0, 1)$. From there, many $100(1 - \alpha)\%$ confidence intervals may be constructed for θ with the use of the standard normal distribution function Φ . Defining $\Phi(x) = P(X \leq x)$ if $X \sim N(0, 1)$ then z_c is such that $\Phi(z_c) = 1 - c$, $0 < c < 1$ and intervals can be constructed from:

1. $1 - \alpha = Pr(\sqrt{n}(\theta - \bar{x})/\sigma \leq z_\alpha | \mathbf{x})$ which implies that $\theta \leq z_\alpha\sigma/\sqrt{n} + \bar{x}$ with posterior probability $1 - \alpha$. Hence, the interval $C_1 = (-\infty, \bar{x} + z_\alpha\sigma/\sqrt{n}]$ is a $100(1 - \alpha)\%$ Bayesian confidence interval for θ . The length of C_1 however is infinity which is not very useful for our summarization purposes. As previously mentioned, one would like to have C as small as possible. The problem with this interval is that it includes (infinitely) many values that have very negligible probability around them.
2. Let z_β and z_γ be numbers such that

$$1 - \alpha = P(-z_\beta \leq \sqrt{n}(\theta - \bar{x})/\sigma \leq z_\gamma | \mathbf{x}).$$

Using the symmetry of the normal distribution,

$$\Phi(-z_\beta) = P(X \leq -z_\beta) = P(X \geq z_\beta) = 1 - P(X < z_\beta) = \beta$$

and the probability of the above interval is given by $\Phi(z_\gamma) - \Phi(-z_\beta) =$

$1 - (\gamma + \beta)$ and therefore $\gamma + \beta = \alpha$. With this assumption,

$$\begin{aligned} 1 - \alpha &= Pr\left(z_\beta \leq \sqrt{n}\frac{(\theta - \bar{x})}{\sigma} \leq z_\gamma | \mathbf{x}\right) \\ &= Pr\left(-\frac{\sigma}{\sqrt{n}}z_\beta + \bar{x} \leq \theta \leq z_\gamma \frac{\sigma}{\sqrt{n}} + \bar{x} | \mathbf{x}\right). \end{aligned}$$

The interval $C = [c_1, c_2]$ where

$$c_1 = \bar{x} + \frac{\sigma}{\sqrt{n}}z_\beta \quad \text{and} \quad c_2 = \bar{x} + \frac{\sigma}{\sqrt{n}}z_\gamma$$

is a $100(1 - \alpha)\%$ Bayesian confidence interval for θ . Note that it has length $(z_\gamma + z_\beta)\sigma/\sqrt{n}$. There still remains the point about minimization of the length of the interval subject to $\gamma + \beta = \alpha$.

Note that if $\phi = \phi(\theta)$ is a monotonically increasing transformation of θ , then $[\phi(c_1), \phi(c_2)]$ is also a $100(1 - \alpha)\%$ Bayesian confidence interval for ϕ . If $\phi = \phi(\theta)$ is a monotonically decreasing transformation of θ , then $[\phi(c_2), \phi(c_1)]$ is a $100(1 - \alpha)\%$ Bayesian confidence interval for ϕ .

Consider without loss of generality that $z_\gamma \leq z_{\alpha/2} \leq z_\beta$ and define $a = z_{\alpha/2} - z_\gamma \geq 0$, $b = z_\beta - z_{\alpha/2} \geq 0$ and A and B as the areas between $z_{\alpha/2}$ and z_γ and between z_β and $z_{\alpha/2}$ respectively. The length of the confidence interval becomes $2z_{\alpha/2} + b - a$ and $A = B$. It is clear from Figure 4.4 that the density over the first interval is strictly larger than under the second interval. Therefore, $b \geq a$ and

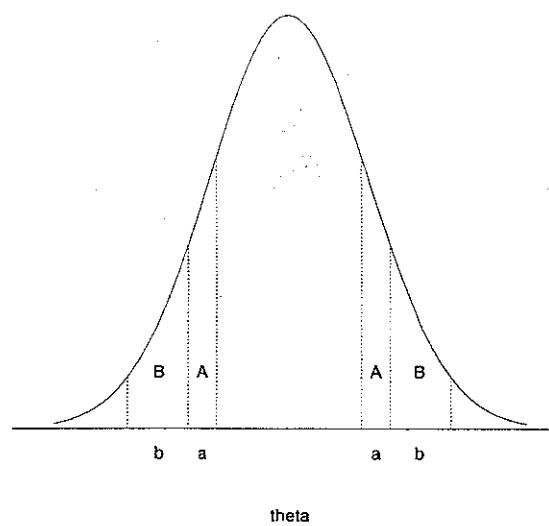


Fig. 4.4 Density of the standard normal distribution.

$b - a \geq 0$. The shortest possible interval is then obtained by taking $b = a = z_{\alpha/2}$. Therefore, the symmetric interval is the shortest one and every value of θ inside it has larger density than any point lying outside the interval.

This simple example provides the key to finding intervals of shortest length. It indicates that the length of the interval is inversely proportional to the density height. Shortest intervals are then provided by inclusion of points of higher density. This idea is mathematically expressed in the definition below and represented graphically in Figure 4.5.

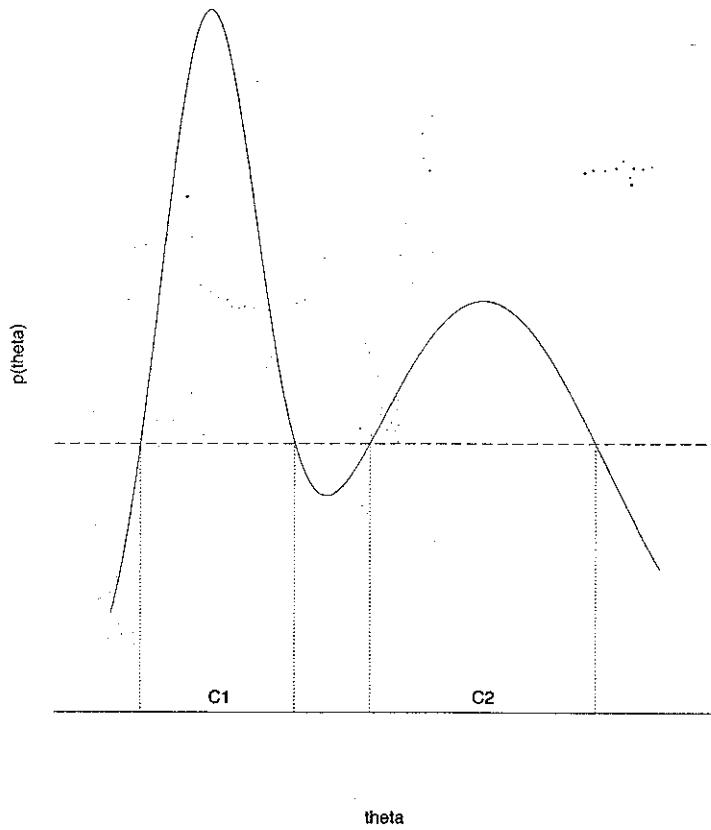


Fig. 4.5 The HPD interval for the density above is given by $C_1 \cup C_2$.

Definition. A $100(1 - \alpha)\%$ Bayesian interval of highest posterior density (HPD, for short) for θ is the $100(1 - \alpha)\%$ Bayesian interval C given by $C = \{\theta \in \Theta : p(\theta | x) \geq k(\alpha)\}$ where $k(\alpha)$ is the largest constant such that $P(\theta \in C | x) \geq 1 - \alpha$.

Example (Berger, 1985). Let $X = (X_1, \dots, X_n)$ be a random sample from the $\text{Cauchy}(\theta, 1)$ distribution and let θ have non-informative prior $p(\theta) \propto k$. The posterior density of θ is

$$p(\theta | x) \propto \prod_{i=1}^n \frac{1}{1 + (x_i - \theta)^2}.$$

Assume now that the observed sample was $x = (4.0, 5.5, 7.5, 4.5, 3.0)$, with sampling average $\bar{x} = 4.9$. Then, the 95% HPD confidence interval for θ can be numerically obtained as $[3.10, 6.06]$. Had we assumed a $N(\theta, 1)$ sampling distribution, the 95% HPD interval would be $[4.02, 5.86]$ which is more affected by the suspect value 7.5. In both cases, the intervals are easily obtained with the help of a computer. It will be seen in the sequel that the exercise is far from trivial for the Cauchy case in the frequentist approach. The main reasons are the absence of a univariate sufficient statistic for θ and the absence of asymptotic results to allow for approximations.

Note however that despite their appeal, HPD regions are not invariant under 1-to-1 transformations. The main reason is the existence of the Jacobian required when obtaining the density of any parametric transformation. Because of the invariance, transformation of HPD regions remain valid confidence intervals with the same confidence level. All they lose is the HPD property.

Finally, assume now that $Pr(\theta_i \in C_i) \geq 1 - \alpha_i, i = 1, \dots, r$, and let $C = C_1 \times \dots \times C_r$. If the θ_i 's are independent a posteriori then

$$Pr(\theta \in C) = \prod_{i=1}^r Pr(\theta_i \in C_i) \geq \prod_{i=1}^r (1 - \alpha_i)$$

and C is a $100(1 - \alpha)\%$ confidence region for θ if $\prod_i (1 - \alpha_i) \geq 1 - \alpha$. If the θ_i 's are not independent then

$$Pr(\theta \in C) \geq 1 - \sum_{i=1}^r Pr(\theta_i \notin C_i).$$

This result is also known as the Bonferroni inequality. As $Pr(\theta_i \notin C_i) \leq \alpha_i$, if one takes $\sum_i \alpha_i = \alpha$, taking for example $\alpha_i = \alpha/r$, then C is a $100(1 - \alpha)\%$ confidence region for θ .

4.4.2 Classical approach

In the case of classical confidence intervals, only sampling distributions can be used since parameters are unknown but fixed. Therefore, they are not liable to the probabilistic description they get under the Bayesian treatment. That is why the concept of confidence instead of probability intervals becomes relevant. Before describing the general formulation, it is useful to see it applied in an example.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution, with σ^2 known. To draw a classical inference one should ideally base calculations on a minimal sufficient statistic for θ . In this case, we have seen that \bar{X} is such a statistic and

$$\bar{X} \sim N\left(\theta, \frac{\sigma^2}{n}\right) \quad \text{or} \quad U = \frac{\bar{X} - \theta}{\sigma/\sqrt{n}} \sim N(0, 1).$$

Observe that U is a function of the sample and of the parameter θ , the parameter of interest and its distribution does not depend on θ . It can be said that

$$P(-z_{\alpha/2} \leq U \leq z_{\alpha/2}) = 1 - \alpha$$

and isolating θ yields

$$P\left(\bar{X} - z_{\alpha/2} \frac{\sigma}{\sqrt{n}} \leq \theta \leq \bar{X} + z_{\alpha/2} \frac{\sigma}{\sqrt{n}}\right) = 1 - \alpha.$$

So, even though an interval for θ was obtained, it cannot be understood as a probability interval for θ as in Bayesian intervals. It can only be interpreted in the sampling framework by saying that if the same experiment were to be repeated many times, in approximately $100(1 - \alpha)\%$ of them, its random limits $\bar{X} - z_{\alpha/2}\sigma/\sqrt{n}$ and $\bar{X} + z_{\alpha/2}\sigma/\sqrt{n}$ would include the value of θ . Also, this assertion is useless from a practical perspective since it is based on an unobserved sample. What can be done is to replace the observed value of \bar{X} in the expression and state that one can have $100(1 - \alpha)\%$ confidence, instead of probability, that the so-formed numerical interval contains θ .

The general procedure to obtain confidence intervals in the frequentist framework is based on a generalization of the steps of the above example to any statistical problem. These are:

1. A statistic $U = G(\mathbf{X}, \theta) \in \mathcal{U}$ with distribution that does not depend on $\theta \in \Theta$ must be found. Ideally, this statistic must depend on \mathbf{X} through minimal sufficient statistics and have a known distribution. Both requirements were met in the example since U depended on the sample through \bar{X} and had a standard normal distribution.
2. With knowledge of the distribution of U , find a region $A \subset \mathcal{U}$ such that $Pr(U \in A) = 1 - \alpha$. When θ is scalar, then a scalar U can be found in many cases with $Pr(a_1 \leq U \leq a_2) = 1 - \alpha$ and, in this case, $A = [a_1, a_2]$.
3. The confidence interval $C \subset \Theta$ is obtained by isolating θ in the above expression and replacing sample values.

A useful complementary text on the subject is Silvey (1970).

Definition. Let θ be an unknown quantity defined in Θ , U be a function $U = G(\mathbf{X}, \theta)$ with values in \mathcal{U} and A be a region in \mathcal{U} such that $Pr(U \in A) \geq 1 - \alpha$. A region $C \subset \Theta$ is a $100(1 - \alpha)\%$ frequentist confidence region for θ if

$$C = \{\theta : G(\mathbf{x}, \theta) \in A\}.$$

In this case, $1 - \alpha$ is called the confidence level. In the scalar case, the inversion in terms of θ usually leads to an interval, $C = [c_1, c_2]$ say, hence the name.

Once again, the use of the word interval for the general case is an abuse of language. The inequality in the definition is taken as an equality, whenever possible.

The quantity U is usually called a pivot or a pivotal quantity and finding one such quantity is fundamental. The choice of U is crucial to the success of the method. It is not at all obvious that reasonable options are available in any given problem. The effort towards the use of minimal sufficiency is in the direction of shortening intervals as much as possible.

All randomness present here is due to the sample \mathbf{X} leading to a probability interval for U and not for θ . The procedure involves an elaboration that is absent from the Bayesian definition. It provides an interval to which is associated a numerical value, the confidence of the interval. For that reason, it is often interpreted misleadingly as a probability interval, as in the Bayesian framework. Care must be exercised to ensure a correct interpretation of the intervals. The existent symmetry in many canonical situations leads to intervals that coincide numerically when obtained by a frequentist or a non-informative Bayesian approach. This happened in the above example but should not be used to unduly equate the two approaches.

In the case of a parameter vector, the use of Cartesian products of confidence intervals can also be applied to the construction of classical intervals. In particular, approximations such as those from the Bonferroni inequality are used more often in classical inference where the search for the pivotal quantity U does not always lead to independent components. This problem typically does not occur in the Bayesian approach where most problems lie in the computations.

4.5 Estimation in the normal model

This section deals with applications of point and interval estimation of means and variances to problems of one and two normal populations. The Bayesian perspective with both non-informative and proper conjugate priors and frequentist perspective are presented and compared. We will particularly emphasize the similarity between the results with the frequentist and the non-informative Bayesian points of view. The similarity is only numerical since we have seen that the derivations are completely different.

4.5.1 One sample case

Assume initially a single sample $\mathbf{X} = (X_1, \dots, X_n)$ from the $N(\theta, \sigma^2)$ distribution with $\phi = \sigma^{-2}$. If ϕ is known and the prior distribution is $\theta \sim N(\mu_0, \tau_0^2)$, we have already obtained that $\theta|\mathbf{x} \sim N(\mu_1, \tau_1^2)$ where

$$\mu_1 = \frac{n\sigma^{-2}\bar{x} + \tau_0^{-2}\mu_0}{n\sigma^{-2} + \tau_0^{-2}} \quad \text{and} \quad \tau_1^2 = \frac{1}{n\sigma^{-2} + \tau_0^{-2}}.$$

So, posterior mean, median and mode coincide and the posterior precision and curvature of the log posterior are given by τ_1^{-2} .

Credibility intervals can be obtained by noticing that

$$\frac{\theta - \mu_1}{\tau_1} \mid \mathbf{x} \sim N(0, 1)$$

and therefore

$$\begin{aligned} 1 - \alpha &= P(-z_{\alpha/2} < (\theta - \mu_1)/\tau_1 < z_{\alpha/2} \mid \mathbf{x}) \\ &= P(\mu_1 - z_{\alpha/2}\tau_1 < \theta < \mu_1 + z_{\alpha/2}\tau_1 \mid \mathbf{x}) \end{aligned}$$

and, due to the symmetry of the normal, $(\mu_1 - z_{\alpha/2}\tau_1, \mu_1 + z_{\alpha/2}\tau_1)$ is the $100(1 - \alpha)\%$ HPD interval for θ .

A non-informative prior can be obtained by letting $\tau_0^2 \rightarrow \infty$. In this case, $\tau_1^{-2} \rightarrow n\sigma^{-2}$ and $\mu_1 \rightarrow \bar{x}$. Posterior mean, median and mode coincide with the moments estimator and MLE, \bar{X} . It is easy to check that \bar{X} is also an unbiased, efficient and (strongly) consistent estimator for θ . Also, the $100(1 - \alpha)\%$ HPD confidence interval for θ coincides with the classical confidence interval obtained in the previous section.

Assuming now that θ is known and the prior for σ^2 is $n_0\sigma_0^2\phi \sim \chi_{n_0}^2$ leads to the posterior for $(n_0\sigma_0^2 + ns_0^2)\phi \mid \mathbf{x} \sim \chi_{n+n_0}^2$ where

$$s_0^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \theta)^2.$$

The following quantities can be obtained:

$$\begin{aligned} E[\phi \mid \mathbf{x}] &= \frac{n_0 + n}{n_0\sigma_0^2 + ns_0^2} \\ \{E[\phi \mid \mathbf{x}]\}^{-1} &= \frac{n_0\sigma_0^2 + ns_0^2}{n_0 + n} = \frac{n_0}{n_0 + n}\sigma_0^2 + \frac{n}{n_0 + n}s_0^2 \end{aligned}$$

which is a weighted average between the prior estimate σ_0^2 and the maximum likelihood estimate s_0^2 with weights $n_0/(n_0 + n)$ and $n/(n_0 + n)$, respectively. Also,

$$E[\sigma^2 \mid \mathbf{x}] = E[\phi^{-1} \mid \mathbf{x}] = \frac{n_0\sigma_0^2 + ns_0^2}{n + n_0 - 2}$$

which coincides with the inverse of the posterior mode (but not of the mean) of ϕ . The main dispersion measures are

$$V(\phi \mid \mathbf{x}) = \frac{2(n + n_0)}{(n_0\sigma_0^2 + ns_0^2)^2} \quad \text{and} \quad J(\text{mode}) = \frac{(n_0\sigma_0^2 + ns_0^2)^2}{2(n + n_0 - 2)}.$$

Note that once again the expression of $J(\text{mode})$ is very similar to the posterior precision $V^{-1}(\phi \mid \mathbf{x})$.

Confidence intervals can be obtained with the percentiles of the χ^2 that are available in many tables and statistical softwares. Defining $\underline{\chi}_{\alpha, v}^2$ and $\bar{\chi}_{\alpha, v}^2$ as the $100\alpha\%$ and $100(1 - \alpha)\%$ percentiles of the χ^2 distribution with v degrees of freedom, respectively gives

$$\begin{aligned} 1 - \alpha &= P\left(\underline{\chi}_{\alpha/2, n_1}^2 < (n_0\sigma_0^2 + ns_0^2)\phi < \bar{\chi}_{\alpha/2, n_1}^2 \mid \mathbf{x}\right) \quad \text{where } n_1 = n_0 + n \\ &= P\left(\frac{\underline{\chi}_{\alpha/2, n_1}^2}{n_0\sigma_0^2 + ns_0^2} < \phi < \frac{\bar{\chi}_{\alpha/2, n_1}^2}{n_0\sigma_0^2 + ns_0^2} \mid \mathbf{x}\right). \end{aligned}$$

This gives rise to a $100(1 - \alpha)\%$ credibility interval for ϕ . Given the asymmetry of the χ^2 distribution, this is not an HPD interval. As $\sigma^2 = 1/\phi$,

$$\left(\frac{n_0\sigma_0^2 + ns_0^2}{\bar{\chi}_{\alpha/2, n_1}^2}, \frac{n_0\sigma_0^2 + ns_0^2}{\underline{\chi}_{\alpha/2, n_1}^2}\right)$$

is a $100(1 - \alpha)\%$ credibility interval for σ^2 .

The non-informative prior can be obtained by letting $n_0 \rightarrow 0$. In this case, the posterior is $ns_0^2\phi \mid \mathbf{x} \sim \chi_n^2$ and $\{E[\phi \mid \mathbf{x}]\}^{-1} = s_0^2$, which coincides with the maximum likelihood estimate of $\phi^{-1} = \sigma^2$. The MLE is S_0^2 with sampling distribution $nS_0^2/\sigma^2 \mid \sigma^2 \sim \chi_n^2$. Therefore, it is unbiased and since

$$\frac{\partial \log p(\mathbf{X} \mid \sigma^2)}{\partial \sigma^2} = \frac{n}{2\sigma^4}(S_0^2 - \sigma^2)$$

it is also efficient. Its variance is $2\sigma^4/n$ and tends to 0 as $n \rightarrow \infty$ which means that the estimator is also consistent. The $100(1 - \alpha)\%$ credibility interval for σ^2 becomes

$$\left(\frac{ns_0^2}{\bar{\chi}_{\alpha/2, n}^2}, \frac{ns_0^2}{\underline{\chi}_{\alpha/2, n}^2}\right).$$

The pivotal quantity used for the construction of the interval is nS_0^2/σ^2 with a χ_n^2 sampling distribution. The classical confidence interval can be easily obtained and shown to coincide with the above interval, obtained for the non-informative prior.

If θ and σ^2 are both unknown quantities, using the conjugate prior $\theta \mid \phi \sim N(\mu_0, (c_0\phi)^{-1})$ and $n_0\sigma_0^2\phi \sim \chi_{n_0}^2$, gives the marginal posterior distributions $\theta \mid \mathbf{x} \sim t_{n_1}(\mu_1, \sigma_1^2/c_1)$ and $n_1\sigma_1^2\phi \mid \mathbf{x} \sim \chi_{n_1}^2$, where $n_1\sigma_1^2 = n_0\sigma_0^2 + (n - 1)s^2 + c_0n(\mu_0 - \bar{x})^2/(c_0 + n)$ and $s^2 = \Sigma(x_i - \bar{x})^2/(n - 1)$. Once again, the posterior mean, mode and median of θ coincide and are given by μ_1 . Also,

$$V(\theta \mid \mathbf{x}) = \frac{n_1}{n_1 - 2} \frac{\sigma_1^2}{c_1} \quad \text{and} \quad J(\mu_1) = \frac{n_1 + 1}{n_1} \frac{c_1}{\sigma_1^2}.$$

Denoting the $100(1 - \alpha)\%$ percentile of the $t_v(0, 1)$ distribution by $t_{\alpha, v}$, gives by symmetry of the Student- t that

$$1 - \alpha = P\left(-t_{\alpha/2, n_1} < \sqrt{c_1} \frac{\theta - \mu_1}{\sigma_1} < t_{\alpha/2, n_1}\right)$$

$$= P \left(\mu_1 - t_{\alpha/2, n_1} \frac{\sigma_1}{\sqrt{c_1}} < \theta < \mu_1 + t_{\alpha/2, n_1} \frac{\sigma_1}{\sqrt{c_1}} \right)$$

and the above is an HPD interval.

For σ^2 , by analogy with the results for known θ , $E[\phi | \mathbf{x}] = \sigma_1^{-2}$ and

$$E[\sigma^2 | \mathbf{x}] = E[\phi^{-1} | \mathbf{x}] = \frac{n_1 \sigma_1^2}{n_1 - 2}$$

which coincides with the inverse of the posterior mode (but not of the mean) of ϕ .

The main dispersion measures are

$$V[\phi | \mathbf{x}] = \frac{2n_1}{(n_1 \sigma_1^2)^2} \quad \text{and} \quad J(\text{mode}) = \frac{(n_1 \sigma_1^2)^2}{2(n_1 - 2)}.$$

Once again, the expression of $J(\text{mode})$ is very similar to the posterior precision

$$V^{-1}[\phi | \mathbf{x}].$$

Confidence intervals can again be obtained with the χ^2 percentiles leading to

$$\begin{aligned} 1 - \alpha &= P \left(\frac{\chi^2_{\alpha/2, n_1}}{n_1 \sigma_1^2} < \phi < \frac{\chi^2_{\alpha/2, n_1}}{n_1 \sigma_1^2} \mid \mathbf{x} \right) \\ &= P \left(\frac{\chi^2_{\alpha/2, n_1}}{n_1 \sigma_1^2} < \phi < \frac{\bar{X}_{\alpha/2, n_1}^2}{n_1 \sigma_1^2} \mid \mathbf{x} \right). \end{aligned}$$

This gives a $100(1 - \alpha)\%$ confidence interval for ϕ (that is also not an HPD interval). As $\sigma^2 = 1/\phi$,

$$\left(\frac{n_1 \sigma_1^2}{\bar{X}_{\alpha/2, n_1}^2}, \frac{n_1 \sigma_1^2}{\chi^2_{\alpha/2, n_1}} \right)$$

is a $100(1 - \alpha)\%$ confidence interval for σ^2 .

The non-informative prior in this case is $p(\theta, \phi) \propto \phi^{-1}$. This can be seen as a limiting case of the conjugate prior above when $c_0, \sigma_0^2 \rightarrow 0$ and $n_0 = -1$. This gives marginal posterior distributions $\theta \mid \mathbf{x} \sim t_{n-1}(\bar{x}, s^2/n)$ and $(n-1)s^2\phi \mid \mathbf{x} \sim \chi^2_{n-1}$. Again, the posterior mean, mode and median of θ coincide with \bar{x} , which is the maximum likelihood estimate.

The dispersion measures for θ are

$$V(\theta \mid \mathbf{x}) = \frac{n-1}{n-3} \frac{s^2}{n} \quad \text{and} \quad J(\bar{x}) = \frac{n^2}{(n-1)s^2}$$

and since $\sqrt{n}(\theta - \bar{x})/s \mid \mathbf{x} \sim t_{n-1}(0, 1)$, the HPD $100(1-\alpha)\%$ confidence interval for θ is analogously obtained as

$$\left(\bar{x} - t_{\alpha/2, n-1} \frac{s}{\sqrt{n}}, \bar{x} + t_{\alpha/2, n-1} \frac{s}{\sqrt{n}} \right).$$

The classical (moments and maximum likelihood) estimator \bar{X} for θ is unbiased, efficient and consistent. Classical confidence intervals for θ cannot be obtained

with the same pivotal quantity used before because it depends on the unknown σ . A new pivotal quantity depending only on \mathbf{X} and θ and with a distribution that is known and does not depend on any of the unknown parameters must be found. Fortunately, this is possible in the normal case with the following results.

Theorem 4.3. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution and let \bar{X} and S^2 be the sample mean and variance respectively. Then, conditional on θ and σ^2 , \bar{X} and S^2 are independent with respective sampling distributions

$$\sqrt{n} \frac{\bar{X} - \theta}{\sigma} \sim N(0, 1) \quad \text{and} \quad \frac{(n-1)S^2}{\sigma^2} \sim \chi^2_{n-1}.$$

Proof. Define $\mathbf{Z} = (Z_1, \dots, Z_n)$ where $Z_i = \sqrt{n}(X_i - \theta)/\sigma$, $i = 1, \dots, n$. Then, the Z_i 's are iid $N(0, 1)$, $\bar{Z} = \sqrt{n}(\bar{X} - \theta)/\sigma$ and, in matrix notation, $\mathbf{Z} \sim N(\mathbf{0}, \mathbf{I}_n)$, where \mathbf{I}_n is the $n \times n$ identity matrix. Let \mathbf{A} be an orthogonal matrix with first row given by $(1/\sqrt{n})\mathbf{1}_n'$ where $\mathbf{1}_n$ is an n -dimensional vector of 1's. There are many methods in linear algebra available for completing orthogonally the other $n-1$ rows of \mathbf{A} . From the invariance of the multivariate normal distribution under linear transformations (Exercise 1.6), it follows that

$$\mathbf{Y} = \mathbf{A} \mathbf{Z} \sim N(\mathbf{0}, \mathbf{I}_n) \text{ since } \mathbf{A}\mathbf{0} = \mathbf{0} \text{ and } \mathbf{A}\mathbf{A}' = \mathbf{I}_n.$$

Therefore, the Y_i 's are iid $N(0, 1)$ variables,

$$Y_1 = \frac{1}{\sqrt{n}} \mathbf{1}_n' \mathbf{Z} = \sqrt{n} \bar{Z} \quad \text{and} \quad \mathbf{Y}'\mathbf{Y} = \sum_{i=1}^n Y_i^2 \sim \chi^2_n.$$

Also, from the independence of the Y_i 's, $\sum_{i=2}^n Y_i^2 \sim \chi^2_{n-1}$. So,

$$\begin{aligned} (n-1) \frac{S^2}{\sigma^2} &= \sum_{i=1}^n \frac{(X_i - \bar{X})^2}{\sigma^2} \\ &= \sum_{i=1}^n (Z_i - \bar{Z})^2 \\ &= \sum_{i=1}^n Z_i^2 - n\bar{Z}^2 \\ &= \sum_{i=1}^n Z_i^2 - Y_1^2 \\ &= \mathbf{Z}'\mathbf{Z} - Y_1^2. \end{aligned}$$

But,

$$\sum_{i=1}^n Y_i^2 = \mathbf{Y}'\mathbf{Y} = (\mathbf{A} \mathbf{Z})'\mathbf{A} \mathbf{Z} = \mathbf{Z}'\mathbf{A}'\mathbf{A}\mathbf{Z} = \mathbf{Z}'\mathbf{Z} = \sum_{i=1}^n Z_i^2.$$

Therefore $(n - 1)S^2/\sigma^2 = \sum_{i=2}^n Y_i^2 \sim \chi_{n-1}^2$ and is independent of Y_1^2 and, consequently of Y_1 and \bar{X} , completing the proof.

□

Lemma. If $T \sim N(0, 1)$ and $W \sim \chi_v^2$ and T and W are independent then $T/\sqrt{W/v} \sim t_v(0, 1)$.

The proof of the Lemma is an adaptation of results previously shown and is left as an exercise.

Corollary. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution and let \bar{X} and S^2 be the sample mean and variance respectively. Then, conditional on θ and σ^2 , \bar{X} has sampling distribution

$$\sqrt{n} \frac{(\bar{X} - \theta)}{S} \sim t_{n-1}(0, 1).$$

Proof. A straightforward application of the last lemma with $T = \sqrt{n}(\bar{X} - \theta)/\sigma$, $W = (n - 1)S^2/\sigma^2$ and $v = n - 1$. Then, $T/\sqrt{W/v} = \sqrt{n}(\bar{X} - \theta)/S$, completing the proof.

□

The above results indicate how to define pivotal quantities for construction of confidence intervals for θ and σ^2 . In the case of θ , σ is replaced by its estimator S leading to the new pivotal quantity $\sqrt{n}(\bar{X} - \theta)/S$, whose sampling distribution is posterior of θ . It is easy to obtain that the classical interval will coincide with the substitution implies an increase in the uncertainty bounds since $t_{\beta,n} > z_\beta$ for small β .

For σ^2 , we have that $\{E[\phi | \mathbf{x}]\}^{-1} = s^2$, the mode of ϕ is $(n-3)/[(n-1)S^2] = \hat{\phi}$ and the dispersion measures are

$$V(\phi | \mathbf{x}) = \frac{2}{(n-1)s^4} \quad \text{and} \quad J(\hat{\phi}) = \frac{(n-1)^2 s^4}{2(n-3)}.$$

The MLE of σ^2 is $\hat{\sigma}^2 = (n-1)S^2/n$ which is biased and is usually replaced by the unbiased estimator S^2 with sampling distribution $(n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$ and $V(S^2) = 2\sigma^4/(n-1)$. Since S^2 is unbiased and $V(S^2) \rightarrow 0$ as $n \rightarrow \infty$, S^2 is a consistent estimator of σ^2 . The difference between S^2 and $\hat{\sigma}^2$ becomes negligible as n increases which means that $\hat{\sigma}^2$ is also consistent.

Non-informative Bayesian and classical intervals for σ^2 once again coincide and are given by

$$\left(\frac{(n-1)s^2}{\bar{X}_{\alpha/2, n-1}^2}, \frac{(n-1)s^2}{\bar{X}_{\alpha/2, n-1}^2} \right).$$

When making an inference about one of the parameters, the other one becomes a nuisance parameter. In the Bayesian approach it is eliminated by integration. In the frequentist approach it is eliminated by appropriate choices of pivotal quantities. In the normal case, we were able to find suitable quantities given by $\sqrt{n}(\bar{X} - \theta)/S$ and $(n - 1)S^2/\sigma^2$. These are based on minimal sufficient statistics and do not depend on the respective nuisance parameters. This fortunate coincidence does not necessarily occur in all statistical problems. In those cases, alternative approaches based on some form of approximation must be used.

4.5.2 Two samples case

From now on, until the end of the section, we will concentrate on two normal samples where $\mathbf{X}_1 = (X_{11}, \dots, X_{1n_1})$ is a random sample from the $N(\theta_1, \sigma_1^2)$ distribution and $\mathbf{X}_2 = (X_{21}, \dots, X_{2n_2})$ is a random sample from the $N(\theta_2, \sigma_2^2)$ distribution. In addition, the two samples will be assumed to be independent.

If σ_1^2 and σ_2^2 are known, the likelihood is

$$p(\mathbf{x}_1, \mathbf{x}_2 | \theta_1, \theta_2) = p(\mathbf{x}_1 | \theta_1) p(\mathbf{x}_2 | \theta_2) \\ \propto \exp \left\{ -\frac{n_1}{2\sigma_1^2} (\theta_1 - \bar{x}_1)^2 \right\} \exp \left\{ -\frac{n_2}{2\sigma_2^2} (\theta_2 - \bar{x}_2)^2 \right\}$$

which factors out into separate likelihoods for θ_1 and θ_2 . So, if θ_1 and θ_2 are prior independent, they will remain posterior independent. One class of conjugate prior is given by independent $\theta_i \sim N(\mu_i, \tau_i^2)$ distributions for $i = 1, 2$. Another class is given by bivariate normal distributions. It includes the previous class by allowing also non-null prior correlation between θ_1 and θ_2 . The first class will be used here for simplicity.

Combining the adopted prior with the likelihood leads to the independent posteriors $\theta_i | \mathbf{x}_i \sim N(\mu_i^*, \tau_i^{*2})$ where

$$\mu_i^* = \frac{n_i \sigma_i^{-2} \bar{x}_i + \tau_i^{-2} \mu_i}{n_i \sigma_i^{-2} + \tau_i^{-2}} \quad \text{and} \quad \tau_i^{*2} = \frac{1}{n_i \sigma_i^{-2} + \tau_i^{-2}}, \quad i = 1, 2.$$

The analysis is exactly like two separate conjugate analyses and all the results for one sample follow. The same comments are true for non-informative priors and for classical inference. The non-informative prior for θ_1 and θ_2 is $p(\theta_1, \theta_2) \propto k$ and

$$\theta_i | \mathbf{x}_i \sim N\left(\bar{x}_i, \frac{\sigma_i^2}{n_i}\right), \quad i = 1, 2 \text{ independent.}$$

The equivalent sampling result is $\bar{X}_i | \theta_i \sim N(\theta_i, \sigma_i^2/n_i)$, $i = 1, 2$ independent, from where point and interval estimation can be processed in the same way.

An interesting problem absent in the single sample case is comparison of means. This can be done by estimation of $\beta = \theta_1 - \theta_2$. The posterior distribution of β is obtained from the properties of the normal distribution as $N(\mu_1^* - \mu_2^*, \tau_1^{*2} + \tau_2^{*2})$.

This posterior reduces in the non-informative case to the $N(\hat{\beta}, \sigma_1^2/n_1 + \sigma_2^2/n_2)$ where $\hat{\beta} = \bar{x}_1 - \bar{x}_2$. In the case of classical inference, estimation is based on $\bar{X}_1 - \bar{X}_2$ with sampling distribution $N(\beta, \sigma_1^2/n_1 + \sigma_2^2/n_2)$. Observe that all the above distributions are symmetric which eases the calculation of estimators and HPD intervals.

Assume now that σ_1^2 and σ_2^2 are unknown but equal with $\phi = \sigma_1^{-2} = \sigma_2^{-2}$. Then a conjugate prior can be constructed in the following way: $\theta_i \mid \phi \sim N(\mu_i, (c_i\phi)^{-1})$, $i = 1, 2$, are conditionally independent and $n_0\sigma_0^2\phi \sim \chi_{n_0}^2$. The prior density of $(\theta_1, \theta_2, \phi)$ is

$$\begin{aligned}
p(\theta_1, \theta_2, \phi) &= p(\theta_1, \theta_2 | \phi)p(\phi) \\
&= p(\theta_1 | \phi)p(\theta_2 | \phi)p(\phi) \\
&\propto \phi^{1/2} \exp\left\{-\frac{\phi}{2}c_1(\theta_1 - \mu_1)^2\right\} \phi^{1/2} \exp\left\{-\frac{\phi}{2}c_2(\theta_2 - \mu_2)^2\right\} \\
&\quad \times \phi^{\{n_0/2\}-1} \exp\left\{-\frac{\phi}{2}n_0\sigma_0^2\right\} \\
&\propto \phi^{n_0/2} \exp\left\{-\frac{\phi}{2}[n_0\sigma_0^2 + c_1(\theta_1 - \mu_1)^2 + c_2(\theta_2 - \mu_2)^2]\right\}.
\end{aligned}$$

In particular, the prior distribution of $\beta \mid \phi$ is $N(\mu_1 - \mu_2, \phi^{-1}(c_1^{-1} + c_2^{-1}))$. Therefore, using the results from Section 3.4, one can obtain its marginal prior $\beta \sim t_{n_0}(\mu_1 - \mu_2, \sigma_0^2(c_1^{-1} + c_2^{-1}))$. The likelihood is

$$p(\mathbf{x}_1, \mathbf{x}_2 | \theta_1, \theta_2, \phi) = p(\mathbf{x}_1 | \theta_1, \phi)p(\mathbf{x}_2 | \theta_2, \phi)$$

$$\propto \prod_{i=1}^2 \phi^{n_i/2} \exp \left\{ -\frac{\phi}{2} \left[(n_i - 1)s_i^2 + n_i(\theta_i - \bar{x}_i)^2 \right] \right\}$$

where

$$s_i^2 = \frac{1}{n_i - 1} \sum_{j=1}^{n_i} (x_{ij} - \bar{x}_i)^2, \quad i = 1, 2,$$

Combining the likelihood with the prior gives the posterior density of $(\theta_1, \theta_2, \phi)$:

$$\phi^{(n_0+n_1+n_2)/2} \times \exp \left\{ -\frac{\phi}{2} \left[n_0 \sigma_0^2 + \nu s^2 + \sum_{i=1}^2 \frac{c_i n_i}{c_i + n_i} (\mu_i - \bar{x}_i)^2 + (c_i + n_i)(\theta_i - \mu_i^*)^2 \right] \right\}$$

where $\mu_i^* = (c_i\mu_i + n_i\bar{x}_i)/(c_i + n_i)$, $i = 1, 2$, $v = n_1 + n_2 - 2$ and $S^2 = [(n_1 - 1)S_1^2 + (n_2 - 1)S_2^2]/v$, that is a weighted average of the s_i^2 obtained within each sample with weights given by $n_i - 1$, $i = 1, 2$. Note that the posterior and prior densities have the same kernel. Hence, the following results can be established by analogy:

$$\theta_i | \phi, \mathbf{x} \sim N\left(\mu_i^*, \frac{1}{c_i^* \phi}\right) \text{ independent and } \theta_i | \mathbf{x} \sim t_{n_0^*}\left(\mu_i^*, \frac{\sigma_0^*}{c_i^*}\right), \quad i = 1, 2,$$

and $n_0^* \sigma_{\phi}^{*2} | x \sim \chi^2_{n_0^*}$ where $c_i^* = c_i + n_i$, $n_0^* = n_0 + n_1 + n_2$ and

$$n_0^* \sigma_0^{*2} = n_0 \sigma_0^2 + vs^2 + \sum_{i=1}^2 \frac{c_i n_i}{c_i + n_i} (\mu_i - \bar{x}_i)^2.$$

In terms of β , one can obtain the conditional posterior distribution $\beta \mid \phi, x \sim N(\mu_1^* - \mu_2^*, \phi^{-1}(c_1^{*-1} + c_2^{*-1}))$ and marginal posterior distribution $\beta \mid x \sim t_{n_0}^*(\mu_1^* - \mu_2^*, \sigma_0^{*2}(c_1^{*-1} + c_2^{*-1}))$. Once again, the posterior is symmetric and posterior mean, mode and median of β coincide. HPD intervals for β can be obtained using percentiles of the Student t distribution. For ϕ , we have that $\{E(\phi \mid x)\}^{-1} = \sigma_0^{*2}$ and credibility intervals can be constructed using percentiles of the χ^2 distribution.

The non-informative prior distribution can be obtained by noting that this is a location scale model with location parameters θ_1 and θ_2 and scale parameter ϕ . Therefore, the non-informative prior is given by $p(\theta_1, \theta_2, \phi) \propto \phi^{-1}$. This can be seen as a limiting case of the conjugate prior above when $c_1, c_2, \sigma_0^2 \rightarrow 0$ and $n_0 = -2$. Replacing these values in the expression of the conjugate posterior gives $c_i^* = n_i$, $\mu_i^* = \bar{x}_i$, $n_0^* = v$ and $n_0^* \sigma_0^{*2} = vs^2$. Therefore, the posterior mean, mode and median of β are given by $\hat{\beta}$ and the $100(1 - \alpha)\%$ HPD interval for β has limits $\hat{\beta} \pm t_{\alpha/2, v}s^2 \sqrt{n_1^{-1} + n_2^{-1}}$. A possible estimate for σ^2 is given by s^2 . The $100(1 - \alpha)\%$ credibility interval for σ^2 obtained as before is given by

$$\left(\frac{v s^2}{\chi_{\alpha/2,v}^2}, \frac{v s^2}{\chi_{-\alpha/2,v}^2} \right).$$

In the case of classical inference, $\hat{\beta} = \bar{X}_1 - \bar{X}_2$ and $\hat{\sigma}^2 = vS^2/(n_1 + n_2)$ are the MLE of β and σ^2 , respectively. It is not difficult to show that $\hat{\beta}$ is an unbiased, efficient and consistent estimator for β and $\hat{\sigma}^2$ is a consistent but biased estimator for σ^2 . It is usually replaced by S^2 which is an unbiased, efficient and consistent estimator for σ^2 . The relevant sampling distributions are

$$\frac{\hat{\beta} - \beta}{S\sqrt{n_1^{-1} + n_2^{-1}}} \sim t_v(0, 1) \quad \text{and} \quad \frac{\nu S^2}{\sigma^2} \sim \chi^2_v.$$

Note that these variables provide pivotal quantities for the construction of confidence intervals for β and σ^2 respectively. The resulting confidence intervals coincide numerically with those provided by the non-informative prior.

If $\sigma_1^2 = \phi_1^{-1}$ and $\sigma_2^2 = \phi_2^{-1}$ are unknown and unequal, the likelihood factors according to

$$p(\mathbf{x} \mid \theta_1, \theta_2, \sigma_1^2, \sigma_2^2) = p(\mathbf{x}_1 \mid \theta_1, \sigma_1^2) p(\mathbf{x}_2 \mid \theta_2, \sigma_2^2).$$

Adopting independent normal-gamma conjugate priors with parameters $(\mu_i, c_i, v_i, s_{0i}^2)$, $i = 1, 2$, for each of the samples leads to the independent posterior

distributions

$$\theta_i | \mathbf{x} \sim t_{v_i^*} \left(\mu_i^*, \frac{s_{0i}^{*2}}{c_i^*} \right), \quad i = 1, 2 \quad \text{and} \quad v_i^* s_{0i}^{*2} \phi_i | \mathbf{x} \sim \chi_{v_i^*}^2, \quad i = 1, 2$$

where the relevant quantities are obtained by the usual one sample operations associated with the conjugacy of the normal-gamma by the normal observational model.

If interest lies in the comparison of the means, the posterior distribution of β must be obtained. First, let τ and ω be such that

$$\tau = \frac{\beta - (\mu_1^* - \mu_2^*)}{\sqrt{s_{01}^{*2}/c_1^* + s_{02}^{*2}/c_2^*}} \quad \text{and} \quad \tan \omega = \frac{s_{01}^*/\sqrt{c_1^*}}{s_{02}^*/\sqrt{c_2^*}}$$

It follows that

$$\sin \omega = \frac{s_{01}^*/\sqrt{c_1^*}}{\sqrt{s_{01}^{*2}/c_1^* + s_{02}^{*2}/c_2^*}} \quad \text{and} \quad \cos \omega = \frac{s_{02}^*/\sqrt{c_2^*}}{\sqrt{s_{01}^{*2}/c_1^* + s_{02}^{*2}/c_2^*}}$$

and therefore

$$\tau = \frac{\theta_1 - \mu_1^*}{s_{01}^*/\sqrt{c_1^*}} \sin \omega - \frac{\theta_2 - \mu_2^*}{s_{02}^*/\sqrt{c_2^*}} \cos \omega$$

where the fractions in the right-hand side of the equation have independent standard Student- t distributions with respective v_1^* and v_2^* degrees of freedom. A random quantity under these conditions is said to have a Behrens–Fisher distribution with parameters v_1^* , v_2^* and ω . This distribution is similar to the Student t distribution and has been tabulated, enabling easy construction of confidence intervals.

In the case of a non-informative prior $p(\theta_1, \theta_2, \sigma_1^2, \sigma_2^2) \propto \sigma_1^{-2} \sigma_2^{-2}$. This can be seen as a limiting case of the conjugate prior above when $c_i, s_{0i}^2 \rightarrow 0$ and $v_i = -1$, $i = 1, 2$. Replacing these values in the expression of the conjugate posterior gives $c_i^* = n_i$, $\mu_i^* = \bar{x}_i$, $v_i^* = n_i - 1$ and $s_{0i}^* = s_i$, where s_i^2 is the usual unbiased estimate of σ_i^2 , $i = 1, 2$. Estimators and confidence intervals can be obtained accordingly. The problem is more difficult to treat under the classical perspective. No easy pivotal quantity for β can be found with known distribution although approximations based on the Behrens–Fisher distribution were proposed.

Another situation of interest is the comparison of variances. Since variances measure the scale of a distribution and are always positive, it makes more sense to compare them through their ratio instead of their difference as we did for the means. Therefore, we will focus on the posterior distribution of $\psi = \sigma_2^2/\sigma_1^2 = \phi_1/\phi_2$. We have just obtained that ϕ_1 and ϕ_2 are posterior independent with joint density

$$p(\phi_1, \phi_2 | \mathbf{x}) \propto \prod_{i=1}^2 \phi_i^{v_i^*/2-1} \exp \left\{ -\frac{v_i^* s_{0i}^{*2}}{2} \phi_i \right\}.$$

The easiest form to obtain the posterior distribution of ψ is to complete the transformation to ensure a bijection and use results available for densities of 1-to-1 transformations of random quantities. Let $\psi_2 = \phi_2$ complete the transformation. The inverse relation $\phi_1 = \psi \phi_2 = \psi \psi_2$ follows. The Jacobian of the transformation is

$$J = \begin{vmatrix} \partial(\phi_1, \phi_2) \\ \partial(\psi, \psi_2) \end{vmatrix} = \begin{vmatrix} \psi_2 & \psi \\ 0 & 1 \end{vmatrix} = \psi_2 > 0$$

and the posterior density of (ψ, ψ_2) is

$$p(\psi, \psi_2 | \mathbf{x}) \propto \psi^{v_1^*/2-1} \psi_2^{(v_1^*+v_2^*)/2-1} \exp \left\{ -\frac{\psi_2}{2} [v_2^* s_{02}^{*2} + v_1^* s_{01}^{*2} \psi] \right\}.$$

Finally, the marginal density of ψ is

$$\begin{aligned} p(\psi | \mathbf{x}) &= \int p(\psi, \psi_2 | \mathbf{x}) d\psi_2 \\ &\propto \psi^{v_1^*/2-1} \int \psi_2^{(v_1^*+v_2^*)/2-1} \exp \left\{ -\frac{\psi_2}{2} [v_2^* s_{02}^{*2} + v_1^* s_{01}^{*2} \psi] \right\} d\psi_2 \\ &= \psi^{v_1^*/2-1} \frac{\Gamma((v_1^*+v_2^*)/2)}{[(v_2^* s_{02}^{*2} + v_1^* s_{01}^{*2} \psi)/2]^{(v_1^*+v_2^*)/2}} \\ &\propto \psi^{v_1^*/2-1} \left(v_2^* + v_1^* \frac{s_{01}^{*2}}{s_{02}^{*2}} \psi \right)^{-(v_1^*+v_2^*)/2}. \end{aligned}$$

It can then be shown that

$$\frac{s_{01}^{*2}}{s_{02}^{*2}} \psi | \mathbf{x} \sim F(v_1^*, v_2^*).$$

Even though the distribution function of the F distribution cannot be obtained analytically, it is tabulated in many books and computer software. Its percentiles can be used in the construction of confidence intervals. The main properties of the F distribution are given in the list of distributions. An interesting property for probability evaluation with the F distribution can be derived from the fact that if $X \sim F(v_2, v_1)$ then $X^{-1} \sim F(v_1, v_2)$ by simple inversion in the ratio of independent χ^2 distributions. Therefore, denoting the α and $1 - \alpha$ quantiles of the $F(v_1, v_2)$ distribution respectively by $F_\alpha(v_1, v_2)$ and $\bar{F}_\alpha(v_1, v_2)$ gives that $F_\alpha(v_1, v_2) = \bar{F}_\alpha^{-1}(v_2, v_1)$.

Point estimators are given by

$$E[\psi | \mathbf{x}] = \frac{s_{02}^{*2}}{s_{01}^{*2}} \frac{v_2^*}{v_2^* - 2} \quad \text{and mode}(\psi) = \frac{s_{02}^{*2} v_2^* (v_1^* - 2)}{s_{01}^{*2} v_1^* (v_2^* + 2)}.$$

Both estimators converge to $(s_{02}^*/s_{01}^*)^2$ when $v_2^* \rightarrow \infty$.

In the case of a non-informative prior the s_{0i}^{*2} 's are given by s_i^2 , the unbiased estimates of the populational variances and the v_i^* 's are given by $n_i - 1$, $i = 1, 2$.

Therefore, for large samples the estimators will be given by the ratio of the variance estimates.

In classical inference, we have already obtained the independent sampling distributions of $(n_i - 1)S_i^2 \phi_i \sim \chi_{n_i-1}^2$, $i = 1, 2$. Therefore, from the properties of the F distribution, $(S_1^2/S_2^2)\psi \sim F(n_1 - 1, n_2 - 1)$. Once again a pivotal quantity was found for inference about ψ and the resulting confidence interval will be numerically equivalent to the non-informative Bayesian one. The unbiased estimator of $\psi^{-1} = \sigma_1^2/\sigma_2^2$ is

$$\frac{n_2 - 3}{n_2 - 1} \frac{S_1^2}{S_2^2}.$$

Confidence intervals are obtained in either case using quantiles of the F distribution. These are not HPD intervals due to the asymmetry of the F distribution. Using the results listed about the F , a $100(1 - \alpha)\%$ credibility interval for ψ is obtained from

$$\begin{aligned} 1 - \alpha &= P\left(\frac{s_{02}^*}{s_{01}^*} \overline{F}_{\alpha/2}(v_2^*, v_1^*) < \frac{s_{02}^{*2}}{s_{01}^{*2}} \psi < \frac{s_{01}^{*2}}{s_{02}^{*2}} \overline{F}_{\alpha/2}(v_1^*, v_2^*) \mid \mathbf{x}_1, \mathbf{x}_2\right) \\ &= Pr\left[\frac{s_{01}^{*2}}{s_{02}^{*2}} \overline{F}_{\alpha/2}(v_2^*, v_1^*) < \psi < \frac{s_{01}^{*2}}{s_{02}^{*2}} \overline{F}_{\alpha/2}(v_1^*, v_2^*) \mid \mathbf{x}_1, \mathbf{x}_2\right] \end{aligned}$$

and, therefore,

$$\left[\left(s_{01}^{*2}/s_{02}^{*2} \right) \overline{F}_{\alpha/2}^{-1}(v_2^*, v_1^*), \left(s_{01}^{*2}/s_{02}^{*2} \right) \overline{F}_{\alpha/2}^{-1}(v_1^*, v_2^*) \right]$$

is a $100(1 - \alpha)\%$ credibility interval for ψ . In the case of non-informative priors, the modifications previously mentioned apply and the resulting interval is

$$\left[\frac{s_2^2}{s_1^2} \overline{F}_{\alpha/2}^{-1}(n_2 - 1, n_1 - 1), \frac{s_2^2}{s_1^2} \overline{F}_{\alpha/2}(n_1 - 1, n_2 - 1) \right].$$

Finally, the confidence interval derived from classical inference coincides with the above interval, obtained with a non-informative prior.

Exercises

§ 4.1

- Prove that the estimator of θ associated with the absolute loss is the posterior median of the updated distribution of θ .
- Show that if L_1 and L_2 are two proportional loss functions, that is, $L_1(\delta, \theta) = kL_2(\delta, \theta)$, then the Bayes estimators associated with these losses coincide.
- Suppose that X with distribution $N(\theta, \sigma^2)$ with σ^2 known is observed and it is known that $\theta \in (a, b)$ ($a < b$).
 - Obtain the non-informative prior for θ .

- Obtain the complete expression of the resulting posterior.
- Obtain the posterior mean and mode.
- Let X_1, \dots, X_n be a random sample from the Bernoulli distribution with unknown parameter θ having unit uniform prior distribution. One wishes to estimate θ using the loss function $L(d, \theta) = (\theta - d)^2/[\theta(1 - \theta)]$.
 - Calculate the Bayes estimator and obtain its risk.
 - Determine the predictive distribution for the $(n + 1)$ th observation and determine its mean and variance.
 - Generalize items (a) and (b) to k possible values for each X_i with respective probabilities θ_j , $j = 1, \dots, k$, and obtain the Bayes estimator of θ_j , $j = 1, \dots, k$.
- Suppose that the loss function used to estimate θ through δ
 - equals the distance between δ and θ if δ is smaller than θ , and
 - triples the distance between δ and θ if δ is larger than θ .
 - Obtain the mathematical expression of the loss function.
 - Show that the estimator of θ is the first quartile of the updated distribution of θ .
 - Generalize the result in (b) for when the loss in (ii) is p times the distance between δ and θ .
- Suppose that $X \sim \text{bin}(n, \theta)$ and the conjugate prior $\theta \sim \text{beta}(a, b)$ is used.
 - What is the value of X that minimizes the variance of the posterior distribution of θ ?
 - What is the value of X that maximizes it? Interpret the results.
 - Repeat items (a) and (b) for the case of a negative binomial distribution for X .
- Assume that $X \sim U[\theta - 1, \theta + 1]$ is observed and assume a prior $p(\theta) \propto \theta^{-1}$, $\theta > 0$.
 - Prove that $p(\theta|x) = c\theta^{-1}$, $\theta \in (x - 1, x + 1)$, where $c^{-1} = \log[(x + 1)/(x - 1)]$, $x > 1$.
 - Calculate the mean, mode and median of the posterior distribution.
- The income tax policy of a given country establishes that an individual pays tax q iff its income is larger than k . Assume that the income distribution for these individuals follows a $\text{Pa}(k, \theta)$ distribution.
 - Show that $\log(X/k) \sim \text{Exp}(\theta)$.
 - Assuming that little is known a priori about θ , a sample of these individuals is observed and their incomes registered. Show that the posterior distribution of θ is $G(n, n \log(\bar{G}/k))$ where n is the sample size and \bar{G} is the geometric average of the observations.

- (c) Assume a change in policy is under study aiming at taxing the rich more. The threshold would be raised to $l > k$ and tax raised to $r > q$. Show that the expected revenue would only rise if

$$r > q \left(1 + \frac{\log(l/k)}{n \log(G/k)} \right)^n.$$

Hint: first calculate the expected revenue given θ .

§ 4.2

9. Consider a simple linear regression where $E(Y_i | \theta) = \theta_0 + \theta_1 X_i$ with $\theta = (\theta_0, \theta_1)$, $i = 1, \dots, n$. Obtain the sum of squares $S(\theta)$ and show that its minimization leads to the least squares estimator

$$(\hat{\theta}_0, \hat{\theta}_1) = \left(\bar{Y} - \hat{\theta}_1 \bar{X}, \frac{\bar{XY} - \bar{X}\bar{Y}}{\bar{X}^2 - \bar{X}^2} \right) \text{ where } \bar{g}(X, Y) = \frac{1}{n} \sum_{i=1}^n g(X_i, Y_i).$$

10. Show that the MLE of a parameter is a minimal sufficient statistic for this parameter.
11. Suppose that the waiting time in a bank queue has $\text{Exp}(\theta)$ distribution with $\theta > 0$. A sample of n customers is observed over a period of T minutes.
- (a) Suppose the individual waiting times were discarded and only the number X of clients was recorded. Determine the MLE of θ based on X .
 - (b) Determine the maximum likelihood and Bayes estimators for θ , assuming that in a sample of $n = 20$ customers the average serving time was 3.8 minutes and all 20 clients were served.
 - (c) Suppose that, in addition to the observations reported in (b), an additional observation was made but all that is known is that it lasted for more than 5 minutes. Obtain the maximum likelihood and Bayes estimators of θ in this case.
12. Suppose one wishes to test three types of bulbs: normal life, long life and extra long life. The lifetimes of the bulbs have exponential distribution with means θ , 2θ and 3θ , respectively. Assume the test consists of observing one randomly selected bulb of each type.
- (a) Determine the MLE of θ .
 - (b) Determine the method of moments estimator of θ .
 - (c) Let $\psi = 1/\theta$ and assume the prior $\psi \sim G(\alpha, \beta)$. Determine the posterior distribution of θ .
 - (d) Determine the Bayes estimators of θ using 0-1 and quadratic loss functions.

13. Let $(X_1, Y_1), \dots, (X_n, Y_n)$ be bidimensional vectors forming a sample from the bivariate normal with mean vector μ , variances σ_i^2 , $i = 1, 2$, and correlation coefficient ρ . Determine the MLE of all model parameters.

§ 4.3

14. Prove the Cramer–Rao inequality for the multiparameter case.
Hint: Consider the joint covariance matrix of δ and the score function and explore the non-negative definiteness of this matrix.
15. Consider a random sample $X = (X_1, \dots, X_n)$ from an unknown distribution function F and let \hat{F} denote the empirical distribution function. Show that
- (a) \hat{F} is non-decreasing and contained in the interval $[0, 1]$.
 - (b) $\hat{F}(x)$ can be written as \bar{Z} where $Z_i = I_{X_i}(-\infty, x]$, $i = 1, \dots, n$.
 - (c) Show that \hat{F} is an unbiased estimator of F .
 - (d) Show that \hat{F} is a consistent estimator of F .

Hint: obtain the sampling distribution of $\hat{F}(x)$, $\forall x$.

16. Let X_1, X_2, \dots, X_n be a random sample from the $\text{Pois}(\theta)$ distribution and $Y = \sum_{i=1}^n X_i$.
- (a) Determine c such that $\exp(-cY)$ is an unbiased estimator of $\exp(-\theta)$.
 - (b) Obtain a bound for the variance of the estimator obtained in (a).
 - (c) Discuss the efficiency of the estimator obtained in (a).
17. Let X_1, \dots, X_n be a random sample from the uniform distribution over interval $[0, \theta]$.
- (a) Obtain $\hat{\theta}_n$, the MLE of θ , and show it is a biased but consistent estimator of θ .
 - (b) Obtain from (a), an unbiased estimator of θ .
 - (c) Calculate the quadratic risks of the estimators in (a) and (b).
 - (d) Find an estimator of θ with smaller risk than those obtained in (a) and (b).

18. It is common in genetics to obtain samples from the binomial distribution with the impossibility of 0 observations.
- (a) Show that the sampling distribution is given by
- $$f(x | \theta) = \binom{n}{x} \cdot \frac{\theta^x (1-\theta)^{n-x}}{1 - (1-\theta)^n}, \quad x = 1, \dots, n.$$
- (b) Obtain the MLE of θ , assuming that $n=2$.
 - (c) Verify whether the above estimator is unbiased.
19. Let X_1, \dots, X_n be a random sample from the uniform distribution on the interval $[a-b, a+b]$, $a \in \mathbb{R}$ and $b > 0$.
- (a) Verify if a and b are location and/or scale parameters.

- (b) Obtain the MLE of a and b .
 (c) Assume now that $b = 1$ and define

$$T_1 = \bar{X} = \frac{\sum_{i=1}^n X_i}{n} \quad \text{and} \quad T_2 = \frac{1}{2} \left(\max_{1 \leq i \leq n} X_i + \min_{1 \leq i \leq n} X_i \right).$$

- (d) Show that T_1 and T_2 are consistent and unbiased estimators of a .
 (e) Compare T_1 and T_2 specifying a choice between them and justifying it.
20. Let X_1, \dots, X_n be iid with probability function $f(x | \theta) = \theta(1-\theta)^x$, $x = 0, 1, 2, \dots$, $\theta \in (0, 1)$ and define $U_i = I_{X_i}(0)$.
- (a) Show that \bar{U} is an unbiased estimator of θ and calculate its variance.
 - (b) Show that the expected Fisher information for θ is $n/(\theta^2(1-\theta))$ and find the efficiency of \bar{U} .
21. Let X_1 and X_2 be iid with $\text{Exp}(\theta)$ distribution.
- (a) Show that $U = \exp(-X_1)$ is an unbiased estimator of $\psi = \theta/(\theta+1)$.
 - (b) Show that $(1-e^{-T})/T$ is a UMVU estimator of ψ , for $T = X_1 + X_2$.

§ 4.4

22. Show that Bayesian confidence intervals are invariant under 1-to-1 transformations of the parameter. So, if C is a $100(1-\alpha)\%$ confidence interval for θ and $\phi = \phi(\theta)$ is a 1-to-1 transformation of θ then $\phi(C)$, the image of C under ϕ , is a $100(1-\alpha)\%$ confidence interval for ϕ . Show that HPD intervals are not invariant under 1-to-1 transformations. Show also that frequentist confidence intervals are invariant under 1-to-1 transformations.
23. Let $\theta = (\theta_1, \dots, \theta_r)$. Show that if the θ_i 's are not independent then

$$\Pr(\theta \in C) \geq 1 - \sum_{i=1}^r \Pr(\theta_i \notin C_i).$$

Show that if $\Pr(\theta_i \in C_i) \geq \alpha_i = \alpha/r$, then C is a $100(1-\alpha)\%$ confidence region for θ . Obtain the equivalent result from a classical perspective.

Hint: Define pivotal quantities $U_i = G_i(X, \theta_i)$, $i = 1, \dots, r$.

24. Let $X \sim \text{bin}(n, \theta)$ and assume the prior $\theta \sim U[0, 1]$. Suppose that the observed value was $X = n$.
- (a) Show that the $100(1-\alpha)\%$ HPD interval for θ has form $[a, 1]$, $a < 1$.
 - (b) Let $\psi = \theta/(1-\theta)$. Show from (a) that $\Pr[a/(1-a) \leq \psi | x] = 1-\alpha$ and therefore $[a/(1-a), \infty)$ is a $100(1-\alpha)\%$ credibility interval for ψ .
 - (c) Obtain the posterior distribution of ψ and discuss the form of a $100(1-\alpha)\%$ HPD interval for ψ .
 - (d) In particular, is the interval obtained in (b) of HPD?

25. Let X_1, \dots, X_n be a random sample from the uniform distribution over the interval $[0, \theta]$.
- (a) Obtain a pivot based on the sampling distribution of $\hat{\theta}_n$, the MLE of θ .
 - (b) Obtain a $100(1-\alpha)\%$ classical confidence interval for θ based on the pivot used in (a).
 - (c) Which conditions must be satisfied for a minimum length interval?
 - (d) Assuming the non-informative prior $p(\theta) \propto \theta^{-1}$, find the $100(1-\alpha)\%$ HPD interval for θ .
26. (Berger, 1985, p. 134) Let X be such that $p(x | \theta) = \exp[-(x-\theta)]$, $x \geq \theta$ and assume the prior $p(\theta) \propto (1+\theta^2)^{-1}$, $\theta \geq 0$.
- (a) Prove that the posterior distribution of θ given x is monotonically increasing and that the mode of θ is x .
 - (b) Show that the $100(1-\alpha)\%$ HPD interval for θ must have the form $[c(\alpha), x]$, where $c(\alpha)$ is such that $P[c(\alpha) \leq \theta \leq x | x] = 1-\alpha$.
 - (c) Obtain the posterior density of $\eta = \exp(\theta)$ and prove it is a monotonically increasing function of η .
 - (d) Show that the $100(1-\alpha)\%$ HPD interval for η must have the form $[1; d(\alpha)]$, where $d(\alpha)$ is such that $P(1 \leq \eta \leq d(\alpha) | x) = 1-\alpha$.
 - (e) Show that the confidence interval in (d) implies a $100(1-\alpha)\%$ lowest posterior density interval for θ .
27. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from a $\text{Pois}(\theta)$ distribution. The prior distribution for θ is judged to be based on information equivalent to that obtained after observing a lifetimes with $\text{Exp}(\theta)$ distribution with observed mean lifetime b , where $a, b > 0$.
- (a) Show that the prior distribution of θ is $G(a+1, ab)$.
 - (b) Obtain the distribution of $\theta | \mathbf{x}$.
 - (c) Suppose that instead of observing the sample, only $T = \sum_{i=1}^n X_i$ was observed. Obtain the distribution of $\theta | T = t$.
- Hint: if \mathbf{X} consist of iid $\text{Pois}(\theta)$ variables then $T \sim \text{Pois}(n\theta)$.
- (d) Compare the distributions obtained in (b) and (c) when $\sum_{i=1}^n x_i = t$, justifying the result obtained.
 - (e) Obtain a $100(1-\alpha)\%$ confidence interval for θ using the percentiles of the χ^2 distribution.
28. Let $X_i = \theta t_i + \epsilon_i$, $i = 1, \dots, n$, where ϵ_i are iid $N(0, \sigma^2)$, σ^2 known and the non-informative prior for θ is assumed.
- (a) Obtain the posterior distribution of θ .
 - (b) Obtain the $100(1-\alpha)\%$ HPD region for θ .
 - (c) Show that the MLE of θ is the UMVU estimator of θ .
 - (d) Based on the sampling distribution of the MLE of θ , construct a $100(1-\alpha)\%$ confidence interval for θ .

- (e) If $0 \leq t_i \leq 1$, $\forall i$, what values of the t_i 's must be chosen to obtain the confidence intervals in (b) and (d) of shortest possible length?

§ 4.5

29. Show that if $T \sim N(0, 1)$ and $W \sim \chi^2_v$ are independent then $T/\sqrt{W/v} \sim t_v(0, 1)$.

30. Consider the situation with two normal samples where $X_1 = (X_{11}, \dots, X_{1n_1})$ is a random sample from the $N(\theta_1, \sigma_1^2)$ distribution and $X_2 = (X_{21}, \dots, X_{2n_2})$ is a random sample from the $N(\theta_2, \sigma_2^2)$ distribution. In addition, the two samples will be assumed to be independent.

- (a) If σ_1^2 and σ_2^2 are known, show that the class of bivariate normal distributions for θ_1 and θ_2 is conjugate with the above observation model. Also, obtain the posterior correlation between θ_1 and θ_2 and compare it with the prior correlation.

- (b) If $\sigma_1^2 = \sigma_2^2 = \sigma^2$ is unknown, show that

$$\frac{\hat{\beta} - \beta}{S\sqrt{n_1^{-1} + n_2^{-1}}} \sim t_v(0, 1) \quad \text{and} \quad \frac{vS^2}{\sigma^2} \sim \chi^2_v.$$

- (c) Show that the resulting confidence intervals for β and σ^2 obtained with the pivotal quantities of the previous item coincide numerically with the credibility intervals obtained with a non-informative prior.

- (d) If σ_1^2 and σ_2^2 are unequal and unknown, show that

$$\frac{s_{01}^{*2}}{s_{02}^{*2}} \psi \mid x \sim F(v_1^*, v_2^*).$$

31. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution and assume that the normal- χ^2 conjugate prior is used for (θ, ϕ) with $n_0 = 5$ where $\phi = \sigma^{-2}$. What should be the sample size to guarantee that $Pr((\theta - \mu_1)^2 \leq 4V_1 | \mathbf{x}) \geq 0.95$ where $\mu_1 = E(\theta | \mathbf{x})$ and $V_1 = V(\theta | \mathbf{x})$?

32. Let X_1 and X_2 be two observations from the $N(\theta, \sigma^2)$ distribution and assume a non-informative prior for (θ, σ^2) . If Y_1 is the smallest of the observation and Y_2 the largest one, show that a posteriori,

$$Pr(y_1 \leq \theta \leq y_2) = 0.5.$$

Hint: write the posterior distribution of θ as a function of y_1 and y_2 .

33. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta_1, \sigma^2)$ distribution and $\mathbf{Y} = (Y_1, \dots, Y_n)$ be a random sample from the $N(\theta_2, k\sigma^2)$ distribution, k known.

- (a) Assuming a non-informative prior for $(\theta_1, \theta_2, \sigma^2)$, obtain the posterior distribution of $\theta_1 - \theta_2$ and σ^2 .

- (b) Construct a $100(1 - \alpha)\%$ HPD interval for $\theta_1 - \theta_2$.
 (c) Obtain a $100(1 - \alpha)\%$ confidence interval for $\theta_1 - \theta_2$ from a frequentist perspective.

34. Let X_1, \dots, X_n be independent observations from the $N(\theta, \sigma^2/k_i)$, $i = 1, \dots, n$, respectively, where the k_i 's are known positive constants.

- (a) Obtain a family of natural conjugate distributions to the observational model above.

Hint: define $\bar{X} = \sum_{i=1}^n k_i X_i / \sum_{i=1}^n k_i$.

- (b) Construct a $100(1 - \alpha)\%$ HPD interval for θ .

- (c) Obtain a $100(1 - \alpha)\%$ confidence interval for θ from a frequentist perspective.

5

Approximate and computationally intensive methods

As we have seen in Chapter 4, the classical approach to statistics requires the sampling distribution of the estimators to be useful from both a practical and theoretical point of view. In the Bayesian approach, all the information needed is described by the posterior distribution. In both approaches, the evaluation of probabilities or expected values and optimization of some criterion function are often demanded. When the exact solution fails, evaluation of these quantities must involve analytical or computationally intensive approximation techniques. Typically, the accuracy of the analytic procedures depends critically on the sample size and the accuracy of simulation-based techniques depends on the number of simulations.

In this chapter the central problem of inference is stated in Section 5.1 and efficient numerical solutions are discussed. Many approximating techniques used in statistics are described in this chapter. Some optimization methods are presented in Section 5.2 and analytical techniques are discussed in Section 5.3. An introduction to numerical integration is presented in Section 5.4, and methods based on simulation are described in Section 5.5, including Monte Carlo, Monte Carlo with importance sampling, classical and Bayesian bootstrap and some ideas on Markov chain Monte Carlo techniques. A good account of some of the techniques presented in this chapter can be found in the books by Davison and Hinkley (1997), Gamerman (1997), Tanner (1996) and Thisted (1976).

5.1 The general problem of inference

In the development of statistical inference, as we have seen, we are often involved with the optimization of some criterion function or with the evaluation of integrals or expected values. In some cases these problems are not tractable analytically. The classical methods are often based on the maximization of an objective function, such as the likelihood function or the squared error loss. Based on decision theory, the minimization of the expected loss function provides Bayes estimators, where expectation is calculated with respect to the posterior distribution.

Generally speaking we are often faced with one of two basic problems:

1. maximization of a function $q(\theta)$ that can be either a posterior density or a

likelihood function;

2. evaluation of an integral

$$E[g(\theta) | \mathbf{x}] = \int g(\theta) p(\theta | \mathbf{x}) d\theta \quad \text{or} \quad E[g(\mathbf{X}) | \theta] = \int g(\mathbf{x}) p(\mathbf{x} | \theta) d\mathbf{x}$$

where \mathbf{x} represents the observed data, $g(\cdot)$ is an integrable function and θ is a p -dimensional vector.

The first problem is related to the definition of the (generalized) MLE and the second includes many alternatives, which are exemplified below:

1. One of the basic problems in Bayesian inference is to find the value of the normalizing constant k of the posterior density. It is obtained as $k^{-1} = \int l(\theta; \mathbf{x}) p(\theta) d\theta$;
2. If one wishes to ascertain the bias of an estimator $\hat{\theta} = g(\mathbf{X})$ of θ then one must evaluate the sampling expectation above.
3. In order to find the Bayes estimator with respect to the loss function $L(\delta, \theta)$, one must evaluate $\int g(\theta) p(\theta | \mathbf{x}) d\theta$ with $g(\cdot) = L(\delta, \cdot)$.
4. Evaluation of the confidence of a region C involves calculation of $P(G(\mathbf{X}, \theta) \in C) = \int g(\mathbf{x}) p(\mathbf{x} | \theta) d\mathbf{x}$ where $g(\mathbf{x}) = I_{\mathbf{x}}(\{\mathbf{x} : G(\mathbf{x}, \theta) \in C\})$.
5. The predictive density of a future sample \mathbf{Y} with density $p(\mathbf{y} | \theta)$, independent of \mathbf{X} conditionally on θ is given by $\int g(\theta) p(\theta | \mathbf{x}) d\theta$ where $g(\theta) = p(\mathbf{y} | \theta)$.

In the next section some useful optimization methods will be presented. Their importance is to yield classical and Bayesian point estimators and also confidence and HPD intervals.

5.2 Optimization techniques

The optimization of some criterion function is present in many theoretical developments of statistical theory, as seen in Chapter 4, from both the classical and the Bayesian perspectives. For example, if we wish to obtain least square estimators, generalized maximum likelihood estimators or the minimum expected loss then some sort of optimization will be required. As will be seen in Section 5.3, even to calculate the approximate value of some integrals via Laplace methods, maximum values of some functions are needed. Since in many relevant problems the optimum cannot be obtained analytically, some numerical optimization techniques will be reviewed in this section. The main goal in this section is to present and illustrate the use of Newton–Raphson techniques and the Fisher scoring methods to locate the maximum of the likelihood function or of the posterior distribution. Many statistical books include chapters on statistical computing with discussion of optimization techniques, as for example Garthwaite, Jolliffe and Jones (1995).

An algorithm to find the zeros of a twice differentiable function $\mathbf{g} : R^p \rightarrow R$ ($p \geq 1$) is easily obtained from the Taylor expansion of \mathbf{g} around an arbitrary point $\mathbf{x}^{(0)} \in R^p$:

$$\mathbf{g}(\mathbf{x}) = \mathbf{g}(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})' \frac{\partial \mathbf{g}(\mathbf{x}^{(0)})}{\partial \mathbf{x}} + \dots$$

Neglecting higher-order terms in $\mathbf{x} - \mathbf{x}^{(0)}$ for suitably close values of \mathbf{x} and $\mathbf{x}^{(0)}$ gives

$$\mathbf{g}(\mathbf{x}) \simeq \mathbf{g}(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})' \frac{\partial \mathbf{g}(\mathbf{x}^{(0)})}{\partial \mathbf{x}}.$$

If \mathbf{x}^* is a zero of \mathbf{g} then solving the above equation for \mathbf{x}^* gives $\mathbf{x}^* \simeq \mathbf{x}^{(1)} = \mathbf{x}^{(0)} - [\partial \mathbf{g}(\mathbf{x}^{(0)}) / \partial \mathbf{x}]^{-1} \mathbf{g}(\mathbf{x}^{(0)})$.

It follows that, starting with an initial value $\mathbf{x}^{(0)}$ and using the relation stated before, the algorithm provides us with a new value $\mathbf{x}^{(1)}$ closer to the root of the above equation. This new point is the intersection of the tangent line, the linear approximation of \mathbf{g} at \mathbf{x}_0 , with the \mathbf{x} axis. The procedure is then repeated with $\mathbf{x}^{(1)}$ replacing $\mathbf{x}^{(0)}$. This will lead to an even better approximation for \mathbf{x}^* denoted by $\mathbf{x}^{(2)}$. Repeating the process successively gives the recursive relation

$$\mathbf{x}^{(j)} = \mathbf{x}^{(j-1)} - \left[\frac{\partial \mathbf{g}(\mathbf{x}^{(j-1)})}{\partial \mathbf{x}} \right]^{-1} \mathbf{g}(\mathbf{x}^{(j-1)}).$$

The procedure is graphically illustrated in Figure 5.1.

This is the well-known Newton–Raphson algorithm and it must be repeated until some convergence criterion is achieved. Typical criteria are $|\mathbf{x}^{(j)} - \mathbf{x}^{(j-1)}| < \delta$ and $|\mathbf{g}(\mathbf{x}^{(j)})| < \epsilon$ where δ and ϵ are preset precisions determined arbitrarily. Since it is easier to evaluate proximity at the \mathbf{x} level rather than at the \mathbf{g} level, the former is sometimes preferred.

5.2.1 Solution of the likelihood equation

Let $\mathbf{U}(\mathbf{X}; \theta) = \partial \log p(\mathbf{X} | \theta) / \partial \theta$ be the score function. The MLE is the solution of the equation

$$\mathbf{U}(\mathbf{X}; \theta) = \mathbf{0}.$$

Remember that $\mathbf{J}(\theta) = -\partial \mathbf{U}(\mathbf{X}; \theta) / \partial \theta$ is the observed information matrix. Then, replacing relevant terms in the expression of the Newton–Raphson iteration gives

$$\boldsymbol{\theta}^{(j)} = \boldsymbol{\theta}^{(j-1)} + [\mathbf{J}(\boldsymbol{\theta}^{(j-1)})]^{-1} \mathbf{U}(\boldsymbol{\theta}^{(j-1)}).$$

There is a sense, to be made clearer in the next section, in which the observed information \mathbf{J} approaches the expected information \mathbf{I} . This idea can be introduced as a modification to the Newton–Raphson algorithm by replacement of the factor involving \mathbf{J} by another one involving \mathbf{I} . The j th step of the iteration becomes

$$\boldsymbol{\theta}^{(j)} = \boldsymbol{\theta}^{(j-1)} + [\mathbf{I}(\boldsymbol{\theta}^{(j-1)})]^{-1} \mathbf{U}(\boldsymbol{\theta}^{(j-1)}).$$

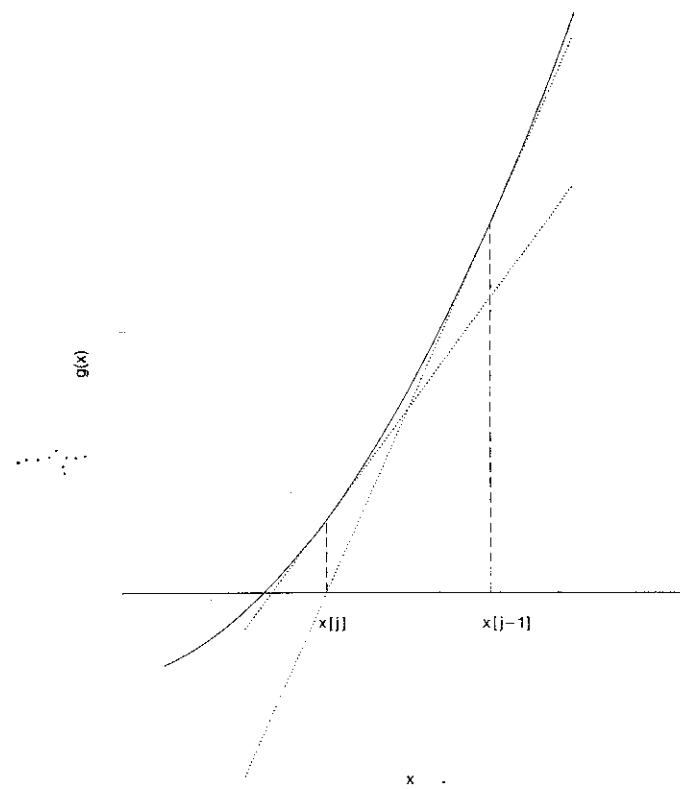


Fig. 5.1 Graphical representation of the iterative method for finding the roots of an equation in the scalar case.

This revised algorithm is usually known as Fisher scoring and enjoys many of the nice properties of the Newton–Raphson algorithm. Under mild regularity conditions found in many applications, it converges to the maximum likelihood estimator. Its strength lies in the fact that considerable reduction in computation is sometimes achieved by replacement of J by I with consequent elimination of a few terms.

Example. In this example an indirect use of the Fisher scoring algorithm is presented. First the profile likelihood defined in Chapter 2 is obtained, reducing the dimension of the parameter space, and then the algorithm is implemented. Let X_1, \dots, X_n be iid $\text{Wei}(\alpha, \beta)$ random variables (see list of distributions). The

log-likelihood function is given by

$$L(\alpha, \beta) = n \log \beta + n \log \alpha + (\alpha - 1) \sum_{i=1}^n \log X_i - \beta \sum_{i=1}^n X_i^\alpha.$$

Substituting β by its MLE $\hat{\beta}(\alpha) = n / \sum_{i=1}^n X_i^\alpha$, the profile log-likelihood function can be written as

$$L(\alpha, \hat{\beta}(\alpha)) = n \log n - n \log \left(\sum_{i=1}^n X_i^\alpha \right) + (\alpha - 1) \sum_{i=1}^n \log X_i + n \log \alpha - n.$$

Differentiating with respect to α and making some algebraic simplifications gives the score function and the observed information matrix as

$$U(\mathbf{X}; \alpha) = \frac{n}{\alpha} - \frac{n\alpha \sum_{i=1}^n X_i^\alpha \log X_i}{\sum_{i=1}^n X_i^\alpha} + \sum_{i=1}^n \log X_i$$

$$J(\alpha) = \frac{n}{\alpha^2} + n \frac{\sum_{i=1}^n X_i^\alpha \sum_{i=1}^n X_i^\alpha (\log X_i)^2 - (\sum_{i=1}^n X_i^\alpha \log X_i)^2}{(\sum_{i=1}^n X_i^\alpha)^2}.$$

Applying the Newton–Raphson with some initial value $\alpha^{(0)}$ gives at the j th step of the iteration the equation

$$\alpha^{(j)} = \alpha^{(j-1)} + J^{-1}(\alpha^{(j-1)}) U(\mathbf{X}; \alpha^{(j-1)}).$$

As soon as convergence in α is reached, β can be estimated using the relationship between α and its conditional MLE given by

$$\hat{\beta}(\alpha) = \frac{n}{\sum_{i=1}^n X_i^\alpha}.$$

Using $n = 25$ observations artificially generated from a Weibull distribution with $\alpha = 1.5$ and $\theta = 2.0$, we obtained via Newton–Raphson the following MLE: $\hat{\alpha} = 1.59$ and $\hat{\theta} = 1.97$. Confidence intervals can be obtained from the asymptotic distribution of the MLE. As will be seen in the next section, this requires the evaluation of the Fisher information. Fortunately, this can be obtained at no extra cost since the Fisher information must be evaluated in the algorithm. Approximate 95% confidence intervals for θ and α are respectively given in this case by (1.80, 2.15) and (1.42, 1.76). Note that they include the true values used in the generation, as expected.

5.2.2 Determining confidence intervals

Recall from Chapter 4 that the HPD interval in the uniparameter case is the collection of all $\theta \in \Theta$ such that $p(\theta|x) > k$, where k is related to γ , the confidence probability, through the equation $P[\theta_l < \theta < \theta_u|x] = \gamma$. Note that

we are assuming here that the confidence region is in fact an interval. Defining $g(\theta) = \log(p(\theta|x)) - \log k$, applying the Newton-Raphson algorithm or the Fisher scoring algorithm twice but with different starting values $\hat{\theta} - \epsilon$ and $\hat{\theta} + \epsilon$, where $\hat{\theta}$ is the posterior mode and $\epsilon > 0$, the convergence of the method is assured as long as starting values are not far from the mode.

Many difficulties can occur with numerical optimization methods. The solution may not be stable when numerical derivatives are used; the choice of a *bad* initial value can lead the solution to a local (instead of global) optimum and so forth. Nevertheless, many solutions to these problems have been proposed in the literature. There are methods that avoid the numerical evaluation of the second derivative, others that allow for different randomly selected initial values and so forth. The interested reader is referred to Thisted (1976) for further discussion.

5.2.3 The EM algorithm

This is a general iterative method useful to obtain the (generalized) MLE when we are faced with an incomplete data set or when it is simpler to maximize the likelihood of a more general problem involving unobserved quantities. Dempster, Laird and Rubin (1977) provided the first unified account of the algorithm. A discussion in the context of data imputation is presented by Little and Rubin (1988). The algorithm is iterative and at each iteration it alternates the operations of expectation (**E**) and maximization (**M**).

Let $X \in R^n$ be the n -dimensional vector of observed quantities and $Z \in R^m$ an m -dimensional vector of unobserved quantities. The complete data is denoted by $Y = (X, Z)' \in R^{n+m}$ and its density function is $p(y|\theta) = p(x, z|\theta)$, $\theta \in \Theta$. On the other hand, let the conditional density of the unobserved data given the observed one be $p(z|x, \theta)$, which also depends on θ . To obtain the MLE of θ the logarithm of the marginal likelihood,

$$L(\theta; x) = \log \left(\int p(x, z; \theta) dz \right)$$

is usually directly maximized. To avoid the highly dimensional integral involved in the marginalization of $p(x, z|\theta)$, the following relationship can be used

$$L(\theta; x) = \log \left(\frac{p(x, z|\theta)}{p(z|x, \theta)} \right) = \log p(x, z|\theta) - \log p(z|x, \theta).$$

Since Z is unobserved, it is necessary to eliminate it before maximizing $L(\theta; x)$. One way to do that is to take expected values with respect to the conditional density, $p(z|x, \theta)$ in the above equation. Noting that $E_{Z|X,\theta}[L(\theta; x)] = L(\theta; x)$, gives

$$L(\theta; x) = Q(\theta; \theta^{(0)}) - H(\theta; \theta^{(0)})$$

where

$$\begin{aligned} Q(\theta; \theta^{(0)}) &= E_{Z|X,\theta^{(0)}}[\log p(X, Z|\theta)] \\ H(\theta; \theta^{(0)}) &= E_{Z|X,\theta^{(0)}}[\log p(Z|X, \theta)] \end{aligned}$$

and $\theta^{(0)}$ is any given starting value for θ . The expectation involved in the definition of Q is based on the likelihood of the complete data set Y . This is usually straightforward because the data augmentation is performed to simplify the problem, as will be shown in the examples below.

If $\theta^{(j)}$ denotes the value of θ in the j th iteration then the EM algorithm is defined through the following two steps:

1. **E** (expectation): evaluation of $Q(\theta, \theta^{(j-1)}) = E_{Z|X,\theta^{(j-1)}}[\log l(\theta; Y)]$;
2. **M** (maximization): evaluation of $\theta^{(j)}$, the value of θ that maximizes $Q(\theta, \theta^{(j-1)})$.

The estimation procedure is iterative and alternates the **E** and **M** operations at each iteration. Dempster, Laird and Rubin (1977) showed that the sequence $\theta^{(j)}$, $j \geq 1$, generated by the EM algorithm satisfies $L(\theta^{(j)}|x) \leq L(\theta^{(j+1)}|x)$ and is monotonically increasing in the likelihood $l(\theta|x)$. Therefore, the sequence $\theta^{(j)}$ converges to $\hat{\theta}$, the MLE, if the likelihood function has only a single maximum. Convergence can be established by criteria such as $|\theta^{(j)} - \theta^{(j-1)}| < \delta$ or $|Q(\theta^{(j)}, \theta^{(j-1)}) - Q(\theta^{(j-1)}, \theta^{(j-1)})| < \epsilon$. The convergence can be slow in some cases specially if the missing information of Z is substantial. The adaptation for posterior mode evaluation involves replacement of the likelihood $l(\theta; Y)$ by the posterior $p(\theta|Y)$ in the **E** step.

Example (Rao, 1973). A classical application in the statistical literature refers to a genetic study stating that the four-dimensional vector of animal counts $X = (X_1, X_2, X_3, X_4)$ has multinomial distribution with parameters n and π , where $\pi = (1/2 + \theta/4, (1 - \theta)/4, (1 - \theta)/4, \theta/4)$. So the probability function of X is

$$p(x|\theta) = \frac{(x_1 + x_2 + x_3 + x_4)!}{x_1!x_2!x_3!x_4!} \left(\frac{1}{2} + \frac{\theta}{4} \right)^{x_1} \left(\frac{1 - \theta}{4} \right)^{x_2+x_3} \left(\frac{\theta}{4} \right)^{x_4}.$$

Direct maximization of the above expression is awkward due to the presence of the term $1/2 + \theta/4$. To avoid it, the EM method described above will be applied.

To do that, let $X_1 = Y_0 + Y_1$ and $X_i = Y_i$, $i \geq 2$, where the augmented vector $Y = (Y_0, Y_1, Y_2, Y_3, Y_4)$ has multinomial distribution with parameters n and $\pi^* = (1/2, \theta/4, (1 - \theta)/4, (1 - \theta)/4, \theta/4)$. To complete the notation, define $Z = Y_0$, so that $Y = (X, Z)$. Therefore,

$$p(y|\theta) = \frac{n!}{y_0!y_1!y_2!y_3!y_4!} \left(\frac{1}{2} \right)^{y_0} \left(\frac{\theta}{4} \right)^{y_1} \left(\frac{1 - \theta}{4} \right)^{y_2+y_3} \left(\frac{\theta}{4} \right)^{y_4}$$

and

$$\begin{aligned} \log p(y|\theta) &= k_1(y) + y_0 \log \left(\frac{1}{2} \right) \\ &\quad + (y_1 + y_4) \log \left(\frac{\theta}{4} \right) + (y_2 + y_3) \log \left(\frac{1 - \theta}{4} \right) \\ &= k_2(y) + (y_1 + y_4) \log \theta + (y_2 + y_3) \log(1 - \theta) \end{aligned}$$

where $k_1(y)$ and $k_2(y)$ are functions of y but not of θ . Therefore

$$\begin{aligned} Q(\theta, \theta^{(j)}) &= E[k_2(Y) + (Y_1 + Y_4) \log \theta + (Y_2 + Y_3) \log(1 - \theta) | X, \theta^{(j)}] \\ &= k(X, \theta^{(j)}) + E(Y_1 + Y_4 | X, \theta^{(j)}) \log \theta \\ &\quad + E(Y_2 + Y_3 | X, \theta^{(j)}) \log(1 - \theta) \\ &= k(X, \theta^{(j)}) + [E(Y_1 | X, \theta^{(j)}) + X_4] \log \theta \\ &\quad + (X_2 + X_3) \log(1 - \theta) \end{aligned}$$

since $Y_i = X_i$, $i = 2, 3, 4$. We only need to evaluate the expectation of Y_1 since $k(X, \theta^{(j)})$ does not depend on θ and will therefore be irrelevant in the M step. From the construction of Y it follows that $(Z|X, \theta) \sim (Z|X_1, \theta) \sim \text{bin}(X_1, p)$ where $p = (1/2)/[(1/2) + (\theta/4)] = 2/(2 + \theta)$. Therefore, $E(Y_1 | X, \theta) = X_1 p$ and

$$Q(\theta, \theta^{(j)}) = E[k_2(Y) | X, \theta^{(j)}] + (X_1 p^{(j)} + X_4) \log \theta + (X_2 + X_3) \log(1 - \theta)$$

where $p^{(j)} = 2/(2 + \theta^{(j)})$. The M step involves finding the value of θ that maximizes $Q(\theta, \theta^{(j)})$. This is easily obtained by differentiation of Q and gives

$$\begin{aligned} \theta^{(j+1)} &= \frac{X_1 p^{(j)} + X_4}{X_1 p^{(j)} + X_2 + X_3 + X_4} \\ &= \frac{(X_1 + X_4)\theta^{(j)} + 2X_4}{(X_1 + X_2 + X_3 + X_4)\theta^{(j)} + 2(X_2 + X_3 + X_4)}. \end{aligned}$$

To illustrate the results, assume it was observed that the counts were $\mathbf{x} = (125, 18, 20, 34)'$ and the EM algorithm was started at $\theta^{(0)} = 0.5$. Then,

$$\theta^{(j+1)} = \frac{159\theta^{(j)} + 68}{197\theta^{(j)} + 144}.$$

The first five iterations of the algorithm give the values 0.608, 0.624, 0.626, 0.627, 0.627.

Example. Randomized response. The proportion θ of individuals belonging to certain stigmatized category must be estimated. To avoid the non-response (and its consequent loss of information), a new sampling scheme is proposed. An alternative question, not related to the main one, with known proportion of YES responses is introduced together with the guarantee that the selected question will be known exclusively by the respondent. The idea is to increase his/her confidence in providing the correct response without revealing its true status. The probability of a YES response will be $\lambda(\theta) = \pi\theta + (1 - \pi)\theta_A$, where θ is the probability of the original question of interest, θ_A is the known probability of a YES answer to the alternative question and π is the probability of selection of the question of interest. In a sample of 150 individuals, 60 YES responses were obtained, based on a procedure with $\pi = 0.7$ and $\theta_A = 0.6$.

Using the EM algorithm we get that the observed data is X , the number of YES responses, and $X|\theta \sim \text{bin}(n, \lambda)$. Also, the unobserved data is Z , the number of individuals that will select the question of interest. Then, $Z|X, \theta \sim \text{bin}(X, p)$, where $p = \pi\theta/\lambda$. The joint density of the observed and unobserved data is given by

$$\begin{aligned} p(x, z | \theta) &= p(z|x, \theta)p(x|\theta) \\ &= \binom{x}{z} \left(\frac{\pi\theta}{\lambda}\right)^z \left(1 - \frac{\pi\theta}{\lambda}\right)^{x-z} \binom{n}{x} \lambda^x (1 - \lambda)^{n-x} \\ &= \frac{n!}{z!(x-z)!(n-x)!} (\pi\theta)^z [(1 - \pi)\theta_A]^{x-z} [1 - \lambda]^{n-x} \end{aligned}$$

$$\log p(x, z | \theta) = k(x, z) + z \log \theta + (n - x) \log(1 - \lambda).$$

Then, the j th iteration of the EM algorithm will have:

1. E step: $Q(\theta, \theta^{(j)}) = E[\log p(X, Z | \theta) | X, \theta^{(j)}] = Xp^{(j)} \log \theta + (n - X) \log(1 - \lambda) + k(X, \theta^{(j)})$.
2. M step: maximization of Q involves finding the solution of

$$\frac{\partial Q(\theta, \theta^{(j)})}{\partial \theta} = 0$$

but

$$\frac{\partial Q(\theta, \theta^{(j)})}{\partial \theta} = \frac{Xp^{(j)}}{\theta} - \frac{n - X}{1 - \lambda}\pi.$$

Solving for θ gives

$$\theta^{(j+1)} = \frac{Xp^{(j)}[1 - (1 - \pi)\theta_A]}{Xp^{(j)} + (n - X)\pi}.$$

In the condition of the example with a sample of 150 individuals, initializing the procedure with $\theta^{(0)} = 0.4$, provides the sequence of estimates 0.338, 0.322, 0.317, 0.315, 0.314, 0.314, ...

5.3 Analytical approximations

Some analytical methods used in statistical inference will be presented in this section. Firstly, results based on asymptotic theory will be presented, followed by the Kullback–Liebler approximation and a technique for numerical integration called the Laplace method.

5.3.1 Asymptotic theory

The behaviour of statistical inference will be studied in this section from the standpoint of an infinite sample size. Clearly, as in practice n is never infinite, the

results presented here can be applied only when n can be thought of as sufficiently large to ensure that the results stated are approximately valid.

The results presented here provide a method to obtain approximate solutions to problems for which the exact solution is not feasible, which is their main utility. There are other ways to obtain approximate solutions but almost all of them are variations on the methods described in this section. The main exception is the class of methods based on the Kullback–Leibler divergence measure which will be developed in the following subsection.

In general, when the sample size increases, its influence on the inference also increases, minimizing the importance of the chosen prior distribution. However, this will only be true if the prior distribution is non-degenerate, that is, $p(\theta) > 0$, $\forall \theta \in \Theta$, as we have seen in Section 3.2. As n increases, the posterior distribution will be more and more concentrated around its mode. Then it follows that

$$\begin{aligned} p(\theta | \mathbf{x}) &\propto l(\theta; \mathbf{x}) p(\theta) \\ &\propto \exp\{L(\theta) + \log p(\theta)\} \quad \text{where } L(\theta) = \log l(\theta; \mathbf{x}). \end{aligned}$$

Since $L(\theta) = \sum_{i=1}^n \log l(\theta; x_i)$, the number of terms in $L(\theta)$ will increase with n , but $p(\theta)$ is fixed and, consequently, the influence of the prior $p(\theta)$ becomes less and less relevant. Then,

$$p(\theta | \mathbf{x}) \approx \exp\{L(\theta)\}.$$

Therefore, the likelihood and the posterior will be approximately equal. Using the Taylor expansion of L around $\hat{\theta}$, the MLE of θ , it follows that

$$L(\theta) = L(\hat{\theta}) + (\theta - \hat{\theta})' \frac{\partial L(\hat{\theta})}{\partial \theta} + \frac{1}{2!} (\theta - \hat{\theta})' \frac{\partial^2 L(\hat{\theta})}{\partial \theta \partial \theta'} (\theta - \hat{\theta}) + R(\theta, \hat{\theta})$$

where $R(\theta, \hat{\theta})$ contains terms of higher order, which can be eliminated when θ is supposed to be close enough to $\hat{\theta}$. As $\partial L(\hat{\theta})/\partial \theta = 0$ it follows that

$$L(\theta) = k - (\theta - \hat{\theta})' \frac{\mathbf{J}(\hat{\theta})}{2} (\theta - \hat{\theta}) \quad \text{where } \mathbf{J}(\theta) = -\frac{\partial^2 L(\theta)}{\partial \theta \partial \theta'}.$$

Therefore, $p(\theta | \mathbf{x}) \propto \exp\left\{-(\theta - \hat{\theta})' \mathbf{J}(\hat{\theta})(\theta - \hat{\theta})/2\right\}$ and so

$$\mathbf{J}^{1/2}(\hat{\theta})(\theta - \hat{\theta}) | \mathbf{x} \sim N(\mathbf{0}, \mathbf{I}_p).$$

Then, if $\mathbf{X}_n = (X_1, \dots, X_n)$ is a random sample from $p(x | \theta)$ and $\hat{\theta}$ is the MLE of θ , under certain mild regularity conditions, it follows that

$$\mathbf{J}^{1/2}(\hat{\theta})(\theta - \hat{\theta}) | \mathbf{x}_n \xrightarrow{\mathcal{D}} N(\mathbf{0}, \mathbf{I}_p) \quad \text{when } n \rightarrow \infty.$$

The result simplifies when θ is a scalar. In this case,

$$\frac{\theta - \hat{\theta}}{\mathbf{J}^{-1/2}(\hat{\theta})} | \mathbf{x}_n \xrightarrow{\mathcal{D}} N(0, 1) \quad \text{when } n \rightarrow \infty$$

and $J(\theta)$ is given by $-\partial^2 L(\theta)/\partial \theta^2$.

The regularity conditions required are basically the same as those involved in the statement of the Cramer–Rao inequality (see Section 4.3). The results can be extended further if the sample is not made up of independent observations. Observe that $\mathbf{J}(\theta) = \sum_{i=1}^n \partial^2 \log f(x_i | \theta)/\partial \theta^2$ and so $J(\theta)$ will typically increase with n .

This result indicates that $\theta | \mathbf{x}_n$ has an approximately normal distribution, $N(\hat{\theta}, \mathbf{J}^{-1}(\hat{\theta}))$, when n is large enough. This means that as long as the regularity conditions are satisfied, it is possible to draw approximate inferences about the parameters. In particular, it is possible to construct approximate credibility regions based on the above results.

Definition. Let θ be an unknown quantity defined in Θ . A region $C \subset \Theta$ is an asymptotic $100(1 - \alpha)\%$ credibility or Bayesian confidence region for θ if $\lim_{n \rightarrow \infty} Pr(\theta \in C | \mathbf{x}_n) \geq 1 - \alpha$. In this case, $1 - \alpha$ is called the asymptotic credibility or confidence level. In the scalar case, the region C is usually given by an interval, $[c_1, c_2]$ say.

More accurate approximations are obtained by retaining the third-order term in the Taylor expansion. Then,

$$p(\theta | \mathbf{x}_n) \simeq p(\hat{\theta} | \mathbf{x}_n) \exp\left[-\frac{1}{2}(\theta - \hat{\theta})' \mathbf{J}(\hat{\theta})(\theta - \hat{\theta})\right] \left[1 + \frac{t(\theta)}{3!}\right]$$

where

$$t(\theta) = \sum_{i,j,k} \frac{\partial^3 L(\hat{\theta})}{\partial \theta_i \partial \theta_j \partial \theta_k} (\theta_i - \hat{\theta}_i)(\theta_j - \hat{\theta}_j)(\theta_k - \hat{\theta}_k).$$

These approximations were studied by Lindley (1980).

Example. The expressions stated above applied for the posterior mean of a scalar θ simplify to

$$E(\theta | \mathbf{x}_n) \simeq \hat{\theta} + \int (\theta - \hat{\theta})^4 \frac{1}{6} \frac{\partial^3 L(\hat{\theta})}{\partial \theta^3} p_N(\theta | \hat{\theta}, \mathbf{J}^{-1}(\hat{\theta})) d\theta$$

where $p_N(\cdot | a, b)$ denotes the density of the $N(a, b)$ distribution. Then

$$E(\theta | \mathbf{x}_n) \simeq \hat{\theta} + \frac{1}{6} \frac{\partial^3 L(\hat{\theta})}{\partial \theta^3} E_N(\theta - \hat{\theta})^4 = \hat{\theta} + \frac{1}{2} \frac{\partial^3 L(\hat{\theta})}{\partial \theta^3} J^{-2}(\hat{\theta}).$$

Example. Let X_1, \dots, X_n be a random sample from $\text{Pois}(\theta)$ and suppose that the prior for θ is a $G(a_0, b_0)$ distribution. So the posterior of θ will be $G(a_1, b_1)$, with $a_1 = a_0 + t$ and $b_1 = b_0 + n$, and $t = \sum_{i=1}^n x_i$.

It is easy to verify that the posterior mode is $\hat{\theta} = (a_1 - 1)/b_1$,

$$\begin{aligned} J(\hat{\theta}) &= \frac{\partial^2 \log p(\hat{\theta} | \mathbf{x}_n)}{\partial \theta^2} = \frac{a_1 - 1}{\hat{\theta}} \\ t(\theta) &= \frac{\partial^3 \log p(\hat{\theta} | \mathbf{x}_n)}{\partial \theta^3} (\theta - \hat{\theta})^3 = \frac{2(a_1 - 1)}{\hat{\theta}^3} (\theta - \hat{\theta})^3. \end{aligned}$$

Then the approximations of second and third order will be

$$\theta|\mathbf{x}_n \sim N\left(\hat{\theta}, \frac{\hat{\theta}^2}{a_1 - 1}\right) \text{ or } \theta|\mathbf{x}_n \sim N\left(\hat{\theta}, \frac{\hat{\theta}^2}{a_1 - 1}\right) \left\{ 1 + \frac{1}{3} \frac{a_1 - 1}{\hat{\theta}^3} (\theta - \hat{\theta})^3 \right\}$$

where the last equation means that the distribution is proportional to the product of a normal distribution and the correction term in brackets. Using the result from the previous example, it follows that $E(\theta|\mathbf{x}_n) \simeq \hat{\theta}$ and $E(\theta|\mathbf{x}_n) \simeq \hat{\theta} + 1/b_1$.

Note that for the validity of the results, $p(\theta)$ must be strictly positive and continuous in Θ . If $p(\theta)$ is proportional to a constant, it follows that the posterior coincides with the likelihood. Also, in order that $t(\theta)$ exist, the posterior must be three times differentiable.

It has just been shown that the distribution of θ converges to the normal distribution when n increases. What can be stated about reparametrizations, that is, transformations of θ ? The answer is not easy because since θ is normal we know that the only transformations that preserves normality are the linear transformations of θ . Since all the above developments do not depend on any special property associated with θ , any transformation of θ preserving the regularity conditions will produce the same results. So, these results are valid for any transformation of θ . The only difference will be on the speed of the convergence to the normal distribution.

The Taylor series expansion applied directly to some transformation of the parameter is another approximation that leads to the same result. Suppose that $E(\mathbf{X}) = \mathbf{a}$, $V(\mathbf{X}) = \mathbf{A}$ and \mathbf{g} is a 1-to-1 transformation of \mathbf{X} with derivatives well defined in the point \mathbf{a} . Then

$$Y = \mathbf{g}(\mathbf{X}) = \mathbf{g}(\mathbf{a}) + (\mathbf{X} - \mathbf{a})' \frac{\partial \mathbf{g}(\mathbf{a})}{\partial \mathbf{X}} + o(\mathbf{X} - \mathbf{a})$$

where $|o(u)|/|u| \rightarrow 0$ when $u \rightarrow 0$. If \mathbf{X} is close to \mathbf{a} then the last term in the right-hand side can be omitted and Y will have an approximately linear relationship with \mathbf{X} where

$$E(Y) \doteq \mathbf{g}(\mathbf{a}) \quad \text{and} \quad V(Y) \doteq \left(\frac{\partial \mathbf{g}(\mathbf{a})}{\partial \mathbf{X}} \right)' \mathbf{A} \frac{\partial \mathbf{g}(\mathbf{a})}{\partial \mathbf{X}}.$$

Also, if \mathbf{X} is normally distributed, so will Y be. This result is commonly known as the delta method.

An interesting question concerns the choice of the optimal reparametrization in the sense of inducing fast convergence. This theme is still under investigation and some preliminary empirical results seem to indicate that the strongest candidates are the parameter transformations that appear in the definition of the exponential family and the transformations leading to constant non-informative priors. It is worth pointing out that the mean of normally distributed data, with known variance, which has a posterior normal distribution for any value of n , belongs to both groups

of transformations. Note also that the first class of transformations encompasses the class of parameters with constant Fisher information and therefore is, in some sense, stable.

Example. Let $X \sim \text{bin}(n, \theta)$ with unknown θ . It then follows that

$$L(\theta) = \log p(x | \theta) = \log \binom{n}{x} + x \log \theta + (n-x) \log(1-\theta)$$

and its derivatives are given by

$$\begin{aligned} \frac{\partial L(\theta)}{\partial \theta} &= \frac{x}{\theta} - \frac{n-x}{1-\theta} \Rightarrow \hat{\theta} = \frac{x}{n} \\ \frac{\partial^2 L(\theta)}{\partial \theta^2} &= -\frac{x}{\theta^2} - \frac{n-x}{(1-\theta)^2} < 0 \end{aligned}$$

and

$$\begin{aligned} J(\hat{\theta}) &= \frac{n\hat{\theta}}{\hat{\theta}^2} + \frac{n-n\hat{\theta}}{(1-\hat{\theta})^2} \\ &= \frac{n}{\hat{\theta}} + \frac{n}{1-\hat{\theta}} = \frac{n}{\hat{\theta}(1-\hat{\theta})}. \end{aligned}$$

So, for large enough n , θ has a posterior distribution approximately $N[\hat{\theta}, \hat{\theta}(1-\hat{\theta})/n]$.

The non-informative prior for this model is given by $p(\theta) \propto \theta^{-1/2}(1-\theta)^{-1/2}$, as seen in Section 3.4, and the transformation producing a constant non-informative prior is

$$\phi \propto \int_0^\theta u^{-1/2}(1-u)^{-1/2} du \propto \text{sen}^{-1}(\sqrt{\theta}).$$

Defining $\phi = \text{sen}^{-1}(\sqrt{\theta})$ it follows that

$$\frac{\partial \phi}{\partial \theta} = \frac{\theta^{-1/2}(1-\theta)^{-1/2}}{2} \quad \text{and} \quad J^{-1}(\hat{\phi}) = \frac{\hat{\theta}(1-\hat{\theta})}{n} \frac{1}{4\hat{\theta}(1-\hat{\theta})} = \frac{1}{4n}.$$

So, the posterior distribution of ϕ for large enough n is approximately $N(\hat{\phi}, 1/4n)$.

The parameter obtained in the definition of the exponential family is $\psi = \log[\theta/(1-\theta)]$. Then,

$$\hat{\psi} = \log\left(\frac{\hat{\theta}}{1-\hat{\theta}}\right) \quad \text{and} \quad \left| \frac{\partial \psi}{\partial \theta} \right| = \frac{1}{\theta(1-\theta)}$$

and so the posterior distribution of ψ is approximately $N(\hat{\psi}, 1/n\hat{\theta}(1-\hat{\theta}))$.

In classical inference, similar calculations can be performed with the asymptotic distribution of the maximum likelihood estimator. Let $\mathbf{X}_n = (X_1, \dots, X_n)$ be a

random sample from $p(x | \theta)$ for a scalar θ and $\hat{\theta}_n$ be the MLE of θ obtained from \mathbf{X}_n . Suppose that the Cramer–Rao regularity conditions are true and also that $|\partial^2 U(\mathbf{X}_n; \theta)/\partial\theta^2| < k$. Under these conditions, it follows that

$$\begin{aligned} \sum_{i=1}^n U(X_i; \hat{\theta}_n) &= \sum_{i=1}^n U(X_i; \theta) \\ &\quad + (\hat{\theta}_n - \theta) \sum_{i=1}^n \frac{\partial U(X_i; \theta)}{\partial\theta} + \frac{(\hat{\theta}_n - \theta)^2}{2} \sum_{i=1}^n \frac{\partial^2 U(X_i; \xi)}{\partial\theta^2} \end{aligned}$$

for ξ between θ and $\hat{\theta}_n$. By the definition of $\hat{\theta}_n$, the term in the left-hand side is null. Isolating the term $(\hat{\theta}_n - \theta)$ it follows that

$$\sqrt{n}(\hat{\theta}_n - \theta) = -\frac{\sqrt{n}/n [\sum_{i=1}^n U(X_i; \theta)]}{(1/n) \left[\sum_{i=1}^n \partial U(X_i; \theta)/\partial\theta + 0.5(\hat{\theta}_n - \theta) \sum_{i=1}^n \partial^2 U(X_i; \xi)/\partial\theta^2 \right]}.$$

Since $E[U(X_i; \theta) | \theta] = 0$ and $V[U(X_i; \theta) | \theta] = I(\theta)$ it follows from the central limit theorem that the numerator converges in distribution to a $N(0, I(\theta))$. The first term of the denominator converges almost surely to $E[\partial U(X_i; \theta)/\partial\theta | \theta] = -I(\theta)$. The second term converges almost surely to zero by the (strong) consistency of the maximum likelihood estimator and by bounds imposed on the second derivatives of $U(\mathbf{X}; \theta)$, by hypothesis.

So, the denominator converges to $-I(\theta)$ with probability one and the fraction converges to the quotient of the limits since if Z_n converges in distribution to Z and W_n converges almost surely to a constant w then Z_n/W_n converges in distribution to Z/w . So,

$$\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{\mathcal{D}} N(0, I^{-1}(\theta)) \quad \text{when } n \rightarrow \infty.$$

The result is equivalent to the Bayesian asymptotic result. It is usually said that for large n the distribution of $\hat{\theta}_n$ is approximately $N(\theta, I^{-1}(\theta)/n)$. It is worth pointing out that the information measure considered is based on a single observation, $p(x | \theta)$.

The above result is also true for multivariate parameters θ . In this case, the asymptotic result is

$$\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{\mathcal{D}} N(\mathbf{0}, \mathbf{I}^{-1}(\theta)) \quad \text{when } n \rightarrow \infty.$$

The result is used to say that when n is large the distribution of $\hat{\theta}_n$ is approximately $N(\theta, \mathbf{I}^{-1}(\theta)/n)$. Approximate classical inference can be made on the basis of those results. Once again, asymptotic confidence regions can be constructed.

Definition. Let θ be an unknown quantity defined in Θ , \mathbf{U} be a function $\mathbf{U} = G(\mathbf{X}_n, \theta)$ with values in \mathcal{U} and A be a region in \mathcal{U} such that $\lim_{n \rightarrow \infty} Pr(\mathbf{U} \in$

$A) \geq 1 - \alpha$. A region $C \subset \Theta$ is an asymptotic $100(1 - \alpha)\%$ frequentist confidence region for θ if

$$C = \{\theta : G(\mathbf{x}_n, \theta) \in A\}.$$

In this case, $1 - \alpha$ is called the asymptotic confidence level. In the scalar case, the inversion in terms of θ usually leads to an interval, $C = [c_1, c_2]$ say.

It is easy to obtain from the asymptotic results above that $E[\hat{\theta}_n | \theta] \rightarrow \theta$, or simply that the MLE is asymptotically unbiased. This means that although the maximum likelihood estimator could be biased for a fixed n , its expected value always tends to the parameter being estimated. Besides, it follows that

$$\lim_{n \rightarrow \infty} nV[\hat{\theta}_n | \theta] \rightarrow \mathbf{I}^{-1}(\theta).$$

So, the variance of the MLE asymptotically reaches the Cramer–Rao lower bound. As $\hat{\theta}_n$ is asymptotically unbiased, the above limit could be thought of as a measure of the asymptotic efficiency of $\hat{\theta}_n$. Estimators satisfying this property are said to be asymptotically efficient.

Sometimes, $\mathbf{I}(\theta)$ depends on θ making the process of inference (point estimation and confidence intervals) harder. It is common in these cases to substitute $\mathbf{I}(\theta)$ by $\mathbf{I}(\hat{\theta}_n)$ or, in cases where the expected value is hard to obtain, by $J(\hat{\theta}_n)$. In the former case the asymptotic distributions coincide and certainly lead to the same numerical results. The consistency of the maximum likelihood estimator and the strong law of large numbers justify the substitutions made to obtain both of the above results.

The convergence in distribution of the score function can be used to make approximate inference about θ . This is particularly useful when there is no explicit form for the maximum likelihood estimator. Then, for example, the asymptotic $100(1 - \alpha)\%$ confidence region for a scalar θ can be constructed using the $z_{\alpha/2}$ percentile of the normal distribution and is given by

$$\left\{ \theta : \left| \frac{1}{\sqrt{n}I(\theta)} \sum_{i=1}^n U(x_i; \theta) \right| < z_{\alpha/2} \right\}.$$

This result can also be extended to the case of parametric vectors.

The same considerations made in the Bayesian case with respect to the invariance over 1-to-1 parametric transformations are again true here since the results are entirely based on the maximum likelihood estimator and on the Fisher information measures, which satisfy the invariance conditions.

Example (continued). If $X \sim \text{bin}(n, \theta)$, then

$$U(X; \theta) = \frac{X}{\theta} - \frac{n - X}{1 - \theta} = \frac{X - n\theta}{\theta(1 - \theta)} \quad \text{and} \quad \hat{\theta}_n = \frac{X}{n}.$$

It is also known that $I(\theta) = 1/\theta(1 - \theta)$. Applying the results developed above gives $\sqrt{n}(\hat{\theta}_n - \theta)$ distributed as a $N(0, \theta(1 - \theta))$, for a large n . Then an asymptotic

confidence region with coverage probability of $100(1 - \alpha)\%$ for θ will be given by

$$\left\{ \theta : \left| \frac{\sqrt{n}(\hat{\theta}_n - \theta)}{\sqrt{\theta(1-\theta)}} \right| < z_{\alpha/2} \right\}.$$

A simple inequality must be solved to get an explicit solution for the confidence region. In this case it could be more convenient to proceed with the suggested modifications of θ by its MLE. In this case, observed and expected information coincide at the MLE and

$$J(\hat{\theta}_n) = \frac{1}{\hat{\theta}_n(1-\hat{\theta}_n)} = I(\hat{\theta}_n).$$

The above confidence region is now replaced by the interval

$$\left(\hat{\theta}_n - z_{\alpha/2} \sqrt{\frac{\hat{\theta}_n(1-\hat{\theta}_n)}{n}}, \hat{\theta}_n + z_{\alpha/2} \sqrt{\frac{\hat{\theta}_n(1-\hat{\theta}_n)}{n}} \right).$$

An alternative can be to use the asymptotic distribution of $\hat{U}(X; \theta)/\sqrt{n}$ given by a $N(0, 1/\theta(1-\theta))$. It is not hard to see that the $100(1-\alpha)\%$ asymptotic confidence region for θ is again given by

$$\left\{ \theta : \left| \frac{\sqrt{n}(\hat{\theta}_n - \theta)}{\sqrt{\theta(1-\theta)}} \right| < z_{\alpha/2} \right\}.$$

In this case, the confidence region coincides with the region obtained from the asymptotic distribution of $\hat{\theta}_n$.

5.3.2 Approximation by Kullback–Liebler divergence

A measure of the divergence between the density $p(\theta)$ and its approximation $p_0(\theta)$ is defined by the expected value (see Bernardo and Smith, 1994)

$$\delta[p(\theta); p_0(\theta)] = \int p(\theta) \log \left(\frac{p(\theta)}{p_0(\theta)} \right) d\theta = E \left[\log \left(\frac{p(\theta)}{p_0(\theta)} \right) \right].$$

In the discrete case, all one needs to do is to substitute the integration with a summation. Smaller values for δ indicate better approximations. Two central questions in the applications are concerned with the determination of the better normal approximation and the most convenient reparametrization to accommodate such approximation.

These two questions will be answered only for particular cases in the exponential family. The first question posed has an easy and general response. The best normal approximation to the distribution $p(\theta)$ is that with mean and variance given respectively by $\mu = E(\theta)$ and $\sigma^2 = V(\theta)$, the mean and variance of the original distribution.

The second question involves finding the best transformation $\zeta = \zeta(\theta)$ to induce normality in the sense of minimizing the divergence measure between $p(\zeta)$ and its normal approximation. This is equivalent to finding ζ that minimizes the expected value

$$\int p(\zeta) \log \left(\frac{p(\zeta)}{p_0(\zeta)} \right) d\zeta,$$

where $p_0(\zeta)$ is the density function of the $N[E(\zeta), V(\zeta)]$ distribution and $p(\zeta) = p(\theta)|d\zeta/d\theta|$. This is a difficult problem for which the solution will be presented only for two particular distributions: the beta and gamma.

1. If $\theta \sim \text{beta}(\alpha, \beta)$ then $\zeta(\theta) = \log[\theta/(1-\theta)]$ is the best transformation to induce normality and the approximating normal will have mean and variance given by

$$\dots \quad E(\zeta) \simeq \log \left(\frac{\mu}{1-\mu} \right) = \log \left(\frac{\alpha}{\beta} \right) \quad \text{where } \mu = \frac{\alpha}{\alpha+\beta}$$

and

$$V(\zeta) \simeq \frac{1}{\mu(1-\mu)} \frac{1}{\alpha+\beta+1} = \frac{(\alpha+\beta)^2}{\alpha\beta(\alpha+\beta+1)}.$$

2. If $\theta \sim G(\alpha, \beta)$ then $\zeta(\theta) = \log \theta$ is the best transformation and the mean and variance of the approximating normal distribution will be

$$E(\zeta) \simeq \log \mu \quad \text{and} \quad V(\zeta) \simeq \frac{1}{\mu\beta} = \frac{1}{\alpha} \quad \text{where } \mu = \frac{\alpha}{\beta}.$$

The delta method was used to get these results: It is worth pointing out that these transformations are exactly the same as those that appear in the exponential family for the Bernoulli and Poisson models, respectively. These are the observational models to which the beta and gamma are conjugate respectively.

Example. Suppose someone wants to elicit the hyperparameters of the conjugate prior for θ in a Poisson model. Assume that $\theta \sim G(\alpha, \beta)$ and now suppose that its mode was assessed as 0.5 and also that $P(\theta < 0.25) = 0.05$. Then $(\alpha-1)/\beta = 0.5$ or $\beta = 2(\alpha-1)$ and $P(\theta < 0.25) = P(\zeta < \log 0.25) \simeq \Phi[(\log 0.25 - E(\zeta))/\sqrt{V(\zeta)}]$ where $E(\zeta) \simeq \log[\alpha/(2\alpha-2)]$ and $V(\zeta) \simeq 1/\alpha$.

So $\beta = 2\alpha - 2$ and $1.96 = [\log 0.25 - \log[\alpha/(2\alpha-2)]]/\sqrt{1/\alpha}$. Solving this implicit function for α gives $\alpha = 3.35$ and $\beta = 4.7$.

5.3.3 Laplace approximation

This class of approximation methods is very useful to evaluate integrals of the type $I = \int f(\theta) d\theta$ by rewriting it as

$$\int g(\theta) \exp[-nh(\theta)] d\theta$$

where $g : R^p \rightarrow R$ and $h : R^p \rightarrow R$ are smooth functions which are at least three times differentiable. Let $\hat{\theta}$ be the value of θ which minimizes h . The Laplace method approximates I by

$$\hat{I} = g(\hat{\theta})(2\pi/n)^{p/2} |\hat{\Sigma}|^{1/2} \exp[-nh(\hat{\theta})] \quad \text{where } \hat{\Sigma} = \left[\frac{\partial^2 h(\hat{\theta})}{\partial \theta \partial \theta'} \right]^{-1}.$$

The Laplace approximation is based on the Taylor approximation for h and g around $\hat{\theta}$. Only the univariate case will be presented here for ease of exposition and $\hat{\Sigma}$ will be denoted by $\hat{\sigma}^2$. As in the last section, it will be supposed that θ and $\hat{\theta}$ are close.

Using a Taylor expansion up to the third order it follows that

$$nh(\theta) = nh(\hat{\theta}) + \frac{n}{2\hat{\sigma}^2}(\theta - \hat{\theta})^2 + \frac{nt(\theta)}{3!} + o(n^{-1})$$

where

$$t(\theta) = \frac{\partial^3 h(\hat{\theta})}{\partial \theta^3} (\theta - \hat{\theta})^3.$$

Exponentiating the last expression and applying a linear expansion to $\exp(-nt(\theta))$, it follows that

$$\begin{aligned} \exp[-nh(\theta)] &= \exp[-nh(\hat{\theta})] \\ &\times \exp\left[-\frac{n}{2\hat{\sigma}^2}(\theta - \hat{\theta})^2\right] \left[1 - \frac{nt(\theta)}{6} + o(n^{-1})\right] [1 + o(n^{-1})]. \end{aligned}$$

The same expansion in Taylor series around $\hat{\theta}$ can be applied to

$$g(\theta) = g(\hat{\theta}) + \frac{\partial g(\theta)}{\partial \theta}(\theta - \hat{\theta}) + o(n^{-1}).$$

Recognizing that

1. $\int \exp[-(n/2\hat{\sigma}^2)(\theta - \hat{\theta})^2] d\theta = (2\pi)^{1/2}(\hat{\sigma}^2/n)^{1/2}$,
2. $\int (\theta - \hat{\theta})^{2k+1} \exp[-(n/2\hat{\sigma}^2)(\theta - \hat{\theta})^2] d\theta = 0$, $\forall k$ integer, and
3. $\int nt(\theta)(\theta - \hat{\theta}) \exp[-(n/2\hat{\sigma}^2)(\theta - \hat{\theta})^2] d\theta = o(n^{-1})$

and applying the expression for I leads to a scalar version of the following proposition.

Proposition. When $n \rightarrow \infty$, $\hat{I} = I[1 + o(n^{-1})]$.

In Bayesian applications, generally $-nh(\theta) = L(\theta) + \log p(\theta)$ which is the expression of the posterior density but for the proportionality constant. If $g(\theta)$ is non-negative, the integral can be redefined by

$$I = \int \exp[-nh^*(\theta)] d\theta \quad \text{where } nh^*(\theta) = nh(\theta) - \log g(\theta).$$

Denoting by $\hat{\theta}^*$ the value that minimizes $h^*(\theta)$ and writing $\hat{\sigma}^{*2} = \partial^2 h^*(\hat{\theta}^*)/\partial \theta^2$, there follows an alternative approximation for I given by

$$\tilde{I} = (2\pi)^{1/2}\hat{\sigma}^* \exp[-nh^*(\hat{\theta}^*)].$$

In the case of a multivariate θ , the expression becomes

$$\tilde{I} = (2\pi)^{p/2} |\hat{\Sigma}^*|^{1/2} \exp[-nh^*(\hat{\theta}^*)]$$

where $\hat{\theta}^*$ is the value of θ that minimizes h^* and $\hat{\Sigma}^*$ is the inverse of the matrix of second derivatives of h^* evaluated at $\hat{\theta}^*$.

Following the same steps as before it is easy to see that $\tilde{I} = I[1 + o(n^{-1})]$. Tierney and Kadane (1986) proposed evaluating

$$E[g(\theta)] = \left[\int g(\theta) \exp[-nh(\theta)] d\theta \right] / \left[\int \exp[-nh(\theta)] d\theta \right]$$

by approximating separately the numerator and the denominator. They have shown that by doing so the $o(n^{-1})$ terms cancel out and an improved approximation of order $o(n^{-2})$ is obtained.

The final expression for their approximation can be obtained by combining the above results to give

$$\hat{E}[g(\theta)] = \frac{g(\hat{\theta}^*) |\hat{\Sigma}^*|^{1/2} \exp[-nh^*(\hat{\theta}^*)]}{|\hat{\Sigma}|^{1/2} \exp[-nh(\hat{\theta})]}.$$

Example. Let X_1, \dots, X_n be a random sample from a $\text{Pois}(\theta)$ and suppose that a conjugate prior $\theta \sim G(a_0, b_0)$ is used. Taking $g(\theta) = \theta$ it follows that

$$\hat{E}[\theta | \mathbf{x}] = \frac{a_1}{b_1} \left(\frac{a_1}{a_1 - 1} \right)^{a_1 - 1/2} e^{-1}, \quad a_1 > 1$$

where $a_1 = a_0 + \sum_{i=1}^n x_i$ and $b_1 = b_0 + n$.

The exact posterior mean in this example is a_1/b_1 and so we can easily evaluate the relative errors $(\hat{E} - E)/E$ involved in the approximation. For example, for $a_1 = 6$ the relative error is 0.0028 and for $a_1 = 10$ it will be only 0.00097.

Example (continued). *Randomized response* (Migon and Tachibana, 1997) A variation of the randomized response model consists in asking as the alternative question the negation of the original one. In this case, $\theta_A = 1 - \theta$ and the probability of a YES response is

$$\lambda = \pi\theta + (1 - \pi)(1 - \theta) = (2\pi - 1)\theta + (1 - \pi).$$

The Laplace approximation for the evaluation of the posterior mean of θ can be applied. The derivatives and points of maxima can all be obtained analytically or numerically. With the same data of the example and with a unit uniform prior for

θ , the exact posterior mode is 0.25. The posterior distribution of θ is a mixture of $n + 1$ beta distributions and the calculations become tedious as the sample size n increases. The performance of the Laplace approximations can be assessed as a function of the sample size by keeping the YES proportion fixed at $60/150 = 0.40$ and letting the sample size n vary. For this example, the exact and approximated posterior means were

n	Exact	Laplace
50	0.28	0.27
150	0.251	0.255
450	0.251	0.251

It is worth noting that even though the exact posterior mean itself depends on the sample size, the Laplace approximation gets better as the value of n increases.

Example. Weibull data Using the same data of the example presented in Section 5.2.1, we obtain the following estimates via Laplace methods with a non-informative prior: $\hat{\alpha} = 1.59$, $V(\hat{\alpha}|x) = 0.064$ and $\hat{\theta} = 2.05$, $V(\hat{\theta}) = 0.079$, which compare well with the maximum likelihood results, as expected.

5.4 Numerical integration methods

Numerical integration, also called the quadrature technique, is a collection of methods convenient to solve some useful problems in inference, mainly when the dimension of the parameter space is moderate. They become useful as soon as the analytical solution fails.

Firstly we will take care of the general problem of obtaining the value of the integral $I = \int_a^b f(x)dx$, where $f : R \rightarrow R$ is a smooth function. Let $w : R \rightarrow R^+$ be a well-defined weight function. Quadrature methods are essentially approximations of I obtained by evaluating f at points x_i , $i = 1, \dots, n$. The simplest solution is given by the weighted sum

$$\hat{I} = \sum_1^n w_i f(x_i) \quad \text{where } w_i = w(x_i), i = 1, \dots, n.$$

The quadrature methods are fully characterized by choosing the points of evaluation or nodes x_1, \dots, x_n in the interval (a, b) and the corresponding weights involved. An integration rule must have easily obtainable weights and nodes. The nodes should lie in the region of integration and the weights should all be positive.

5.4.1 Newton–Cotes type methods

The interval of integration (a, b) with finite a, b is divided into n equal parts, the function $f(x)$ is evaluated in the middle point of each interval and the weights are

then applied. Then

$$\hat{I}_{NC} = h \sum_{i=1}^n f(a + (2i - 1)h/2)$$

approximates I , with $h = (b - a)/n$. These methods are generically named Newton–Cotes rules.

This is an approximation by the area of the rectangles with equal base $(b - a)/n$. Often a good approximation is obtained with n of the order of 10^2 , which seems reasonable for the unidimensional case.

A slight variation is the trapezoidal rule involving unit weights except in the extremes of the interval, when they are set to 1/2. The rule gives the approximation

$$\hat{I}_T = h \left[\frac{f(a)}{2} + \sum_{i=1}^{n-1} f(a + (2i - 1)h/2) + \frac{f(b)}{2} \right].$$

The Simpson rule is another variation described by weights alternating between 4/3 and 2/3, except in the extreme where they assume the value 1/3. In this case the approximation is given by

$$\hat{I}_S = \frac{h}{3} [f(a) + 4 \sum_{i=1}^{n/2} f(a + (4i + 1)h/2) + 2 \sum_{i=1}^{n/2} f(a + (4i + 3)h/2) + f(b)].$$

The p -dimensional case is slightly more demanding. A general solution follows from an iterative application of the cartesian product rule. Let $\mathbf{x} = (x_1, x_2)$ be a bidimensional vector. A quadrature in two dimensions is based on the cartesian product rule ...

$$\int f(\mathbf{x})d\mathbf{x} = \int \left[\int f(x_1, x_2)dx_2 \right] dx_1 = \int f_1(x_1)dx_1.$$

The last term in the right-hand side is obtained by integration with respect to x_2 using the unidimensional quadrature rule

$$\int f_2(x_2)dx_2 \cong \sum_{j=1}^m w_j f_2(x_{2,j}),$$

where $f_2(x_2) = f(x_1, x_2)$. As the last integral is also unidimensional it can be approximated again by quadrature $\int f_1(x_1)dx_1 \cong \sum_{i=1}^n w_i f_1(x_{1,i})$. Joining the two weights it follows that

$$\int f(\mathbf{x}) d\mathbf{x} \cong \sum_{i=1}^n \sum_{j=1}^m w_i w_j f(x_{1,i}, x_{2,j}).$$

This is a bidimensional rule based on $n \times m$ evaluation points $(x_{1,i}, x_{2,j})$ and with weights (w_i, w_j) , $i = 1, \dots, n$, $j = 1, \dots, m$.

In the sequel, a general integration method more adequate to statistical problems will be introduced. From now on we will be concerned with an integral over the whole real line.

5.4.2 Gauss–Hermite rules

The method introduced in this section is specifically useful for integration over the whole real line. If the domain of the integral is $\tau > 0$, which sometimes happens when we are dealing with a precision or a scale parameter, then the reparametrization $\log \tau$ is often convenient.

Assuming that the integrand $f(x)$ can be expressed in the form of $g(x) \exp(-x^2)$, a general unidimensional rule is

$$\int_{-\infty}^{\infty} f(x) dx = \int_{-\infty}^{\infty} g(x) e^{-x^2} dx \simeq \sum_{i=1}^n h_i f(x_i)$$

where the x_i 's are the zeros of a Hermite polynomial of degree n , $H_n(x)$, and the weights h_i depend on n and on the Hermite polynomial $H_{n-1}(x)$, evaluated at x_i , $i = 1, \dots, n$. As long as $g(x)$ is a polynomial of maximum degree $2n - 1$, the formula is exact. There are other Gaussian integration rules that can be applied when the domain of integration is finite. The values of the zeroes and heights are tabulated in many mathematical tables and can be found in Abramowitz and Stegun (1965), for example.

The accuracy of the approximation depends on $f(x)$ being well approximated by a polynomial of degree $2n - 1$ or less times a normal weight function. For the multivariate case some sort of parameter orthogonality must be guaranteed for application of the cartesian rule.

The integration rule introduced before can be extended to a more general approximating function $f(x) = g(x)(2\pi\sigma^2)^{-1/2} \exp\{-0.5[(x - \mu)/\sigma]^2\}$ leading to

$$\int f(x) dx = \int g(x) \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left[-\frac{1}{2}\left(\frac{x - \mu}{\sigma}\right)^2\right] dx.$$

Making the variable transformation $z = (x - \mu)/\sqrt{2\sigma^2}$, it follows that

$$\begin{aligned} \int f(x) dx &= \int g(\sqrt{2\sigma^2}z + \mu) \frac{1}{\sqrt{2\pi\sigma^2}} e^{-z^2} \sqrt{2\sigma^2} dz \\ &= \pi^{-1/2} \int g(\sqrt{2\sigma^2}z + \mu) e^{-z^2} dz \end{aligned}$$

and the Gauss–Hermite approximation will be

$$\begin{aligned} \int f(x) dx &\simeq \pi^{-1/2} \sum_{i=1}^n h_i \sqrt{2\pi\sigma^2} e^{x_i^2} g(\sqrt{2\sigma^2}z_i + \mu) \\ &= \sum_{i=1}^n m_i g(x_i) \end{aligned}$$

where $m_i = \sqrt{2\sigma^2} e^{x_i^2} h_i$ and $x_i = \sqrt{2\sigma^2} z_i + \mu$, $i = 1, \dots, n$.

In general, the mean μ and variance σ^2 of the approximating normal are unknown. The following strategy, proposed by Naylor and Smith (1982), can be used:

1. Let μ_0 and σ_0^2 be initial values for μ and σ^2 .
2. Apply the above expression with $g(\theta) = \theta$ and $g(\theta) = \theta^2$, respectively. This will provide approximating values for the mean and variance.
3. Repeat the last step until the mean and variance values obtained are stable.

In the multiparameter case, the cartesian product rule is applied after some orthogonalization is operated over the parameters. This can be done through a Cholesky decomposition of the (approximated) variance–covariance matrix Σ . Let $\Sigma = \mathbf{H} \mathbf{D} \mathbf{H}'$, where \mathbf{D} is a diagonal matrix and \mathbf{H} is a lower triangular matrix, and define the transformation $z = \mathbf{D}^{-1/2} \mathbf{H}^{-1} (\mathbf{x} - \boldsymbol{\mu})/\sqrt{2}$. Therefore

$$\begin{aligned} \int f(\mathbf{x}) d\mathbf{x} &= \pi^{-p/2} \int g(\boldsymbol{\mu} + \sqrt{2}\mathbf{H} \mathbf{D}^{1/2} \mathbf{z}) \exp(-\mathbf{z}' \mathbf{z}) d\mathbf{z} \\ &\simeq \pi^{-p/2} \sum_{i_p=1}^{n_p} \cdots \sum_{i_1=1}^{n_1} h_{i_1} \cdots h_{i_p} f(\boldsymbol{\mu} + \sqrt{2}\mathbf{H} \mathbf{D}^{1/2} \mathbf{z}) \end{aligned}$$

where n_i is the number of nodes involved in the approximation at the i th coordinate and $\mathbf{z} = (z_{i_1}, \dots, z_{i_p})$, $i_j = 1, \dots, n_j$, $j = 1, \dots, p$.

If the mean vector $\boldsymbol{\mu}$ and variance–covariance matrix Σ of the approximating normal density are unknown, an extension of the iterative method of Naylor and Smith presented before can be applied.

5.5 Simulation methods

In this section we will be discussing a collection of techniques useful to solve many of the relevant statistical problems. From a classical point of view we are interested in simulating a sampling distribution arising from a possibly complex model and in the Bayesian case we are interested in obtaining some characteristics of the posterior distribution. All the techniques described in this section share a common characteristic: they involve the random generation of samples from a distribution of interest. This is the empirical distribution in the classical approach and the posterior distribution in the Bayesian case. The interested reader can complement their study by reading the books by Efron (1982), Gamerman (1997) and Ripley (1987).

5.5.1 Monte Carlo method

The basic idea of the Monte Carlo method consists in writing the desired integral as an expected value with respect to some probability distribution. To motivate

148 Approximate and computationally intensive methods

our discussion we will begin with a very simple problem. Assume we wish to calculate the integral of a smooth function in a known interval (a, b) , that is

$$I = \int_a^b g(\theta) d\theta.$$

The above integral can be rewritten as

$$I = \int_a^b [(b-a)g(\theta)] \frac{1}{b-a} d\theta.$$

This problem can be thought of as the evaluation of the expectation of $[(b-a)g(\theta)]$ with respect to the uniform distribution over (a, b) and

$$I = E_{U(a,b)}[(b-a)g(\theta)]$$

where $U(a, b)$ represents the uniform distribution in (a, b) .

A method of moments estimator of this quantity is

$$\hat{I} = \frac{1}{n} \sum_{i=1}^n (b-a)g(\theta_i)$$

where $\theta_1, \dots, \theta_n$ is a random sample selected from the uniform distribution on (a, b) .

An algorithm can be described by the following steps

1. generate $\theta_1, \theta_2, \dots, \theta_n$ from a $U(a, b)$ distribution;
2. calculate $g(\theta_1), g(\theta_2), \dots, g(\theta_n)$;
3. obtain the sample mean: $\bar{g} = (1/n) \sum_{i=1}^n g(\theta_i)$;
4. finally, determine: $\hat{I} = (b-a)\bar{g}$.

A generalization can be obtained straightforwardly. Let $I = E_p[g(\theta)]$ be the expected value of $g(\theta)$ with respect to a distribution with density $p(\theta)$. The algorithm is similar to that described above with modification of the sampling in step 1 from the $U(a, b)$ to $p(\cdot)$.

The multivariate extension is based on evaluation of

$$I = \int_{a_1}^{b_1} \cdots \int_{a_p}^{b_p} g(\theta) d\theta$$

and the Monte Carlo estimator is

$$\hat{I} = \frac{1}{n} \sum_{i=1}^n g(\theta_i)$$

where $\theta_1, \dots, \theta_n$ is a random sample selected from the uniform distribution on $(a_1, b_1) \times \cdots \times (a_p, b_p)$.

Some questions need to be answered to implement these techniques. How large must n be? How do Monte Carlo methods compare with quadrature rules? Or, to pose it in another way: when is Monte Carlo preferred to numerical integration? An elementary example will be useful to motivate further developments.

Example. The evaluation of $I = \int_0^1 e^x dx$ is desired. As we have just seen, the simple Monte Carlo estimator of I is $\bar{I} = (1/n) \sum_{i=1}^n \exp(X_i)$, $X_i \sim U(0, 1)$, $j = 1, \dots, n$. Since \bar{I} is a sample mean, its precision to estimate I can be measured by its variance. It is given by

$$V(\bar{I}) = \frac{1}{n} V(e^X) = \frac{1}{n} \left(\int_0^1 e^{2x} dx - I^2 \right) = \frac{0.242}{n}.$$

5.5.2 Monte Carlo with importance sampling

The Monte Carlo method with importance sampling is a technique developed to reduce the variance of the estimator. Consider now explicitly that the integral I of interest is the expectation of a given function g with respect to a density p . It can then be rewritten as

$$I = \int g(x) p(x) dx = \int g(x) \frac{p(x)}{h(x)} h(x) dx$$

where $h(x)$ is a positive function for all x where $p(x) > 0$ and $\int h(x) dx = 1$. It is therefore a density. An alternative method of moments estimator for I can be obtained as

$$\bar{I} = \frac{1}{n} \sum_{i=1}^n g(X_i) w(X_i),$$

where

$$w(X_i) = \frac{p(X_i)}{h(X_i)}$$

and

$$X_i \sim h(x), i = 1, \dots, n.$$

where h is called the importance sampling density. The only difference with respect to simple Monte Carlo is the first step where sampling from the uniform distribution is replaced by sampling from h and the third step where the values of $g \times w$ instead of values of g are averaged. Therefore, $V(\bar{I}) = (1/n) \int (g(x)w(x) - I)^2 h(x) dx$. Choosing $g(x)w(x)$ approximately constant can make $V(\bar{I})$ as small as we want. Therefore, whenever possible the importance sampling density should be roughly similar to $g(x)p(x)$.

The multivariate extension is again trivial. Assume we wish to obtain the value of the expectation of g with respect to p given by

$$I = \int g(x) p(x) dx,$$

and use the multivariate density h as the importance density. Once again, h must be positive whenever p is. The Monte Carlo estimator is given by

$$\hat{I} = \frac{1}{n} \sum_{i=1}^n g(\mathbf{X}_i) w(\mathbf{X}_i),$$

where

$$w(\mathbf{X}_i) = \frac{p(\mathbf{X}_i)}{h(\mathbf{X}_i)}$$

and

$$X_i \sim h(x), i = 1, \dots, n.$$

Example (continued). If we take as importance sampling density $h(x) = \frac{2}{3}(1+x)$, $x \in (0, 1)$ then $w(x) = p(x)/h(x) = 3/[2(1+x)]$ and $g(x)w(x) = (3/2)e^x/(1+x) \propto k$, for $x \in (0, 1)$. Then

$$I = \int_0^1 e^x \left[\frac{3}{2(1+x)} \right] \left[\frac{2}{3}(1+x) \right] dx = \int_0^1 \frac{3}{2} \frac{e^x}{1+x} h(x) dx.$$

Therefore, $\bar{I} = (1/n) \sum_{i=1}^n 3e^{X_i}/[2(1+X_i)]$ where $X_i \sim h(x)$, $i = 1, \dots, n$, and $V(\bar{I}) = (1/n)[(3/2)^2 \int_0^1 e^{2x}/(1+x)^2 dx - I^2] = 0.027/n$. The variance reduction is quite substantial. It was reduced to approximately one tenth of the value obtained with a simple Monte Carlo using the same number of replications.

The implementation of the algorithm in this case depends on sampling from the density h . The importance sampling distribution function is given by

$$H(x) = \begin{cases} 0 & \text{if } x < 0 \\ \frac{2}{3}(x + \frac{x^2}{2}) & \text{if } 0 \leq x < 1 \\ 1 & \text{if } x \geq 1. \end{cases}$$

Using the probability integral transform, one simply has to generate a unit uniform random variable U and solve for $U = (2/3)(X + X^2/2)$. The unique solution satisfying the equation is $X = [(4 + 12U)^{1/2} - 2]/2$.

Example. Let $\theta = P[X > 2]$ where X has a standard Cauchy distribution with density

$$p(x) = \frac{1}{\pi(1+x^2)}, \quad x \in R.$$

In this example, $g(x) = I_x[(2, \infty)]$ and $\theta = \int_{-\infty}^{\infty} g(x)p(x) dx$.

Let X_1, \dots, X_n be a random sample from the Cauchy distribution. It is easy to obtain that

$$\hat{\theta} = \frac{1}{n} \sum_{i=1}^n I_{X_i}[(2, \infty)] = \frac{\#(X_i > 2)}{n}$$

and $n\hat{\theta} \sim \text{bin}(n, \theta)$. Note that $P(x) = \int_{-\infty}^x p(t)dt = 0.5 + \frac{1}{\pi} \arctan x$, so that $\theta = 1 - P(2) = 0.1476$. Then, $V(\hat{\theta}) = \theta(1-\theta)/n = 0.126/n$.

Let h be a density defined by $h(x) = 2/x^2 I_x[(2, \infty)]$. It is easy to obtain that the distribution function is $H(x) = 1 - 2/x$, if $x \geq 2$. For any $U \in U(0, 1)$ it follows that $X = 2/(1-U)$ has density h and a sample X_1, \dots, X_n from which h can easily be obtained. The importance sampling estimator of θ is given by

$$\tilde{\theta} = \frac{1}{n} \sum_{i=1}^n w(X_i), \quad \text{where } w(x) = \frac{1}{2\pi} \frac{x^2}{1+x^2}.$$

It can be shown that the variance of this estimator is smaller than that of $\hat{\theta}$ (see Exercise 5.16).

Therefore, the Monte Carlo algorithm can be used to solve any of the basic inference problems cited in the introduction of this chapter that can be written as an expectation. In the Bayesian case, when one wants to evaluate $E[g(\theta)|\mathbf{x}]$, the algorithm can be summarized as follows:

1. Generate $\theta_1, \dots, \theta_n$ from the posterior density $p(\theta|\mathbf{x})$ (or the importance density $h(\theta)$).
2. Calculate $g_i = g(\theta_i)$ (or $g_i = g(\theta_i)p(\theta_i|\mathbf{x})/h(\theta_i)$), $i = 1, \dots, n$.
3. Obtain the estimator $\hat{E}[g(\theta)] = (1/n) \sum_{i=1}^n g_i$.

5.5.3 Resampling methods

In this section we will be concerned with some sampling and resampling techniques from a classical and a Bayesian point of view. Firstly, the classical bootstrap, which essentially consists in resampling from the empirical distribution function, will be presented. A weighted version of the bootstrap will be useful to implement the Bayesian argument. Intuitively the argument follows as: a sample is generated from the prior distribution and a resample is taken using some well-defined weights. It is not difficult to show that the points in the resample constitute an approximate sample from the posterior distribution. This approximation becomes better as sample sizes increase. Another classical resample technique, named jackknife, will be presented and exemplified. Its Bayesian version will be developed for the exponential family. A general account of these resampling methods from a classical perspective is provided by Davison and Hinkley (1997) and Efron (1982).

The main objective of jackknife and bootstrap is to obtain a measure of the accuracy of some complex statistics. From a classical point of view the question arises from the fact that the sampling distribution is hard to be determined in some cases. As some examples we can mention the robust statistics, like trimmed means, the correlation coefficient, concordance measures in probabilistic classification and so on.

On the Bayesian side the interest in techniques like jackknife and bootstrap is slightly different. One important application of leave-one-out methods is to obtain information about the influence of particular observations or, more generally, to define diagnostic measures. One may also use the bootstrap as a resample technique useful to implement the Bayesian paradigm, as will be shown later in this section.

This class of procedures is characterized by its computational demand, although they are often very easy to implement even in complex situations or high-dimensional problems.

152 Approximate and computationally intensive methods

Jackknife

The jackknife is a useful technique to build up confidence intervals and it works as well as a bias reduction technique as will be shown in an example. This is a generic tool and so in specific problems it could provide less accurate results. The basic idea of splitting the sample was introduced by Quenouille (1949, 1956) to eliminate estimation bias.

Suppose that X_1, \dots, X_n is a random sample from $p(x|\theta)$ and that $\hat{\theta}(X)$ is an estimator of θ . Denote by $\tilde{\theta}_i$ the estimator based on the original sample without the i th observation. Let $\tilde{\theta}_i = n\hat{\theta} - (n-1)\hat{\theta}_i$ be a sequence of pseudo-values and define the jackknife estimator of θ as

$$\hat{\theta}_J = \frac{1}{n} \sum_{i=1}^n \tilde{\theta}_i.$$

The name pseudo-value derives from the fact that for the special case where $\hat{\theta}(X) = \bar{X}$, the pseudo-value coincides with the i th observation, that is $\tilde{\theta}_i = \sum_j^n X_j - \bar{X}_i$. It is not difficult to show that $\hat{\theta}_J$ is unbiased if $\hat{\theta}$ and $\tilde{\theta}_i$ are also unbiased. Besides that, the jackknife estimator has the property of eliminating terms of order $1/n$ on the bias of the estimator.

Example. Let X_1, \dots, X_n be iid observations from the uniform distribution on $(0, \theta)$. It is well known that $T = \max_i X_i$ is a sufficient statistic for θ with $E(T) = (1 - 1/n)\theta, \forall \theta$. A jackknife estimator is given by

$$\hat{\theta}_J = \frac{1}{n} \sum_{i=1}^n \tilde{\theta}_i, \quad \text{where } \tilde{\theta}_i = nT - (n-1)\hat{\theta}_i,$$

and $\hat{\theta}_i = \max\{X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_n\}$. Then, $E(\hat{\theta}_J) = n(1 - 1/n)\theta - (n-1)[1 - 1/(n-1)]\theta = \theta$.

Let $\tilde{\theta}_i, i = 1, \dots, n$, represent random variables approximately independent and identically distributed with mean θ . A jackknife estimate of the sample variance will be given by

$$\hat{\sigma}_J^2 = \frac{1}{n-1} \sum_{i=1}^n (\tilde{\theta}_i - \hat{\theta}_J)^2$$

and therefore the statistic

$$\frac{\hat{\theta}_J - \theta}{(\hat{\sigma}_J^2/n)^{1/2}}$$

has approximately a standard Student t distribution with $n-1$ degrees of freedom. An approximate $100(1-\alpha)\%$ confidence interval for θ is given by

$$\left(\tilde{\theta}_J - t_{n-1, \alpha/2} \frac{\hat{\sigma}_J}{n^{1/2}}, \tilde{\theta}_J + t_{n-1, \alpha/2} \frac{\hat{\sigma}_J}{n^{1/2}} \right).$$

Example. Correlation coefficient. A very popular data set in statistics was given by Fisher (1936) and contains measurements on three species of iris. We concentrate on the sepal length and sepal width of the species *iris setosa*. The correlation coefficient of length and width is calculated in a sample of 50 observations. The sample correlation coefficient is estimated as $\hat{\rho} = 0.742$. Using the jackknife methodology we obtain the point estimation 0.743 and the 95% confidence interval based on the Student t distribution was (0.63, 0.84).

From a Bayesian perspective the notion of jackknife corresponds to obtaining the posterior or predictive distribution leaving one of the observations out and is very useful for model checking. For example, the influence of an observation can be assessed by this procedure. Using a divergence measure like Kullback-Liebler, the posterior or predictive distributions based on the whole sample could be compared with that leaving one observation out to evaluate its influence in the analysis.

For the exponential family with one parameter, defined in Chapter 2, the conjugate analysis leads to the posterior density

$$p(\theta|x) \propto \exp\{\alpha_1 \phi(\theta) + \beta_1 b(\theta)\}$$

where $\alpha_1 = \alpha + T(x)$, $\beta_1 = \beta + n$ and $T(x) = \sum_{i=1}^n u(x_i)$ and $p(y|x) = a(y)k(\alpha_1, \beta_1)/k(\alpha_1 + u(y), \beta_1 + n + 1)$. So, leaving one observation out would make the posterior density

$$p(\theta|x_i) \propto \exp\{\alpha'_1 \phi(\theta) + \beta'_1 b(\theta)\}$$

where $\alpha'_1 = \alpha + T(x_i)$, $\beta'_1 = \beta + n - 1 = \beta_1 - 1$, $T(x_i) = \sum_{j \neq i}^n u(x_j)$.

Bootstrap

The concept of bootstrap was introduced by Efron (1979). The method of bootstrap consists in generating a large number of samples based on the empirical distribution obtained from the original sampled data. Confidence intervals with some pre-specified coverage probability can be built up easily under mild assumptions.

Let X_1, \dots, X_n be the observed data from a random sample of a distribution $p(x|\theta)$, where $\theta \in \Theta$ is the unknown parameter. Let $\hat{\theta}(x)$ be an estimator of θ . The empirical distribution function is defined by $\hat{F}_n(x) = (1/n)\#(X_i \leq x)$, $\forall x \in R$, as seen in Chapter 4.

A resample procedure consists of the selection of samples with replacement from a finite population using equal probability. This corresponds to selecting a sample from the empirical distribution $\hat{F}_n(x)$. These sampled values will be denoted by $\{X_1^*, \dots, X_n^*\}$ and the bootstrap estimator of θ by $\hat{\theta}^*(x^*)$. The inference will be based on B replications of the above procedure and in the evaluation of the statistic of interest, in this case the estimator $\hat{\theta}^* = \hat{\theta}^*(x^*)$ for each of the B replications. Denote the resulting values by $\hat{\theta}_1^*, \dots, \hat{\theta}_B^*$. The bootstrap distribution of $\hat{\theta}^*$ is given

by the empirical distribution formed by the resampled values. Summarizing, the bootstrap distribution of the $\hat{\theta}^*$ is used in the place of the sampling distribution of $\hat{\theta}$ in order to make the inferences about θ .

A central assumption in the method is that \hat{F}_n is a good approximation of F , that is, the bootstrap distribution of $\hat{\theta}^*$ is similar to that of $\hat{\theta}$ or, that the distribution of $\hat{\theta}^* - \tilde{\theta}$ is similar to that of $\hat{\theta} - \theta$, where $\tilde{\theta}$ is the value of the parameter for \hat{F}_n .

The mean and variance of these B replications will be denoted by $\bar{\theta}^* = (1/n) \sum_{i=1}^B \hat{\theta}_i^*$ and $\hat{\sigma}^2(\hat{\theta}^*) = [1/(B-1)] \sum_{i=1}^B (\hat{\theta}_i^* - \bar{\theta}^*)^2$. From the above suppositions it follows that $V(\hat{\theta}) = V(\hat{\theta}^*) \simeq \hat{\sigma}^2(\hat{\theta}^*)$ and $E[\hat{\theta} - \theta] = \bar{\theta}^* - \tilde{\theta}$. A bias adjusted estimate of θ will be $\hat{\theta} = \hat{\theta} - [\bar{\theta}^* - \tilde{\theta}]$.

Confidence intervals for θ can be built from the percentiles of the bootstrap distribution. Let $\theta^*(\alpha)$ be the $100(\alpha)\%$ percentile of the bootstrap distribution of $\hat{\theta}^*$, that is, $P[\hat{\theta}^* \leq \theta^*(\alpha)] = \alpha$. The interval $(\theta_{\alpha/2}^*, \theta_{1-\alpha/2}^*)$ obtained as described before is named the $100(1-\alpha)\%$ bootstrap confidence interval.

Example (continued). Returning to the evaluation of the correlation between length and width of samples of *iris setosa* and applying the bootstrap with different values of B gives the following results

B	$L_{5\%}$	Mean	$U_{95\%}$
100	0.64	0.736	0.81
400	0.71	0.741	0.82
1600	0.71	0.742	0.82

It seems that $n = 400$ is a reasonable number of replications to accurately describe the bootstrap distribution. It is interesting to note that this is the number of replication usually recommended in the literature. The point estimates for the bootstrap and jackknife are almost the same although the confidence intervals are shorter for the bootstrap.

Weighted bootstrap

Sometimes we are not able to sample directly from the distribution of interest, $p(x)$. A useful strategy is to sample from an approximation of this distribution and use the accept-reject scheme:

1. Generate x from an auxiliary density $h(x)$.
2. Generate u independently from a uniform distribution on $(0, 1)$.
3. If $u \leq p(x)/Ah(x)$, where $A = \max p(x)/h(x)$, accept x , otherwise return to step 1.

The probability of accepting a value x generated from $h(x)$ is

$$P(\text{accept } x) = \int \int I_u(0, p(x)/Ah(x)]h(x)dxdu = \frac{1}{A}.$$

The expected number of accepted values in n independent runs of the algorithm will be n/A . So, the algorithm is improved by decreasing the value of A as much as possible.

If the determination of A is difficult, the following modification of the algorithm can be applied:

1. Take a sample from x_1, \dots, x_n from $h(x)$.
2. Evaluate the weights $w(x_i) = p(x_i)/h(x_i)$, $i = 1, \dots, n$.
3. Select a new sample $x_1^*, x_2^*, \dots, x_m^*$ from the set $\{x_1, \dots, x_n\}$ with respective probabilities given by $w_i / \sum_{i=1}^n w_i$, $i = 1, \dots, n$ with replacement.

Note that

$$P(x^* \leq a) = \sum_{i=1}^n \frac{w_i}{\sum_{j=1}^n w_j} I_{x_i}(-\infty, a).$$

Taking the limit as $n \rightarrow \infty$,

$$\sum_{i=1}^n \frac{w_i}{\sum_{j=1}^n w_j} I_{x_i}(-\infty, a) \rightarrow \int_{-\infty}^a p(x)dx.$$

It is interesting to note that the algorithm allows approximate sampling from $p(x)$ even when p is known up to an arbitrary constant. This is particularly useful for Bayesian inference where in many cases the proportionality constant of the posterior distribution is not known.

The above algorithm is known in the literature as the weighted bootstrap. As before, many questions deserve consideration in the applications. For example, how big must n , the initial sample size, be? And m , the resampling size? Is this approximation efficient? Note that if values of x were not generated in some regions then these values will never be resampled even if the weights were large.

We shall concentrate here on a modification of the algorithm to solve Bayes' theorem numerically. Remember that

$$p(\theta | \mathbf{x}) = kp(\theta)l(\theta; \mathbf{x}), \quad \theta \in \Theta$$

where $p(\theta)$ is the prior distribution, $l(\theta; \mathbf{x})$ is the likelihood function and \mathbf{x} denotes the available data.

Taking $h(x) = p(\theta)$ in the algorithm gives $w(x) = p(\theta | \mathbf{x})p(\theta) = kl(\theta; \mathbf{x})$. Therefore, the algorithm simplifies to

1. Take a sample $\theta_1, \dots, \theta_n$ from the prior distribution $p(\theta)$.
2. Evaluate the weights $w_i = p(\theta | \mathbf{x})/p(\theta) = kl(\theta_i; \mathbf{x})$, $i = 1, \dots, n$.
3. Sample $\theta_1^*, \theta_2^*, \dots, \theta_m^*$ with replacement from the finite population $\{\theta_1, \dots, \theta_n\}$ with respective weights $w_i / \sum_{i=1}^n w_i$, where $l_i = l(\theta_i; \mathbf{x})$, $i = 1, \dots, n$.

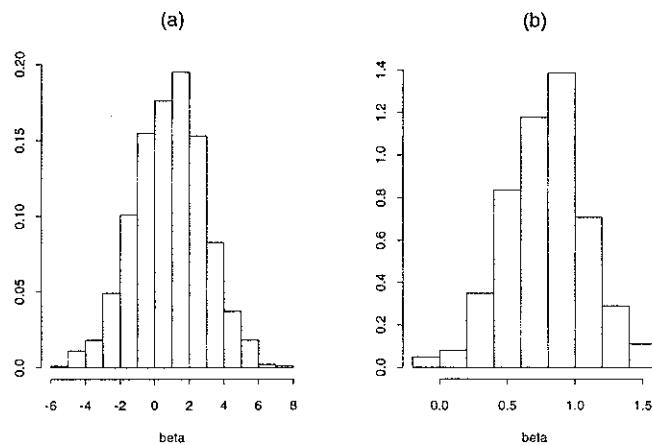


Fig. 5.2 Summary of inference for β in the example: (a) initial sample; (b) final resample.

It is worth noting that often it is necessary to make some adjustments to take care of numerical difficulties. One must make sure to sample over the relevant region of the parameter space. It may well be that the prior is concentrated over a region of low posterior probability. In this case, the prior is not a suitable candidate for initial sampling and other distributions must be used.

Example. Let $Y_i \sim N(\mu_i, \sigma^2)$, where $\mu_i = \beta x_i$ and $\sigma^2 = 1$, $i = 1, \dots, 5$. The observed data is $y = (-2, 0, 0, 0, 2)$ and $x = (-2, -1, 0, 1, 2)$. Using a reasonably vague prior $N(0, 4)$ for β and simulating $n = 1000$ samples and $m = 500$ resamples we can easily obtain numerical summaries. The estimated mean is 0.801 and the estimated variance is 0.087. The exact MLE in this example is $\hat{\beta} = 0.8$, which is in complete agreement with the Bayesian bootstrap resample depicted in Figure 5.2.

5.5.4 Markov chain Monte Carlo methods

The central idea behind the Markov chain Monte Carlo (MCMC, in short) method is to build up a Markov chain that is easy to simulate and has equilibrium distribution given by the distribution of interest. These techniques are often more powerful than the quadrature rules and simple Monte Carlo because they can be successfully applied to highly dimensional problems. A general discussion about this topic can be found in Gamerman (1997).

Let X_1, \dots, X_p have the joint density $p(x) = p(x_1, \dots, x_p)$ defined in the space $\mathcal{X} \subset R^p$. In fact, the derivations below are also valid for the more general case where the X_i 's are vector variables. Suppose that a homogeneous, irreducible

and aperiodic Markov chain with state-space \mathcal{X} and equilibrium distribution $p(x)$ can be constructed. Denote by $q(x, y)$ the transition kernel of the chain, which means that $q(x, \cdot)$ defines a conditional distribution governing the transitions from state x .

In other words, it is possible to build a chain with transition probabilities invariant in time, where each state can be reached from any other state with a finite number of iterations and also without absorbing states. Assume further that it is easy to generate values from these transition probabilities. This means that for any given initial stage, a trajectory from the chain can be generated. For a sufficiently large number of iterations, this trajectory will eventually produce draws from the equilibrium distribution $p(x)$. By constructing a suitable Markov chain, one is able to perform a Monte Carlo simulation of values from p , hence the name MCMC.

There are many possible ways to construct such a chain. One scheme is provided by the Gibbs sampler algorithm, proposed by Geman and Geman (1984) and popularized to the statistical community by Gelfand and Smith (1990). Let $p_i(x_i | \mathbf{x}_{-i})$ denote the conditional density function of X_i given values of all the other X_j 's ($j \neq i$) and assume that it is possible to generate from these distributions for each $i = 1, \dots, p$. The algorithm starts by arbitrarily choosing initial values $\mathbf{x}^0 = (x_1^0, \dots, x_p^0)$. If in the j th iteration we have the chain at state $\mathbf{x}^{(j)}$, then the position of the chain at iteration $j + 1$ will be denoted by $\mathbf{x}^{(j+1)}$ and will be given after:

- generating a random quantity $x_1^{(j+1)}$ from

$$p_1(x_1 | \mathbf{X}_{-1} = (x_2^{(j)}, \dots, x_p^{(j)}));$$

- generating a random quantity $x_2^{(j+1)}$ from

$$p_2(x_2 | \mathbf{X}_{-2} = (x_1^{(j+1)}, x_3^{(j)}, \dots, x_p^{(j)}));$$

- successively repeating the procedure for $i = 3, \dots, p$ where at the last step a random quantity $x_p^{(j+1)}$ from $p_p(x_p | \mathbf{X}_{-p} = (x_1^{(j+1)}, \dots, x_{p-1}^{(j+1)}))$.

This way, a vector $\mathbf{x}^{(j+1)} = (x_1^{(j+1)}, \dots, x_p^{(j+1)})$ is formed. Under suitable regularity conditions the limiting distribution of $\mathbf{x}^{(j)}$, as $j \rightarrow \infty$, is just $p(x)$.

Another scheme is provided by the Metropolis–Hastings algorithm initially proposed by Metropolis et al. (1953) and later extended by Hastings (1970). A clear introductory explanation of the algorithm is presented by Chib and Greenberg (1995). It is based on the same idea of using an auxiliary distribution, previously used for importance sampling, accept-reject schemes and weighted bootstrap. Let $q^*(x, \cdot)$ denote an arbitrary transition kernel and assume that at iteration j the chain is at state $\mathbf{x}^{(j)}$. Then, the position of the chain at iteration $j + 1$ will be denoted by $\mathbf{x}^{(j+1)}$ and will be given after

- proposing a move to \mathbf{x}^* according to $q^*(\mathbf{x}^{(j)}, \cdot)$;

- accepting the proposed move with probability

$$\alpha(\mathbf{x}^{(j)}, \mathbf{x}^*) = \min \left\{ 1, \frac{p(\mathbf{x}^*)/q^*(\mathbf{x}^{(j)}, \mathbf{x}^*)}{p(\mathbf{x}^{(j)})/q^*(\mathbf{x}^*, \mathbf{x}^{(j)})} \right\}$$

thus setting $\mathbf{x}^{(j+1)} = \mathbf{x}^*$ or rejecting the move with probability $1 - \alpha(\mathbf{x}^{(j)}, \mathbf{x}^*)$ thus setting $\mathbf{x}^{(j+1)} = \mathbf{x}^{(j)}$ otherwise.

It is not difficult to show that the Metropolis–Hastings chain has equilibrium distribution given by $p(\mathbf{x})$. Note that the move is made in block for all model parameters. In practice, with highly dimensional models it is very difficult to find suitable kernels q^* for such spaces that ensure large enough acceptance probabilities. A commonly used variation of the algorithm incorporates the blocking strategy used in the Gibbs sampler and performs moves componentwise by defining transition kernels q_i^* , for $i = 1, \dots, p$. A transition is then completed after cycling through all p components of \mathbf{x} and the generations of the components are made according to the Metropolis–Hastings scheme.

Whatever the scheme used to generate the chain, a stream of values $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$ is formed. Although consecutive values $\mathbf{x}^{(j)}$ and $\mathbf{x}^{(j+1)}$ are correlated, a random sample of size n from $p(\mathbf{x})$ can be formed by retaining n successive values after convergence has been ascertained. If approximately independent observations are required one might hold only the n observations lagged by l units, for example $\mathbf{x}^{(m)}, \mathbf{x}^{(m+l)}, \dots, \mathbf{x}^{(m+(n-1)l)}$, where m is large enough to ensure convergence has been achieved and l is large enough to carry only residual correlation over the chain. This will be a random sample of n approximately iid elements of the joint distribution $p(\mathbf{x})$. This sample is valid for any positive value of m and l ; in particular, $l = 1$ is a common choice since independence is not really required.

After generating a large random sample the inference about each x_i can be done as in any Monte Carlo method. For example, the mean of the i th component of \mathbf{x} is estimated by $(1/n) \sum_{k=0}^{n-1} X_i^{(m+kl)}$. This idea can be applied in the Bayesian context to obtain a sample from the posterior distribution of a parametric vector θ or in the frequentist context to obtain a sample from the sampling distribution of an estimator or of a test statistic $T(\mathbf{X})$. Nevertheless, most of the work and applications in the area are geared towards the Bayesian approach.

Estimates using the known conditional distribution can also be obtained. In the case of the mean of X_i , the Rao–Blackwellized estimator of $E_p(X_i)$ is

$$\hat{E}_p(X_i) = \frac{1}{n} \sum_{k=0}^{n-1} E_p(X_i | \mathbf{X}_{-i}^{(m+kl)}).$$

This estimator of $E_p(X_i)$ is usually better than \bar{X}_i , which is based only on the generated values. The improvement is justified by a more efficient use of the (probabilistic) information available. A very similar idea was used in the Rao–Blackwell theorem (see Section 4.3.2) to prove that conditioning of sufficient

statistics improves the estimator, hence the name. It is worth pointing out that inferences about any quantity related to the X 's are easily done. For example the mean value of $g(\mathbf{X})$, given by $E_p[g(\mathbf{X})]$, is estimated as $(1/n) \sum_{k=0}^{n-1} g(\mathbf{x}^{(m+kl)})$, where the $\mathbf{x}^{(j)}, \forall j \geq 1$, are the values generated from the chain.

There are many practical problems that are easily handled by the combination of Bayesian methods and Gibbs sampling or some other MCMC methods, but are difficult to handle by other means.

Example. Let X_1, \dots, X_n be a random sample from the $N(\theta, \sigma^2)$ distribution and assume independent prior distributions $\theta \sim N(\mu_0, \tau_0^2)$ and $\phi \sim G(n_0/2, n_0\sigma_0^2/2)$. Note that this distribution is different from the usual conjugate prior used so far in this book but may be a suitable representation of the prior knowledge in some situations. Then the joint posterior is

$$\begin{aligned} p(\theta, \phi | \mathbf{x}) &\propto l(\theta, \sigma^2; \mathbf{x}) p(\theta) p(\phi) \\ &\propto \phi^{n/2} \exp \left\{ -\frac{\phi}{2} [ns^2 + n(\bar{x} - \theta)^2] \right\} \\ &\quad \times \exp \left\{ -\frac{1}{2\tau_0^2} (\theta - \mu_0)^2 \right\} \phi^{n_0/2-1} \exp \left(-\frac{n_0\sigma_0^2\phi}{2} \right) \\ &\propto \phi^{[(n+n_0)/2]-1} \\ &\quad \times \exp \left\{ -\frac{1}{2} [\phi(n_0\sigma_0^2 + ns^2 + n(\theta - \bar{x})^2) + \tau_0^{-2}(\theta - \mu_0)^2] \right\}. \end{aligned}$$

This distribution has no known form and it is not possible to perform the analytic integration to obtain the proportionality constant. The kernel of the marginal distributions can be obtained but are of no known form which prevents the exact evaluation of their mean, variance and so forth. Nevertheless, the posterior conditional distributions of $\theta|\phi$ and $\phi|\theta$ are easy to obtain. They are given by the posterior density once the terms that do not depend on the quantity of interest are incorporated into the proportionality constant. So,

$$\begin{aligned} p(\theta|\phi, \mathbf{x}) &\propto \exp \left\{ -\frac{1}{2} [n\phi(\theta - \bar{x})^2 + \tau_0^{-2}(\theta - \mu_0)^2] \right\} \\ &\propto \exp \left\{ -\frac{1}{2} [\theta^2(\tau_0^{-2} + n\phi) - 2\theta(\tau_0^{-2}\mu_0 + n\phi\bar{x})] \right\} \\ p(\phi|\theta, \mathbf{x}) &\propto \phi^{[(n+n_0)/2]-1} \exp \left\{ -\frac{\phi}{2} [n_0\sigma_0^2 + ns^2 + n(\theta - \bar{x})^2] \right\} \end{aligned}$$

from which it is clear that $(\theta|\phi, \mathbf{x}) \sim N[\mu_1(\phi), \tau_1^2(\phi)]$ and $(\phi|\theta, \mathbf{x}) \sim G(n_1/2, n_1\sigma_1^2(\theta)/2)$ where

$$\mu_1(\phi) = \frac{\tau_0^{-2}\mu_0 + n\phi\bar{x}}{\tau_0^{-2} + n\phi}, \quad \tau_1^2(\phi) = \frac{1}{\tau_0^{-2} + n\phi}, \quad n_1 = n_0 + n$$

and $n_1\sigma_1^2(\theta) = n_0\sigma_0^2 + ns^2 + n(\theta - \bar{x})$. Therefore, it is easy to generate from the conditionals and the Gibbs sampler becomes easy to implement.

One difficult problem in the applications of MCMC schemes is to ensure the convergence of the chain. Some theoretical results and many diagnostic statistics are available. The practical recommendation is to monitor the trajectory of the chain using output diagnostics. A common approach is to plot the averages of selected quantities such as the components of the vector \mathbf{X} and assess by visual inspection whether the convergence has occurred. More formal diagnostic tools have already been derived and should also be used in addition to visual inspection. Convergence of the chain can be slow for many reasons. For example, if the components of \mathbf{x} are highly correlated, or if the joint density has multiple modes with regions of low probability between some of them, then the chain may take a large number of iterations to converge. The interested reader is referred to the books by Gamerman (1997) and Gilks et al. (1996) for some theoretical and practical aspects of MCMC methodology.

Exercises

§ 5.2

- Consider the genetic application of Section 5.2 where a four-dimensional vector of counts $\mathbf{X} = (X_1, X_2, X_3, X_4)$ has multinomial distribution with parameters n and π , where $\pi = (1/2 + \theta/4, (1 - \theta)/4, (1 - \theta)/4, \theta/4)$. Assume that the observed data was (125, 18, 20, 34).
 - Obtain the equations required for calculation of the MLE via the Newton–Raphson algorithm and apply it to obtain the maximum likelihood estimate for the given data set.
 - Obtain the equations required for calculation of the MLE via the Fisher scoring algorithm and apply it to obtain the maximum likelihood estimate for the given data set.
 - Use the expressions given for the successive iterates in the EM algorithm to show that the likelihood is monotonically increasing through the steps of the algorithm.
 - Compare the three different algorithms for finding the MLE in terms of computational complexity and time.
 - Assume now a prior $\theta \sim \text{beta}(a, b)$. Repeat the exercise to obtain the (generalized) MLE. Specify numerical values for a and b and obtain the corresponding posterior modes.
- Consider a random sample X_1, \dots, X_n from the $G(\alpha, \beta)$ distribution with both parameters unknown. Obtain the maximum likelihood equations and describe the use of an iterative scheme to obtain the MLE of α and β .
- Consider the randomized response example. Show that the MLE of θ can be calculated directly using invariance properties of the MLE and evaluate its value with the figures provided in the example.

- Consider the EM algorithm with a sequence of iterated values $\theta^{(j)}$, $j \geq 1$. Show that the sequence satisfies $L(\theta^{(j)}|\mathbf{x}) \leq L(\theta^{(j+1)}|\mathbf{x})$ and is therefore monotonically increasing in the likelihood $l(\theta|\mathbf{x})$.

§ 5.3

- Let $\mathbf{X}_n = (X_1, \dots, X_n)$ be a random sample from the $N(0, \theta^2)$ distribution.
 - Obtain the asymptotic posterior distribution of θ as $n \rightarrow \infty$.
 - Obtain the asymptotic posterior mean and variance of θ^2 .
Hint: $X \sim N(0, 1) \Rightarrow X^2 \sim \chi_1^2$.
 - Obtain the asymptotic distribution of θ^2 based on the delta method and compare it with the results obtained in (b).
- Let $X \sim \text{bin}(20, \theta)$ and assume that $X = 7$ was observed. Obtain a 90% confidence interval for θ using a uniform prior and
 - the fact that if $z \sim \text{beta}(a, b)$ then

$$\frac{b}{a} \frac{z}{1-z} \sim F(2a, 2b);$$
 - an asymptotic approximation for $\psi = \theta/1 - \theta$;
 - an asymptotic approximation for $\phi = \sin^{-1}(\sqrt{\theta})$.
 - Compare the results.
- Let $\mathbf{X}_n = (X_1, \dots, X_n)$ be a vector of independent random variables where $X_i \sim \text{Pois}(\theta t_i)$, $i = 1, \dots, n$, and t_1, \dots, t_n are known times.
 - Prove that the MLE of θ is $\hat{\theta} = \bar{X}/\bar{t}$ where $\bar{X} = \sum_{i=1}^n X_i/n$ and $\bar{t} = \sum_{i=1}^n t_i/n$.
 - Obtain the asymptotic posterior distribution of $\theta | \mathbf{x}_n |$ and construct an asymptotic $100(1 - \alpha)\%$ confidence interval for θ assuming that n is large.
 - Obtain the asymptotic posterior distribution of $\theta^{1/2} | \mathbf{x}_n |$ and, based on it, construct an asymptotic $100(1 - \alpha)\%$ confidence interval for θ assuming that n is large.
 - Compare the confidence intervals obtained in (b) and (c), considering especially their lengths.
- Let X_1, \dots, X_n be a random sample from the distribution with density

$$f(x | \theta) = \theta x^{\theta-1} I_x([0, 1]).$$
 - Verify which function(s) of θ (up to linear transformations) can be estimated with highest efficiency and determine its (their) corresponding estimator(s).
 - Obtain the asymptotic $100(1 - \alpha)\%$ confidence interval for θ based on approximations for the posterior distribution of θ .

- (c) Repeat item (b) basing calculations now on the asymptotic distribution of the score function $U(\mathbf{X}; \theta)$.
- (d) Repeat item (b) basing calculations now on the central limit theorem applied to the sample $\mathbf{X}_n = (X_1, \dots, X_n)$.
9. Let X_1, \dots, X_n be a random sample from the $\text{Pois}(\theta)$ distribution and define $\lambda = \theta^{1/a}$, $a \neq 0$.
- Obtain the likelihood function $l(\lambda; \mathbf{X})$.
 - Obtain the Jeffreys non-informative prior for λ .
 - Obtain the Taylor expansion of $L(\lambda) = \log l(\lambda)$ around the MLE of λ and determine the value(s) of a for which the third-order term vanishes.
 - Discuss the importance of the result obtained in the previous item in terms of asymptotic theory.
10. Let X_1, \dots, X_n be a random sample from the uniform distribution over the interval $[0, \theta]$ and let $\hat{\theta}_n$ be the MLE of θ .
- Obtain a non-degenerate asymptotic distribution for $\hat{\theta}_n$, or in other words, find functions $h(n), a(\theta)$ and $b(\theta)$ and a non-degenerate asymptotic distribution P such that
- $$h(n) \frac{\hat{\theta}_n - a(\theta)}{b(\theta)} \xrightarrow{\mathcal{D}} P \quad \text{when } n \rightarrow \infty.$$
- Hint: use the density of $\hat{\theta}_n$ to obtain the form of h , a and b and use the result $(1 + s/n)^n \rightarrow e^s$ when $n \rightarrow \infty$ for $s \in R$.
- Comment on the convergence rate found.
 - Obtain the asymptotic $100(1 - \alpha)\%$ confidence interval for θ of smallest length based on the results of item (a).
 - Show that the parameter
- $$h(n) \frac{\hat{\theta}_n - a(\theta)}{b(\theta)}$$
- converges in distribution to P where h , a , b and P are the same ones obtained in item (a). Therefore, the asymptotic result obtained with the Bayesian inference is similar to the result obtained with the classical inference.
- Obtain the asymptotic $100(1 - \alpha)\%$ HPD confidence interval for θ .
 - Compare the intervals obtained in items (c) and (e) with the exact interval.
11. Show that for any distribution $p(\theta)$ in the exponential family, the best normal approximation in the Kullback–Leibler sense has mean and variance given respectively by $\mu = E(\theta)$ and $\sigma^2 = V(\theta)$, the mean and variance of the original distribution.

12. Consider again the variation of the randomized response model which consists in asking as the alternative question the negation of the original one.
- Show that the posterior distribution of θ is a mixture of $n + 1$ beta distributions.
 - Obtain the relevant derivatives and points of maxima required for the evaluation of the posterior mean of θ analytically or numerically.
 - Apply the results of the previous items to reproduce the table of exact and approximated posterior means for θ given in the text.

§ 5.4

13. Apply the Gauss–Hermite integration rules to obtain approximations for the posterior expectation of α and θ given the observed values already provided in the Weibull example of Section 5.3, and compare the results with the approximations from the Laplace method.

§ 5.5

14. Use the simple Monte Carlo method to evaluate $\int_{-\infty}^{\infty} e^{-x^2/2} dx$ and compare it with the known answer $\sqrt{2\pi}$. Also, evaluate the variance of the estimator. Hint: make a transformation to take the line into the interval $[0, 1]$ and then proceed as before.
15. Show that if an integral $I = \int g(x)p(x)dx$ is estimated by importance sampling then its estimator

$$\bar{I} = \frac{1}{n} \sum_{i=1}^n g(x_i)w(x_i),$$

where

$$w(x_i) = \frac{p(x_i)}{h(x_i)}$$

and

$$x_i \sim h(x), i = 1, \dots, n,$$

is unbiased and has variance given by $V(\bar{I}) = (1/n) \int (g(x)w(x) - I)^2 h(x)dx$.

16. Let $\theta = P(X > 2)$ where X has a standard Cauchy distribution with density

$$p(x) = \frac{1}{\pi(1+x^2)}, \quad x \in R.$$

Let h be an importance sampling density defined by

$$h(x) = 2I_x[(2, \infty)]/x^2.$$

Show that use of this sampling density reduces the variance of the estimator of θ over the simple Monte Carlo estimator.

164 Approximate and computationally intensive methods

17. Show that the Rao–Blackwellized estimator of $E_p(X_i)$ given by

$$\hat{E}_p(X_i) = \frac{1}{n} \sum_{k=0}^{n-1} E_p(X_i | \mathbf{X}_{-i}^{(m+kl)})$$

provides an unbiased and consistent estimator of $E_p(X_i)$. Generalize the result to obtain the Rao–Blackwellized estimator of the marginal density of X_i and show that it is also an unbiased and consistent estimator.

18. Let X_1, \dots, X_n be a random sample from a Poisson distribution with mean that is either θ or ϕ . The mean is θ up to an unknown break point m from where it becomes ϕ .

- (a) Obtain the likelihood of the unknown parameters θ, ϕ and m .
- (b) Suggest a reasonable family of conjugate prior distributions for θ, ϕ and m .
- ... Hint: to simplify matters, assume independent priors for θ, ϕ and m .
- (c) Obtain the full conditional distributions required for implementation of the Gibbs sampler.
- (d) Generate data (X_1, \dots, X_n) for given values of θ, ϕ and m and apply the Gibbs sampler to draw inference about them.

19. (Casella and George, 1992) Let π denote the following discrete distribution over $S = \{0, 1\}^2$

		X_2	
		0	1
X_1	0	π_{00}	π_{01}
	1	π_{10}	π_{11}

where $\pi_{00} + \pi_{01} + \pi_{10} + \pi_{11} = 1$ and $\pi_{ij} > 0$, for $i, j = 0, 1$. Assume that instead of drawing samples directly from π , one decides to draw values from π through the Gibbs sampler.

- (a) Show that the transition probabilities for X_1 are given by the conditional distribution π_1 of $X_1 | X_2 = j$,

$$\pi_1(0|j) = \frac{\pi_{0j}}{\pi_{+j}} \quad \text{and} \quad \pi_1(1|j) = \frac{\pi_{1j}}{\pi_{+j}}$$

where $\pi_{+j} = \pi_{0j} + \pi_{1j}$, $j = 0, 1$.

- (b) Show that the transition probabilities for X_2 are given by the conditional distribution π_2 of $X_2 | X_1 = i$,

$$\pi_2(0|i) = \frac{\pi_{i0}}{\pi_{i+}} \quad \text{and} \quad \pi_2(1|i) = \frac{\pi_{i1}}{\pi_{i+}}$$

where $\pi_{i+} = \pi_{i0} + \pi_{i1}$, $i = 0, 1$.

Exercises 165

- (c) Show that the 4×4 transition matrix P of the chain formed by the Gibbs sampler has elements

$$\begin{aligned} P((i, j), (k, l)) &= Pr((X_1, X_2)^{(n)} \\ &= (k, l) | (X_1, X_2)^{(n-1)} = (i, j)) \\ &= \frac{\pi_{kl}}{\pi_{k+} \pi_{+j}} \end{aligned}$$

for $(i, j), (k, l) \in S$.

- (d) Show that π is the only stationary distribution of this chain.
- (e) Extend the results for cases when X_1 can take n_1 values and X_2 can take n_2 values.

20. Show that the Metropolis–Hastings chain has equilibrium distribution given by $p(\mathbf{x})$.

6

Hypothesis testing

6.1 Introduction

In this chapter we still consider statistical problems involving an unknown quantity θ belonging to a parametric space Θ . In many instances, the inferential process may be summarized in the verification of some assertions or conjectures about θ . For example, one may be interested in verifying whether a coin is fair, a collection of quantities is independent or if distinct populations are probabilistically equal. Each one of the assertions above constitutes an hypothesis and can be associated with a model. This means here that it can be parametrized in some form. Considering the simple case of two alternative hypotheses, two disjoint subsets Θ_0 and Θ_1 belonging to Θ are formed.

Denote by H_0 the hypothesis that $\theta \in \Theta_0$ and by H_1 the hypothesis that $\theta \in \Theta_1$. A new statistical problem is to decide whether H_0 or H_1 is accepted, or in other words, whether θ is in Θ_0 or Θ_1 . If the subset of the parameter space defining an hypothesis contains a single element, the hypothesis is said to be simple. Otherwise, it is said to be composite. Under a simple hypothesis, the observational distribution is completely specified whereas under a composite hypothesis it is only specified that the observational distribution belongs to a family. From now on, the hypotheses H_0 and H_1 will be uniquely associated with disjoint subsets Θ_0 and Θ_1 of the parameter space. Whenever they are simple, the notation $\Theta_0 = \{\theta_0\}$ and $\Theta_1 = \{\theta_1\}$ will be used. Note that in some cases, only one of the hypotheses is simple.

Typically, a test of hypotheses is a decision problem with a number of possible actions. If the researcher makes the wrong decision he incurs a penalty or suffers a loss. Once again, his/her objective is to minimize his/her loss in some form. For example, under the Bayesian approach he/she would try to minimize the expected loss. A rule to decide the hypothesis to be accepted is called a test procedure or simply a test and will be denoted by ψ . One may define for example that $\psi = i$ if the hypothesis accepted is H_i .

From the Bayesian perspective, one may have many alternative hypotheses H_1, \dots, H_k that can be compared through $P(H_i \mid \mathbf{x})$, $i = 1, \dots, k$. Under the classical perspective, it is important to have only two hypotheses H_0 and H_1 .

168 Hypothesis testing

The theory of this chapter is developed for this case in order to ease comparisons between the two approaches.

Usually there is an hypothesis that is more important. This will be denoted by H_0 and called the null hypothesis. The other hypothesis is denoted by H_1 and called the alternative hypothesis. In general, H_0 and H_1 are mutually exclusive which means that at most one of the hypotheses is true. If, in addition, H_0 and H_1 exhaust all possibilities then necessarily one of them must be true. In this case, rejecting one of them necessarily implies accepting the other one.

In this chapter, we start the presentation of the classical procedures and later present the Bayesian procedure. We then move on to establish a connection between hypothesis tests and confidence intervals. Finally, tests based on asymptotic theory results are described from both classical and Bayesian perspectives. A systematic introductory account of this topic is presented in Bickel and Doksum (1977) and DeGroot (1970). The classical theory is presented at a more formal level in Lehmann (1986).

Assume that before deciding which hypothesis to accept, the statistician is offered the choice of observing a sample X_1, \dots, X_n from a distribution that depends on the unknown parameter θ . In a problem of this kind, the statistician can specify a test procedure by splitting the sample space into two subsets. Under the frequentist viewpoint, this is the only available option as the only source of information comes from the data. One subset of the sample space will contain the values of X that will lead to acceptance of H_0 and the other one will lead to rejection of H_0 . This latter set is called the critical region and a test procedure gets completely specified by the critical region. Of course, the complement of the critical region contains sample results that lead to the acceptance of H_0 .

6.2 Classical hypothesis testing

The general theory of classical hypothesis testing comes from the pioneering work of Neyman and Pearson (1928). The probabilistic characteristics of a classical test can be described by specification of $\pi(\theta)$, the probability that the test leads to the rejection of H_0 , for each value of $\theta \in \Theta$. The function π is called the power of the test. If C is the critical region then π is defined by

$$\pi(\theta) = P(X \in C \mid \theta), \quad \forall \theta \in \Theta.$$

Some textbooks define the power function only for $\theta \notin \Theta_0$. The size or significance level α of a test procedure is defined as

$$\alpha \geq \sup_{\theta \in \Theta_0} \pi(\theta).$$

Just as in the case of confidence levels seen in Section 4.4, the inequality above is a technical requirement. It is more useful in discrete sample spaces where not all

values in $[0, 1]$ are possible probabilities. As will shortly be seen, one wishes to use as small a value for α as possible. In practice, this means that an equality is used.

For any given test procedure ψ , two types of error can be committed. A type I error is committed when the test indicates rejection of H_0 when it is true. Note that the largest possible value for the probability of this error is α . Similarly, the type II error is committed when the test indicates acceptance of H_0 when it is false. The probability of a type II error is usually denoted by β . Note that $\beta(\theta) = 1 - \pi(\theta)$, for $\theta \in \Theta_1$. In the case of simple hypotheses, the probability of a type I error is $\alpha = \pi(\theta_0)$ and the probability of a type II error is $\beta = 1 - \pi(\theta_1)$.

6.2.1 Simple hypotheses

It is useful to start the study of the theory with the case of two simple hypotheses $H_0 : \theta = \theta_0$ and $H_1 : \theta = \theta_1$. Ideally, one wishes to find a test procedure for which the two error probabilities are as small as possible. In practice, it is impossible to find a test for which these probabilities are simultaneously minimized. As an alternative, one may seek to construct a test that minimizes linear combinations of α and β .

Theorem (optimal test). Assume that $X = (X_1, \dots, X_n)$ is a random sample from $p(x|\theta)$, $H_0 : \theta = \theta_0$ and $H_1 : \theta = \theta_1$. Let ψ^* be a test of H_0 versus H_1 such that H_0 is accepted if $p_0/p_1 > k$ and H_0 is rejected if $p_0/p_1 < k$, where $p_i = p(x \mid \theta_i)$, $i = 0, 1$ and $k > 0$. (If $p_0/p_1 = k$, nothing can be decided.) Then, any other test ψ will be such that

$$a\alpha(\psi^*) + b\beta(\psi^*) \leq a\alpha(\psi) + b\beta(\psi),$$

where $\alpha(\psi)$ and $\beta(\psi)$ respectively denote the probabilities of errors of type I and II of test ψ , for any $a, b \in R^+$.

Proof. Let C be the critical region of any arbitrary test ψ and define $p_i = p(x \mid \theta_i)$, $i = 0, 1$. Then, for $a, b \in R^+$,

$$\begin{aligned} a\alpha(\psi) + b\beta(\psi) &= a \int_C p(x \mid \theta_0) dx + b \int_{C^c} p(x \mid \theta_1) dx \\ &= a \int_C p(x \mid \theta_0) dx + b \left[1 - \int_C p(x \mid \theta_1) dx \right] \\ &= b + \int_C (ap_0 - bp_1) dx. \end{aligned}$$

So, minimization of $a\alpha(\psi) + b\beta(\psi)$ is equivalent to choosing the critical region C in such a way that the value of the integral be minimal. This will occur if the integration is performed over a set that includes every point x such that $ap_0 - bp_1 < 0$ and does not include points x such that $ap_0 - bp_1 > 0$.

Therefore, minimization of $a\alpha(\psi) + b\beta(\psi)$ is achieved by having the critical region \mathbf{C} including only points \mathbf{x} such that $ap_0 - bp_1 < 0$. (If the sampling distribution is continuous and $ap_0 - bp_1 = 0$, this point has 0 contribution and is irrelevant.) This completes the demonstration because $ap_0 - bp_1 < 0$ iff $p_0/p_1 < k = b/a$, which corresponds to the description of the test ψ^* . \square

The ratio p_0/p_1 is called the likelihood ratio (LR). The theorem establishes that a test that minimizes $a\alpha(\psi) + b\beta(\psi)$, rejects H_0 when the LR is small and accepts H_0 when the LR is large. Usually, the null hypothesis H_0 and error of type I are privileged. Therefore, one considers only tests ψ such that $\alpha(\psi)$ cannot be larger than a pre-specified level α_0 and, among them search for the one that minimizes $\beta(\psi)$. This is a variation of the problem solved with the theorem and the solution is provided by the following lemma.

Lemma (Neyman–Pearson). Assume that $\mathbf{X} = (X_1, \dots, X_n)$ is a random sample from $p(x|\theta)$, $H_0 : \theta = \theta_0$ and $H_1 : \theta = \theta_1$. Let ψ^* be a test of H_0 versus H_1 such that H_0 is accepted if $p_0/p_1 > k$ and H_0 is rejected if $p_0/p_1 < k$, where $p_i = p(\mathbf{x} | \theta_i)$, $i = 0, 1$. (If $p_0/p_1 = k$, nothing can be decided.) Then, for any other test ψ such that $\alpha(\psi) \leq \alpha(\psi^*)$, $\beta(\psi) \geq \beta(\psi^*)$. Also, $\alpha(\psi) < \alpha(\psi^*)$ implies $\beta(\psi) > \beta(\psi^*)$.

Proof. Following the definition of the optimal test ψ^* in the theorem, it follows that for any other test ψ

$$\alpha(\psi^*) + k\beta(\psi^*) \leq \alpha(\psi) + k\beta(\psi),$$

for $k > 0$. If $\alpha(\psi) \leq \alpha(\psi^*)$ then necessarily $\beta(\psi) \geq \beta(\psi^*)$. Also, if $\alpha(\psi) < \alpha(\psi^*)$, it follows that $\beta(\psi) > \beta(\psi^*)$, completing the demonstration. \square

Special attention must be given to the wording of the lemma. It only considers acceptance or rejection of H_0 with no reference to H_1 . This is consistent with the preferential status given to H_0 and also to the fact that H_0 and H_1 do not exhaust the parameter space. This point will be readdressed below.

In the lemma, α_0 plays the role of significance level. Recalling that $\pi(\theta_1) = 1 - \beta(\psi)$, minimization of β implies maximization of π . Hence, the Neyman–Pearson lemma shows that of all tests with a given significance level, that based on the LR has largest power or is more powerful.

Example. In the $N(\theta, \sigma^2)$ with known σ^2 , consider the test of $H_0: \theta = \theta_0$ versus $H_1: \theta = \theta_1$ with $\theta_0 < \theta_1$. Then

$$\frac{p_0}{p_1} = \frac{p(\mathbf{x}|\theta_0)}{p(\mathbf{x}|\theta_1)} = \frac{(2\pi\sigma^2)^{-n/2} \exp\left\{-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \theta_0)^2\right\}}{(2\pi\sigma^2)^{-n/2} \exp\left\{-\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \theta_1)^2\right\}}$$

$$\begin{aligned} &= \exp\left\{\frac{1}{2\sigma^2} \left[-2\theta_1 \sum_{i=1}^n x_i + n\theta_1^2 + 2\theta_0 \sum_{i=1}^n x_i - n\theta_0^2 \right]\right\} \\ &= \exp\left\{\frac{1}{2\sigma^2} \left[2(\theta_0 - \theta_1) \sum_{i=1}^n x_i \right]\right\} \exp\left\{\frac{1}{2\sigma^2} n(\theta_1^2 - \theta_0^2)\right\} \\ &\propto \exp\left\{\frac{n(\theta_0 - \theta_1)\bar{x}}{\sigma^2}\right\} \end{aligned}$$

where the proportionality constant involves the constants θ_0 , θ_1 and σ^2 .

The LR test accepts H_0 when $p_0/p_1 > k$. Then,

$$\frac{p_0}{p_1} > k \iff \exp\left\{\frac{n(\theta_0 - \theta_1)\bar{x}}{\sigma^2}\right\} > c_3 \iff \frac{n(\theta_0 - \theta_1)\bar{x}}{\sigma^2} > c_2 \iff \bar{x} < c_1$$

since $\theta_0 < \theta_1$, for constants c_1 , c_2 and c_3 . As the best estimator of θ is \bar{X} , the sample mean, when testing H_0 against H_1 , one expects the test to accept H_0 for small values of \bar{X} . That is exactly the result of the optimal, LR test. The next step is to determine the value of c_1 . To do this, note that the test has level α and therefore $\alpha = P(\text{rejection of } H_0 | \theta = \theta_0) = P(\bar{X} > c_1 | \theta = \theta_0)$. But

$$\bar{X} | \theta_0 \sim N\left(\theta_0, \frac{\sigma^2}{n}\right) \quad \text{or} \quad Z = \frac{\bar{X} - \theta_0}{\sigma/\sqrt{n}} \sim N(0, 1).$$

Since $\alpha = P(Z > (c_1 - \theta_0)\sqrt{n}/\sigma)$,

$$\frac{(c_1 - \theta_0)}{\sigma}\sqrt{n} = z_\alpha \Rightarrow c_1 = \theta_0 + z_\alpha \frac{\sigma}{\sqrt{n}}.$$

The test with significance level α accepts H_0 if $\bar{X} < \theta_0 + \sigma z_\alpha / \sqrt{n}$.

Note that the test is completely specified and does not depend on the value of θ_1 but for the fact that $\theta_1 > \theta_0$. This means that the test is the same for any value of θ_1 such that $\theta_1 > \theta_0$. Therefore the LR test is also more powerful to test H_0 versus $H_1: \theta_1 > \theta_0$.

The fact that the test does not depend on the value of θ_1 may also cause problems. Consider the case when H_0 is rejected but \bar{x} is much closer to θ_0 than to θ_1 , that is, $\bar{x} - \theta_0 \ll \theta_1 - \bar{x}$. In this case, common sense suggests that given the choices of H_0 and H_1 one should choose H_0 . This is another reason to avoid commitments towards acceptance or rejection of H_1 since the test does not provide information about it. Similar comments apply if the distances of \bar{x} to θ_0 and θ_1 are much larger than the distance between θ_0 and θ_1 . The intuitive reasoning based on the frequentist argument is that if sampling of \mathbf{X} is repeated many times, in only 100 α % of them, H_0 will be erroneously rejected.

The power of the test $\pi(\theta) = P(\text{rejection of } H_0 | \theta)$ is given by

$$\pi(\theta) = P\left[\bar{X} > \theta_0 + \frac{\sigma}{\sqrt{n}}z_\alpha | \theta > \theta_0\right].$$

But $\bar{X} | \theta > \theta_0 \sim N(\theta, \sigma^2/n)$ and therefore $\sqrt{n}(\bar{X} - \theta)/\sigma \sim N(0, 1)$. Then,

$$\begin{aligned}\pi(\theta) &= P\left[\sqrt{n}\frac{(\bar{X} - \theta)}{\sigma} > \frac{\theta_0 + (\sigma/\sqrt{n})z_\alpha - \theta}{\sigma/\sqrt{n}} \mid \theta > \theta_0\right] \\ &= P\left[\sqrt{n}\frac{(\bar{X} - \theta)}{\sigma} > \frac{(\theta_0 - \theta)}{\sigma}\sqrt{n} + z_\alpha \mid \theta > \theta_0\right] \\ &= 1 - \Phi\left(z_\alpha - \frac{(\theta - \theta_0)}{\sigma}\sqrt{n}\right),\end{aligned}$$

which is an increasing function of θ . So, the more distant is the parameter value in the alternative, the smaller are the chances of a type II error.

This test is not only the most powerful test of H_0 versus $H_1: \theta > \theta_0$, but also to test $H_0: \theta \leq \theta_0$ versus $H_1: \theta > \theta_1$ because the level of the test with the new hypothesis is

$$\max_{\theta \leq \theta_0} \pi(\theta) = \max_{\theta \leq \theta_0} P\left[\bar{X} > \theta_0 + \frac{\sigma}{\sqrt{n}}z_\alpha\right].$$

As just seen, π is an increasing function of θ and the maximum in the region $\{\theta : \theta \leq \theta_0\}$ is given at the value θ_0 and the value of π at this point is $1 - \Phi(z_\alpha) = \alpha$.

Also, in the example we could evaluate the size at which we would reject H_0 after observing $\mathbf{X} = \mathbf{x}$, that is, we could evaluate γ such that $\theta_0 - \sigma z_\gamma/\sqrt{n} = \bar{x}$. This gives a more precise account of the strength of the data evidence in favour of or against the hypotheses. Smaller values of γ indicate a lower probability or type I error and therefore more evidence in favour of H_0 . Likewise, larger values of γ indicate a higher probability or type I error and therefore more evidence against H_0 .

More generally, suppose we have a test where H_0 is rejected when a test statistic T belongs to a region of the form $[T > c]$ and let t be the observed value of T . Then, evaluation of $Pr(T > t | H_0)$ gives an idea of how extreme the observed value is under H_0 . This probability is usually known as the *p*-value. In the previous example, the *p*-value is given by $1 - \Phi(\sqrt{n}(\bar{x} - \theta_0)/\sigma)$. The notion of a *p*-value is useful for determining the size at which one would reject H_0 based on the information actually obtained for

$$H_0 \text{ is rejected} \iff p\text{-value} < \alpha,$$

where α is a pre-specified level of the test. It should be stressed that under the frequentist treatment, no probabilities can be associated to the hypothesis as θ is not random. Therefore, no association between the *p*-value and the probability of H_0 can be made because such a probability simply cannot be defined. The notion of *p*-value can be placed in a general setting whenever it makes sense to specify the border of the critical region in terms of observed values of the sample.

Returning now to the case of discrete populations, it is not always possible to obtain tests of any pre-specified level exactly. By exact, we mean to have $\alpha = \eta$

where $\eta = \sup_{\theta \in \Theta_0} P(\text{rejection of } H_0 | \theta \in \Theta_0)$. The notion of *p*-value becomes even more important here.

There is an alternative approach that allows one to obtain tests of an exact level even for discrete distributions. This alternative is known as randomized tests where any pre-specified level is obtained after realization of an additional independent Bernoulli experiment with success probability conveniently chosen to complete the difference between α and η .

6.2.2 Composite hypotheses

Consider again the test ψ of $H_0: \theta \in \Theta_0$ versus $H_1: \theta \in \Theta_1$. Let α be the fixed significance level and $\pi_\psi(\theta)$ the power function of ψ . Then, $\pi_\psi(\theta) \leq \alpha$ for every $\theta \in \Theta_0$.

Definition. A test ψ^* is uniformly more powerful (UMP, in short) for $H_0: \theta \in \Theta_0$ versus $H_1: \theta \in \Theta_1$ at the significance level α if

1. $\alpha(\psi^*) \leq \alpha$;
2. $\forall \psi$ with $\alpha(\psi) \leq \alpha$, $\pi_\psi(\theta) \leq \pi_{\psi^*}(\theta)$, $\forall \theta \in \Theta_1$.

The test of the example above is UMP to test $\theta = \theta_0$ versus $\theta > \theta_0$ and also to test $\theta \leq \theta_0$ versus $\theta > \theta_0$.

Theorem 6.1. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from $p(x|\theta)$ and $p(x|\theta)$ belong to the one-parameter exponential family with density

$$p(\mathbf{x} | \theta) = a(\mathbf{x}) \exp\{\phi(\theta)T(\mathbf{x}) + b(\theta)\}$$

and let ϕ be a strictly increasing function of θ . Then the UMP test of level α to test $H_0: \theta \leq \theta_0$ versus $H_1: \theta > \theta_0$ is given by the critical region $T(\mathbf{X}) > c$ where c is such that $\alpha \geq P(T(\mathbf{X}) > c | \theta_0)$ (with equality in the continuous case). The power of this test is an increasing function of θ . If the hypotheses are interchanged or ϕ is a strictly decreasing function of θ , then the UMP test of level α rejects H_0 if $T(\mathbf{X}) < c$ where c is such that $\alpha \geq P(T(\mathbf{X}) < c | \theta_0)$ and the power of this test is again an increasing function of θ . If the two conditions above are simultaneously true, the UMP test remains unaltered.

Proof. Consider the standard case where ϕ is strictly increasing and the hypotheses are $H_0: \theta = \theta_0$ and $H_1: \theta = \theta_1 > \theta_0$. In this case, the Neyman–Pearson lemma ensures that the most powerful test rejects H_0 when

$$\begin{aligned}\frac{p(\mathbf{x}|\theta_0)}{p(\mathbf{x}|\theta_1)} < c_4 &\iff \frac{\exp\{\phi(\theta_0)T(\mathbf{x}) + b(\theta_0)\}}{\exp\{\phi(\theta_1)T(\mathbf{x}) + b(\theta_1)\}} < c_3 \\ &\iff \exp\{\phi(\theta_0) - \phi(\theta_1)\}T(\mathbf{x}) < c_2 \\ &\iff [\phi(\theta_0) - \phi(\theta_1)]T(\mathbf{x}) < c_1 \\ &\iff T(\mathbf{x}) > c\end{aligned}$$

where c_4, \dots, c_1, c are constants such that $\alpha = P(T(\mathbf{X}) > c | \theta_0)$. For the most powerful test, it must be true that $\pi(\theta_0) \leq \pi(\theta_1)$ (see the exercises). So, the power function is an increasing function of θ and

$$\pi(\theta_0) = \sup_{\{\theta: \theta < \theta_0\}} \pi(\theta).$$

So, the test is UMP for $H_0 : \theta \leq \theta_0$. As in the calculation above only the condition $\theta_1 > \theta_0$ was used, so the results must be equally true for any such value of θ_1 . Therefore, the test is UMP for $H_1 : \theta > \theta_0$.

In the case of a strictly decreasing ϕ ,

$$[\phi(\theta_0) - \phi(\theta_1)]T(\mathbf{x}) < c_1 \iff T(\mathbf{x}) < c$$

and the critical region becomes $\{\mathbf{x} : T(\mathbf{x}) < c\}$.

In the case of interchanged hypotheses, all inequalities must be reversed because one must work with the ratio $p(\mathbf{x}|\theta_1)/p(\mathbf{x}|\theta_0)$ instead of $p(\mathbf{x}|\theta_0)/p(\mathbf{x}|\theta_1)$. This leads to the critical region in the form $\{\mathbf{x} : T(\mathbf{x}) < c\}$.

Finally, in the case of a strictly decreasing ϕ and interchanged hypotheses, double reversal of the inequalities preserves it as it was and the critical region remains in the form $\{\mathbf{x} : T(\mathbf{x}) < c\}$. \square

Example. Let X_1, \dots, X_n be a random sample from the $\text{Ber}(\theta)$ distribution. From Section 2.5, we know that $\phi(\theta) = \log[\theta/(1-\theta)]$ which is an increasing function of θ and that $T(\mathbf{X}) = \sum_{i=1}^n X_i$. Therefore, the UMP test for $H_0 : \theta \leq \theta_0$ versus $H_1 : \theta > \theta_0$ has critical region of the form $\sum_{i=1}^n X_i > c$.

The property that guarantees the existence of UMP tests in the exponential family is in fact more general. It finds an appropriate setting under families with monotone LR.

Definition. The family of distributions $\{p(\mathbf{x} | \theta), \theta \in \Theta\}$ is said to have monotone likelihood ratio if there is a statistic $T(\mathbf{X})$ such that $\forall \theta_1, \theta_2 \in \Theta$ with $\theta_1 < \theta_2$ the likelihood ratio

$$\frac{p(\mathbf{X} | \theta_2)}{p(\mathbf{X} | \theta_1)}$$

is a monotone function of $T(\mathbf{X})$.

The uniform distribution over the interval $[0, \theta]$ does not belong to the exponential family. Nevertheless, it has monotone LR because the ratio of sampling densities is a monotonically decreasing function of $T(\mathbf{X}) = \max_i X_i$.

The results just proved for exponential families can be extended for families with monotone likelihood ratio. So, if the LR is an increasing function of $T(\mathbf{X})$, then the UMP test of level α for $H_0 : \theta \leq \theta_0$ versus $H_1 : \theta > \theta_0$ is given by the critical

region of the form $T(\mathbf{X}) < c$ where c is such that $\alpha = P(T(\mathbf{X}) < c | \theta_0)$ and the power of this test is an increasing function of θ . Likewise, if the hypotheses are interchanged or the LR is a decreasing function of $T(\mathbf{X})$, the UMP test of level α rejects H_0 if $T(\mathbf{X}) > c$ where c is such that $\alpha \geq P(T(\mathbf{X}) > c | \theta_0)$ and the power of this test is again an increasing function of θ . If the two conditions above are simultaneously true, the UMP test remains unaltered.

These results make intuitive sense. The larger the LR, the more plausible is the value θ_0 relative to θ_1 . If the LR is an increasing function of $T(\mathbf{X})$, the same reasoning is true for $T(\mathbf{X})$. Therefore, a reasonable rejection region for H_0 would be given by small values of $T(\mathbf{X})$.

Example (continued). Let X_1, \dots, X_n be a random sample from the $\text{Ber}(\theta)$ distribution. Then, $p(\mathbf{X} | \theta) = \theta^T(1-\theta)^{n-T}$ with $T = \sum_i X_i$. If $\theta_1 < \theta_2$, the LR is

$$\frac{\theta_2^T(1-\theta_2)^{n-T}}{\theta_1^T(1-\theta_1)^{n-T}} = \left[\frac{\theta_2(1-\theta_1)}{\theta_1(1-\theta_2)} \right]^T \left(\frac{1-\theta_2}{1-\theta_1} \right)^n = \xi^T \eta^n$$

with

$$\xi = \frac{\theta_2(1-\theta_1)}{\theta_1(1-\theta_2)} \quad \text{and} \quad \eta = \frac{1-\theta_2}{1-\theta_1}.$$

Since $1-\theta_1 > 1-\theta_2$, $\xi > 0$, and the LR is increasing in T , confirming results obtained earlier in the example. \square

So far, only one-sided tests have been considered. These are tests where the parametric regions defining the hypotheses are given by a single strict inequality. An example of interest of an hypothesis that is not one-sided is $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$. This test may be useful when comparing two competing treatments.

Assuming now the observation of a sample from the $N(\theta, \sigma^2)$ distribution with known σ^2 , consider the three tests of $H_0: \theta = \theta_0$ below based on \bar{X} :

1. reject H_0 if $|\bar{X} - \theta_0| > 1.645\sigma/\sqrt{n}$;
2. reject H_0 if $|\bar{X} - \theta_0| > 1.282\sigma/\sqrt{n}$;
3. reject H_0 if $|\bar{X} - \theta_0| < 0.126\sigma/\sqrt{n}$.

Calculation of the rejection probability of H_0 shows that the three tests have level 0.1. The next step is to proceed with evaluation of the power of each of the tests. It can be easily seen from Figure 6.1 that none of the tests is UMP over the other two. Nevertheless, the first test is the only one with $\min_{\Theta_1} \pi(\theta) > \pi(\theta_0)$. This means that the rejection probability is larger under the alternative than under the null hypothesis. This way, one guarantees that the chances of rejecting H_0 are larger when H_0 is false. This seems like a reasonable property to require from tests.

Definition. A test ψ for $H_0 : \theta \in \Theta_0$ against $H_1 : \theta \in \Theta_1$ is said to be unbiased if for every pair (θ, θ') where $\theta \in \Theta_0$ and $\theta' \in \Theta_1$, then $\pi_\psi(\theta) \leq \pi_\psi(\theta')$. The power function is at least as large in Θ_1 as it is in Θ_0 . If the test does not satisfy the condition above, it is said to be biased.

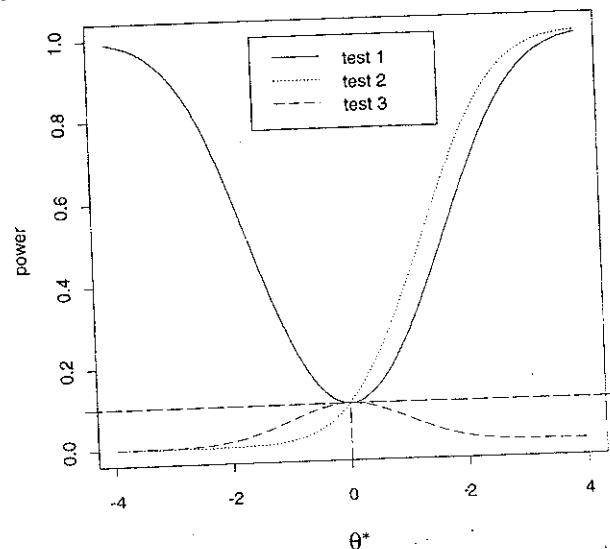


Fig. 6.1 Power functions for tests 1, 2 and 3 as functions of $\theta^* = \sqrt{n}(\theta - \theta_0)/\sigma$.

One can then try to construct UMP tests for $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$ within the class of unbiased tests. In the one-parameter exponential family, it can be shown that if ϕ is a strictly increasing function of θ , the UMP unbiased test of level α for $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$ accepts H_0 when $c_1 < T(\mathbf{X}) < c_2$ with $P(c_1 < T(\mathbf{X}) < c_2 | \theta_0) = 1 - \alpha$ and (c_1, c_2) is an interval of highest sampling density of $T(\mathbf{X})$. The sampling distributions of $T(\mathbf{X})$ are not necessarily symmetric (around θ_0) as in the normal case above. For such tests, we are led to two distinct p -values. In general, the smallest one is used.

It may not be possible in general to find unbiased tests. A general procedure for testing $H_0: \theta \in \Theta_0$ versus $H_1: \theta \in \Theta_1$ is based on the maximum likelihood ratio (MLR, in short) statistic given by

$$\lambda(\mathbf{X}) = \frac{\sup_{\theta \in \Theta_0} p(\mathbf{X} | \theta)}{\sup_{\theta \in \Theta_1} p(\mathbf{X} | \theta)}.$$

The most common case for use of this procedure is when Θ_0 and Θ_1 are exclusive and exhaustive and Θ_0 is of smaller dimension than Θ_1 . In these cases, the denominator is replaced by the supremum over the whole parametric space Θ , which is easier to evaluate. Formally, the statistic $\lambda(\mathbf{X})$ is being replaced by $\max\{\lambda(\mathbf{X}), 1\}$. In any case, $\lambda(\mathbf{X})$ is a random variable depending on the sample. The maximum (or generalized) LR test for H_0 of level α accepts H_0 if $\lambda(\mathbf{X}) > c$ where c satisfies

$$\alpha \geq \sup_{\theta \in \Theta_0} P(\lambda(\mathbf{X}) < c | \theta).$$

Once again, α is taken as equal to the supremum of the above probabilities whenever possible. The power of the test is given by $\pi(\theta) = P(\lambda(\mathbf{X}) < c | \theta)$. This test rejects the null hypothesis H_0 if the maximized value of the likelihood under H_0 is distant from the global maximized value. This is an indication that there is great improvement in likelihood by consideration of points outside H_0 and this hypothesis does not provide a good description of the data. In this case, it makes sense to reject H_0 .

It is important to distinguish between the above test and the test based on monotone likelihood ratios. Although the maximum likelihood ratio test enjoys good asymptotic properties, it is not always unbiased or UMP. The main difficulties associated with it are the calculation of the maximized likelihood in closed form and determination of its sampling distribution. The first point was dealt with in Chapter 5 and the second one will be addressed below when asymptotic tests are treated in Section 6.5.

Other desirable properties in test procedures are similarity and invariance. Consider the problem of testing $H: \theta = \theta_0$ in the presence of a nuisance parameter ϕ . A test is said to be similar if the level of the test is the same whatever the value of the disturbance parameter. An example of a similar test is the t test, to be seen later in this section. A test is said to be invariant (under a specific family of transformations) if the distribution of any transformation of the observations inside the family remains in the same family. This will allow the hypotheses to remain unaltered with any of the transformations operated over the data. The tests presented below in this section are all invariant under linear transformations of the observations.

6.2.3 Hypothesis testing with the normal distribution

The most common tests for samples from a normal distribution are presented here. Once again, let X_1, \dots, X_n be iid with $X_i \sim N(\theta, \sigma^2)$ and suppose one wishes to test $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$. Assume initially that σ^2 is known. In this case, we have shown that the UMP unbiased test of level α is given by the critical region $\sqrt{n}(\bar{X} - \theta_0)/\sigma > z_{\alpha/2}$. We will now obtain the MLR test.

Since $\Theta_0 = \{\theta_0\}$ then

$$\sup_{\theta \in \Theta_0} p(\mathbf{x} | \theta) = p(\mathbf{x} | \theta_0).$$

For $\theta \in \Theta_1 = \Theta - \{\theta_0\}$, the maximum of $p(\mathbf{x} | \theta)$ is obtained at $\hat{\theta}$, the MLE of θ , which in this case is the sample mean \bar{X} . Note that for every value of θ , $P(\hat{\theta} = \theta_0) = 0$ and therefore to consider maximization over Θ instead of over Θ_1 in the expression of the MLR does not produce any change. Then the MLR statistic is given by

$$\lambda(\mathbf{X}) = \frac{p(\mathbf{X} | \theta_0)}{p(\mathbf{X} | \hat{\theta})}$$

$$\begin{aligned}
 &= \frac{(2\pi\sigma^2)^{-n/2} \exp\left\{-[\sum(X_i - \bar{X})^2 + n(\bar{X} - \theta_0)^2]/2\sigma^2\right\}}{(2\pi\sigma^2)^{-n/2} \exp\left\{-\sum(X_i - \bar{X})^2/2\sigma^2\right\}} \\
 &= \exp\left\{-\frac{n}{2\sigma^2}(\bar{X} - \theta_0)^2\right\}.
 \end{aligned}$$

Observe that under H_0 , $\bar{X} \sim N(\theta_0, \sigma^2/n)$ and therefore $Z = \sqrt{n}(\bar{X} - \theta_0)/\sigma \sim N(0, 1)$. The MLR is given by $\lambda(\mathbf{X}) = \exp(-Z^2/2)$. The MLR test rejects H_0 if

$$\lambda(\mathbf{X}) < c_2 \iff Z^2/2 > c_1 \iff |Z| > c.$$

Given a significance level α , $\alpha = P(|Z| > c \mid H_0)$ or $c = z_{\alpha/2}$. The power of the test is $\pi(\theta) = P(|Z| > z_{\alpha/2} \mid H_1)$.

Under H_1 , $\bar{X} \sim N(\theta, \sigma^2/n)$. So,

$$\bar{X} - \theta_0 \sim N\left(\theta - \theta_0, \frac{\sigma^2}{n}\right) \iff W = \frac{\sqrt{n}}{\sigma}[\bar{X} - \theta_0 - (\theta - \theta_0)] \sim N(0, 1).$$

Then, $W = Z - \sqrt{n}(\theta - \theta_0)/\sigma$ and therefore the power of the test is

$$\begin{aligned}
 \pi(\theta) &= 1 - P(-z_{\alpha/2} < Z < z_{\alpha/2} \mid \theta) \\
 &= 1 - P\left(-z_{\alpha/2} - \sqrt{n}\frac{(\theta - \theta_0)}{\sigma} < W < z_{\alpha/2} - \sqrt{n}\frac{(\theta - \theta_0)}{\sigma} \mid \theta\right) \\
 &= 1 + \Phi\left(-z_{\alpha/2} - \sqrt{n}\frac{(\theta - \theta_0)}{\sigma}\right) - \Phi\left(z_{\alpha/2} - \sqrt{n}\frac{(\theta - \theta_0)}{\sigma}\right) > \alpha.
 \end{aligned}$$

The MLR test is unbiased as $\pi(\theta) > \alpha, \forall \theta \neq \theta_0$. The rate of increase of the power depends on σ . The smaller is σ (more precise distribution, more concentrated population), the faster is the rate of growth of the power towards 1.

In the case where σ^2 is unknown, $\Theta_0 = \{(\theta, \sigma^2) : \theta = \theta_0, \sigma^2 > 0\}$ and $\Theta = \{(\theta, \sigma^2) : \theta \in R, \sigma^2 > 0\}$. Since the dimension of Θ_0 is smaller than the dimension of Θ_1 and of Θ , we will work with the latter. As before, this change can be proved to affect only a zero probability set. As seen in Section 4.5.1,

$$\sup_{(\theta, \sigma^2) \in \Theta_0} p(\mathbf{X} \mid \theta, \sigma^2) = p(\mathbf{X} \mid \theta_0, \hat{\sigma}_0^2)$$

where

$$\hat{\sigma}_0^2 = \sum(X_i - \theta_0)^2/n$$

and

$$\sup_{(\theta, \sigma^2) \in \Theta} p(\mathbf{X} \mid \theta, \sigma^2) = p(\mathbf{X} \mid \hat{\theta}, \hat{\sigma}^2)$$

where

$$\hat{\theta} = \bar{X} \text{ and } \hat{\sigma}^2 = \sum(X_i - \bar{X})^2/n.$$

Therefore, the MLR statistic becomes

$$\begin{aligned}
 \lambda(\mathbf{X}) &= \frac{p(\mathbf{X} \mid \theta_0, \hat{\sigma}_0^2)}{p(\mathbf{X} \mid \hat{\theta}, \hat{\sigma}^2)} \\
 &= \frac{(2\pi\hat{\sigma}_0^2)^{-n/2} \exp\left\{-\frac{1}{2\hat{\sigma}_0^2} \sum(X_i - \theta_0)^2\right\}}{(2\pi\hat{\sigma}^2)^{-n/2} \exp\left\{-\frac{1}{2\hat{\sigma}^2} \sum(X_i - \bar{X})^2\right\}} \\
 &= \left(\frac{\hat{\sigma}^2}{\hat{\sigma}_0^2}\right)^{n/2} \frac{\exp\left(-\frac{1}{2\hat{\sigma}_0^2} n\hat{\sigma}_0^2\right)}{\exp\left(-\frac{1}{2\hat{\sigma}^2} n\hat{\sigma}^2\right)} \\
 &= \left(\frac{\hat{\sigma}^2}{\hat{\sigma}_0^2}\right)^{n/2}.
 \end{aligned}$$

One can also write

$$\begin{aligned}
 \frac{\hat{\sigma}_0^2}{\hat{\sigma}^2} &= \frac{\sum(X_i - \bar{X})^2 + n(\bar{X} - \theta_0)^2}{\sum(X_i - \bar{X})^2} \\
 &= 1 + \frac{n(\bar{X} - \theta_0)^2}{(n-1)\hat{\sigma}^2} \\
 &= 1 + \frac{T^2}{n-1} \quad \text{where } T = \sqrt{n} \frac{(\bar{X} - \theta_0)}{\hat{\sigma}}
 \end{aligned}$$

and the MLR can be rewritten as $\lambda(\mathbf{X}) = (1 + T^2/(n-1))^{-n/2}$. Therefore, the MLR test accepts H_0 if $T^2 < c_1$ or $|T| < c$. As $T \mid H_0 \sim t_{n-1}(0, 1)$, the value of c of the level α test is $t_{\alpha/2, n-1}$. This test is usually known as the t test, and is possibly the most used test in statistics. It can be shown that the power of the test is a strictly increasing function of $|\theta - \theta_0|$ (see Exercise 6.5). An immediate consequence is the unbiasedness of the test since the smallest value of the power occurs when $\theta = \theta_0$.

This test is similar because none of the properties of the test is affected by the value of the nuisance parameter σ^2 . This was achieved by replacement of σ^2 by its estimator $\hat{\sigma}^2$ and the existence of the pivotal quantity T . This test is also invariant under linear transformations.

Another common test is the test of equality of two means. Consider two random samples $\mathbf{X}_1 = (X_{11}, \dots, X_{1n_1})$ from the $N(\theta_1, \sigma^2)$ distribution and $\mathbf{X}_2 = (X_{21}, \dots, X_{2n_2})$ from the $N(\theta_2, \sigma^2)$ distribution. It was assumed also for simplicity that the variances are equal. Then, the hypothesis of interest is defined by the parameter space $\Theta_0 = \{(\theta_1, \theta_2, \sigma^2) : \theta_1 = \theta_2 = \theta \in R, \sigma^2 > 0\}$. Observe that under H_0 the problem contains only a single sample of size $n_1 + n_2$ from the $N(\theta, \sigma^2)$ distribution. The parameter vector can be more usefully described by 3 new parameters of interest $\theta_2 - \theta_1$, any other transformation of θ_1 and θ_2 parameters not multiple of $\theta_2 - \theta_1$ and σ^2 . The last two are nuisance parameters.

Once again, the MLR is based on maximizations over Θ and Θ_0 . These operations give

$$\begin{aligned} \sup_{(\theta_1, \theta_2, \sigma^2) \in \Theta_0} p(\mathbf{X}_1, \mathbf{X}_2 | \theta, \sigma^2) \\ = \left(\frac{1}{2\pi\hat{\sigma}_0^2} \right)^{n_1+n_2} \exp \left\{ -\frac{1}{2\hat{\sigma}_0^2} \left[\sum_{i=1}^2 \sum_{j=1}^{n_i} (X_{ij} - \hat{\theta})^2 \right] \right\} \end{aligned}$$

where

$$\hat{\theta} = \frac{1}{n_1 + n_2} \left[\sum_{i=1}^{n_1} X_{1i} + \sum_{i=1}^{n_2} X_{2i} \right]$$

and

$$\hat{\sigma}_0^2 = \frac{1}{n_1 + n_2} \left[\sum_{i=1}^{n_1} (X_{1i} - \hat{\theta})^2 + \sum_{i=1}^{n_2} (X_{2i} - \hat{\theta})^2 \right]$$

and

$$\begin{aligned} \sup_{(\theta_1, \theta_2, \sigma^2) \in \Theta} p(\mathbf{X}_1, \mathbf{X}_2 | \theta_1, \theta_2, \sigma^2) \\ = \left(\frac{1}{2\pi\hat{\sigma}^2} \right)^{n_1+n_2} \exp \left\{ -\frac{1}{2\hat{\sigma}^2} \left[\sum_{i=1}^2 \sum_{j=1}^{n_i} (X_{ij} - \hat{\theta}_i)^2 \right] \right\} \end{aligned}$$

where

$$\hat{\theta}_i = \bar{X}_i, \quad i = 1, 2 \quad \text{and} \quad \hat{\sigma}^2 = \frac{1}{n_1 + n_2} \left[\sum_{i=1}^{n_1} (X_{1i} - \hat{\theta}_1)^2 + \sum_{i=1}^{n_2} (X_{2i} - \hat{\theta}_2)^2 \right].$$

This gives $\lambda(\mathbf{X}_1, \mathbf{X}_2) = (\hat{\sigma}^2 / \hat{\sigma}_0^2)^{(n_1+n_2)/2}$. Note that

$$\bar{X}_1 - \hat{\theta} = \frac{n_2}{n_1 + n_2} (\bar{X}_1 - \bar{X}_2) \quad \text{and} \quad \bar{X}_2 - \hat{\theta} = \frac{n_1}{n_1 + n_2} (\bar{X}_2 - \bar{X}_1)$$

which implies that

$$\begin{aligned} \sum_{i=1}^2 \sum_{j=1}^{n_i} (X_{ij} - \hat{\theta})^2 &= \sum_{i=1}^2 \sum_{j=1}^{n_i} (X_{ij} - \bar{X}_i)^2 + n_i (\bar{X}_i - \hat{\theta})^2 \\ &= \sum_{i=1}^2 \sum_{j=1}^{n_i} (X_{ij} - \bar{X}_i)^2 + \frac{n_1 n_2}{n_1 + n_2} (\bar{X}_1 - \bar{X}_2)^2. \end{aligned}$$

Therefore

$$\begin{aligned} \lambda(\mathbf{X}_1, \mathbf{X}_2) &= \left(1 + \frac{(n_1^{-1} + n_2^{-1})(\bar{X}_1 - \bar{X}_2)^2}{\sum_{i=1}^2 \sum_{j=1}^{n_i} (X_{ij} - \bar{X}_i)^2} \right)^{-(n_1+n_2)/2} \\ &= \left(1 + \frac{T^2}{n_1 + n_2 - 2} \right)^{-(n_1+n_2)/2} \end{aligned}$$

where

$$T = \frac{\bar{X}_1 - \bar{X}_2}{S\sqrt{n_1^{-1} + n_2^{-1}}} \quad \text{with} \quad S^2 = \frac{n_1 + n_2}{n_1 + n_2 - 2} \hat{\sigma}^2.$$

So, the MLR test accepts H_0 when $\lambda(\mathbf{X}_1, \mathbf{X}_2) > c_1$ or when $|T| < c$. As seen in Section 4.5.2, $T|H_0 \sim t_{n_1+n_2-2}(0, 1)$. For a level α test, $c = t_{\alpha/2, n_1+n_2-2}$. Analogously, it can be shown that the power of this test is a strictly increasing function of $|\theta_1 - \theta_2|$. Therefore, this test is unbiased for the same reasons as the previous t test. It is also similar because none of its properties are affected by the two nuisance parameters and invariant under linear transformations of the observations. \square

6.3 Bayesian hypothesis testing

In the Bayesian context, the problem of deciding about which hypothesis to accept is conceptually simpler. Typically, one would compare the hypotheses H_1, \dots, H_k through their respective posterior probabilities, obtained via Bayes' theorem as

$$p(H_i | \mathbf{x}) \propto p(\mathbf{x} | H_i)p(H_i).$$

Once again, this setup can be framed as a decision problem. In addition to the (posterior) probabilities attached to the hypotheses (or states of nature), a loss structure associated with the possible actions can be incorporated.

Returning to the special cases of two hypotheses, suppose one wishes to test $H_0: \theta \in \Theta_0$ versus $H_1: \theta \in \Theta_1$. It suffices to examine the posterior probabilities $p(H_0 | \mathbf{x})$ and $p(H_1 | \mathbf{x})$. If $p(H_0 | \mathbf{x}) > p(H_1 | \mathbf{x})$, then H_0 should be accepted as the most plausible hypothesis for θ . In this case, it can be said that H_0 is preferable to H_1 . Otherwise, H_1 is preferred to H_0 . There is a clear-cut rule for the choice between the hypotheses, which is not always true under the frequentist framework.

As

$$\begin{aligned} p(H_0 | \mathbf{x}) &\propto p(\mathbf{x} | H_0)p(H_0) \\ p(H_1 | \mathbf{x}) &\propto p(\mathbf{x} | H_1)p(H_1) \end{aligned}$$

and recalling that the proportionality constant is the same in both expressions,

$$\frac{p(H_0 | \mathbf{x})}{p(H_1 | \mathbf{x})} = \frac{p(H_0)}{p(H_1)} \frac{p(\mathbf{x} | H_0)}{p(\mathbf{x} | H_1)}.$$

The ratio $p(H_0)/p(H_1)$ is called the prior odds between H_0 and H_1 and the ratio $p(H_0 | \mathbf{x})/p(H_1 | \mathbf{x})$ is called the posterior odds between H_0 and H_1 . The ratio

$$\frac{p(\mathbf{x} | H_0)}{p(\mathbf{x} | H_1)}$$

is called the Bayes factor and is denoted by $BF(H_0; H_1)$. This concept was introduced by Jeffreys (1961). Note that it is in some way a ratio of likelihoods. So, the posterior odds is given by product of the prior odds and the Bayes factor. Once again, the likelihood ratio introduces the influence of the observations in the setting of hypothesis testing. In general, the likelihoods here are marginal in the sense that they are obtained after integrating out some of the parameters not associated with the specification of the hypotheses.

Despite their notational simplicity, it is not easy in many cases to specify $p(H_j)$ when H_j is a simple hypothesis, $j = 0, 1$ and θ is continuous. If a prior density f is specified for $\theta \in \Theta$, one will have that $p(H_j) = p(H_j | \mathbf{x}) = 0$. In these cases, one solution is to attribute a lump prior probability π to the simple hypothesis, say H_0 , for $\pi \in (0, 1)$. So, if H_1 is the complement of a simple hypothesis H_0 , then $p(H_1) = 1 - \pi$ and this probability is distributed over the different values of θ under H_1 , according to the prior distribution for $\theta | H_1$. This distribution will have density f over Θ_1 .

As H_0 is a simple hypothesis, it follows that $p(\mathbf{x} | H_0) = p(\mathbf{x} | \theta_0)$, the marginal density of \mathbf{X} given H_0 . This can also be referred to as the marginal likelihood of H_0 based on \mathbf{X} . The marginal likelihood of H_1 based on \mathbf{X} is

$$\begin{aligned} p(\mathbf{x} | H_1) &= \int_{\Theta - \{\theta_0\}} p(\mathbf{x}, \theta | H_1) d\theta \\ &= \int_{\Theta - \{\theta_0\}} p(\mathbf{x} | \theta, H_1) p(\theta | H_1) d\theta \\ &= \int_{\Theta} p(\mathbf{x} | \theta) f(\theta) d\theta \end{aligned}$$

observing that the last integration is performed over the entire parameter space Θ because a single point does not alter its value. The Bayes factor is reduced to

$$BF(H_0; H_1) = \frac{p(\mathbf{x} | \theta_0)}{\int p(\mathbf{x} | \theta) f(\theta) d\theta}.$$

Note that it provides the relative odds between H_0 and H_1 without taking into account the prior odds. It is a Bayesian measure of the goodness of fit of a given model to the data set. A Bayes factor larger than 1 indicates that H_0 fits the data better than H_1 and thus has larger likelihood. Also, the marginal prior for θ is mixed and therefore the marginal distribution of \mathbf{X} is

$$\begin{aligned} p(\mathbf{x}) &= \int p(\mathbf{x} | \theta) dF(\theta) \\ &= \pi p(\mathbf{x} | \theta_0) + (1 - \pi) \int p(\mathbf{x} | \theta) f(\theta) d\theta \\ &= \pi p(\mathbf{x} | \theta_0) + (1 - \pi) p(\mathbf{x} | H_1). \end{aligned}$$

Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution and assume one wishes to test $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$. Consider initially, the

case of known variance σ^2 . Then, with a prior probability π for H_0 and assuming $\theta | H_1 \sim N(\mu, \tau^2)$ gives

$$\begin{aligned} p(\mathbf{x} | H_1) &= \int p(\mathbf{x} | \theta) p(\theta) d\theta \\ &= \left(\frac{1}{2\pi\sigma^2} \right)^{n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \bar{x})^2 \right\} \\ &\quad \times \int \exp \left\{ -\frac{n}{2\sigma^2} (\theta - \bar{x})^2 \right\} \frac{1}{\sqrt{2\pi}\tau} \exp \left\{ -\frac{1}{2\tau^2} (\theta - \mu)^2 \right\} d\theta. \end{aligned}$$

After a few substitutions and algebraic transformations

$$\begin{aligned} p(\mathbf{x} | H_1) &= \left(\frac{1}{2\pi\sigma^2} \right)^{n/2} \\ &\quad \times \exp \left\{ -\frac{ns^2}{2\sigma^2} \right\} \left(\frac{\sigma^2/n}{\tau^2 + \sigma^2/n} \right)^{1/2} \exp \left\{ -\frac{1}{2} \frac{(\bar{x} - \mu)^2}{\tau^2 + \sigma^2/n} \right\} \end{aligned}$$

where $s^2 = \sum_{i=1}^n (x_i - \bar{x})^2 / n$. The Bayes factor is

$$\begin{aligned} BF(H_0; H_1) &= \frac{\left(\frac{1}{2\pi\sigma^2} \right)^{n/2} \exp \left\{ -\frac{ns^2}{2\sigma^2} \right\} \exp \left\{ -\frac{n}{2\sigma^2} (\bar{x} - \theta_0)^2 \right\}}{\left(\frac{1}{2\pi\sigma^2} \right)^{n/2} \exp \left\{ -\frac{ns^2}{2\sigma^2} \right\} \left(\frac{\sigma^2/n}{\tau^2 + \sigma^2/n} \right)^{1/2} \exp \left\{ -\frac{1}{2} \frac{(\bar{x} - \mu)^2}{\tau^2 + \sigma^2/n} \right\}} \\ &= \left(\frac{\tau^2 + \sigma^2/n}{\sigma^2/n} \right)^{1/2} \exp \left\{ \frac{n}{2} \left[\frac{(\bar{x} - \mu)^2}{\sigma^2 + n\tau^2} - \frac{(\bar{x} - \theta_0)^2}{\sigma^2} \right] \right\}. \end{aligned}$$

As expected, the Bayes factor only depends on the sample through \bar{x} . To obtain the sample value that maximizes the Bayes factor, it suffices to solve the maximization problem for \bar{x} . Taking the logarithm of the previous expression and differentiating with respect to \bar{x} gives

$$\begin{aligned} \frac{\partial \log BF}{\partial \bar{x}} &= \frac{n}{2} \left[\frac{2(\bar{x} - \mu)}{\sigma^2 + n\tau^2} - \frac{2(\bar{x} - \theta_0)}{\sigma^2} \right] = 0 \\ \frac{\partial^2 \log BF}{\partial \bar{x}^2} &= \frac{n}{2} \left[\frac{2}{\sigma^2 + n\tau^2} - \frac{2}{\sigma^2} \right] = \frac{n}{2} \left[-\frac{2n\tau^2}{\sigma^2(\sigma^2 + n\tau^2)} \right] < 0 \end{aligned}$$

and solving the first equation provides the maximum. The solution is

$$\bar{x}_{\max} = \theta_0 + \frac{\sigma^2}{n\tau^2}(\theta_0 - \mu)$$

and the maximized value of the Bayes factor is

$$\exp \left\{ \frac{(\theta_0 - \mu)^2}{2\tau^2} \right\} \left(1 + \frac{n\tau^2}{\sigma^2} \right)^{1/2} > 1.$$

The larger the sample value n , the larger the chances of H_0 and the larger is τ^2 , the larger is the maximized value of the Bayes factor. The prior uncertainty plays a crucial role in the Bayesian comparison of sharp null hypotheses. In the limit, when $\tau^2 \rightarrow \infty$, the Bayes factor $BF(H_0; H_1)$ also increases indefinitely. This limit is the result obtained when a non-informative prior for $\theta \mid H_1$ is used. This was first noted by Lindley (1957) and is known as Lindley's paradox. It has been the object of study of comparisons between the Bayesian and frequentist approaches to hypothesis testing.

To ease comparisons between the approaches, take $\mu = \theta_0$ and $\tau^2 = \sigma^2$. The first assumption centres the alternative prior distribution over the single value of H_0 and the second takes the prior variance in the alternative as equal to the observational variance. Then,

$$\begin{aligned} BF(H_0; H_1) &= \left(\frac{\sigma^2 + \sigma^2/n}{\sigma^2/n} \right)^{1/2} \exp \left\{ \frac{n}{2} \left[\frac{(\bar{x} - \theta_0)^2}{\sigma^2 + n\sigma^2} - \frac{(\bar{x} - \theta_0)^2}{\sigma^2} \right] \right\} \\ &= (n+1)^{1/2} \exp \left\{ -\frac{n}{2} \left[\frac{n(\bar{x} - \theta_0)^2}{\sigma^2(n+1)} \right] \right\} \\ &= (n+1)^{1/2} \exp \left\{ -\frac{n}{n+1} \frac{z^2}{2} \right\} \quad \text{where } z = \sqrt{n} \frac{\bar{x} - \theta_0}{\sigma}. \end{aligned}$$

If $p(H_0) = \pi$, then

$$\begin{aligned} \frac{p(H_0 \mid \mathbf{x})}{p(H_1 \mid \mathbf{x})} &= \frac{\pi}{1-\pi} BF(H_0; H_1) \\ \iff p(H_0 \mid \mathbf{x}) &= \left\{ 1 + \left[\frac{\pi}{1-\pi} BF(H_0; H_1) \right]^{-1} \right\}^{-1}. \end{aligned}$$

Assuming prior indifference ($p(H_0) = p(H_1) = 1/2$) gives

$$p(H_0 \mid \mathbf{x}) = \left\{ 1 + \left[(1+n)^{1/2} \exp \left\{ -\frac{n}{n+1} \frac{z^2}{2} \right\} \right]^{-1} \right\}^{-1}.$$

The posterior probability of H_0 can be calculated for different values of n and z since both classical and Bayesian tests are based on z^2 . Working with the most common values, associated with the p -values 0.05 and 0.01 gives Table 6.1.

The probabilities of H_0 are smaller for the larger value of $|z|$ which is reasonable. What is less reasonable is the values that are obtained for these probabilities but they reinforce the idea that p -values should not be taken as probabilities that can be associated with H_0 . Another interesting result is that for any given value of z , posterior probabilities can vary substantially from a very low value that would lead to rejection of H_0 to a very large value that would lead to acceptance of H_0 . One possible way to reconcile these findings with the significance level is that levels should also be changed with sample size. A large sample size should call for a

Table 6.1 Values of $P(H_0 \mid \mathbf{x})$

n	$ z $ (p -value)	
	1.96 (0.05)	2.576 (0.01)
1	0.35	0.21
10	0.37	0.14
100	0.60	0.27
1000	0.80	0.53

reduced level and vice versa. A more detailed discussion of the subject can be found in Berger and Sellke (1987).

Assume now the same situation but with an unknown observational variance σ^2 and take $\phi = \sigma^{-2}$. A prior for ϕ must also be specified. Since the hypotheses do not involve ϕ , it is reasonable to take the same marginal prior for ϕ under both hypotheses. Adopting conjugate priors under both hypotheses gives

$$n_0 \sigma_0^2 \phi \sim \chi_{n_0}^2 \text{ under } H_0 \text{ and } H_1 \text{ and } \theta \mid \phi, H_1 \sim N(\mu, (c\phi)^{-1}).$$

The relevant quantities for the calculation of the Bayes factor are

$$\begin{aligned} p(\mathbf{x} \mid H_0) &= \int p(\mathbf{x} \mid \theta_0, \phi) p(\phi \mid H_0) d\phi \\ p(\mathbf{x} \mid H_1) &= \int \int p(\mathbf{x} \mid \theta, \phi) p(\theta \mid \phi, H_1) p(\phi \mid H_1) d\theta d\phi \end{aligned}$$

where all the above densities are known. Substituting their expressions gives

$$\begin{aligned} BF(H_0; H_1) &= \left(\frac{c+n}{c} \right)^{1/2} \left\{ \frac{n_0 \sigma_0^2 + (n-1)s^2 + [cn/(c+n)](\bar{x} - \mu)^2}{n_0 \sigma_0^2 + (n-1)s^2 + n(\bar{x} - \theta_0)^2} \right\}^{(n_0+n)/2} \end{aligned}$$

where now $s^2 = \sum (x_i - \bar{x})^2 / (n-1)$.

It is interesting to study the behaviour of the Bayes factor in extreme situations such as with non-informative priors. Taking $n_0 \rightarrow 0$ and assuming as before that $\mu = \theta_0$, gives

$$\begin{aligned} BF(H_0; H_1) &\rightarrow \left(\frac{c+n}{c} \right)^{1/2} \left\{ \frac{(n-1)s^2 + [c/(c+n)]n(\bar{x} - \theta_0)^2}{(n-1)s^2 + n(\bar{x} - \theta_0)^2} \right\}^{n/2} \\ &= k^{-1/2} \left\{ \frac{n-1+kt^2}{n-1+t^2} \right\}^{n/2} \quad \text{where } k = \frac{c}{c+n} \text{ and } t = \sqrt{n} \frac{\bar{x} - \theta_0}{s}. \end{aligned}$$

The above expression is graphed in Figure 6.2. The Bayes factor is a symmetric function of the sample values through the statistic t . It varies from the highest value of support of H_0 when $t = 0$ to a minimum value when $|t| \rightarrow \infty$, as expected.

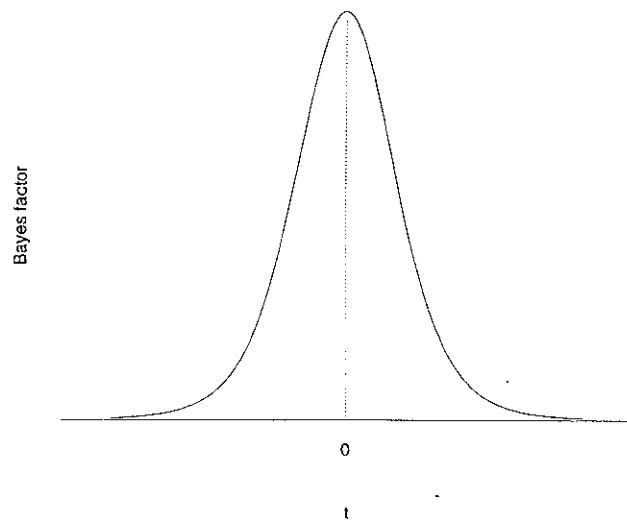


Fig. 6.2 Bayes factor for $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$ as a function of t .

In the general case, the hypotheses can be incorporated into the prior distribution and the problem of hypothesis testing can be thought of as comparison of possible alternative distributions for θ . For example, in the above test, the priors corresponding to H_0 and H_1 would be $P(\theta = \theta_0 | \phi) = 1$ and $\theta | \phi \sim N(\mu, (c\phi)^{-1})$. The prior corresponding to H_0 is degenerate. This will be the case whenever one of the hypotheses corresponds to a parameter space with smaller dimension than the complete parameter space Θ .

It will be seen in Chapter 8 that it is common to test if a few of the components of θ are null. Writing $\theta = (\theta_1, \theta_2)$, where θ_1 is q -dimensional and θ_2 is $p - q$ -dimensional, one may wish to test whether $\theta_2 = \mathbf{0}$. In this case, $P(\theta | H_0)$ will be concentrated on a q -dimensional subspace of R^p passing through the point $\theta_2 = \mathbf{0}$.

6.4 Hypothesis testing and confidence intervals

You may have noticed the strong connection between hypothesis testing and confidence intervals. This connection is clearer with the frequentist framework but it is also relevant under the Bayesian approach as will shortly be seen.

Starting with the classical approach, consider a level α test for the hypothesis that $\theta = \theta_0$. Assume that, based on a given sample size, the test yields an acceptance region $A = A(\theta_0)$ where it is important to explicitly denote the dependence of the

region on θ_0 . So,

$$P(\mathbf{X} \in A(\theta_0) | \theta_0) = 1 - \alpha.$$

Varying the value of θ_0 in Θ leads to a family of regions, all depending on θ and with probability $1 - \alpha$. This automatically implies that, after observing the value \mathbf{x} for \mathbf{X} , the region

$$\{\theta : \mathbf{x} \in A(\theta)\}$$

will have confidence $1 - \alpha$.

Example. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $N(\theta, \sigma^2)$ distribution with known σ^2 . The UMP unbiased test of level α to test $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$ has acceptance region $\{\mathbf{X} : \theta_0 - z_{\alpha/2}\sigma/\sqrt{n} \leq \bar{X} \leq \theta_0 + z_{\alpha/2}\sigma/\sqrt{n}\}$. This interval can be rewritten as $\{\mathbf{X} : \bar{X} - z_{\alpha/2}\sigma/\sqrt{n} \leq \theta_0 \leq \bar{X} + z_{\alpha/2}\sigma/\sqrt{n}\}$. Replacing now θ_0 by θ gives the $1 - \alpha$ confidence interval for θ :

$$\left\{ \theta : \bar{X} - \frac{\sigma}{\sqrt{n}}z_{\alpha/2} \leq \theta \leq \bar{X} + \frac{\sigma}{\sqrt{n}}z_{\alpha/2} \right\}.$$

Of course, this is exactly the same interval as that obtained in Section 4.4.

This relation is more heavily explored in the next section when asymptotic tests are introduced. In fact, this relation can also be used in the reverse direction with hypothesis tests obtained from confidence regions. To see this, suppose that $\{\theta : G(\mathbf{x}, \theta) \in C\}$ is a $100(1 - \alpha)\%$ confidence region for θ . Then, for every value θ_0 of θ , it is true that $P(G(\mathbf{X}, \theta_0) \in C | \theta_0) = 1 - \alpha$. A level α test for the hypothesis $\theta = \theta_0$ can be defined by the critical region $\{\mathbf{X} : G(\mathbf{X}, \theta_0) \notin C\}$.

This brings us naturally to an alternative definition of Bayesian hypothesis testing. The method consists in constructing a credibility region for θ with probability $1 - \alpha$ and accepting the hypothesis $\theta = \theta_0$ if the above region contains the value θ_0 . As before, one should aim to construct HPD regions (intervals, in the scalar case) or at least with smallest possible volume (length, in the scalar case). The value of α is typically low but there is no prescription about the value to be adopted in any given problem. This method was proposed and extensively used by Lindley (1965).

Example. Taking again a random sample $\mathbf{X} = (X_1, \dots, X_n)$ from the $N(\theta, \sigma^2)$ distribution with known σ^2 and prior $\theta \sim N(\mu, \tau^2)$ gives the posterior $\theta | \mathbf{x} \sim N(\mu_1, \tau_1^2)$. The $100(1 - \alpha)\%$ confidence interval for θ is given by $[\mu_1 - \tau_1 z_{\alpha/2}, \mu_1 + \tau_1 z_{\alpha/2}]$. The hypothesis $\theta = \theta_0$ can be accepted if θ_0 belongs to the interval above. In the non-informative case, $\mu_1 \rightarrow \bar{X}$ and $\tau_1 \rightarrow \sigma/\sqrt{n}$ and the hypothesis is accepted if θ_0 belongs to the interval $[\bar{X} - z_{\alpha/2}\sigma/\sqrt{n}, \bar{X} + z_{\alpha/2}\sigma/\sqrt{n}]$, coinciding with the classical test.

In the previous section, Bayesian tests of a simple hypothesis in the form $\theta = \theta_0$ were defined by attributing a lump prior probability π to this hypothesis and the remaining $1 - \pi$ was distributed according to some probability distribution of

188 Hypothesis testing

$\theta_0 | H_1$. This specification is criticized for giving a sometimes unjustified special, different status to the value θ_0 .

This discussion is in fact wider and goes beyond the boundaries of Bayesian thinking. It has to do with the capacity of a single hypothesis to represent adequately the situation under study. Some authors suggest that these hypotheses are always at most a useful approximation of more realistic hypotheses where θ actually belongs to a neighbourhood of θ_0 . We will not pursue this discussion here further than acknowledging that the spectrum of thoughts on this matter goes from total rejection to total acceptance of this formulation. Good references for this discussion are Berger and Delampady (1987) and Lindley (1993).

6.5 Asymptotic tests

Asymptotic tests are those based on asymptotic approximation for the distribution of the test quantity, irrespective of whether it is under the Bayesian or frequentist paradigm. This is a very broad definition including results based on the central limit theorem. We shall concentrate here on tests based on the MLE, score function and MLR. In most cases we are typically led to a limiting χ^2 distribution. The use of asymptotic theory to develop useful test statistics was carried out by Bartlett, Wald and Wilks among others in the 1940s.

In many cases, it is not possible to analytically obtain the exact distribution of the MLR $\lambda(\mathbf{X})$ and asymptotic methods are frequently used. Alternative approximating methods were described in the previous chapter. Suppose that $\theta \in \Theta \subset R^p$ and one wishes to test the hypothesis $H_0: \theta = \theta_0$. In this case, a Taylor series expansion of the function $L(\theta; \mathbf{X}) = \log p(\mathbf{X} | \theta)$ around $\hat{\theta}$ gives

$$\begin{aligned} L(\theta_0; \mathbf{X}) &\simeq L(\hat{\theta}; \mathbf{X}) + [\mathbf{U}(\mathbf{X}; \hat{\theta})]'(\theta_0 - \hat{\theta}) - \frac{1}{2}(\theta_0 - \hat{\theta})' \mathbf{J}(\hat{\theta})(\theta_0 - \hat{\theta}) \\ &= L(\hat{\theta}; \mathbf{X}) - \frac{1}{2}(\theta_0 - \hat{\theta})' \mathbf{J}(\hat{\theta})(\theta_0 - \hat{\theta}) \end{aligned}$$

since $\mathbf{U}(\mathbf{X}; \hat{\theta}) = \mathbf{0}$. The higher-order terms are neglected since, under H_0 , θ_0 and $\hat{\theta}$ are close for n large. Therefore,

$$\begin{aligned} -2 \log \lambda(\mathbf{X}) &= -2 \log \left(\frac{p(\mathbf{X} | \theta_0)}{p(\mathbf{X} | \hat{\theta})} \right) \\ &= -2[L(\theta_0; \mathbf{X}) - L(\hat{\theta}; \mathbf{X})] \\ &\simeq (\theta_0 - \hat{\theta})' \mathbf{J}(\hat{\theta})(\theta_0 - \hat{\theta}). \end{aligned}$$

Since the MLE is asymptotically normal and $\mathbf{J}(\hat{\theta})/n$ converges almost surely to its expectation $\mathbf{I}(\theta_0)/n$ under H_0 , the quadratic form on the right-hand side of the equation has a χ_p^2 asymptotic distribution. Therefore, the asymptotic distribution of $-2 \log \lambda(\mathbf{X})$ is χ_p^2 and the test with asymptotic level α accepts H_0 if $-2 \log \lambda(\mathbf{X}) < \chi_{\alpha, p}^2$. If the null hypothesis is of the form $H_0: \theta \in \Theta_0$ with

$\dim \Theta_0 = p - q > 0$, then the asymptotic distribution of $-2 \log \lambda(\mathbf{X})$ is χ_q^2 and the test with asymptotic level α accepts H_0 if $-2 \log \lambda(\mathbf{X}) < \chi_{\alpha, q}^2$.

These results are also useful in the construction of confidence intervals. The idea is once again to explore the relation between confidence intervals and tests. Suppose that $H_0: \theta = (\theta_1, \dots, \theta_{p-q}, \theta_{p-q+1, 0}, \dots, \theta_{p, 0})$. In other words, the last q components of θ are fixed. Let the MLR now be denoted by $\lambda(\mathbf{X} | \theta_{p-q+1, 0}, \dots, \theta_{p, 0})$. A $100(1 - \alpha)\%$ confidence region for $(\theta_{p-q+1}, \dots, \theta_p)$ is given by

$$\{(\theta_{p-q+1}, \dots, \theta_p) : -2 \log \lambda(\mathbf{x} | \theta_{p-q+1}, \dots, \theta_p) < \chi_{\alpha, q}^2\}.$$

Similarly, from the Bayesian point of view, the asymptotic posterior distribution of $-2 \log \lambda(\mathbf{x} | \theta_{p-q+1}, \dots, \theta_p)$ is χ_q^2 . Tests and confidence regions can be constructed as described before. Note that we have chosen to test for known values of the last q components for simplicity. The same Bayesian and frequentist results hold for a test on any q components of θ .

The asymptotic distributions of the score function and of the MLE lead to two classes of classical tests. Assume that $\theta \in \Theta \subset R^p$ and one wishes to test $H_0: \theta = \theta_0$. Defining then

$$W_E(\theta_0) = (\hat{\theta} - \theta_0)' \mathbf{I}(\theta_0)(\hat{\theta} - \theta_0)$$

and

$$W_U(\theta_0) = [\mathbf{U}(\mathbf{X}; \theta_0)]' \mathbf{I}^{-1}(\theta_0) \mathbf{U}(\mathbf{X}; \theta_0)$$

we have that both statistics have a χ_p^2 asymptotic distribution under H_0 . So, the tests of asymptotic level α reject H_0 if $W_E(\theta_0) > \chi_{\alpha, p}^2$ and $W_U(\theta_0) > \chi_{\alpha, p}^2$, respectively. Given the almost sure convergence of $J(\hat{\theta})$ and $J(\theta_0)$ to $I(\theta_0)$, these replacements can be made in the definitions of W_E and W_U and the same results are obtained. Equivalent Bayesian results can be obtained that the asymptotic posterior distributions of $W_E(\theta)$ and $W_U(\theta)$ are χ_p^2 . These results were partially presented in Section 5.3. Hypothesis testing then can be made as described in the previous section.

The tests based on W_E and W_U can also be applied to situations where $\dim \Theta_0 = p - q > 0$. In this case, the value of θ_0 is replaced in their expressions by $\hat{\theta}_0$, the estimator of θ under H_0 . Their asymptotic distribution becomes a χ_q^2 , just like the MLR test.

Note that three general tests have been defined and they all have asymptotic χ^2 distribution under H_0 where the number of degrees of freedom depends on the difference between the dimensions of Θ and Θ_0 . The test based on the MLR depends on maximization on both hypotheses whereas the score test requires maximization under the null hypothesis and the test based on the MLE requires maximization only under the alternative hypothesis. In most but not all cases, the null hypothesis provides a simplification to the model. It is then simpler to estimate parameters under the null hypothesis which favours the score test. When estimation under the null hypothesis is harder, it is simpler to perform the test based on the MLE. A

good reference for comparison between and interpretation of these tests is Buse (1982).

Example. Assume a specific model for observations Y_1, \dots, Y_n is of the form $Y_i \sim N(f_i(\theta), \sigma^2)$ where $\theta = (\theta_0, \theta_1, \theta_2)$ and $f_i(\theta) = \theta_0 + \theta_1 \exp(\theta_2 z_i)$, $i = 1, \dots, n$. Let the null hypothesis be in the form $H_0: \theta_2 = \theta_{2,0}$. Under the null hypothesis the model becomes a simple linear regression with known regressor variable $x_i = \exp(\theta_2 z_i)$, $i = 1, \dots, n$. It is then simpler to estimate the model under the null hypothesis. If, however, the null hypothesis is of the form $H_0: \partial f_i / \partial z_i = 1/3$ then estimation under the null hypothesis becomes a non-linear problem with restriction whereas estimation under the alternative is only a non-linear problem. In those cases, it is easier to use the test based on the MLE.

Another interesting question concerns the appropriateness of the asymptotic approximation. It can be shown that the approximation of $\lambda(\mathbf{X})$ to the χ_p^2 is of order n^{-1} . Bartlett (1947) showed that $P(\lambda(\mathbf{X}) \leq x) = P(Z \leq x) + O(n^{-1})$, where $Z \sim \chi_p^2$, and obtained a corrected MLR statistic $\lambda^*(\mathbf{X}) = \lambda(\mathbf{X})[1 + b(\theta_0)/n]^{-1}$ such that $E[\lambda^*(\mathbf{X})] = p + O(n^{-2})$ for many multivariate problems. Lawley (1956) showed that all moments of λ^* agree with those of a χ_p^2 distribution to order n^{-2} . Cordeiro (1987) proved that this approximation order is also valid for the distribution function of λ^* . This result was extended further to any test statistic with asymptotic χ_p^2 distribution by Cordeiro and Ferrari (1991).

A particular case of special interest is when n items are observed and classified independently into one of p possible groups. The parameter of interest is the vector with the group probabilities $\theta = (\theta_1, \dots, \theta_{p-1})$ where the probability for the p th group is obtained from the unit sum restriction. Suppose one wishes to test $H_0: \theta = \theta_0 = (\theta_{1,0}, \dots, \theta_{p-1,0})$. Then it can be shown (see Exercise 6.21) that the test statistics W_E and W_U are given by

$$W_E(\theta_0) = \sum_{i=1}^p \frac{(N_i - n\hat{\theta}_{i,0})^2}{N_i} \quad \text{and} \quad W_U(\theta_0) = \sum_{i=1}^p \frac{(N_i - n\hat{\theta}_{i,0})^2}{n\hat{\theta}_{i,0}}.$$

Both have an asymptotic χ_{p-1}^2 distribution under H_0 . The only difference between them is the replacement of N_i by $n\hat{\theta}_{i,0}$. But, under H_0 , $N_i/n \rightarrow \theta_{i,0}$ almost surely when $n \rightarrow \infty$, by the strong law of large numbers. This is an indication of an asymptotic equivalence between the two statistics. The tests reject H_0 when the values of the statistics W_E and W_U are larger than the $1 - \alpha$ quantile of the χ_{p-1}^2 distribution.

These tests try to measure how well a given hypothesis fits the data. For that reason, they are known as goodness-of-fit tests and are heavily used in statistics whenever a situation can be represented in p mutually exclusive categories. Of course the assessment of the fit of a model is much more general and leads to a variety of other tests. Analogously, the Bayesian asymptotic result is that when the above quantities are written as functions of θ (instead of θ_0) they will have an asymptotic χ_{p-1}^2 posterior distribution.

When H_0 is no longer a simple hypothesis but has instead dimension $p - q - 1 > 0$, the goodness-of-fit statistics are given by

$$W_E(\hat{\theta}_0) = \sum_{i=1}^p \frac{(N_i - n\hat{\theta}_{i,0})^2}{N_i} \quad \text{and} \quad W_U(\hat{\theta}_0) = \sum_{i=1}^p \frac{(N_i - n\hat{\theta}_{i,0})^2}{n\hat{\theta}_{i,0}}$$

where $\hat{\theta}_{i,0}$ is the MLE of θ_i , $i = 1, \dots, p - 1$ under H_0 . Both test statistics have asymptotic χ_q^2 distribution under H_0 and reject H_0 if and only if the value of the statistic is larger than the $1 - \alpha$ quantile of the χ_q^2 distribution.

An important application of goodness-of-fit tests is the test of the fit of a given distribution to a data set. In this case, a random sample $\mathbf{X} = (X_1, \dots, X_n)$ from an unknown distribution is observed and one wishes to test whether this distribution is of a given known form. Partitioning the line into intervals I_i , $i = 1, \dots, p$, the number of observations in each interval can be counted. These counts have jointly multinomial distribution with probabilities $\theta_{i,0} = P(X \in I_i | H_0)$ and the test of the fit is given as above. If the null hypothesis specifies a class of distributions with $q > 0$ unknown parameters then one should first estimate the unknown parameters, by maximum likelihood say, under H_0 . Then the statistics W_E and W_U can be evaluated with the estimates $\hat{\theta}_{i,0} = \hat{P}(X \in I_i | H_0)$. The test statistics will now have an asymptotic χ_{p-q-1}^2 distribution and the level α test rejects H_0 if the value of the test statistic is larger than the $1 - \alpha$ quantile of the χ_{p-q-1}^2 distribution. Another important application of these tests to contingency tables is left as an exercise.

All these tests have the property that their power function converges to 1 when $n \rightarrow \infty$ for any parameter value in the alternative. This result can be obtained from the Taylor series expansion of the log-likelihood around the point in the alternative. This follows essentially from the consistency of MLE. Tests with this property are said to be consistent.

Exercises

§ 6.2

1. Assume that X_1, \dots, X_n are iid with density $p(x|\theta) = \theta x^{\theta-1} I_x([0, 1])$, and $\theta > 0$ is unknown. Determine the UMP test of level 0.05 for $H_0: \theta \leq 1$ against $H_1: \theta > 1$.
2. Let $X \sim \text{bin}(n, \theta)$ and suppose one wishes to test $H_0: \theta = 1/2$ versus $H_1: \theta \neq 1/2$.
 - (a) Show that the MLR test statistic is $|2X - n|$.
 - (b) Find the critical region for a test of level 0.05 when $n = 25$ using the normal approximation.
3. Suppose that k independent tests about the same hypothesis $H: \theta = \theta_0$ have been performed using different data sets and were based on independent statistics T_1, \dots, T_k with continuous distributions under H_0 . Let $\alpha(T_1), \dots, \alpha(T_k)$ be their respective p -values.

- (a) Show that $\alpha(T_1), \dots, \alpha(T_k)$ form a random sample from the uniform distribution on $(0, 1)$.
- (b) Define $F = -2 \sum_{i=1}^k \log \alpha(T_i)$ as the test statistic for a combined test of H . Which values of F should lead to the rejection of the hypothesis H ?
- (c) Show that $F \sim \chi_{2k}^2$, under H and specify the critical region of a level α test.
4. Let $\mathbf{X} = (X_1, \dots, X_n)$ and $\mathbf{Y} = (Y_1, \dots, Y_n)$ be independent samples from the exponential distributions with means θ_1 and θ_2 respectively and suppose we wish to test the hypothesis H_0 of equality of the distributions.

- (a) Show that the MLR test rejects H_0 if

$$\left| \frac{\bar{X}}{\bar{X} + \bar{Y}} - \frac{1}{2} \right| > c.$$

- (b) Show that the value of c is given by

$$\frac{1}{1 + F_\beta(2n, 2n)} - \frac{1}{2}$$

for the level α test and β is such that $P[F_\beta(2n, 2n) < F(2n, 2n) < F_\beta(2n, 2n) + 2] = 1 - \alpha$.

Hint: show that under H_0 , $\bar{X}/\bar{Y} \sim F(2n, 2n)$.

5. Show that the power of the t test is a strictly increasing function of $|\theta - \theta_0|$. What is the expression of the p -value for this test?
6. Show that the uniform distribution over the interval $[0, \theta]$ has monotone LR because the ratio of sampling densities is a monotonically non-increasing function of $T(\mathbf{X}) = \max_i X_i$.
7. Consider a random sample X_1, \dots, X_n from the $N(\theta, \sigma^2)$ distribution with both parameters unknown and define the hypotheses $H_0 : \theta = \theta_0$ and $H_1 : \theta \neq \theta_0$. Show that working with the full parameter space Θ instead of parameter space Θ_1 under H_1 amounts to a zero probability change in the MLR statistic.
8. Let X_1, \dots, X_p be independent random variables with respective $\text{Pois}(\mu_i)$ distributions, $i = 1, \dots, p$ and we wish to test the hypothesis H_0 of the equality of the distributions.
- (a) Obtain the MLR test for H_0 .
- (b) Show that $(X_1, \dots, X_p) | X_1 + \dots + X_p = s$ has a multinomial distribution with parameters s and $(\theta_1, \dots, \theta_p)$ where $\theta_i = \mu_i / (\mu_1 + \dots + \mu_p)$, $i = 1, \dots, p$.
- (c) Justify the following test (commonly used in such situations): reject H_0 if $\sum_{i=1}^p (X_i - \bar{X})^2 / \bar{X} \geq \chi_{p-1, \alpha}^2$.

- (d) Are the test in (a) and the test in (c) similar?
9. Prove that the tests in Section 6.2.3 are unbiased, similar and invariant under linear transformations.

§ 6.3

10. Let $X \sim \text{Exp}(\theta)$ and assume one wishes to test $H_0: \theta = \theta_0$ versus $H_1: \theta = \theta_1$, where $\theta_1 < \theta_0$.
- (a) Prove that the level α likelihood test accepts H_0 if $X < -\theta_0^{-1} \log \alpha$.
- (b) Obtain the p -value associated with $X = 3$, when $\theta_0 = 1$.
- (c) Assuming that $\theta_1 = 1/2$, calculate the BF for $X = 3$ and $\theta_0 = 1$.
- (d) Supposing $p(H_0) = p(H_1)$, calculate the posterior probability of H_0 .
- (e) Compare the results of the classical test obtained in item (a) with those from the Bayesian test.
11. Suppose that $X \sim \text{Cauchy}(\theta, 1)$ and that one wishes to test $H_0: \theta = 0$ versus $H_1: \theta \neq 0$. To do that, set a mass probability $\pi > 0$ to H_0 and the remaining probability over H_1 distributed according to a density $p(\theta)$.
- (a) Prove that $p(H_0 | x) \rightarrow \pi$ when $|x| \rightarrow \infty$.
- (b) What conclusions can be drawn from this result?
12. The data set in Table 6.2 represents the number of vehicles that travel through a section of a highway during a 15 minute interval in the afternoons of six consecutive days of two consecutive weeks.

Table 6.2 Number of vehicles

Week	Day of the week					
	Mon	Tue	Wed	Thu	Fri	Sat
1st.	50	65	52	63	84	102
2nd.	56	49	60	45	112	90

Assume that the number of vehicles follows a Poisson distribution. So, the average number of vehicles traveling from Monday to Thursday is λ and μ for Friday and Saturday.

- (a) Obtain 95% credibility intervals for λ and μ based on non-informative priors and comment on the results.
- (b) Test the hypothesis that $2\lambda = \mu$.
13. Suppose that independent random variables X_1 and X_2 where $P(X_i = 0) = \theta_i$ and $P(X_i = 1) = 1 - \theta_i$, $i = 1, 2$, are observed and one wishes to test $H_0: \theta_1 = \theta_2 = \theta$ versus $H_1: \theta_1 \neq \theta_2$. Assume also that all prior distributions are uniform, that is, under H_0 , θ is uniform over the unit interval, and under H_1 , (θ_1, θ_2) is uniform over the unit square.

- (a) Show that the MLE of θ under H_0 is $(X_1 + X_2)/2$ and the MLE's of θ_i under H_1 are X_i , $i = 1, 2$.
- (b) Obtain the posterior distributions under the two hypotheses, that is, the distributions of $\theta | \mathbf{x}, H_0$ and $(\theta_1, \theta_2) | \mathbf{x}, H_1$.
- (c) Obtain the GMLE's under H_0 and H_1 and compare them with the MLE's obtained in item (a).
- (d) Obtain the predictive distributions of (X_1, X_2) under H_0 and under H_1 .
- (e) Show that $BF(H_0, H_1) = 4/3$, if $x_1 = x_2$, and $4/6$ if $x_1 \neq x_2$. Interpret the result.
14. Suppose one wishes to verify if a quantity θ is smaller than a prespecified value θ_0 . Therefore, assume $\theta \sim N(\mu, \tau^2)$ and observe $X \stackrel{\text{d}}{\sim} N(\theta, \sigma^2)$, σ^2 known.
- (a) Obtain the prior probability of $H_0: \theta < \theta_0$.
 - (b) Obtain the posterior probability of H_0 .
 - (c) Prove that the probability of H_0 increases after observing $X = x$ iff $x < \theta_0(1 - 1/\sqrt{2})$ in the case $\sigma^2 = \tau^2$ and $\mu = 0$.

§ 6.4

15. Construct a level α test for the hypothesis $\theta = \theta_0$ from confidence intervals assuming that independent $X_i \sim \text{Pois}(\theta t_i)$, $i = 1, \dots, n$, are observed with t_i , $i = 1, \dots, n$, known. Can the hypothesis be accepted if $\sum_{i=1}^n x_i = 10$, $\sum_{i=1}^n t_i = 5$, $\theta_0 = 1$ and $\alpha = 0.05$? And if $\alpha = 0.01$?
16. Let $X_1 \sim N(\theta_1, 1)$ and $X_2 \sim N(\theta_2, 1)$ be independent and define $\rho = \theta_1/\theta_2$.
- (a) Show that the test that rejects the hypothesis $\rho = \rho_0$ when $|X_1 - \rho_0 X_2| > (1 + \rho_0^2)^{1/2} z_{\alpha/2}$ has level α .
 - (b) Obtain a $100(1 - \alpha)\%$ confidence region for ρ from the test described in (a). Draw a graph of the region and interpret the result.

§ 6.5

17. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from the $\text{Exp}(\theta)$ distribution and suppose one wishes to test $H: \theta \leq 1$.
- (a) Show that the likelihood ratio test rejects H when $\sum_{i=1}^n X_i < c$.
 - (b) What is the value of c for the test with level α ?
 - (c) Show that the power of this test is a strictly monotonic function of θ .
 - (d) Draw a graph of the power for $\alpha = 0.05$ and $n = 15$.
 - (e) Construct a test based on asymptotic results and compare it with the test based on exact results.
18. Let $\mathbf{X} = (X_1, \dots, X_{10})$ be a random sample from an unknown distribution. After observing $\mathbf{X} = (1, 0.7, 0.2, -1.3, -0.5, 1.52, -0.85, 0.25,$

- $0.47, -0.67)$, test whether one can assume that the data comes from a standard normal distribution. Check also whether one can assume that the data comes from a normal distribution. Obtain the p -values of the two tests and compare them.
19. Show that the t test of $H_0: \theta = \theta_0$ versus $H_1: \theta \neq \theta_0$ based on a random sample of size n from the $N(\theta, \sigma^2)$, θ and σ^2 unknown, is consistent.
20. Suppose that X_1, \dots, X_n are iid with $N(\theta, \sigma^2)$ and one wishes to test $H_0: \sigma = \sigma_0$ versus $H_1: \sigma \neq \sigma_0$.
- (a) Show that the MLR test with asymptotic level α accepts H_0 if $\sum_{i=1}^n (X_i - \bar{X})^2 / \sigma_0^2 \in [c_1, c_2]$ where c_1 and c_2 are such that $F(c_2) - F(c_1) = 1 - \alpha$ and F is the distribution function of the χ_{n-1}^2 distribution.
 - (b) What condition must be satisfied by c_1 and c_2 for an unbiased test?
 - (c) Show that the normal approximation gives $c_1 = n - \sqrt{2n}z_{\alpha/2}$ and $c_2 = n + \sqrt{2n}z_{\alpha/2}$ where $z_{\alpha/2}$ is the $1 - \alpha/2$ quantile of the $N(0, 1)$ distribution.
 - (d) Show that the equal tail test, where $F(c_2) = 1 - \alpha/2$ and $F(c_1) = \alpha/2$ is asymptotically unbiased, i.e. that the test is unbiased when $n \rightarrow \infty$.
21. Prove that in the case of multinomial observations, the statistics W_E and W_U are respectively given by
- $$W_E(\theta_0) = \sum_{i=1}^p \frac{(N_i - n\theta_{i,0})^2}{N_i} \quad \text{and} \quad W_U(\theta_0) = \sum_{i=1}^p \frac{(N_i - n\theta_{i,0})^2}{n\theta_{i,0}}$$
22. A contingency table is a table of multiple classification of observational units into cells. In the simplest case of double entry, the cells are defined by the intersection of two factors: A with levels A_1, \dots, A_p and B with levels B_1, \dots, B_r . Define the probabilities $\theta_{ij} = P(A = A_i, B = B_j)$, $\forall(i, j)$, and assume that n independent observations are made and each one of them is classified into a single cell (i, j) and the counts N_{ij} associated with the cells registered.
- (a) Show that $\{N_{ij}\}$ has multinomial distribution with parameters n and $\{\theta_{ij}\}$. Define the parametric space Θ and obtain the MLE's of $\{\theta_{ij}\}$, $\forall(i, j)$.
 - (b) The factors A and B are said to be independent if $P(A = A_i, B = B_j) = P(A = A_i)P(B = B_j)$. Define the parametric space Θ_0 under independence between A and B and calculate its dimension.
Hint: define $\theta_{i+} = \sum_{j=1}^r \theta_{ij}$ and $\theta_{+j} = \sum_{i=1}^p \theta_{ij}$.
 - (c) Calculate the MLE's of $\{\theta_{ij}\}$ under Θ_0 .
 - (d) Obtain the goodness-of-fit and MLR tests of level α , specifying the critical regions and the distributions involved.

7

Prediction

All that has been done up to now concerns estimation, that is, understanding a phenomenon through probabilistic assertions that relate directly or indirectly to unobserved quantities of interest. That means we can never be refuted. The quantities we deal with do not exist from a practical point of view. Their usefulness is only associated with the valuable, but insufficient on its own, help in describing in the best possible way the random process under study. An improvement was obtained when we criticized the adopted models by putting them under hypotheses tests but we are still restricted to the observed data. A real test of a theory or model is obtained when its assertions are applied to future experiences and observations. In this chapter, we will only deal with this topic: prediction.

Statistical prediction has a well-defined meaning and is an integral part of the inferential procedure as will be seen. It deals with making probabilistic statements about quantities to be observed in the future. Note that at the moment inference is made, the problem is similar to that already studied of parametric estimation. So, much of the material described in this chapter is a mere adaptation of the material from previous chapters. The big and fundamental difference here is that all statements will be confronted with reality and are subject to approval or dismissal without dispute.

Even though notions of decision theory were introduced when we dealt with Bayesian estimation, we think the appropriate moment for decision taking is when we make predictions. We have a clear view of the consequences and respective losses associated with our acts when we face our positions about quantities that will become known.

7.1 Bayesian prediction

The typical situation of the prediction problem is that in which a quantity X related to an unobserved quantity θ through $P_1(x | \theta)$ is observed and we are interested in producing statements about another random quantity Y that is related to X and θ through $P_2(Y | \theta, X)$. So, after observing $X = x$ we have updated information to make an inference about Y and this information is contained in the distribution

of $\mathbf{Y} | \mathbf{X}$. Therefore, the distribution of $\mathbf{Y} | \mathbf{X} = \mathbf{x}$ is

$$\begin{aligned} p(\mathbf{y} | \mathbf{x}) &= \int p(\mathbf{y}, \theta | \mathbf{x}) d\theta \\ &= \int p(\mathbf{y} | \theta, \mathbf{x}) p(\theta | \mathbf{x}) d\theta \\ &= \int p(\mathbf{y} | \theta) p(\theta | \mathbf{x}) d\theta \end{aligned}$$

with the last equality valid in the common case of conditional independence between \mathbf{X} and \mathbf{Y} given θ . That happens, for instance, when we sample future and past observations from the same population. In the trivial case in which it is possible to directly specify the distribution of $\mathbf{Y} | \mathbf{x}$, the above calculations involving the removal of unobserved quantities are not needed.

In the case of a random sample $\mathbf{X} = (X_1, \dots, X_n)$ from $p(\cdot | \theta)$ and a single future observation Y that also comes from the same population, that is, Y has density $p(\cdot | \theta)$, then

$$p(\mathbf{y} | \mathbf{x}) = \int p(\mathbf{y} | \theta) p(\theta | \mathbf{x}) d\theta \quad \text{where } p(\theta | \mathbf{x}) \propto p(\theta) \prod_{i=1}^n p(x_i | \theta)$$

and $p(\mathbf{y} | \theta)$ and $p(x_i | \theta)$ have the same form. One can then write

$$p(\mathbf{y} | \mathbf{x}) = E_{\theta | \mathbf{x}} [p(\mathbf{y} | \theta)].$$

In the case of prediction in samples from the one-parameter exponential family with conjugate prior, we have, using the notation of Section 3.2 that

$$\begin{aligned} p(x | \theta) &= a(x) \exp\{u(x)\phi(\theta) + b(\theta)\} \\ p(\theta) &= k(\alpha, \beta) \exp\{\alpha\phi(\theta) + \beta b(\theta)\} \end{aligned}$$

and so

$$\begin{aligned} p(\theta | \mathbf{x}) &= k \left(\alpha + \sum_{i=1}^n u(x_i), \beta + n \right) \exp \left\{ \left[\alpha + \sum_{i=1}^n u(x_i) \right] \phi(\theta) + [\beta + n] b(\theta) \right\} \end{aligned}$$

and

$$p(\mathbf{x}) = \prod_{i=1}^n a(x_i) \frac{k(\alpha, \beta)}{k(\alpha + \sum_{i=1}^n u(x_i), \beta + n)}.$$

The expression of $p(\mathbf{y} | \mathbf{x})$ is obtained similarly to $p(\mathbf{x})$, using the posterior density $p(\theta | \mathbf{x})$ instead of $p(\theta)$. So

$$p(\mathbf{y} | \mathbf{x}) = a(y) \frac{k(\alpha + \sum_{i=1}^n u(x_i), \beta + n)}{k(\alpha + \sum_{i=1}^n u(x_i) + u(y), \beta + n + 1)}.$$

Example. Let X_1, \dots, X_n be a sample from the exponential distribution with parameter θ . The conjugate distribution is $G(\gamma, \delta)$. The densities can be written in the form

$$\begin{aligned} p(x | \theta) &= \exp\{-x\theta + \log \theta\}, \quad x > 0 \\ p(\theta) &= \frac{\delta^\gamma}{\Gamma(\gamma)} \exp\{-\delta\theta + (\gamma - 1)\log \theta\}, \quad \theta > 0. \end{aligned}$$

Identifying with the above notation, we have

$$\begin{aligned} \alpha &= -\delta \\ \beta &= \gamma - 1 \\ k(\alpha, \beta) &= (-\alpha)^{\beta+1} / \Gamma(\beta + 1) \\ a(x) &= 1 \\ u(x) &= -x \end{aligned}$$

from which we get

$$\begin{aligned} p(\mathbf{y} | \mathbf{x}) &= \frac{(-\alpha - \sum_{i=1}^n u(x_i))^{\beta+n+1} / \Gamma(\beta + n + 1)}{(-\alpha - \sum_{i=1}^n u(x_i) - u(y))^{\beta+n+2} / \Gamma(\beta + n + 2)}, \quad y > 0 \\ &= \frac{\beta + n + 1}{-\alpha - \sum_{i=1}^n u(x_i) - u(y)} \\ &\quad \times \left(1 + \frac{u(y)}{\alpha + \sum_{i=1}^n u(x_i)} \right)^{-(\beta+n+1)}, \quad y > 0. \end{aligned}$$

Rewriting as a function of γ and δ leads to

$$p(\mathbf{y} | \mathbf{x}) = \frac{\gamma + n}{\delta + \sum_{i=1}^n x_i + y} \left(1 + \frac{y}{\delta + \sum_{i=1}^n x_i} \right)^{-(\gamma+n)}, \quad y > 0.$$

Note that this density is strictly decreasing with y .

With the predictive distribution, we can proceed to an inference as previously seen in Chapter 4. In particular, we can make an inference by point prediction. To do this, we put the problem into the framework of decision theory whose elements are:

1. States of nature – here, represented by the possible values of the quantity \mathbf{Y} to be observed in the future and that we wish to predict.
2. Space of possible actions – containing possible actions to be taken. Here, taking an action is to choose a value for \mathbf{Y} , its point predictor δ .
3. Loss function – to each possible value of \mathbf{Y} and to each predictor δ we have a loss $L(\mathbf{Y}, \delta)$.

Put in this form, the problem is mathematically identical to that from Section 4.1, that is, we choose the predictor δ so as to minimize the expected loss

$$\int L(y, \delta) p(y | x) dy.$$

This decision theoretic approach to prediction is pursued by Aitchison and Dunsmore (1975) and Geisser (1993). The difference with respect to parameter estimation is not in the equation but in the interpretations of its elements. We can objectively quantify the loss we will incur by predicting Y by δ because Y is observable. This quantification is less clear in the case of estimation and only makes sense when related to observed quantities associated with the parameters. The example of John and his doctor in Section 4.1 illustrates this point. It was only possible to construct the loss table after referring to observed quantities in the future such as death and definition of the disease state of the patient.

That is due to the fact that we do not take decisions against values of theoretical objects used to understand a phenomenon (parameters) but only after evaluating the consequences these values will have upon observables. This point is the basis of the predictivist approach to inference which itself is not free from arguments. Its advantage however is that it allows judgments that are not ambiguous with a clear and unquestionable meaning. A prediction can always be confronted against reality while estimation never is.

So, a point predictor can be chosen according to the loss function we incur. Using results from Section 4.1, we have that:

1. The predictor associated with quadratic loss is the mean of the predictive distribution.
2. The predictor associated with absolute loss is the median of the predictive distribution.
3. The predictor associated with the 0-1 loss is the mode of the predictive distribution.

Example (continued). The predictors described above are given by:

(a)

$$\begin{aligned} E(Y | x) &= E_{\theta|x}[E(Y | \theta)] \\ &= E_{\theta|x}[1/\theta] \\ &= \frac{\delta + \sum_{i=1}^n x_i}{\gamma + n - 1}. \end{aligned}$$

(b) The solution med of the equation

$$\frac{1}{2} = \int_0^{\text{med}} p(y | x) dy$$

given by $\text{med} = (\delta + \sum_{i=1}^n x_i)(2^{1/\gamma+n} - 1)$.

(c) by 0, the mode of the predictive distribution of $Y | x$.

Observe that for large n , the predictive mean is approximately equal to \bar{x} and for the non-informative prior ($\gamma, \delta \rightarrow 0$) the predictive mean is

$$E(Y | x) = \frac{\sum_{i=1}^n x_i}{n - 1}.$$

In analogy to the estimation problem, 100(1 - α)% predictive regions may be obtained for Y . All that is needed is to find a region C such that $P(Y \in C | x) \geq 1 - \alpha$. In the case of a scalar Y , the region C may be reduced to an interval $[a_1, a_2]$ satisfying $P(a_1 < Y < a_2 | x) \geq 1 - \alpha$. The same comments are still valid here that for a given value of α one wishes to find the interval with smallest possible length, leading naturally to the choice of regions where their predictive density is higher. This is formalized by the concept of regions of highest predictive density (HPRD) C given by

$$C = \{y : p(y | x) \geq k(\alpha)\}$$

where $k(\alpha)$ is the highest constant guaranteeing that $P(Y \in C | x) \geq 1 - \alpha$.

Example (continued). As the predictive density is strictly decreasing, the HPRD interval for Y must be in the form $[0, a]$ where a is such that

$$\int_0^a p(y | x) dy = 1 - \alpha.$$

Solving for a gives $(\delta + \sum_{i=1}^n x_i)((1 - \alpha)^{-1/\gamma+n} - 1)$.

In the case of a large sample, the posterior distribution of $\theta | x$ gets concentrated around $\hat{\theta}$, the MLE of θ . So,

$$p(y | x) = \int p(y | \theta) p(\theta | x) d\theta \doteq p(y | \hat{\theta}).$$

This approximation neglects the variability of the parameter and leaves only the sampling variability of the quantity to be predicted. In the general case, both forms of variability are important and should be taken into account.

7.2 Classical prediction

There are no clear rules as to how to proceed to make prediction of future observations in classic inference. One of the most frequently used procedures is to substitute the value of the parameter appearing in the sampling distribution of future observations by some estimate based on past data.

Specifically, assuming that Y with sampling distribution $p(y | \theta)$ must be predicted based on observations X from a sample of $p(x | \theta)$, one uses the distribution $p(y | \hat{\theta})$ where $\hat{\theta}$ is an estimator of θ . The most common choice is the MLE, which takes us back to the discussion in the final paragraph of the last section. The drawback of this procedure is not to take into account the variability associated with the estimation of θ .

As an example, assume that $E(\mathbf{Y} | \theta) = \mathbf{b}(\theta)$ and that $V(\mathbf{Y} | \theta) = \mathbf{B}(\theta)$. It is common practice to take $\mathbf{b}(\hat{\theta})$ as a point predictor of \mathbf{Y} and $\mathbf{B}(\hat{\theta})$ as a measure of the variability associated with that prediction. This procedure underestimates the variability of the prediction by not taking into account the variability of the estimation of θ .

Example (continued). The maximum likelihood estimator of θ is $1/\bar{X}$ and the mean of \mathbf{Y} is $1/\theta$. After doing the substitutions we get the point predictor \bar{X} for \mathbf{Y} . Note that the estimator of the variance of \mathbf{Y} , $1/\bar{X}^2$, does not consider the variability of $\hat{\theta}$ with respect to θ .

One approach avoiding this problem consists in obtaining a pivotal quantity $G(\mathbf{Y}, \mathbf{X})$ whose distribution does not depend on θ . This approach was used in Section 4.4 in interval estimation. As before, the function G must depend on the sample \mathbf{X} in an optimal form (for instance, through minimal sufficient statistics for θ). Once the pivot is obtained, one can make probabilistic statements about it. In particular, for a given value of α ($\alpha \in (0, 1)$), one can obtain that $P(G(\mathbf{Y}, \mathbf{X}) \in C) \geq 1 - \alpha$. Then, it becomes possible to construct a $100(1 - \alpha)\%$ predictive region for \mathbf{Y} given by

$$\{y : G(y, \mathbf{x}) \in C\}.$$

The problem with this approach is that it is not always possible to find such a function G . The example we are dealing with in this chapter is one of the exceptions.

Example (continued). We now want to find a function $G(\mathbf{Y}, \mathbf{X})$ whose distribution does not depend on θ . Preferably, this function would depend on \mathbf{X} through \bar{X} which is a minimal sufficient statistic for θ . Fortunately, in this case, this is possible since

$$Y \sim \text{Exp}(\theta) = G(1, \theta) = \frac{\chi_2^2}{2\theta} \quad \text{and} \quad \sum_{i=1}^n X_i \sim G(n, \theta) = \frac{\chi_{2n}^2}{2\theta}$$

and they are independent. So, $Y/\bar{X} \sim F(2, 2n)$, which does not depend on θ . Using the properties of the F distribution, we have that $E(Y/\bar{X}) = 2n/(2n - 2)$ from where we can take as a point predictor for Y

$$\bar{X} \frac{2n}{2n - 2} = \frac{\sum_{i=1}^n X_i}{n - 1}$$

which differs from the predictor obtained above but coincides with the Bayesian predictor with non-informative prior. The $100(1 - \alpha)\%$ confidence interval for Y can be constructed with the percentiles $\underline{F}_{\alpha/2}(2, 2n)$ and $\bar{F}_{\alpha/2}(2, 2n)$ as $P(\underline{F}_{\alpha/2}(2, 2n) < Y/\bar{X} < \bar{F}_{\alpha/2}(2, 2n)) = 1 - \alpha$. The prediction interval is given by $(\bar{X}\underline{F}_{\alpha/2}(2, 2n), \bar{X}\bar{F}_{\alpha/2}(2, 2n))$.

7.3 Prediction in the normal model

7.3.1 Bayesian approach

The structure of normal models allows many easy calculations especially in the case of linearity. The application of these rules to prediction reduces the calculation considerably. Basically consider an observation $\mathbf{Y} | \theta \sim N(\theta, \sigma^2)$ where σ^2 is known. One can rewrite this model as $\mathbf{Y} = \theta + \epsilon$ where $\epsilon \sim N(0, \sigma^2)$ can be regarded as an observation error. Observe that the distribution of ϵ does not depend on θ and therefore ϵ and θ are independent. Assume now that the updated distribution of θ (possibly after the observation of a sample) is $N(\mu, \tau^2)$. \mathbf{Y} is therefore the sum of independent normal quantities and has a $N(\mu, \sigma^2 + \tau^2)$ distribution. This is the predictive distribution of \mathbf{Y} . Point and HPRD interval predictions can be made as described in Section 7.1.

Example. Assume the updated distribution of θ is the posterior distribution relative to a sample \mathbf{X} from a $N(\theta, \sigma^2)$ distribution. This leads to $\mathbf{Y} | \mathbf{x} \sim N(\mu_1, \sigma^2 + \tau_1^2)$ where μ_1 and τ_1^2 are given by Theorem 2.1. The point predictor of \mathbf{Y} in this case will be μ_1 and the $100(1 - \alpha)\%$ HPRD interval for \mathbf{Y} is of the form $(\mu_1 - z_{\alpha/2}\sqrt{\sigma^2 + \tau_1^2}, \mu_1 + z_{\alpha/2}\sqrt{\sigma^2 + \tau_1^2})$.

In the case of a non-informative prior, $\mu_1 = \bar{x}$ and $\tau_1^2 = \sigma^2/n$, leading to the predictive distribution

$$\mathbf{Y} | \mathbf{x} \sim N\left(\bar{x}, \sigma^2 \left(1 + \frac{1}{n}\right)\right).$$

The point predictor of \mathbf{Y} in this case will be \bar{x} and the $100(1 - \alpha)\%$ HPRD interval for \mathbf{Y} is of the form $(\bar{x} - z_{\alpha/2}\sigma\sqrt{1 + n^{-1}}, \bar{x} + z_{\alpha/2}\sigma\sqrt{1 + n^{-1}})$.

The above example can be generalized in various forms. If \mathbf{Y} is a vector with sampling distribution $N(\theta, \Sigma)$ and the updated distribution of θ is $N(\mu, \tau)$, we have by the same reasoning that the predictive distribution of \mathbf{Y} is $N(\mu, \Sigma + \tau)$. Assume now that the sampling mean of \mathbf{Y} is given by the linear relation $\mathbf{X}\theta$ where \mathbf{X} is a matrix of known constants and that θ remains with the same distribution. If the matrix \mathbf{X} is square, the dimension of θ and the hyperparameters of its distribution remain unaltered. This restriction is unnecessary and we can consider any matrix \mathbf{X} with the correspondent change in the dimension and distribution of θ . We shall see in the next chapter that the cases of interest involve a lower dimension of θ than the dimension of \mathbf{Y} implying a reduction in the dimensionality of the problem. Schematically,

$$\mathbf{Y} = \mathbf{X}\theta + \epsilon$$

where θ and ϵ are independent. So, the predictive distribution of \mathbf{Y} is given by the sum of two independent normal quantities, the first relative to θ given by a $N(\mathbf{X}\mu, \mathbf{X}\tau\mathbf{X}' + \Sigma)$ distribution and so

$$\mathbf{Y} | \mathbf{X} \sim N(\mathbf{X}\mu, \mathbf{X}\tau\mathbf{X}' + \Sigma).$$

An example is the case of prediction of a series of future observations from a population from which a sample had previously been observed.

Up to now, the sampling variance was assumed known. Generally it is not and the usual approach in this case is to try to specify a conjugate distribution for the variance. Although a conjugate analysis when the sampling covariance matrix is totally unknown is possible, we will consider here only the case where the covariance matrix is totally known but for an unknown multiplicative scalar σ^2 . Without loss of generality we will assume that $V(\mathbf{Y} | \theta, \sigma^2) = \sigma^2 \mathbf{I}_p$, where p is the dimension of \mathbf{Y} . For a conjugate analysis, it is necessary that the updated covariance matrix of θ be proportional to σ^2 as seen in Chapter 3. So, $\theta | \sigma^2 \sim N(\mu, \sigma^2 \tau)$ and $\mathbf{Y} | \sigma^2 \sim N(\mathbf{X}\mu, \sigma^2(\mathbf{X}\tau\mathbf{X}^T + \mathbf{I}_p))$. Assuming as before that the updated distribution of σ^2 is $\eta_0 \sigma_0^2 \phi \sim \chi_{n_0}^2$ with $\phi = \sigma^{-2}$,

$$\begin{aligned} p(\mathbf{y}) &= \int p(\mathbf{y} | \phi) p(\phi) d\phi \\ &\propto \phi^{p/2} \exp\{-\phi Q(\mathbf{y})/2\} \phi^{(n_0/2)-1} \exp\{-\phi n_0 \sigma_0^2 / 2\} d\phi \\ &= \int \phi^{(n_0+p/2)-1} \exp\{-\phi[n_0 \sigma_0^2 + Q(\mathbf{y})]/2\} d\phi \\ &\propto [n_0 \sigma_0^2 + Q(\mathbf{y})]^{-(n_0+p)/2} \end{aligned}$$

where $Q(\mathbf{y}) = (\mathbf{y} - \mathbf{X}\mu)^T(\mathbf{X}\tau\mathbf{X}^T + \mathbf{I}_p)^{-1}(\mathbf{y} - \mathbf{X}\mu)$.

It is easy to obtain then that

$$\mathbf{Y} \sim t_{n_0}(\mathbf{X}\mu, \sigma_0^2(\mathbf{X}\tau\mathbf{X}^T + \mathbf{I}_p)).$$

Note that the only changes with respect to the known variance case are the substitutions of the normal by the Student t distribution with n_0 degrees of freedom and of σ^2 by its updated estimator σ_0^2 .

The main properties of the multivariate Student t distribution that are relevant here are, if $\mathbf{U} \sim t_v(m, \mathbf{C})$ then

- (i) the marginal distribution of any q -dimensional subvector from \mathbf{U} ($q < p$), say \mathbf{U}_1 , is also Student t with v degrees of freedom and parameters \mathbf{m}_1 and \mathbf{C}_1 obtained from the components of the vector \mathbf{m} and of the matrix \mathbf{C} corresponding to the components of the vector \mathbf{U}_1 . In particular, the j th component of \mathbf{U} , U_j , has (univariate) $t_v(m_j, C_{jj})$ distribution.
- (ii) $\mathbf{LU} \sim t_v(\mathbf{L}\mathbf{m}, \mathbf{L}\mathbf{C}\mathbf{L}')$, for any $r \times p$ matrix \mathbf{L} of full rank.

Example (continued). The same results are valid with the modifications cited above. As seen in Section 3.3, the posterior distribution of (θ, ϕ) is given by the normal- χ^2 with parameters $(\mu_1, c_1, n_1, \sigma_1)$. A future observation \mathbf{Y} will have predictive distribution

$$\mathbf{Y} | \mathbf{x} \sim t_{n_1}(\mu_1, \sigma_1^2(1 + c_1^{-1})).$$

In the case of a non-informative prior, the predictive distribution is given by a $t_{n-1}(\bar{x}, s^2(1 + n^{-1}))$ distribution where s^2 is the unbiased estimator of σ^2 . The HPRD intervals coincide with those given for known σ^2 but for the substitutions of the percentiles of the $N(0, 1)$ for those of the $t_{n-1}(0, 1)$ distribution and of σ by s .

7.3.2 Classical approach

For the frequentist inference, one should only work with the sampling distributions of \mathbf{Y} and \mathbf{X} in search of a pivot whose distribution does not depend on the parameters. Again, we write $\mathbf{Y} - \theta \sim N(0, \sigma^2)$ and all that is left is to find suitable estimators for θ and σ . Assuming initially that σ^2 is known we have the estimator $\hat{\theta} = \bar{X}$. As \mathbf{Y} and \bar{X} are both independent normal with same mean, $\mathbf{Y} - \bar{X} \sim N(0, \sigma^2(1 + n^{-1}))$ is the pivot with distribution identical to that obtained with a non-informative prior. Confidence intervals for \mathbf{Y} will also coincide. The difference in the approaches is theoretical: the Bayesian result is conditional on \mathbf{X} which in fact is the way it will be used in classical inference too.

In the case where σ^2 is unknown, another pivot must be found because the distribution of $\mathbf{Y} - \bar{X}$ depends on σ^2 . In the normal case, this is easy because S^2 is independent from $\mathbf{Y} - \bar{X}$ and $(n-1)S^2/\sigma^2 \sim \chi_{n-1}^2$. So,

$$\frac{\mathbf{Y} - \bar{X}}{S\sqrt{1 + 1/n}} \sim t_{n-1}(0, 1)$$

is the required pivot. Again, we get the same results previously obtained with a non-informative prior.

The results of the example above were obtained only for the simplest case of a single future observation following the observation of a sample from the same population. They can be extended to more general situations. Some of these extensions will be seen in the next chapter.

In the purely predictive case without unknown parameters, we have \mathbf{Y} and \mathbf{X} with joint distribution

$$\begin{pmatrix} \mathbf{Y} \\ \mathbf{X} \end{pmatrix} \sim N \left[\begin{pmatrix} \mu_Y \\ \mu_X \end{pmatrix}, \begin{pmatrix} \Sigma_Y & \Sigma_{YX} \\ \Sigma_{XY} & \Sigma_X \end{pmatrix} \right]$$

and the prediction of \mathbf{Y} given the observation of \mathbf{X} is based on the distribution of $\mathbf{Y} | \mathbf{X} = \mathbf{x}$ given by

$$N \left(\mu_Y + \Sigma_{XY} \Sigma_X^{-1}(\mathbf{x} - \mu_X), \Sigma_Y - \Sigma_{XY} \Sigma_X^{-1} \Sigma_{YX} \right).$$

If Y and X are scalar quantities

$$Y | X = x \sim N \left(\mu_Y + \frac{\sigma_{XY}}{\sigma_X^2} (x - \mu_X), \sigma_Y^2 - \frac{\sigma_{XY}^2}{\sigma_X^2} \right)$$

where $\sigma_X^2 = V(X)$, $\sigma_Y^2 = V(Y)$ and $\sigma_{XY} = \text{Cov}(X, Y)$. We can see from the results above that the knowledge of the value of X reduces the variance of Y . The larger the correlation between X and Y in absolute value the greater the reduction. More important than that is the fact that the optimal predictor (under any reasonable criteria) is a linear function of X . This result is explored in the next section.

7.4 Linear prediction

The problem of linear prediction takes place in the absence of complete distributional information. Assume the presence of two random vectors \mathbf{X} and \mathbf{Y} from which it is only known that

$$\begin{aligned} E\left(\begin{array}{c} \mathbf{Y} \\ \mathbf{X} \end{array}\right) &= \left(\begin{array}{c} \mu_Y \\ \mu_X \end{array}\right) \\ V\left(\begin{array}{c} \mathbf{Y} \\ \mathbf{X} \end{array}\right) &= \left(\begin{array}{cc} \Sigma_Y & \Sigma_{YX} \\ \Sigma_{XY} & \Sigma_X \end{array}\right). \end{aligned}$$

The complete distribution of \mathbf{X} and \mathbf{Y} may not be known for a number of reasons. Our problem here is to establish a form of predicting \mathbf{Y} based on \mathbf{X} . A few restrictions must be imposed to solve it. The most natural one is to restrict the class of predictors to linear functions of \mathbf{X} . This restriction is not only justified by rendering the problem tractable. Linear solutions are obtained as first-order approximations. The linear predictor thus serves as a preliminary predictor even in the case where the joint distribution is completely specified.

To choose the predictor among all predictors of the form $\hat{\mathbf{Y}}(\mathbf{X}) = \mathbf{a} + \mathbf{B} \mathbf{X}$ requires the definition of an optimality criterion. It is reasonable that this criterion be based on the quadratic risk (or expected loss) given by

$$\mathbf{R}(\mathbf{Y}, \hat{\mathbf{Y}}) = E[\mathbf{Y} - \hat{\mathbf{Y}}(\mathbf{X})][\mathbf{Y} - \hat{\mathbf{Y}}(\mathbf{X})]'$$

and the optimality obtained by the minimization of the trace of the matrix \mathbf{R} . The search of the optimal predictor reduces to the search of optimal constants \mathbf{a} and \mathbf{B} . The predictor so obtained is called a linear predictor and the associated risk, the linear risk. Observe that the predictor that (globally) minimizes the risk is the conditional expectation $E(\mathbf{Y} | \mathbf{X})$ and the associated risk is the conditional variance. For that reason, the linear predictor is also called the linear expectation and its risk, the linear variance. In the specific case of the normal distribution, the conditional expectation is a linear function of \mathbf{X} . So, the linear predictor minimizes the quadratic expected loss among all possible predictors (linear or not).

In the special case where Y and X are scalar, the risk is given by

$$R(Y, \hat{Y}) = E[Y - \hat{Y}(X)]^2 = E[Y - (a + bX)]^2.$$

In this case, $\text{trace}(\mathbf{R}) = R$, and to find the values of a and b that minimize R we calculate

$$\frac{\partial R}{\partial a} = E[2(Y - a - bX)(-1)] = 2[a + b\mu_X - \mu_Y]$$

and

$$\frac{\partial R}{\partial b} = E[2(Y - a - bX)(-X)] = 2[a\mu_X + bE(X^2) - E(XY)].$$

Equating the derivatives to zero gives the system of equations

$$\hat{a} = \mu_Y - \hat{b}\mu_X$$

and

$$\hat{b} = \frac{E(XY) - \hat{a}\mu_X}{E(X^2)}.$$

Replacing the first in the second equation gives

$$\begin{aligned} \hat{b} &= \frac{E(XY) - (\mu_Y - \hat{b}\mu_X)\mu_X}{E(X^2)} \\ &= \frac{E(XY) - \mu_X\mu_Y + \hat{b}\mu_X^2}{E(X^2)} \\ &= \frac{\sigma_{XY} + \hat{b}\mu_X^2}{E(X^2)} \end{aligned}$$

which has solution $\hat{b} = \sigma_{XY}/\sigma_X^2$ where $\sigma_{XY} = \text{Cov}(X, Y)$. So the optimal linear predictor is

$$\hat{Y} = \mu_Y + \frac{\sigma_{XY}}{\sigma_X^2}(X - \mu_X)$$

and its risk is given by

$$\begin{aligned} R(Y, \hat{Y}) &= E[Y - \mu_Y - \frac{\sigma_{XY}}{\sigma_X^2}(X - \mu_X)]^2 \\ &= E(Y - \mu_Y)^2 + \left(\frac{\sigma_{XY}}{\sigma_X^2}\right)^2 E(X - \mu_X)^2 \\ &\quad - 2\frac{\sigma_{XY}}{\sigma_X^2} E[(Y - \mu_Y)(X - \mu_X)] \\ &= \sigma_Y^2 + \left(\frac{\sigma_{XY}}{\sigma_X^2}\right)^2 \sigma_X^2 - 2\frac{\sigma_{XY}}{\sigma_X^2} \sigma_{XY} \\ &= \sigma_Y^2 - \frac{\sigma_{XY}^2}{\sigma_X^2}. \end{aligned}$$

In the multivariate case, it can be shown that the linear predictor is

$$\mu_Y + \Sigma_{XY} \Sigma_X^{-1}(\mathbf{x} - \mu_X)$$

and its risk is given by

$$\Sigma_Y - \Sigma_{XY} \Sigma_X^{-1} \Sigma_{YX}$$

generalizing the result obtained when Y and X are scalar. As mentioned above, the linear predictor is given by the expression of the conditional expectation in the normal case. This implies that the linear predictor is as close to the global optimum (in terms of minimization of the quadratic risk) as the normal distribution is close to the joint distribution of \mathbf{X} and \mathbf{Y} .

The method of linear prediction was developed in the 1970s by Hartigan and Goldstein among others. It does not depend on the point of view adopted for inference. However, it generally appears in the Bayesian literature where it receives the name of linear Bayes methodology.

Exercises

§ 7.1

1. Let $\mathbf{X} = (X_1, \dots, X_n)$ be a random sample from a $U[0, \theta]$ distribution and $\theta \sim \text{Pa}(\alpha, \beta)$.
 - (a) Obtain the predictive distribution for a new observation from the same population based on all the information.
 - (b) Assuming that $\max x_i > \beta$, what is the probability of $\mathbf{Y} > \max x_i$?
 - (c) Repeat (a) and (b) for $\alpha \rightarrow 0$ and $\beta \rightarrow 0$.
2. Show that the problem of choice of a point predictor for Y may be reformulated as that of finding the predictor δ that minimizes

$$\int V(\theta, \delta) p(\theta | \mathbf{x}) d\theta$$

where

$$V(\theta, \delta) = E_{Y|\theta}[L(Y, \delta)].$$

3. Consider the problem of choosing the point predictor of $\mathbf{Y} \in R^p$ based on the quadratic loss

$$L(\mathbf{Y}, \delta) = (\mathbf{Y} - \delta)' \mathbf{M} (\mathbf{Y} - \delta)$$

where \mathbf{M} is a known positive definite matrix. Show that the point predictor is given by $\delta = E(\mathbf{Y} | \mathbf{x})$ and that the expected loss is given by the trace of the matrix $\mathbf{M} V(\mathbf{Y} | \mathbf{x})$.

§ 7.2

4. Assume that one wishes to make a prediction for $Y \sim \text{bin}(m, \theta)$ based on observation of $X \sim \text{bin}(n, \theta)$. How can that be done from a frequentist point of view?
5. A geologist wishes to study the incidence of seismic movements in a given region. He then selects m independent but geologically similar observation points and counts the number of movements in a specific time interval. The observational model is $X_i \sim \text{Pois}(\theta)$, where $X_i, i = 1, \dots, m$, is the

number of occurrences in the i th observation point and θ is the average rate of seismic movements. From his previous experience, the researcher assumes that $E[\theta] = 2$ movements per time interval and that $V[\theta] = 0.36$ and uses these values to specify a conjugate prior.

- (a) Assuming that $x = (2, 3, 0, 0, 1, 0, 2, 0, 3, 0, 1, 2)$ was observed, what is the posterior distribution?
- (b) He wishes to predict the expected number of seismic movements and its precision in an $(m+1)$ th site based on the observation he had made. Establish the necessary hypotheses and perform the calculations using the data above.

§ 7.3

6. Consider the observation of independent samples $\mathbf{Y} = (Y_1, \dots, Y_m)$ and $\mathbf{X} = (X_1, \dots, X_n)$ from a $N(\theta, \sigma^2)$ population. Indicate how to obtain point and confidence interval predictors for the components of \mathbf{Y} using:

- (a) a conjugate prior for (θ, σ^2) ;
- (b) a non-informative prior for (θ, σ^2) ;
- (c) frequentist inference.

§ 7.4

7. Prove that, in the multivariate case, the linear predictor is given by

$$\mu_Y + \Sigma_{XY} \Sigma_X^{-1} (\mathbf{x} - \mu_X)$$

and its risk is given by

$$\Sigma_Y - \Sigma_{XY} \Sigma_X^{-1} \Sigma_{YX}.$$

8

Introduction to linear models

In this chapter one of the most important problems in statistics will be considered. The scope is introductory because it deserves a whole book to be considered in depth. A complete approach to the subject can be found in books such as Draper and Smith (1966) and Broemeling (1985).

The class of normal linear models is characterized in the first section. In this same section the class of generalized linear models will also be introduced, as an extension of the normal linear models to the exponential family. In the following sections, classical and Bayesian inferences are developed for these models. Two other broad classes of models, natural extensions of the normal linear models, are described in Sections 8.4 (hierarchical linear models) and 8.5 (dynamic linear models).

8.1 The linear model

In this section, the problem of observing a random variable Y with values affected by other variables will be discussed. For example, the income of a firm is affected by its capital and by the number of persons it employs, the production of a machine is influenced by the maintenance scheme implemented and how well trained its operator is, the arterial blood pressure depends upon the age of the patient, and so on.

In these cases, the variability of Y is explained by some other variables. As a first approximation only a linear relationship will be used to describe how these variables influence Y . Let X_1, \dots, X_p denote the set of p explanatory variables. It follows that

$$E(Y) = \beta_1 X_1 + \dots + \beta_p X_p.$$

If a model with an intercept is required, all we need to do is to specify $X_1 = 1$. The model introduced above is called a linear model or linear regression model. The case with $p = 1$, where only one explanatory variable is involved, is named simple linear regression. Note that the expectation of Y is calculated conditionally on the values of the variables X_1, \dots, X_p . Throughout this chapter it will be assumed that the values of the explanatory variable are known.

If, in addition, a common $V(Y_i) = \sigma^2$ is assumed for all i , the errors, after observing the sample Y_1, \dots, Y_n can be specified as

$$e_i = Y_i - (\beta_1 x_{1i} + \dots + \beta_p x_{pi}), i = 1, \dots, n$$

and the parameters $\beta = (\beta_1, \dots, \beta_p)'$ can easily be estimated by least squares. In this case, the estimator is given by the value of β that minimizes $\sum_{i=1}^n e_i^2$. Note that no hypothesis about the distribution of Y or the errors was needed. If the hypotheses of normality and independence of the distribution of the Y_i 's are assumed, then it is easy to obtain the likelihood function as

$$l(\beta, \sigma^2; Y_1, \dots, Y_n) \propto \sigma^{-n} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n e_i^2 \right\}.$$

Then, the value of β that maximizes the likelihood is equivalent to the value that minimizes $\sum_{i=1}^n e_i^2$, that is, the maximum likelihood and the least squares estimators coincide.

It is useful, for further development, to adopt a matrix notation. Let us define

$$\mathbf{Y} = \begin{pmatrix} Y_1 \\ \vdots \\ Y_n \end{pmatrix} \quad \text{and} \quad \mathbf{X} = \begin{pmatrix} \mathbf{x}'_1 \\ \vdots \\ \mathbf{x}'_n \end{pmatrix} = \begin{pmatrix} x_{11} & \cdots & x_{p1} \\ \vdots & \ddots & \vdots \\ x_{1n} & \cdots & x_{pn} \end{pmatrix}$$

from which it follows that $\mathbf{Y} | \beta, \sigma^2 \sim N(\mathbf{X}\beta, \sigma^2 \mathbf{I}_p)$ where \mathbf{I}_p is the $p \times p$ identity matrix. The likelihood equation can be rewritten as $\sigma^{-n} \exp\{-S(\beta)/2\sigma^2\}$, where

$$\begin{aligned} S(\beta) &= \sum_{i=1}^n (Y_i - \mathbf{x}'_i \beta)^2 \\ &= (\mathbf{Y} - \mathbf{X}\beta)'(\mathbf{Y} - \mathbf{X}\beta). \end{aligned}$$

The model presented before can be rewritten through the following equations

$$Y_i \sim N(\mu_i, \sigma^2), i = 1, \dots, n, \text{ independent}$$

$$\mu_i = \lambda_i, i = 1, \dots, n$$

$$\lambda_i = \mathbf{x}'_i \beta$$

where the (apparently unnecessary) second equation states the relationship between the mean of the observations and the explanatory structure in the model. Stating the main objective in this form it is clear that there is no reason to be restricted to the normal distribution and to the class of linear relationships. One of the most relevant extensions in the model presented before constitutes the class of generalized linear models, which allows us to model observations described by any member of the exponential family and to relate the mean of the distribution

to the explanatory relationship through any differentiable function. Therefore the Y_i 's have density given by

$$\begin{aligned} p(y_i | \theta) &= a(y_i) \exp\{u(y_i)\phi(\theta) + b(\theta)\}, i = 1, \dots, n, \text{ independent} \\ g(\mu_i) &= \lambda_i \text{ where } g \text{ is differentiable and } \mu_i = E(Y_i | \theta) \\ \lambda_i &= \mathbf{x}'_i \beta. \end{aligned}$$

This class is broad enough to include many of the most frequently used models in statistics. A complete description of these models, including many inferential aspects, can be found in McCullagh and Nelder (1989).

8.2 Classical linear models

Classical estimators for β and σ^2 can be obtained from the likelihood function exhibited before. Beginning with the estimation of β we see that the least square and maximum likelihood estimators do coincide and are given by the value of β that minimizes the quadratic form $S(\beta)$. Differentiating this expression with respect to the elements of the parameter vector β , it follows that

$$\frac{\partial S(\beta)}{\partial \beta} = 2(\mathbf{X}'\mathbf{X}\beta - \mathbf{X}'\mathbf{Y})$$

(see Exercise 8.1). Since the matrix of second derivatives is positive definite, the solution of $\partial S(\beta)/\partial \beta = \mathbf{0}$ provides the point of minimum $\hat{\beta}$ of $S(\beta)$ and must satisfy

$$\mathbf{X}'\mathbf{X}\hat{\beta} = \mathbf{X}'\mathbf{Y}.$$

The above p equations are known as the normal equations. If the matrix $\mathbf{X}'\mathbf{X}$ is of full rank, or if the columns of \mathbf{X} are linearly independent, then $\mathbf{X}'\mathbf{X}$ has an inverse and the maximum likelihood estimator of β is given by

$$\hat{\beta} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y}.$$

It will be assumed that the matrix $\mathbf{X}'\mathbf{X}$ has full rank, from here on. This restriction does not seem to be that serious since \mathbf{X} not of full rank means that there are some redundancies in its specification or in the model specification. These redundancies can be useful for understanding the model but can be eliminated without any loss from the following inferences. The quadratic form can be expressed as

$$S(\beta) = (\beta - \hat{\beta})'\mathbf{X}'\mathbf{X}(\beta - \hat{\beta}) + S_e$$

where

$$\begin{aligned} S_e &= \mathbf{Y}'\mathbf{Y} - \hat{\beta}'\mathbf{X}'\mathbf{X}\hat{\beta} \\ &= (\mathbf{Y} - \mathbf{X}\hat{\beta})'(\mathbf{Y} - \mathbf{X}\hat{\beta}). \end{aligned}$$

The maximum likelihood estimator of σ^2 is obtained as the solution of the equation

$$\frac{\partial \log l(\hat{\beta}, \hat{\sigma}^2; \mathbf{Y})}{\partial \sigma^2} = -\frac{n-1}{2\hat{\sigma}^2} + \frac{S(\hat{\beta})}{2\hat{\sigma}^4} = 0$$

and will be denoted by $\hat{\sigma}^2$. Since $S(\hat{\beta}) = S_e$, by the above development, the maximum likelihood estimator of σ^2 is given by S_e/n .

The sampling distribution of these estimators can easily be obtained. Since $\hat{\beta}$ is a linear function of \mathbf{Y} , its sampling distribution is a multivariate normal with mean and variance given by

$$\begin{aligned} E(\hat{\beta} | \beta, \sigma^2) &= (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'E(\mathbf{Y} | \beta, \sigma^2) \\ &= (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{X}\beta \\ &= \beta \\ V(\hat{\beta} | \beta, \sigma^2) &= (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'V(\mathbf{Y} | \beta, \sigma^2)\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1} \\ &= \sigma^2(\mathbf{X}'\mathbf{X})^{-1} \end{aligned}$$

and so it is an unbiased estimator. Since the score function is linear in $\hat{\beta}$, it also has minimum variance. The quadratic form given by $[S(\beta) - S_e]/\sigma^2$ has χ_p^2 sampling distribution. On the other hand, it can be shown that S_e is independent of $\hat{\beta}$ and so of $S(\beta) - S_e$. Since $S(\beta)/\sigma^2 \sim \chi_n^2$, $S_e/\sigma^2 \sim \chi_{n-p}^2$ and, therefore, $s^2 = S_e/(n-p)$ is an unbiased estimator of σ^2 . So, it follows that $(\hat{\beta} - \beta)/s$ has a multivariate Student t sampling distribution with $n-p$ degrees of freedom and location parameter $\mathbf{0}$ and scale parameter matrix $(\mathbf{X}'\mathbf{X})^{-1}$.

The above statistic can be used as a pivot to obtain confidence intervals for β or its components. In particular, it is easy to obtain the sampling distribution of $(\hat{\beta}_j - \beta_j)/s$ which is $t_{n-p}(0, c_{jj})$ where c_{jj} is the (j, j) th element of the matrix $(\mathbf{X}'\mathbf{X})^{-1}$, $j = 1, \dots, p$. The $100(1-\alpha)\%$ confidence interval for β_j is given by

$$[\hat{\beta}_j - t_{\alpha/2, n-p} s c_{jj}^{1/2}, \hat{\beta}_j + t_{\alpha/2, n-p} s c_{jj}^{1/2}].$$

Alternatively, from the independence between $S(\beta) - S_e$ and S_e , it follows that $[S(\beta) - S_e]/ps^2 \sim F(p, n-p)$ and confidence regions for β can be obtained.

Example 1. Simple linear regression. Consider the model $Y_i = \beta_0 + \beta_1(x_i - \bar{x}) + e_i$, $i = 1, \dots, n$ with only one explanatory variable and where \bar{x} is the mean of the observed values x_i 's. Then $\mathbf{X} = (x_1, \dots, x_n)'$, $\mathbf{X} = (\mathbf{1}_n, \mathbf{x} - \bar{x}\mathbf{1}_n)$,

$$\mathbf{X}'\mathbf{X} = \begin{pmatrix} n & 0 \\ 0 & \sum_{i=1}^n (x_i - \bar{x})^2 \end{pmatrix} \quad \text{and} \quad \hat{\beta} = \begin{pmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \end{pmatrix}$$

with

$$\hat{\beta}_0 = \bar{Y}, \hat{\beta}_1 = \frac{\sum_{i=1}^n (Y_i - \bar{Y})(x_i - \bar{x})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

and

$$s^2 = \frac{1}{n-2} \sum_{i=1}^n (Y_i - \hat{\beta}_0 - \hat{\beta}_1(x_i - \bar{x}))^2$$

where $\mathbf{1}_n$ is an n -dimensional vector of 1's and $\mathbf{x}' = (x_1, \dots, x_n)$. Due to the centring of the values of the x_i 's, the covariance matrix of the sampling distribution of $\hat{\beta}$ is diagonal. As this is a multivariate normal, this is equivalent to saying that $\hat{\beta}_0$ and $\hat{\beta}_1$ are independent.

Each of these estimators will provide separately information for each parameter. Confidence intervals can be built up based on the independent sampling distributions

$$\frac{\hat{\beta}_0 - \beta_0}{s} \sim t_{n-2}(0, 1/n) \quad \text{and} \quad \frac{\hat{\beta}_1 - \beta_1}{s} \sim t_{n-2} \left(0, \left[\sum_{i=1}^n (x_i - \bar{x})^2 \right]^{-1} \right).$$

Example 2. Analysis of variance with one factor of classification. Let $Y_{ji} = \beta_j + e_{ji}$, $i = 1, \dots, n_j$, $j = 1, \dots, k$, be the model, that is, the n_j observations in group j have the same mean, $j = 1, \dots, k$. This model is frequently referred to as *One-way ANOVA*. The total number of observations n is given by $\sum_{j=1}^k n_j$. Other parametrizations are possible, with the most usual given by $\beta_j = \mu + \alpha_j$ where μ is the global mean and α_j the deviation of the group mean with respect to the global mean. This parametrization is not used here due to the redundancy $\sum_{j=1}^k \alpha_j = 0$, which we are trying to avoid.

The model is completely characterized with parameter vector defined by $\beta = (\beta_1, \dots, \beta_k)'$, the observation vector

$$\mathbf{Y} = \begin{pmatrix} Y_{11} \\ \vdots \\ Y_{1n_1} \\ \vdots \\ Y_{kk} \\ \vdots \\ Y_{kn_k} \end{pmatrix} \quad \text{and the matrix } \mathbf{X} = \begin{pmatrix} 1 & 0 & \cdots & \cdots & 0 \\ \vdots & \vdots & & & \vdots \\ 1 & 0 & \cdots & \cdots & 0 \\ \vdots & & & & \vdots \\ 0 & \cdots & \cdots & 0 & 1 \\ \vdots & & & \vdots & \vdots \\ 0 & \cdots & \cdots & 0 & 1 \end{pmatrix}.$$

As in the previous example, the matrix $(\mathbf{X}'\mathbf{X})^{-1}$ has a diagonal form with (j, j) element given by n_j^{-1} and $\hat{\beta}_j = \bar{Y}_j$ where \bar{Y}_j is the observation mean in group j . Besides that,

$$S_e = \sum_{j=1}^k \sum_{i=1}^{n_j} (Y_{ji} - \bar{Y}_j)^2$$

and $s^2 = S_e/(n-k)$ is the unbiased estimator of σ^2 . Therefore, $(\hat{\beta}_j - \beta_j)/s \sim t_{n-k}(0, n_j^{-1})$, $j = 1, \dots, k$.

216 Introduction to linear models

Hypothesis tests based on the likelihood ratio can be obtained. One of the most useful is the model validation test, that is, the test of the hypothesis $H_0: \beta_2 = \dots = \beta_p = 0$ in the model $Y_i = \beta_1 + \beta_2 x_{2i} + \dots + \beta_p x_{pi} + e_i, i = 1, \dots, n$. Under H_0 , none of the explanatory variables have influence over the value of Y . As we have seen in Section 4.2, the maximum likelihood estimator of (β_1, σ^2) under H_0 is $(\bar{Y}, \hat{\sigma}_0^2)$ where $\hat{\sigma}_0^2 = \Sigma(Y_i - \bar{Y})^2/n$. The maximized likelihood under H_0 is given by $\hat{\sigma}_0^{-n} e^{-n/2}$. So, the maximum likelihood ratio is given by

$$\left(\frac{n\hat{\sigma}_0^2}{S_e} \right)^{n/2} = \left(1 + \frac{n\hat{\sigma}_0^2 - S_e}{S_e} \right)^{n/2} = \left(1 + \frac{p-1}{n-p} F \right)^{n/2}$$

where $F = (n\hat{\sigma}_0^2 - S_e)/[(p-1)s^2]$.

The likelihood ratio test of significance level α rejects H_0 if $F > c$ where c is implicitly given by the equation $\alpha = Pr(F > c | H_0)$. Using results about quadratic forms of normal random variables it is possible to show that $n\hat{\sigma}_0^2 - S_e$ is independent of S_e . As $n\hat{\sigma}_0^2/\sigma^2 \sim \chi_{n-1}^2$, then it follows that $(n\hat{\sigma}_0^2 - S_e)/\sigma^2 \sim \chi_{p-1}^2$ and so $F \sim F(p-1, n-p)$, under H_0 . Therefore, $c = \bar{F}_\alpha(p-1, n-p)$. Confidence regions for $(\beta_2, \dots, \beta_p)$ follows from the above test statistic.

Example 1 (continued). The model validation test in this case corresponds to the test of $H_0: \beta_1 = 0$. The expression $n\hat{\sigma}_0^2$ is given by

$$\begin{aligned} \sum_{i=1}^n (Y_i - \bar{Y})^2 &= \sum_{i=1}^n [(Y_i - \bar{Y} - \hat{\beta}_1(x_i - \bar{x})) + \hat{\beta}_1(x_i - \bar{x})]^2 \\ &= \sum_{i=1}^n [Y_i - \bar{Y} - \hat{\beta}_1(x_i - \bar{x})]^2 + \sum_{i=1}^n \hat{\beta}_1^2(x_i - \bar{x})^2 \\ &\quad + \sum_{i=1}^n [Y_i - \bar{Y} - \hat{\beta}_1(x_i - \bar{x})]\hat{\beta}_1(x_i - \bar{x}) \\ &= S_e + \hat{\beta}_1 \sum_{i=1}^n (Y_i - \bar{Y})(x_i - \bar{x}) \\ &= S_e + \hat{\beta}_1^2 \sum_{i=1}^n (x_i - \bar{x})^2. \end{aligned}$$

So, $F = \hat{\beta}_1^2/(s^2 / \sum_{i=1}^n (x_i - \bar{x})^2) \sim F(1, n-2) \sim t_{n-2}^2$ and the model validation test rejects the null hypothesis H_0 if $\bar{F} > \bar{F}_\alpha(1, n-2)$ or equivalently if

$$|\hat{\beta}_1| > t_{\alpha/2, n-2} s \sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}.$$

This result can be extended to multiple regression in the following sense. Consider the model $Y_i = \beta_1 + \beta_2(x_{2i} - \bar{x}_2) + \dots + \beta_p(x_{pi} - \bar{x}_p) + e_i, i = 1, \dots, n$ where

all the explanatory variables were conveniently centred. The likelihood ratio test for $H_j: \beta_j = 0, j = 2, \dots, p$ with level α rejects H_j if

$$|\hat{\beta}_j| > t_{\alpha/2, n-p} s \sqrt{\sum_{i=1}^n (x_{ji} - \bar{x}_j)^2}.$$

Example 2 (continued). Now the model validation test is concerned with the following hypothesis, $H_0: \beta_1 = \dots = \beta_k$. The expression of $n\hat{\sigma}_0^2$ is

$$\sum_{j=1}^k \sum_{i=1}^{n_j} (Y_{ji} - \bar{Y}_j)^2 = \sum_{j=1}^k \sum_{i=1}^{n_j} [(Y_{ji} - \bar{Y}_j) + (\bar{Y}_j - \bar{Y})]^2$$

where

$$\begin{aligned} \bar{Y} &= \frac{1}{n} \sum_{j=1}^k n_j \bar{Y}_j \text{ and} \\ n\hat{\sigma}_0^2 &= \sum_{j=1}^k \sum_{i=1}^{n_j} (Y_{ji} - \bar{Y}_j)^2 + \sum_{j=1}^k n_j (\bar{Y}_j - \bar{Y})^2 \\ &\quad + 2 \sum_{j=1}^k \sum_{i=1}^{n_j} (Y_{ji} - \bar{Y}_j)(\bar{Y}_j - \bar{Y}) \\ &= S_e + \sum_{j=1}^k n_j (\bar{Y}_j - \bar{Y})^2 \text{ because the third term vanishes.} \end{aligned}$$

We have already seen that squared deviations around the mean for a normally distributed variable are independent of the sample mean. Then, S_e is independent of the means $\bar{Y}_1, \dots, \bar{Y}_k$ and so, independent of the second term in the above expression. This term is also in the form of squared deviations of k normal observations with respect to its mean and so have the sampling distribution given by $\sigma^2 \chi_{k-1}^2$. Therefore,

$$F = \frac{\sum_{j=1}^k n_j (\bar{Y}_j - \bar{Y})^2 / (k-1)}{s^2} \sim F(k-1, n-k)$$

confirming the result described above for the general linear model. F is the MLR test statistic for H_0 and the test of significance level α rejects H_0 if $F > \bar{F}_\alpha(k-1, n-k)$.

Finally, we will address the relevant question of how to make predictions using the classical linear model. The problem here is how to make an inference about an m -dimensional vector of future observations Y^* with explanatory variables gathered into a $m \times p$ matrix x^* . Once again, the trick is to find a convenient pivot,

which in this case is given by $(\mathbf{Y}^* - \mathbf{x}^* \hat{\beta})/s$. To verify this result, it suffices to observe that \mathbf{Y}^* , $\hat{\beta}$ and s^2 are independent and so $\mathbf{Y}^* - \mathbf{x}^* \hat{\beta}$ is independent of s^2 . The numerator is distributed as normal with mean and variance given by

$$\begin{aligned} E[\mathbf{Y}^* - \mathbf{x}^* \hat{\beta} | \beta, \sigma^2] &= \mathbf{x}^* \beta - \mathbf{x}^* \beta = 0 \\ V[\mathbf{Y}^* - \mathbf{x}^* \hat{\beta} | \beta, \sigma^2] &= \sigma^2 + \mathbf{x}^* \sigma^2 (\mathbf{X}' \mathbf{X})^{-1} \mathbf{x}^* \\ &= \sigma^2 \mathbf{C}^* \text{ where } \mathbf{C}^* = \mathbf{I}_m + \mathbf{x}^* (\mathbf{X}' \mathbf{X})^{-1} \mathbf{x}^*. \end{aligned}$$

Since $(n-p)s^2/\sigma^2 \sim \chi_{n-p}^2$, the pivotal quantity is distributed as a $t_{n-p}(0, \mathbf{C}^*)$ independently of the parameters in the model. So, confidence intervals for \mathbf{Y}^* with confidence level $100(1-\alpha)\%$ can be built. These results are easily extended to the prediction of a vector of future observations following the same steps as before.

8.3 Bayesian linear models

In Bayesian inference, a prior distribution for the parameters must be specified in addition to the likelihood function. Firstly, the results involving proper priors (in fact natural conjugate, as will be shown) will be presented. The analysis with non-informative prior will be presented in the sequel and some comparisons with classical inference will be made. The examples presented in the previous section will be revisited.

The prior distribution adopted for the parameters is a multivariate generalization of the normal- χ^2 presented in Section 3.3. Assume that the parametric vector β has a conditional prior distribution $N(\mu_0, \phi^{-1} \mathbf{C}_0^{-1})$ where $\phi = \sigma^{-2}$ and that $n_0 \sigma_0^2 \phi \sim \chi_{n_0}^2$. In this way, the prior distribution is fully specified with density given by

$$\begin{aligned} p(\beta, \phi) &= (2\pi)^{-p/2} |\phi \mathbf{C}_0|^{1/2} \exp \left\{ -\frac{\phi}{2} (\beta - \mu_0)' \mathbf{C}_0 (\beta - \mu_0) \right\} \\ &\quad \times \frac{(n_0 \sigma_0^2 / 2)^{n_0/2}}{\Gamma(n_0/2)} \phi^{(n_0/2)-1} \exp \left\{ -\frac{n_0 \sigma_0^2}{2} \phi \right\} \\ &\propto \phi^{[(n_0+p)/2]-1} \exp \left\{ -\frac{\phi}{2} [n_0 \sigma_0^2 + (\beta - \mu_0)' \mathbf{C}_0 (\beta - \mu_0)] \right\}. \end{aligned}$$

Then, as in the univariate case, the conditional distribution of $\phi | \beta$ can be obtained from the joint prior distribution of β and ϕ collecting only the terms involving ϕ . It is given by $[n_0 \sigma_0^2 + (\beta - \mu_0)' \mathbf{C}_0 (\beta - \mu_0)] \phi | \beta \sim \chi_{n_0+p}^2$. The marginal distribution of β can be obtained dividing $p(\beta, \phi)$ by $p(\phi | \beta)$ or, as done before, integrating the joint distribution with respect to ϕ . Its density is given by

$$p(\beta) \propto [n_0 \sigma_0^2 + (\beta - \mu_0)' \mathbf{C}_0 (\beta - \mu_0)]^{-(n_0+p)/2}$$

which corresponds to the density of a $t_{n_0}(\mu_0, \sigma_0^2 \mathbf{C}_0^{-1})$, as seen in Chapter 4. The normalizing constant is

$$\frac{\Gamma[(n_0+p)/2]}{\Gamma(n_0/2) n_0^{p/2}} (n_0 \sigma_0^2)^{n_0/2} |\mathbf{C}_0|^{1/2}.$$

On the other hand, the likelihood of β and ϕ is given by

$$\phi^{n/2} \exp \left\{ -\frac{\phi}{2} [S_e + (\beta - \hat{\beta})' \mathbf{X}' \mathbf{X} (\beta - \hat{\beta})] \right\}$$

and has the same form as the prior density. Therefore, the posterior is given by

$$\begin{aligned} p(\beta, \phi | \mathbf{y}) &\propto \phi^{(n+n_0+p)/2-1} \\ &\quad \times \exp \left\{ -\frac{\phi}{2} [n_0 \sigma_0^2 + S_e + (\beta - \mu_0)' \mathbf{C}_0 (\beta - \mu_0) + (\beta - \hat{\beta})' \mathbf{X}' \mathbf{X} (\beta - \hat{\beta})] \right\}. \end{aligned}$$

It can be shown that

$$\begin{aligned} &(\beta - \mu_0)' \mathbf{C}_0 (\beta - \mu_0) + (\beta - \hat{\beta})' \mathbf{X}' \mathbf{X} (\beta - \hat{\beta}) \\ &= (\beta - \mu_1)' \mathbf{C}_1 (\beta - \mu_1) + \mu_0' \mathbf{C}_0 \mu_0 + \hat{\beta}' \mathbf{X}' \mathbf{X} \hat{\beta} + \mu_1' \mathbf{C}_1 \mu_1 \end{aligned}$$

where $\mu_1 = \mathbf{C}_1^{-1} (\mathbf{C}_0 \mu_0 + \mathbf{X}' \mathbf{y})$ and $\mathbf{C}_1 = \mathbf{C}_0 + \mathbf{X}' \mathbf{X}$. Note that even if \mathbf{X} does not have full rank, \mathbf{C}_1 will have and can always be inverted. We still have to deal with the terms

$$\begin{aligned} S_e + \mu_0' \mathbf{C}_0 \mu_0 + \hat{\beta}' \mathbf{X}' \mathbf{X} \hat{\beta} + \mu_1' \mathbf{C}_1 \mu_1 \\ &= \mathbf{y}' \mathbf{y} - \hat{\beta}' \mathbf{X}' \mathbf{X} \hat{\beta} + \mu_0' \mathbf{C}_0 \mu_0 + \hat{\beta}' \mathbf{X}' \mathbf{X} \hat{\beta} + \mu_1' (\mathbf{C}_0 \mu_0 + \mathbf{X}' \mathbf{y}) \\ &= (\mathbf{y} - \mathbf{X} \mu_1)' \mathbf{y} + (\mu_0 - \mu_1)' \mathbf{C}_0 \mu_0. \end{aligned}$$

Then, the posterior density of β and ϕ can be written as

$$p(\beta, \phi | \mathbf{y}) \propto \phi^{p/2} \exp \left\{ -\frac{\phi}{2} (\beta - \mu_1)' \mathbf{C}_1 (\beta - \mu_1) \right\} \phi^{(n_1/2)-1} \exp \left\{ -\frac{\phi}{2} n_1 \sigma_1^2 \right\}$$

where $n_1 = n+n_0$ and $n_1 \sigma_1^2 = n_0 \sigma_0^2 + (\mathbf{y} - \mathbf{X} \mu_1)' \mathbf{y} + (\mu_0 - \mu_1)' \mathbf{C}_0 \mu_0$. This density has the same form of the prior and, so, is natural conjugate to the normal linear models. In particular, $\beta | \mathbf{y} \sim t_{n_1}(\mu_1, \sigma_1^2 \mathbf{C}_1^{-1})$ and $\beta_j | \mathbf{y} \sim t_{n_1}(\mu_{1j}, \sigma_1^2 (\mathbf{C}_1^{-1})_{jj})$, $j = 1, \dots, p$. The posterior mean and variance-covariance matrix of β are given, respectively, by

$$\mu_1 \quad \text{and} \quad \frac{n_1}{n_1 - 2} \sigma_1^2 \mathbf{C}_1^{-1}, \quad n_1 > 2.$$

The posterior distribution of ϕ is $n_1 \sigma_1^2 \phi | \mathbf{y} \sim \chi_{n_1}^2$ with mean σ_1^{-2} . The point estimators of β and ϕ are given by μ_1 and σ_1^{-2} , respectively. Confidence intervals for β_j and ϕ are obtained from the percentiles of t_{n_1} and $\chi_{n_1}^2$ distributions,

respectively. It is also possible to make inference about the joint distribution of the β based on the fact that $(\beta - \mu_1)'C_1(\beta - \mu_1)/\sigma_1^2 \mid y \sim F(p, n_1 - p)$.

The non-informative prior can be used to represent, in some sense, the absence of initial information. Using the same approach as in Section 3.4, as β is essentially a (multivariate) location parameter and ϕ is a scale parameter, it follows that the joint non-informative prior distribution is given by $p(\beta, \phi) \propto \phi^{-1}$. This prior is a particular, degenerated case of the natural conjugate prior. Making the convenient substitutions, the posterior density is

$$\begin{aligned} p(\beta, \phi \mid y) &\propto \phi^{(n/2)-1} \exp \left\{ -\frac{\phi}{2} [S_e + (\beta - \hat{\beta})'X'X(\beta - \hat{\beta})] \right\} \\ &\propto \phi^{p/2} \exp \left\{ -\frac{\phi}{2} (\beta - \hat{\beta})'X'X(\beta - \hat{\beta}) \right\} \phi^{((n-p)/2)-1} \exp \left\{ -\frac{\phi}{2} (n-p)s^2 \right\}. \end{aligned}$$

Therefore, the posterior will remain in the same class with only changes to the values of the hyperparameters of the relevant distributions. So, $\beta \mid y \sim t_{n-p}(\bar{\beta}, s^2(X'X)^{-1})$ and $(n-p)s^2\phi \mid y \sim \chi_{n-p}^2$ and the quadratic form in β will be reduced to $(\beta - \hat{\beta})'X'X(\beta - \hat{\beta})/s^2$ with posterior distribution $F(p, n-p)$. These distributions provide the parallel Bayesian results to those obtained in the previous section using the classical approach.

Example 1 (continued). Suppose that in the prior distribution, β_0 and β_1 are conditionally independent given ϕ with distributions $N(\mu_0, (c_0\phi)^{-1})$ and $N(\mu_1, (c_1\phi)^{-1})$, respectively and $n_0\sigma_0\phi \sim \chi_{n_0}^2$. (The case where β_0 and β_1 are not conditionally independent is left as an exercise.) Since $X'X$ is a diagonal matrix, the quadratic form $(\beta - \hat{\beta})'X'X(\beta - \hat{\beta})$ is reduced to $S(\beta_0) + S(\beta_1)$ where

$$S(\beta_0) = n(\beta_0 - \bar{y})^2 \text{ and } S(\beta_1) = \sum_{i=1}^n (x_i - \bar{x})^2(\beta_1 - \hat{\beta}_1)^2$$

and the likelihood function is given by

$$l(\beta_0, \beta_1, \phi; y) \propto \phi^{n/2} \exp \left\{ -\frac{\phi}{2} [S_e + S(\beta_0) + S(\beta_1)] \right\}.$$

Combining with the prior, the posterior distribution follows:

$$\begin{aligned} p(\beta_0, \beta_1, \phi \mid y) &\propto \phi^{(n+n_0+2)/2-1} \\ &\times \exp \left\{ -\frac{\phi}{2} [n_0\sigma_0^2 + S_e + c_0(\beta_0 - \mu_0)^2 + S(\beta_0) + c_1(\beta_1 - \mu_1)^2 + S(\beta_1)] \right\} \end{aligned}$$

from where it is easy to see that, conditional on ϕ , β_0 and β_1 stay independent with distributions $N(\mu_j^*, (c_j^*\phi)^{-1})$, $j = 0, 1$, respectively where $c_0^* = c_0 + n$, $c_1^* = c_1 + \sum_{i=1}^n (x_i - \bar{x})^2$,

$$\mu_0^* = \frac{c_0\mu_0 + n\bar{y}}{c_0 + n} \quad \text{and} \quad \mu_1^* = \frac{c_1\mu_1 + \sum_{i=1}^n (x_i - \bar{x})^2\hat{\beta}_1}{c_1 + \sum_{i=1}^n (x_i - \bar{x})^2}.$$

The posterior distribution of ϕ is $n_1\sigma_1^2\phi \mid y \sim \chi_{n_1}^2$ with $n_1 = n + n_0$ and

$$n_1\sigma_1^2 = \sum_{i=1}^n y_i^2 - n\mu_0^*\hat{\beta}_0 - \mu_1^*\hat{\beta}_1 \sum_{i=1}^n (x_i - \bar{x})^2 + c_0(\beta_0 - \mu_0)^2 + c_1(\beta_1 - \mu_1)^2.$$

The marginal posterior distributions of β_0 and β_1 are Student t with n_1 degrees of freedom and parameters μ_j^* and σ_1^2/c_j^* , $j = 0, 1$, respectively.

In the case of a non-informative prior $p(\beta_0, \beta_1, \phi) \propto \phi^{-1}$, the posterior is

$$\begin{aligned} p(\beta_0, \beta_1, \phi \mid y) &\propto \phi^{-1}\phi^{n/2} \exp \left\{ -\frac{\phi}{2} [(n-2)s^2 + S(\beta_0) + S(\beta_1)] \right\} \\ &\propto \phi^{1/2} \exp \left\{ -\frac{\phi}{2} S(\beta_0) \right\} \phi^{1/2} \exp \left\{ -\frac{\phi}{2} S(\beta_1) \right\} \\ &\quad \times \phi^{((n-2)/2)-1} \exp \left\{ -\frac{\phi}{2} (n-2)s^2 \right\}. \end{aligned}$$

Examining the above expression, it is easy to identify that $\beta_0 \mid \phi, y \sim N(\bar{y}, (\phi\bar{y})^{-1})$, $\beta_1 \mid \phi, y \sim N(\hat{\beta}_1, (\phi\sum_{i=1}^n (x_i - \bar{x})^2)^{-1})$ and $(n-2)s^2\phi \mid y \sim \chi_{n-2}^2$. Then, $\beta_0 \mid y \sim t_{n-2}(\bar{y}, s^2/n)$ and $\beta_1 \mid y \sim t_{n-2}(\hat{\beta}_1, s^2/\sum_{i=1}^n (x_i - \bar{x})^2)$ and so the inference coincides numerically with the one obtained following the classical approach.

Example 2 (continued). Assume that $\beta_j \mid \phi \sim N(\mu_j, (c_j\phi)^{-1})$, $j = 1, \dots, k$, are conditionally independent and that $n_0\sigma_0^2\phi \sim \chi_{n_0}^2$. An alternative to the conditional independence of the β_j 's will be considered in the next section. Then, as in Example 1, $X'X$ is a diagonal matrix and the quadratic form $(\beta - \hat{\beta})'X'X(\beta - \hat{\beta})$ reduces to

$$\sum_{j=1}^k n_j(\beta_j - \bar{y}_j)^2.$$

The likelihood is given by

$$l(\beta_1, \dots, \beta_k, \phi; y) \propto \phi^{n/2} \exp \left\{ -\frac{\phi}{2} \left[(n-k)s^2 + \sum_{j=1}^k n_j(\beta_j - \bar{y}_j)^2 \right] \right\}.$$

Combining with the prior distribution, the following posterior is obtained

$$\begin{aligned} p(\beta_1, \dots, \beta_k, \phi \mid y) &\propto \phi^{(n+n_0+k)/2-1} \\ &\times \exp \left\{ -\frac{\phi}{2} \left[n_0\sigma_0^2 + (n-k)s^2 + \sum_{j=1}^k c_j(\beta_j - \mu_j)^2 + \sum_{j=1}^k n_j(\beta_j - \bar{y}_j)^2 \right] \right\} \\ &\propto \left\{ \prod_{j=1}^k \phi^{1/2} \exp \left[-\frac{\phi}{2} c_j^*(\beta_j - \mu_j^*)^2 \right] \right\} \phi^{(n_1/2)-1} \exp \left\{ -\frac{\phi}{2} n_1\sigma_1^2 \right\} \end{aligned}$$

where $c_j^* = c_j + n_j$, $\mu_j^* = (c_j \mu_j + n_j \bar{y}_j)/c_j^*$, $j = 1, \dots, k$, $n_1 = n + n_0$ and

$$n_1 \sigma_1^2 = n_0 \sigma_0^2 + (n - k)s^2 + \sum_{j=1}^k \frac{n_j c_j}{n_j + c_j} (\bar{y}_j - \mu_j)^2.$$

It is easy to obtain that $\beta_j \mid \phi, \mathbf{y} \sim N[\mu_j^*, (c_j^* \phi)^{-1}]$, thus retaining the prior independence and $n_1 \sigma_1^2 \phi \mid \mathbf{y} \sim \chi_{n_1}^2$ and, therefore, the marginal posterior distributions of β_j are $t_{n_1}(\mu_j^*, s^2/c_j^*)$, $j = 1, \dots, k$. The point estimators of β_j and ϕ are given by μ_j^* and σ_1^{-2} , $j = 1, \dots, k$, respectively. Confidence intervals for β_j and ϕ (or σ^2) can be obtained using the percentiles of the t_{n_1} and $\chi_{n_1}^2$ distributions, respectively.

In the case of a non-informative prior, we will have the posterior

$$\begin{aligned} p(\beta_1, \dots, \beta_k, \phi \mid \mathbf{y}) &\propto \phi^{(n/2)-1} \exp \left\{ -\frac{\phi}{2} \left[(n - k)s^2 + \sum_{j=1}^k n_j (\beta_j - \bar{y}_j)^2 \right] \right\} \\ &\propto \left\{ \prod_{j=1}^k \phi^{1/2} \exp \left[-\frac{\phi}{2} n_j (\beta_j - \bar{y}_j)^2 \right] \right\} \phi^{((n-k)/2)-1} \\ &\quad \times \exp \left\{ -\frac{\phi}{2} (n - k)s^2 \right\}. \end{aligned}$$

Making the same identifications as Example 1, the marginal posterior distributions $\beta_j \mid \mathbf{y} \sim t_{n-k}(\bar{y}_j, s^2/n_j)$, $j = 1, \dots, k$, and $(n - k)s^2 \phi \sim \chi_{n-k}^2$ are easily obtained. These distributions are similar to the ones obtained in classical inference.

Predictive distributions are needed to perform hypothesis testing and prediction from a Bayesian point of view. As we have already seen, these distributions can be obtained by integrating the sample distribution with respect to the distribution of the parameters. Fortunately, in the normal case, this job is greatly simplified by the linear structure of the distribution. Consider the model

$$\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \mathbf{e}, \quad \mathbf{e} \sim N(\mathbf{0}, \phi^{-1}\mathbf{I}_n) \quad \text{where } \phi = 1/\sigma^2$$

and suppose that the conditional prior distribution of $\boldsymbol{\beta}$ is

$$\boldsymbol{\beta} \mid \phi \sim NM(\boldsymbol{\mu}, \phi^{-1}\mathbf{C}^{-1}).$$

Combining these two results, we get

$$\begin{aligned} \mathbf{Y} \mid \phi &\sim XN(\boldsymbol{\mu}, \phi^{-1}\mathbf{C}^{-1}) + N(\mathbf{0}, \phi^{-1}\mathbf{I}_n) \\ &\sim N[\mathbf{X}\boldsymbol{\mu}, \phi^{-1}(\mathbf{I}_n + \mathbf{X}\mathbf{C}^{-1}\mathbf{X}')] \end{aligned}$$

since the above normal distributions are independent. Supposing now that $\nu \sigma_0^2 \phi \sim \chi_v^2$ leads to

$$\mathbf{Y} \sim t_v[\mathbf{X}\boldsymbol{\mu}, \sigma_0^2(\mathbf{I}_n + \mathbf{X}\mathbf{C}^{-1}\mathbf{X}')].$$

The marginal distribution of Y_i is $t_v[\mathbf{x}_i \boldsymbol{\mu}, \sigma_0^2(1 + \mathbf{x}_i \mathbf{C}^{-1} \mathbf{x}'_i)]$, $i = 1, \dots, n$.

To compare two hypotheses H_0 and H_1 , evaluation of the distributions of $(\mathbf{Y} \mid H_l)$, $l = 0, 1$ is required. The model validation test is based on the Bayes factor

$$BF(H_0; H_1) = \frac{p(\mathbf{y} \mid H_0)}{p(\mathbf{y} \mid H_1)}$$

where the denominator is the density of the above distribution given by

$$\begin{aligned} &\frac{\Gamma[(n_0 + n)/2]}{\Gamma(n_0/2)n_0^{n/2}} (n_0 \sigma_0^2)^{n/2} |\mathbf{I}_n + \mathbf{X}\mathbf{C}_0^{-1}\mathbf{X}'|^{-1/2} \\ &\quad \times [n_0 \sigma_0^2 + (\mathbf{y} - \mathbf{X}\boldsymbol{\mu}_0)'(\mathbf{I}_n + \mathbf{X}\mathbf{C}_0^{-1}\mathbf{X}')^{-1}(\mathbf{y} - \mathbf{X}\boldsymbol{\mu}_0)]^{-(n_0+n)/2}. \end{aligned}$$

If it is desirable to test the model validation hypothesis, we build up H_0 : $\beta_2 = \dots = \beta_p = 0$. Under H_0 , the model simplifies to $\mathbf{Y} = \mathbf{1}_n \boldsymbol{\beta}_1 + \mathbf{e}$ and $\boldsymbol{\beta}_1 \mid \phi \sim N[\boldsymbol{\mu}_{00}, (c_{00}\phi)^{-1}]$ and therefore $\mathbf{Y} \mid \phi \sim N[\mathbf{1}_n \boldsymbol{\mu}_{00}, \phi^{-1}(\mathbf{I}_n + c_{00}^{-1}\mathbf{1}_n \mathbf{1}'_n)]$. Then, the density $p(\mathbf{y} \mid H_0)$ is given by

$$\begin{aligned} &\frac{\Gamma[(n_0 + n)/2]}{\Gamma(n_0/2)n_0^{n/2}} (n_0 \sigma_0^2)^{n/2} |\mathbf{I}_n + c_{00}^{-1}\mathbf{1}_n \mathbf{1}'_n|^{-1/2} \\ &\quad \times [n_0 \sigma_0^2 + (\mathbf{y} - \mathbf{1}_n \boldsymbol{\mu}_{00})'(\mathbf{I}_n + c_{00}^{-1}\mathbf{1}_n \mathbf{1}'_n)^{-1}(\mathbf{y} - \mathbf{1}_n \boldsymbol{\mu}_{00})]^{-(n_0+n)/2}. \end{aligned}$$

The Bayes factor is, then, given by

$$\begin{aligned} &\frac{\left\{ |\mathbf{I}_n + \mathbf{X}\mathbf{C}_0^{-1}\mathbf{X}'| \right\}^{\frac{1}{2}}}{\left\{ |\mathbf{I}_n + c_{00}^{-1}\mathbf{1}_n \mathbf{1}'_n| \right\}^{\frac{1}{2}}} \\ &\quad \times \left\{ \frac{|n_0 \sigma_0^2 + (\mathbf{y} - \mathbf{X}\boldsymbol{\mu}_0)'(\mathbf{I}_n + \mathbf{X}\mathbf{C}_0^{-1}\mathbf{X}')^{-1}(\mathbf{y} - \mathbf{X}\boldsymbol{\mu}_0)|}{|n_0 \sigma_0^2 + (\mathbf{y} - \mathbf{1}_n \boldsymbol{\mu}_{00})'(\mathbf{I}_n + c_{00}^{-1}\mathbf{1}_n \mathbf{1}'_n)^{-1}(\mathbf{y} - \mathbf{1}_n \boldsymbol{\mu}_{00})|} \right\}^{\frac{n_0+n}{2}} \end{aligned}$$

In the case of the prediction of an m -dimensional vector of future observations \mathbf{Y}^* with explanatory variable matrix \mathbf{x}^* , the same result used above can be applied as far as the prediction is based on the predictive distribution of $\mathbf{Y}^* \mid \mathbf{y}$ with density given by

$$\begin{aligned} p(\mathbf{y}^* \mid \mathbf{y}) &= \int \int p(\mathbf{y}^* \mid \boldsymbol{\beta}, \phi, \mathbf{y}) p(\boldsymbol{\beta}, \phi \mid \mathbf{y}) d\boldsymbol{\beta} d\phi \\ &= \int \left\{ \int p(\mathbf{y}^* \mid \boldsymbol{\beta}, \phi) p(\boldsymbol{\beta} \mid \phi, \mathbf{y}) d\boldsymbol{\beta} \right\} p(\phi \mid \mathbf{y}) d\phi \end{aligned}$$

and the calculus is similar to that involved in the evaluation of the Bayes factor. An important difference is that the marginalizations are with respect to the posterior distribution of the parameters while the Bayes factor is obtained using marginalizations with respect to the prior distribution. Then, using the adopted notation it follows that

$$\mathbf{Y}^* \mid \mathbf{y} \sim t_{n_1}(\mathbf{x}^* \boldsymbol{\mu}_1, \sigma_1^2(\mathbf{I}_m + \mathbf{x}^* \mathbf{C}_1^{-1} \mathbf{x}'^*))$$

and so point predictions and confidence intervals for \mathbf{Y}^* can be easily obtained. The analysis using non-informative priors leads to $\mu_1 \rightarrow \hat{\beta}$, $C_1 \rightarrow \mathbf{X}'\mathbf{X}$, $n_1 \rightarrow n$ and $\sigma_1^2 \rightarrow s^2$ and the predictive distribution of \mathbf{Y}^* reduces to

$$\mathbf{Y}^* | \mathbf{y} \sim t_{n-p}(\mathbf{x}^*\hat{\beta}, s^2(I_m + \mathbf{x}^*(\mathbf{X}'\mathbf{X})^{-1}\mathbf{x}^*)').$$

This distribution coincides with the predictive distribution of the classical approach providing the same predictions.

8.4 Hierarchical linear models

In Section 3.5 we have seen how to combine structural information with strictly subjective information to build up the prior distribution in stages. This strategy is explored within the context of linear models in this section. The linear structure of the model combines very well with the hierarchical modelling in the context of linear normal models and this structure can be explored in more detail. Specifically, the hierarchical structure described in Chapter 3 is used with the additional assumption of linearity and normality. This setup preserves the model linearity as a whole and the conjugacy at any stage of the model. This area was organized systematically and received great research development following a paper by Lindley and Smith (1972).

In order to specify the model and its prior, we can rewrite it as

$$\begin{aligned} \mathbf{Y} | \beta_1, \phi &\sim N(\mathbf{X}_1\beta_1, \phi^{-1}\mathbf{I}_n) \\ \beta_1 | \beta_2, \phi &\sim N(\mathbf{X}_2\beta_2, \phi^{-1}\mathbf{C}_1^{-1}) \\ \beta_2 | \phi &\sim N(\mu, \phi^{-1}\mathbf{C}_2^{-1}) \\ n_0\sigma_0^2\phi &\sim \chi_{n_0}^2 \end{aligned}$$

where the matrix with explanatory variables is renamed as \mathbf{X}_1 due to the presence of another matrix including the second stage explanatory variables, \mathbf{X}_2 . This matrix includes the values which explain the variations in the β_1 and the coefficient of this explanation is given by β_2 . The model specified above is the simplest in the class of hierarchical models containing only two stages. More stages can be included in the model specification depending on the structure of the problem. The form of the model is not modified. Only some extra equations are included, each in the form

$$\beta_j | \beta_{j+1}, \phi \sim N(\mathbf{X}_{j+1}\beta_{j+1}, \phi^{-1}\mathbf{C}_j^{-1}).$$

Generally, for higher stages, it is more difficult to specify the equations as the level of elaboration involved gets deeper and deeper. The dependence of the variances on ϕ is not by chance. It allows the use of results about conjugacy developed in Sections 3.3 and 8.3. All the derivations that follow are done only for the model with two stages to keep the notation simple although there is no technical problem to extend the results to models with K stages, $K > 2$.

Firstly it is worth mentioning that the analysis conditional on β_2 is completely similar to the developments made in the last section. If β_2 is known, the prior does not depend on its probabilistic specification. So, all that is needed is to apply the results of the last section substituting μ_0 by $\mathbf{X}_2\beta_2$. In particular, it is obtained that $\beta_1 | \mathbf{y} \sim t_{n_1}(\mu_1, \sigma_1^2\mathbf{C}^{*-1})$ and $n_1\sigma_1^2\phi \sim \chi_{n_1}^2$ where $\mu_1 = \mathbf{C}^{*-1}(\mathbf{C}_1\mathbf{X}_2\beta_2 + \mathbf{X}_1'\mathbf{y})$, $\mathbf{C}^* = \mathbf{C}_1 + \mathbf{X}_1'\mathbf{X}_1$, $n_1 = n + n_0$ and $n_1\sigma_1^2 = n_0\sigma_0^2 + (\mathbf{y} - \mathbf{X}_1\mathbf{X}_2\beta_2)'\mathbf{y} + (\mathbf{X}_2\beta_2 - \mu_1)'\mathbf{C}_0\mathbf{X}_2\beta_2$.

Considering again the case with β_2 unknown, it is worth noting that the distributions of $\beta_1 | \phi$ and β_1 can be obtained via marginalization as was done at the end of the last section to obtain the predictive distribution of the observations. Therefore, combining the distributions of $\beta_1 | \beta_2, \phi$ and $\beta_2 | \phi$ it follows that

$$\begin{aligned} \beta_1 | \phi &\sim \mathbf{X}_2N(\mu, \phi^{-1}\mathbf{C}_2^{-1}) + N(\mathbf{0}, \phi^{-1}\mathbf{C}_1^{-1}) \\ &\sim N(\mathbf{X}_2\mu, \phi^{-1}\mathbf{C}_1^{*-1}) \quad \text{where } \mathbf{C}_1^{*-1} = \mathbf{C}_1^{-1} + \mathbf{X}_2\mathbf{C}_2^{-1}\mathbf{X}_2' \end{aligned}$$

and, by integration, we obtain that $\beta_1 \sim t_{n_0}(\mathbf{X}_2\mu, \sigma_0^2\mathbf{C}_1^{*-1})$.

Example 2 (continued). In the last section, it was assumed that the means β_j , $j = 1, \dots, k$ were independently distributed a priori. A reasonable alternative for the case in which similarity among the k groups can be assumed is to suppose that the means are a random sample from a population of means. This population is fictitious and, to fix ideas, taken as homogeneous. To keep the structure presented above, it is assumed that this population is normal. So, it follows that β_1, \dots, β_k is a sample from the $N(\mu, (c_1\phi)^{-1})$ where c_1 measures the precision of the population of means relatively to the precision of the likelihood. The model is completed with the specification of the distributions of $\mu | \phi$ and ϕ . It follows that the prior is completely specified by

$$\begin{aligned} \text{1st level : } \beta | \mu, \phi &\sim N[\mathbf{1}_k\mu, (c_1\phi)^{-1}\mathbf{I}_k] \\ \text{2nd level : } \mu | \phi &\sim N[\mu_0, (c_2\phi)^{-1}] \\ n_0\sigma_0^2\phi &\sim \chi_{n_0}^2 \end{aligned}$$

and the distribution of the $\beta | \phi$ is $N[\mathbf{1}_k\mu_0, \phi^{-1}(c_1^{-1}\mathbf{I}_k + c_2^{-1}\mathbf{1}_k\mathbf{1}_k')]$ and the components of β are no longer independent. In particular, the prior correlation between any two distinct components is given by $(1 + c_2/c_1)^{-1}$. This may be helpful in the specification of the constants c_1 and c_2 . A larger value for the constant c_2/c_1 leads to a smaller prior correlation and vice versa.

It is interesting to see that the distributions of β_1 and β_2 conditional on ϕ are multivariate normal and therefore the theory developed in the last section is directly applicable and the predictive and posterior distributions can easily be obtained. So, the predictive distribution for $\mathbf{Y} | \phi$ is given by

$$\mathbf{X}_1N(\mathbf{X}_2\mu, \phi^{-1}\mathbf{C}_1^{*-1}) + N(\mathbf{0}, \phi^{-1}\mathbf{I}_n) \sim N(\mathbf{X}_1\mathbf{X}_2\mu, \phi^{-1}(\mathbf{I}_n + \mathbf{X}_1\mathbf{C}_1^{*-1}\mathbf{X}_1'))$$

and the marginal for \mathbf{Y} is obtained from the above expression substituting the N by t_{n_0} and ϕ^{-1} by σ_0^2 . The posterior distribution of (β_1, ϕ) is a multivariate normal- χ^2 with parameters μ_1 , \mathbf{C}^* , n_1 and σ_1^2 given by

$$\begin{aligned}\mu_1 &= \mathbf{C}_1^{*-1}(\mathbf{C}_1^*\mathbf{X}_2\mu_0 + \mathbf{X}'_1\mathbf{y}) \\ \mathbf{C}^* &= \mathbf{C}_1^* + \mathbf{X}'_1\mathbf{X}_1 \\ n_1 &= n + n_0 \\ n_1\sigma_1^2 &= n_0\sigma_0^2 + (\mathbf{y} - \mathbf{X}_1\mu_1)' \mathbf{y} + (\mathbf{X}_2\mu_0 - \mu_1)' \mathbf{C}_1^{*-1}\mathbf{X}_2\mu_0.\end{aligned}$$

Example 2 (continued). Consider again the analysis of variance model with one classification factor with a hierarchical prior. Applying the above results, it follows that

$$\begin{aligned}\mathbf{C}_1^{*-1}\mathbf{X}_2\mu_0 + \mathbf{X}'_1\mathbf{y} &= (c_1^{-1} + kc_2^{-1})\mu_0\mathbf{1}_k + (n_1\bar{y}_1, \dots, n_k\bar{y}_k)' \\ \cdots \mathbf{C}^* &= \text{diag}(c_1 + n_1, \dots, c_1 + n_k) - \frac{c_1^2}{kc_1 + c_2}\mathbf{1}_k\mathbf{1}'_k.\end{aligned}$$

It can still be shown in the case of a non-informative prior in the second level ($c_2 \rightarrow 0$) and equal number of observations in each group that the posterior mean of β_j is in the form

$$w_j\bar{y}_j + (1 - w_j)\bar{\bar{y}}, \quad \text{with } 0 \leq w_j \leq 1, j = 1, \dots, k,$$

where $\bar{\bar{y}}$ is the average of group averages. This type of estimator is known as a shrinkage estimator since it brings all the group estimates closer to the global mean, shrinking the distance among the means. Shrinkage estimators and their properties were studied by many authors including Copas, James, Morris and Stein, as cited in O'Hagan (1994).

To obtain a posterior distribution of β_2 it is first necessary to write the likelihood function of β_2 and ϕ . Using again the results involving combinations of the normal, it follows that

$$\begin{aligned}\mathbf{Y} | \beta_2, \phi &\sim \mathbf{X}_2 N(\mathbf{X}_2\beta_2, \phi^{-1}\mathbf{C}_2^{-1}) + N(\mathbf{0}, \phi^{-1}\mathbf{I}_n) \\ &\sim N(\mathbf{X}_1\mathbf{X}_2\beta_2, \phi^{-1}\mathbf{C}_0^{-1})\end{aligned}$$

where $\mathbf{C}_0^{-1} = \mathbf{I}_n + \mathbf{X}_1\mathbf{C}_1^{-1}\mathbf{X}'_1$. Therefore, likelihood and prior are normal as before with a slight difference since the observational variance is not proportional to an identity matrix. The same results are still valid with the respective substitutions, that is, the posterior distribution of β_2 and ϕ is multivariate normal- χ^2 with parameters μ_2 , \mathbf{C}_2^* , n_2 and σ_2^2 given by

$$\begin{aligned}\mu_2 &= \mathbf{C}_2^{*-1}(\mathbf{C}_2^{-1}\mu_0 + \mathbf{X}'_1\mathbf{C}_0\mathbf{y}) \\ \mathbf{C}_2^* &= \mathbf{C}_2 + \mathbf{X}_1\mathbf{X}_2\mathbf{C}_0\mathbf{X}'_1\mathbf{X}_2 \\ n_2 &= n + n_0 \\ n_2\sigma_2^2 &= n_0\sigma_0^2 + (\mathbf{y} - \mathbf{X}_1\mathbf{X}_2\mu_2)' \mathbf{C}_0\mathbf{y} + (\mu_0 - \mu_2)' \mathbf{C}_2^{*-1}\mu_0.\end{aligned}$$

8.5 Dynamic linear models

This is a broad class of models with time-varying parameters, useful to model time series data and regression. It was introduced by Harrison and Stevens (1976) and is very well documented in the book by West and Harrison (1997). In this section some fundamental aspects of dynamic models will be introduced and some examples in time series as well as in regression will be addressed.

8.5.1 Definition and examples

Dynamic linear models are parametric models where the parameter variation with time and the available data information are described probabilistically. They are characterized by a pair of equations, named observational equation and parameters evolution or system equation. These are given by

$$\begin{aligned}Y_t &= \mathbf{x}'_t \beta_t + \epsilon_t, \quad \epsilon_t \sim N(0, \sigma_t^2) \\ \beta_t &= \mathbf{G}_t \beta_{t-1} + \omega_t, \quad \omega_t \sim N(0, \mathbf{W}_t)\end{aligned}$$

where Y_t is a time sequence of observations, independent conditionally on the sequence of parameters β_t , \mathbf{x}_t is a vector of explanatory variables as described in Section 8.1, β_t is a $p \times 1$ vector of parameters, \mathbf{G}_t is a $p \times p$ matrix describing the parameter evolution and, finally, σ_t^2 and \mathbf{W}_t are the variances of the errors associated with the unidimensional observation and with the p -dimensional vector of parameters, respectively.

Summarizing, a dynamic linear model is completely specified by the quadruple $\{\mathbf{x}_t, \mathbf{G}_t, \sigma_t^2, \mathbf{W}_t\}$. Two special cases are, respectively, time series models characterized by $\mathbf{x}_t = \mathbf{x}$ and $\mathbf{G}_t = \mathbf{G}, \forall t$ and dynamic regression models, described by $\mathbf{G}_t = \mathbf{I}_p$.

Example 3. The simplest model in time series is the first-order polynomial model, which corresponds to a first-order Taylor series approximation of a smooth time function, named the time series trend. This model is completely defined by the quadruple $\{1, 1, \sigma_t^2, W_t\}$. The above equations specialize to

$$\begin{aligned}Y_t &= \beta_t + \epsilon_t, \quad \epsilon_t \sim N(0, \sigma_t^2) \\ \beta_t &= \beta_{t-1} + \omega_t, \quad \omega_t \sim N(0, W_t)\end{aligned}$$

where β_t is unidimensional.

Although this model is very simple, it can be applied in many short-term forecasting systems involving a large number of time series such as in stock control or production planning. The observational and parameter variance can also evolve in time, offering a broad scope for modelling.

A slightly more elaborated model, named the linear growth model (LGM, in short), is derived after including an extra parameter $\beta_{2,t}$ to describe the underlying

growth of the process. Then, it follows that

$$\begin{aligned} Y_t &= \beta_{1,t} + \epsilon_t, \quad \epsilon_t \sim N(0, \sigma_t^2) \\ \beta_{1,t} &= \beta_{1,t-1} + \beta_{2,t-1} + \omega_{1,t} \\ \beta_{2,t} &= \beta_{2,t-1} + \omega_{2,t}, \quad \omega_t = (\omega_{1,t}, \omega_{2,t})' \sim N(\mathbf{0}, \mathbf{W}_t). \end{aligned}$$

The parameter $\beta_{1,t}$ is interpreted as the current level of the process and it is easy to verify that $\mathbf{x}_t = (1, 0)$ and $\mathbf{G}_t = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$, $\forall t$ characterizing a time series model.

Example 1 (continued). Simple dynamic linear regression. Suppose, in this example, that pairs of values (x_t, Y_t) are observed through time and that it is wished to model the existing relationship between x_t and Y_t . Assuming that the linear model is a good approximation for the relationship between these values, a simple linear regression model can be set. Its parameters can be estimated via classical methods or through the use of the Bayesian argument as described in Section 8.3. Since the linear relationship is only a local approximation for the true functional dependence involving x and Y , a model with varying parameters is appropriate. In many applications time-varying parameters are more adequate. For example, the omission of some variables can justify the parameter oscillation, the non-linearity of the functional relationship connecting x and Y or some structural changes occurring in the process under investigation can also be responsible for the parameter instability. Then, these situations can be modelled as

$$\begin{aligned} Y_t &= \mathbf{x}'_t \boldsymbol{\beta}_t + \epsilon_t \\ \boldsymbol{\beta}_t &= \boldsymbol{\beta}_{t-1} + \omega_t \end{aligned}$$

where $\mathbf{x}_t = (1, x_t)'$ and $\omega_t \sim N(\mathbf{0}, \mathbf{W}_t)$. Note that, in this case, $\mathbf{G}_t = \mathbf{I}_2$.

As we can observe, the choice of \mathbf{x}_t and \mathbf{G}_t depends on the model and the nature of the data that is being analysed. To complete the model specification the variances σ_t^2 and \mathbf{W}_t must be set. The observational variance is usually supposed time invariant, as in the previous sections and \mathbf{W}_t describes the speed of the parameter evolution. In applications σ_t^2 is, often, larger than the elements of \mathbf{W}_t . In what follows the parameter estimation method, including the observational variance, will be described. To make it easier for the conjugate analysis, \mathbf{W}_t is scaled by σ_t^2 and the conditional variances of the ω_t become $\sigma_t^2 \mathbf{W}_t$. Therefore, the matrix \mathbf{W}_t can be interpreted as a matrix of relative weights with respect to the observational variance. The parameter evolution variance matrix must be assessed subjectively by the user of the method and, in order to do that, the notion of discount factor will be useful. Alternatively, it can be estimated by one of the approximating methods described in Chapter 5.

The equations presented before can be rewritten as

$$\begin{aligned} Y_t | \boldsymbol{\beta}_t &\sim N(\mathbf{x}'_t \boldsymbol{\beta}_t, \sigma_t^2) \\ \boldsymbol{\beta}_t | \boldsymbol{\beta}_{t-1} &\sim N(\mathbf{G}_t \boldsymbol{\beta}_{t-1}, \sigma_t^2 \mathbf{W}_t). \end{aligned}$$

Let $D_t = \{D_{t-1}, y_t\}$ with D_0 describing the initial available information, including the values of \mathbf{x}_t and \mathbf{G}_t , $\forall t$, which are supposed to be known.

It is worth noting that it is assumed that for any time t , the current observation Y_t is independent of the past observations given the knowledge of $\boldsymbol{\beta}_t$. This means that the temporal dynamics is summarized in the state parameter evolution. This linear structure for modelling data observed through time combines very well with the principles of Bayesian inference by the possibility to describe subjectively the involved probabilities and by its sequential nature. Therefore, subjective information is coherently combined with past information to produce convenient inferences.

8.5.2 Evolution and updating equations

The equations described before enable a joint description of $(Y_t, \boldsymbol{\beta}_t)$ given the past observed data D_{t-1} via

$$p(y_t, \boldsymbol{\beta}_t | D_{t-1}) = p(y_t | \boldsymbol{\beta}_t) p(\boldsymbol{\beta}_t | D_{t-1}).$$

This leads to the predictive distribution after integrating out $\boldsymbol{\beta}_t$.

One of the main characteristics of the dynamic linear model is that, at each instant of time, all the information available is used to describe the posterior distribution of the state vector. The theorem that follows shows how to evolve from the posterior distribution at time $t - 1$ to the posterior at t .

Theorem 8.1. Consider a normal dynamic linear model with $\sigma_t^2 = \sigma^2$, $\forall t$. Denote the posterior distribution at $t - 1$ by $(\boldsymbol{\beta}_{t-1} | D_{t-1}, \sigma^2) \sim N(\mathbf{m}_{t-1}, \sigma^2 \mathbf{C}_{t-1})$ and the marginal posterior distribution of $\phi = \sigma^{-2}$ as

$$\phi | D_{t-1} \sim G(n_{t-1}/2, n_{t-1}s_{t-1}/2).$$

Then,

1. Conditionally on σ^2 , it follows that

(a) Evolution – the prior distribution at t will be

$$\boldsymbol{\beta}_t | \sigma^2, D_{t-1} \sim N(\mathbf{a}_t, \sigma^2 \mathbf{R}_t)$$

with $\mathbf{a}_t = \mathbf{G}_t \mathbf{m}_{t-1}$ and $\mathbf{R}_t = \mathbf{G}_t \mathbf{C}_{t-1} \mathbf{G}_t' + \mathbf{W}_t$.

(b) The one-step-ahead predictive distribution will be

$$y_t | \sigma^2, D_{t-1} \sim N(f_t, \sigma^2 Q_t)$$

with $f_t = \mathbf{x}'_t \mathbf{a}_t$ and $Q_t = \mathbf{x}'_t \mathbf{R}_t \mathbf{x}_t + 1$.

(c) Updating – the posterior distribution at t will be

$$\beta_t | \sigma^2, D_t \sim N(\mathbf{m}_t, \sigma^2 \mathbf{C}_t)$$

with $\mathbf{m}_t = \mathbf{a}_t + \mathbf{A}_t e_t$ and $\mathbf{C}_t = \mathbf{R}_t - \mathbf{A}_t \mathbf{A}'_t Q_t$, where $\mathbf{A}_t = \mathbf{R}_t \mathbf{x}'_t / Q_t$ and $e_t = y_t - f_t$.

2. The precision ϕ is updated by the relation

$$\phi | D_t \sim G(n_t/2, n_t s_t/2)$$

with $n_t = n_{t-1} + 1$ and $n_t s_t = n_{t-1} s_{t-1} + e_t^2 / Q_t$.

3. Unconditionally on σ^2 , we will have

- (a) $\beta_t | D_{t-1} \sim t_{n_{t-1}}(\mathbf{a}_t, s_{t-1} \mathbf{R}_t)$;
- (b) $Y_t | D_{t-1} \sim t_{n_{t-1}}(f_t, Q_t^*)$, with $Q_t^* = s_{t-1} Q_t$;
- (c) $\beta_t | D_t \sim t_{n_t}(\mathbf{m}_t, s_t \mathbf{C}_t)$.

Proof. (1) Item (a) follows immediately using the parameter evolution equation and standard facts from the normal theory. With respect to (b), using the prior distribution in (a), it follows that

$$\begin{aligned} f_t &= E[E(Y_t | \beta_t) | \sigma^2, D_{t-1}] = E[\mathbf{x}'_t \beta_t | \sigma^2, D_{t-1}] = \mathbf{x}'_t \mathbf{a}_t \\ Q_t &= V[E(Y_t | \beta_t) | \sigma^2, D_{t-1}] + E[V(Y_t | \beta_t) | \sigma^2, D_{t-1}] \\ &= V[\mathbf{x}'_t \beta_t | \sigma^2, D_{t-1}] + \sigma^2 = \sigma^2 (\mathbf{x}'_t \mathbf{R}_t \mathbf{x}_t + 1) \end{aligned}$$

and the normality is a consequence of the fact that all the distributions involved are normal.

To prove part (c), suppose that the posterior distribution at $t-1$ is as given in the theorem. We wish to show that (c) follows from the application of Bayes' theorem, that is,

$$p(\beta_t | \sigma^2, D_t) \propto p(\beta_t | \sigma^2, D_{t-1}) p(y_t | \beta_t, \sigma^2).$$

This can be shown using Theorem 2.1 and the identity

$$\mathbf{C}_t^{-1} = \mathbf{R}_t^{-1} + \mathbf{x}'_t \mathbf{x}_t \sigma^{-2}.$$

If σ^2 is unknown and defining $\phi = \sigma^{-2}$ it will follow that

- By hypothesis, $\phi | D_{t-1} \sim G(n_{t-1}/2, n_{t-1} s_{t-1}/2)$, and $y_t | \phi, D_{t-1} \sim N(f_t, Q_t/\phi)$. Then, by Bayes' theorem,

$$p(\phi | D_t) \propto \phi^{(n_{t-1}+1)/2-1} \exp \left\{ -\frac{\phi}{2} \left(n_{t-1} s_{t-1} + \frac{e_t^2}{Q_t} \right) \right\}$$

and therefore, $\phi | D_t \sim G(n_t/2, n_t s_t/2)$.

- Finally, for part 3 of the theorem, the proofs of items (a)–(c) follow from the results about conjugacy of the normal- χ^2 to the normal model and from the marginal distributions obtained in Sections 3.3 and 8.3.

□

8.5.3 Special aspects

Among the special aspects involved in the dynamic Bayesian modelling, the fact that it is possible to model the observational variance deserves special attention. For example, it can be modelled as a power law, $\sigma_t^2 = \sigma^2 \mu_t^b$, where $\mu_t = \mathbf{x}'_t \beta_t$ is the process mean level. The constant b can be chosen in parallel to the well-known Box–Cox family of transformations. The scale factor σ^2 can be sequentially estimated as stated in the theorem. The main advantage in this form of modelling is to avoid data transformation, thus leaving original data and interpretation of parameters unchanged. This can be useful for times where one wishes to perform subjective intervention in the series. Also, when analysing similar time series, it is possible to incorporate the hierarchical structure into the dynamic framework. This idea is formalized with dynamic hierarchical models (Gamerman and Migon, 1993).

To avoid directly setting the state parameter evolution matrix, the use of discount factors is proposed. These are fixed numbers between zero and one describing subjectively the loss of information through time. Remember that $\mathbf{R}_t = \mathbf{P}_t + \mathbf{W}_t$, where $\mathbf{P}_t = \mathbf{G}_t \mathbf{C}_{t-1} \mathbf{G}'_t$. Denoting the discount factor by δ , we can rewrite $\mathbf{R}_t = \mathbf{P}_t / \delta$, showing clearly that there is a relationship between \mathbf{W}_t and δ . This is given by $\mathbf{W}_t = (\delta^{-1} - 1) \mathbf{P}_{t-1}$, showing that the loss of information is proportional to the posterior variance of the state parameters. For example, if $\delta = 0.9$, only about 90% of the information passes through time.

Other relevant aspects of dynamic linear models are to easily take care of missing observations and to automatically implement subjective interventions. In the first case, it suffices not to use the updating equations at the time the observations are missing. In this way, the uncertainties increase with the evaluation of the new prior distribution and the recurrence equation continues to be valid without any additional problem. From the intervention point of view, the simplest proposal is to use a small discount factor, close to zero, at the time of announced structural changes in the data generation process. In this way the more recent observations will be strongly considered in the updating of the prior distribution and the system can be more adaptive to possible changes.

Finally, it is worth mentioning that parameter distribution at time t can be revised with the arrival of the new observations. We can generically state the parameter distributions $p(\beta_t | D_{t+k})$, $\forall k$ integer. If $k > 0$, this is named the smoothed distribution, if $k = 0$, it is just the posterior and if $k < 0$, it is the prior distribution. In a dynamic model it is common to use the distributions $p(\beta_t | D_n)$, $\forall t = 1, \dots, n$, where n is the size of the series, to retrospectively analyse the parameter behaviour. For example, one may want to quantify the effect of a behaviour change induced by some measure of economic policy. The future data would inform about the change occurring in any particular parameter of the model describing the behaviour of the economic agents involved.

232 Introduction to linear models

Exercises

§ 8.2

1. Show that the derivative of $S(\beta)$ with respect to β is given by $2(\mathbf{X}'\mathbf{X}\beta - \mathbf{X}'\mathbf{y})$. Verify that $\hat{\beta}$ does in fact minimize $S(\beta)$ by calculating the second derivative of $S(\beta)$ with respect to β .
2. Prove that $S(\beta) = (\beta - \hat{\beta})'\mathbf{X}'\mathbf{X}(\beta - \hat{\beta}) + S_e$ where

$$\begin{aligned} S_e &= \mathbf{y}'\mathbf{y} - \hat{\beta}'\mathbf{X}'\mathbf{X}\hat{\beta} \\ &= (\mathbf{y} - \mathbf{X}\hat{\beta})'(\mathbf{y} - \mathbf{X}\hat{\beta}). \end{aligned}$$

3. Obtain a $100(1 - \alpha)\%$ confidence region for β based on $[S(\beta) - S_e]/ps^2$. What is the form of the region?
4. Consider the model $Y_i = \beta_1 + \beta_2x_{2i} + \dots + \beta_px_{pi} + e_i$, $i = 1, \dots, n$ with the e_i 's iid $N(0, \sigma^2)$. Construct a $100(1 - \alpha)\%$ confidence region for $(\beta_2, \dots, \beta_p)$ from the MLR test of validity of the model with level α .
5. Consider the model $Y_i = \beta_1 + \beta_2(x_{2i} - \bar{x}_2) + \dots + \beta_p(x_{pi} - \bar{x}_p) + e_i$, $i = 1, \dots, n$ with the e_i 's iid $N(0, \sigma^2)$ and $\bar{x}_j = (1/n) \sum_{i=1}^n x_{ji}$, $j = 2, \dots, p$. Show that the MLR test of the hypothesis $H_j: \beta_j = 0$, $j = 2, \dots, p$, of level α rejects H_j if

$$|\hat{\beta}_j| > t_{\alpha/2, n-p}s / \sqrt{\sum_{j=1}^n (x_{ji} - \bar{x}_j)^2}.$$

6. Obtain the expression of the $100(1 - \alpha)\%$ confidence interval for the prediction of a future observation Y^* :
 - (a) in the simple linear regression with explanatory variable x^* .
 - (b) in the analysis of variance model with a single classification factor, for an observation for group j , $j = 1, \dots, k$.

§ 8.3

7. Show that

$$\begin{aligned} (\beta - \mu_0)'C_0(\beta - \mu_0) + (\beta - \hat{\beta})'\mathbf{X}'\mathbf{X}(\beta - \hat{\beta}) \\ = (\beta - \mu_1)C_1(\beta - \mu_1) + \mu_0' C_0 \mu_0 + \hat{\beta}'\mathbf{X}'\mathbf{X}\hat{\beta} + \mu_1' C_1 \mu_1 \end{aligned}$$

where $\mu_1 = C_1^{-1}(C_0\mu_0 + \mathbf{X}'\mathbf{y})$ and $C_1 = C_0 + \mathbf{X}'\mathbf{X}$.

8. Obtain the expressions of μ_1 , C_1 , n_1 and σ_1^2 for the particular case of a simple linear regression, showing that they coincide with the expressions obtained in Example 1.
9. Obtain the joint posterior distribution of β_0 and β_1 in the simple linear regression model. Are these parameters independent a posteriori? (Note that they are conditionally independent given ϕ .) Repeat the exercise for the one-way analysis of variance model.

10. Perform Bayesian inference for the simple linear regression model with the same prior as before but with prior correlation ρ ($0 < |\rho| < 1$) conditional on ϕ , instead of 0, as assumed before.

11. Prove that if

- (a) $\mathbf{Y} = \mathbf{X}\beta + \mathbf{e}$, and $\mathbf{e} \sim N(\mathbf{0}, \phi^{-1}\mathbf{I}_n)$ where $\phi = 1/\sigma^2$,
- (b) $\beta | \phi \sim N(\mu, \phi^{-1}\mathbf{C})$ and
- (c) $v\sigma_0\phi \sim \chi_v^2$,

then $\mathbf{Y} \sim t_v(\mathbf{X}\mu, \sigma_0^2(\mathbf{I}_n + \mathbf{X}\mathbf{C}\mathbf{X}'))$, using only the formula $p(\mathbf{z}) = \int p(\mathbf{z} | \mathbf{w})p(\mathbf{w})d\mathbf{w}$.

12. Obtain the expression of the Bayes factor to test the hypothesis of validity of the model for
 - (a) the simple linear regression model.
 - (b) the analysis of variance model with a single classification factor with k levels.

§ 8.4

13. Consider the hierarchical model with K stages given by

$$\begin{aligned} \mathbf{Y} | \beta_1, \phi &\sim N(\mathbf{X}_1\beta_1, \phi^{-1}\mathbf{I}_n) \\ \beta_k | \beta_{k+1}, \phi &\sim N(\mathbf{X}_{k+1}\beta_{k+1}, \phi^{-1}\mathbf{C}_k^{-1}), k = 1, \dots, K-1 \\ \beta_K | \phi &\sim N(\mu, \phi^{-1}\mathbf{C}_K^{-1}) \\ n_0\sigma_0^2\phi &\sim \chi_{n_0}^2. \end{aligned}$$

- (a) Obtain the prior distribution of $\beta_k | \phi$, $k = 1, \dots, K-1$.
- (b) Obtain the marginal prior distribution of β_k , $k = 1, \dots, K-1$.
- (c) Obtain the marginal distribution of \mathbf{Y} .

14. Show that in the hierarchical analysis of variance model with a single classification factor with k levels, $\text{Cov}(\beta_j, \beta_{j'}) | \phi = \text{Cov}(\beta_j, \beta_{j'}) = c_2^{-1}$ and $\text{Cor}(\beta_j, \beta_{j'} | \phi) = \text{Cor}(\beta_j, \beta_{j'}) = (1 + c_2/c_1)^{-1}$, $\forall(j, j')$, where Cor denotes the correlation.
15. Consider the hierarchical model with two stages. Show that the posterior distribution of β_2 and ϕ is multivariate normal- χ^2 with parameters μ_2 , \mathbf{C}_2^* , n_2 and σ_2^2 given by

$$\begin{aligned} \mu_2 &= \mathbf{C}_2^{*-1}(\mathbf{C}_2^{-1}\mu_0 + \mathbf{X}_1'\mathbf{C}_0\mathbf{y}) \\ \mathbf{C}_2^* &= \mathbf{C}_2 + \mathbf{X}_1\mathbf{X}_2\mathbf{C}_0\mathbf{X}_1' \\ n_1 &= n + n_0 \\ n_1\sigma_1^2 &= n_0\sigma_0^2 + (\mathbf{y} - \mathbf{X}_1\mathbf{X}_2\mu_2)' \mathbf{C}_0 \mathbf{y} + (\mu_0 - \mu_2)' \mathbf{C}_2^{*-1} \mu_0. \end{aligned}$$

234 Introduction to linear models

16. Show that in the hierarchical analysis of variance model with a single classification factor, m observations in each group and non-informative prior at the second level,

$$E(\beta_j | \phi, \mathbf{y}) = w_j \bar{y}_j + (1 - w_j) \bar{\bar{y}}$$

with $0 \leq w_j \leq 1$, $j = 1, \dots, k$ and $\bar{\bar{y}}$ is the average of group averages. Obtain the expressions for w_j , $j = 1, \dots, k$.

§ 8.5

17. Consider the first-order polynomial dynamic model with $\sigma_t^2 = \sigma^2$ and $W_t = W$, $\forall t$. Derive the predictive distribution for a horizon $k > 0$ and show that

$$y_{t+k} | D_t \sim N(m_t, C_t + kW + \sigma^2).$$

Obtain the distribution of $(y_{t+1} + \dots + y_{t+k} | D_t)$ from the joint predictive distribution of $(y_{t+1}, \dots, y_{t+k} | D_t)$. For example, take $k = 2$, and prove that

$$y_{t+1} + y_{t+2} | D_t \sim N(2m_t, 4C_t + 2\sigma^2 + 5W).$$

18. Consider again the first-order polynomial dynamic model. Show that

$$\beta_{t-1} | D_t \sim N(m_{t-1}^*, C_{t-1}^*)$$

with

$$E(\beta_{t-1} | D_t) = m_{t-1} + C_{t-1}/R_t[m_t - m_{t-1}]$$

and

$$V(\beta_{t-1} | D_t) = C_{t-1} - (C_{t-1}/R_t)^2(R_t - C_t).$$

19. Suppose the series was not observed at time t so that Y_t was missing and therefore, $D_t = D_{t-1}$. Obtain the distributions of $\beta_t | D_t$ and $y_{t+1} | D_t$ assuming knowledge of m_{t-1} and C_{t-1} for a first-order polynomial dynamic model.

Sketched solutions to selected exercises

Chapter 2

2. The test X is such that $P(X = 1 | \theta = 1) = 0.95$, $P(X = 1 | \theta = 0) = 0.40$ and the disease prevalence is $P(\theta = 1) = 0.70$.
 - (a) From the example on page 26, $P(\theta = 1 | X = 1) = 0.847$. By Bayes' theorem, $P(\theta = 1 | X = 0) \propto l(\theta = 1; X = 0)P(\theta = 1) = 0.035$ and $P(\theta = 0 | X = 0) \propto l(\theta = 0; X = 0)P(\theta = 0) = 0.180$. Therefore, $P(\theta = 1 | X = 0) = 0.035/(0.035 + 0.180) = 0.163$. The result $X = 1$ makes the doctor more certain about John's illness because $P(\theta = 1 | X = 1) = 0.847 > 0.163 = P(\theta = 1 | X = 0)$.
 - (b) Using $P(\theta | X = 1)$ as the prior for the second experiment, it follows from Bayes' theorem that $P(\theta = 1 | X_1 = 1, X_2 = 1) \propto 0.95 \times 0.847 = 0.805$ and $P(\theta = 0 | X_1 = 1, X_2 = 1) \propto 0.40 \times 0.153 = 0.063$. So, the probability that John is ill is $P(\theta = 1 | X_1 = 1, X_2 = 1) = 0.805/(0.063 + 0.805) = 0.927$.
 - (c) The likelihood for n positive results is $(0.95)^n$. Therefore the solution of $P(\theta = 1 | X_1 = 1, \dots, X_n = 1) = (0.95)^n 0.70 / ((0.4)^n 0.3 + (0.95)^n 0.7) \geq 0.99$ is obtained by trial and error as $n = 8$.
5. Let θ represent the event *the driver is drunk* with $P(\theta) = 0.75$. The test X_1 will be positive ($=1$) if the level of alcohol in his/her blood is high and zero otherwise. It is known that $P(X_1 = 1 | \theta = 1) = 0.8$ and for a second test $P(X_2 = 0 | \theta = 0) = 1$ and $P(X_2 = 0 | \theta = 1) = 0.10$.
 - (a) $P(X_1 = 1) = (0.8)(0.25) + (0.2)(0.75) = 0.35$ can be interpreted as the proportion of drivers stopped that have to be submitted to a second test.
 - (b) By Bayes' theorem, $P(\theta = 1 | X_1 = 1) \propto (0.8)(0.25)/0.35 = 0.571$. Also, $P(X_2 = 1 | X_1 = 1) = 0(1 - 0.571) + (0.9)(0.571) = 0.514$. Therefore, $P(\theta = 1 | X_1 = 1, X_2 = 1) = (0.9)(0.571)/0.514 = 1$.
 - (c) Obviously, $P(X_1 = 0) = 1 - P(X_1 = 1) = 0.65$.
8. (a) $p(\theta | x, \mu) \propto p(\theta, \mu, x) = p(x | \theta, \mu)p(\theta | \mu)p(\mu)$. But, $p(\theta | x, \mu) \propto \exp\{-(0.5)[(x - \theta)^2/\sigma^2 + (\theta - \mu)^2/\tau^2 + \mu^2]\}$. After some algebra

- we get $(\theta|x, \mu) \sim N(\mu_1, \tau_1^2)$ with $\mu_1 = \tau_1^2(x/\sigma^2 + \mu/\tau^2)$ and $\tau_1^{-2} = (\sigma^{-2} + \tau^{-2})$. This is in fact an application of Theorem 2.1.
- (b) In order to apply Bayes' theorem we need the likelihood $l(\mu; x) = p(x|\mu) = \int p(x|\theta, \mu)p(\theta|\mu)d\theta \propto \int \exp\{-(0.5)[(x-\theta)^2/\sigma^2 - (\theta - \mu)^2/\tau^2]\}d\theta$. After some calculations we get

$$p(x|\mu) \propto \exp[-(0.5)(x-\mu)^2/(\tau^2 + \sigma^2)].$$

- Using Theorem 2.1 once again, it follows that $\mu|x \sim N(\mu_1, \tau_1^2)$ with $\mu_1 = x/(\sigma^2 + \tau^2 + 1)$ and $\tau_1^2 = (\sigma_2 + \tau^2)/(\sigma^2 + \tau^2 + 1)$.
- (c) $p(\theta) \propto \int p(\theta|\mu)p(\mu)d\mu$ or $\theta \sim N(0, 1+\tau^2)$. It follows immediately by Bayes' theorem that $\theta|x \sim N(\theta_1, \tau_2^2)$ where $\theta_1 = \tau_2^2 x \sigma^2$ and $\tau_2^2 = \sigma_2(\tau^2 + 1)/(\sigma^2 + \tau^2 + 1)$.

In fact the above results can be easily obtained after rewriting the exercise as $X = \theta + \epsilon$, $\epsilon \sim N(0, \sigma^2)$, $\theta = \mu + \omega$, $\omega \sim N(0, \tau^2)$ and $\mu \sim N(0, 1)$, where ϵ , ω and μ are independent. Using results about linear combinations of normal variates we get $X = \mu + \omega + \epsilon$ or $X|\mu \sim N(\mu, \tau^2 + \sigma^2)$. Analogously, $\theta = \mu + \omega$ or $\theta \sim N(0, \tau^2 + 1)$.

15. From exchangeability of the X_i 's and Theorem 2.2, any sequence of n X_i 's having k values of 1 and $n-k$ values 0 has probability given by $\int_0^1 \theta^k (1-\theta)^{n-k} dF(\theta)$, $\forall k \leq n, n > 0$.

- (a) Since $T = \sum X_i$, $P(T = t) = \sum_{X \in A} \int_0^1 \theta^k (1-\theta)^{n-k} dF(\theta)$ where $A = \{X | \sum X_i = t\}$. The number of n -tuples in A is $\binom{n}{t}$ and they all have the same probability. Then, $P(T = t) = \binom{n}{t} \int \theta^t (1-\theta)^{n-t} dF(\theta)$.
- (b) $E[T] = \sum_{i=1}^n E(X_i) = \sum_{i=1}^n E(X_1) = nE(X_1) = nE[E(X_1|\theta)] = nE[\theta]$ since $X_1|\theta \sim \text{Ber}(\theta)$.

21. (a) The sample space depends on the unknown parameter, hence the distribution does not belong to the exponential family.
- (c)

$$\begin{aligned} p(x|\theta) &= (2\pi)^{-1/2} \exp\{-0.5[x^2 - 2\theta x + \theta^2]/\theta\} \\ &= (2\pi)^{-1/2} \exp(x) \exp(-0.5[x^2/\theta - \theta - \log(\theta)]). \end{aligned}$$

Identifying $a(x) = (2\pi)^{-1/2} \exp(x)$, $u(x) = x^2/2$, $\phi(\theta) = \theta^{-1}$ and $b(\theta) = -(\theta + \log(\theta))/2$, we conclude that it is a member of the one-parameter exponential family and, by the Neyman factorization criterion, X^2 is a sufficient statistic.

- (e) $p(x|x \neq 0, \theta) = p(x|\theta)/P[X \neq 0|\theta] = \binom{n}{x} \theta^x (1-\theta)^{n-x}/[1 - (1-\theta)^n]$, $x = 1, \dots, n$ or $p(x|x \neq 0, \theta) = \binom{n}{x} \exp[x \log(\theta)/(1-\theta)] + n \log(1-\theta) - \log[1 - (1-\theta)^n]$. So, $a(x) = \binom{n}{x}$, $u(x) = x$, $\phi(\theta) = \log(\theta/(1-\theta))$ and $b(\theta) = n \log(1-\theta) - \log[1 - (1-\theta)^n]$. It is clear that X is a sufficient statistic for θ .

- (g) $p(x|\theta) = \theta^x \log(\theta)/(\theta - 1) = \exp[x \log \theta + \log[\log \theta/(\theta - 1)]]$. So, $a(x) = 1$, $u(x) = x$, $\phi(\theta) = \log(\theta)$ and $b(\theta) = \log[\log(\theta)/(\theta - 1)]$. So it is a member of the exponential family and $T(X) = X$ is a minimal sufficient statistic.

24. Remember that $p(x|\theta) = a(x) \exp[u(x)\phi(\theta) + b(\theta)]$ and $\int p(x|\theta)dx = 1$.

- (a) Differentiating both sides with respect to θ it follows that

$$\int [u(x)\phi'(\theta) + b'(\theta)]p(x|\theta) dx = 0.$$

It is then easy to get $E[u(X)] = -b'(\theta)/\phi'(\theta)$.

- (b) Differentiating again $\int [u(x)\phi'(\theta) + b'(\theta)]p(x|\theta)dx = 0$, we get

$$\begin{aligned} &\int [u(x)\phi''(\theta)p(x|\theta) + u(x)\phi'(\theta)p'(x|\theta) + b''(\theta)p(x|\theta) \\ &+ b'(\theta)p'(x|\theta)]dx \end{aligned}$$

where $p'(x|\theta) = [u(x)\phi'(\theta) + b'(\theta)]p(x|\theta)$. After some calculation it follows that $E[u^2(X)] = [\phi''(\theta)b'(\theta) - \phi'(\theta)b''(\theta)]/[\phi'(\theta)]^3$. The variance follows from the above calculations.

32. (a) It is easy to get $p(\psi, \xi) = p_{\theta, \phi}(\psi, \xi)|J| \propto \xi$, because the Jacobian of the transformation is just ξ .

- (b) $p(\psi, \xi|x, y) \propto \theta^x \exp(-\theta)\phi^y \exp(-\phi)$ where $\theta = \psi\xi$ and $\phi = \xi(1-\psi)$. So, $p(\psi, \xi|x, y) \propto \xi^{x+y+1} \exp(-\xi)\psi^x(1-\psi)^y$ or proportional to the product of a $G(x+y+2, 1)$ and a beta $(x+1, y+1)$ distribution.

- (c) First of all, it is easy to see that $X + Y \sim \text{Pois}(\xi)$, where $\xi = \theta + \phi$, since X and Y are independent Poisson distributed with parameters θ and ϕ respectively. Since $p(x|x+y, \psi, \xi) = p(x|\psi, \xi)p(x+y|x, \psi, \xi)/p(x+y|\psi, \xi)$ it follows that, $p(x|x+y, \psi, \xi) = \binom{x+y}{x} \psi^x (1-\psi)^y$ which is only a function of ψ .

- (d) Using the distributions obtained in (c) and the factorization criterion it is simple to obtain the results.

- (e) The marginal likelihoods of ψ and ϕ are given by $l(\psi; x, y) = p(x, y|\psi) = \int p(x, y|\psi, \xi)p(\xi|\psi)d\xi$ and $l(\xi; x, y) = p(x, y|\xi) = \int p(x, y|\psi, \xi)p(\psi|\xi)d\psi$. Since $p(\xi|\psi) \propto \xi$ and $p(\psi|\xi) \propto k$, simple integrations lead to $l(\psi; x, y) \propto \psi^x(1-\psi)^y$ and $l(\xi; x, y) \propto \xi^{x+y} \exp(-\xi)$.

35. (a) The inverse transformations are $\theta_1 = \lambda\psi$ and $\theta_2 = \psi(1-\lambda)$ and the Jacobian is $J = \psi$. The distribution of (λ, ψ) follows as $p(\lambda, \psi) \propto \psi$, since $p(\theta) \propto k$.

- (b) $p(x|\lambda, \psi) = \lambda^{x_1}(1-\lambda)^{x_2}\psi^{x_1+x_2}(1-\psi)^{1-x_1-x_2}$. Then the marginal likelihood for ψ is $p(x|\psi) = \int p(x|\lambda, \psi)p(\lambda|\psi)d\lambda$. Since $p(\lambda|\psi) \propto k$, $p(x|\psi) \propto \psi^{x_1+x_2}(1-\psi)^{1-x_1-x_2}$.

- (c) By the Neyman factorization criterion with $f(x_1 + x_2, \psi) = \psi^{x_1+x_2} (1-\psi)^{1-x_1+x_2}$ and $g(\mathbf{x}) = k$ it follows that $X_1 + X_2$ is sufficient for ψ .

Chapter 3

1. (a) My experience of living in Rio gives me confidence to assess the first quartile as 20°C and the median as 28°C .
- (b) A good approximation for the standard deviation is $\sigma \approx 1.25|Q_1 - Q_2|$, where Q_i is the i th quartile. So my assessment is that the temperature is normally distributed with mean 28°C and standard deviation equal to 10°C .
- (c) The 0.1 quantile determined from the normal distribution corresponds to 2.7°C conflicting with the subjective assessment of 8°C . The assessment of normality must be revised.
5. (a) $I(\theta; \mathbf{y}) = \prod_{i=1}^4 \binom{n_i}{y_i} \theta^{y_i} (1-\theta)^{n_i-y_i} \propto \theta^9 (1-\theta)^3$.
- (b) The class of the beta distributions can be used as prior, say with $a = 1$ and $b = 1$ representing vague initial information. So, $p(\theta|\mathbf{y}) \propto \theta^9 (1-\theta)^3$ or a $\theta|\mathbf{y} \sim \text{beta}(10, 4)$.
- (c) Since conditions are similar, it is natural to assume that $y_5|\theta \sim \text{bin}(n_5, \theta)$ with $n_5 = 3$. The predictive distribution will be a beta-binomial distribution with parameter $(3, 10, 4)$ providing the following probabilities

y_5	0	1	2	3
$p(y_5 \theta)$	0.036	0.178	0.393	0.393

10. (a) Assuming that $\theta \sim G(a, b)$ then $\mu = E[\theta] = a/b = 4$ and the coefficient of variation $CV(\theta) = \sigma/|\mu| = 1/a^{1/2} = 1/2$. It follows that $a = 4$ and $b = 1$. The posterior variance will be less than or equal to 0.01 if and only if $(t+a)/(n+b)^2 \leq 0.01$, where $t = \sum x_i$. Solving the quadratic inequality in n it follows that $n \geq 10(4+t)^{1/2} - 1$.
- (b) The posterior mean can be written as $\mu_1 = (a+t)/(b+n) = \gamma_n \bar{x}_n + (1-\gamma_n)\mu_0$ where $\mu_0 = a/b$ and $\gamma_n = n/(n+b)$. The limit of γ_n , when $n \rightarrow \infty$, will be 1.
- (c) The posterior for θ is $\text{beta}(a+t, b+n-t)$ where $t = \sum x_i$. This distribution has mean $(a+t)/(a+b+n) = \gamma_n \bar{x}_n + (1-\gamma_n)\mu_0$ where $\mu_0 = a/(a+b)$ and $\gamma_n = n/(a+b+n)$. The limit of γ_n , when $n \rightarrow \infty$, will be 1.
16. (a) The Pareto distribution is a member of the one-parameter exponential family with $a(\mathbf{x}) = 1/\prod_i x_i$, $\phi(\theta) = -\theta$, $U(\mathbf{x}) = \sum \log(x_i)$ and $b(\theta) = n \log \theta + n\theta \log b$. So the sufficient statistic for θ is $U(\mathbf{X})$.

- (b) The observed information is $-d^2 \log(p(\mathbf{x}|\theta)) / d\theta^2 = n/\theta^2$ which coincides with the Fisher information for a sample of size n . The Jeffreys prior will be $p(\theta) \propto I(\theta)^{1/2} \propto \theta^{-1}$. It is obviously improper since $\int p(\theta)d\theta = \int k\theta^{-1}d\theta$ diverges.
- (c) The posterior distribution is $p(\theta|\mathbf{x}) \propto \theta^{n-1} b^{n\theta} / (\{\prod_{i=1}^n x_i\}^{1/n})^{n\theta}$. Let $z = \{\prod_{i=1}^n x_i\}^{1/n}$ be the sample geometric mean. It is clear that $p(\theta|\mathbf{x}) \propto \theta^{n-1} \exp(-n\theta \log(z/b))$ which corresponds to a $G(n, n \log(z/b))$.
20. The non-informative prior for θ is $p(\theta) \propto k$. The density for $\phi = a\theta + b$, $a \neq 0$, is $p(\phi) = p_\theta[(\phi - b)/a] |d\theta/d\phi| \propto k d[(\phi - b)/a]/d\phi \propto k$. If $p(\theta) \propto \theta^{-1}$, $\theta > 0$ then $p(\phi) = p_\theta(\phi^{1/a}) |d\theta/d\phi|$ with $\phi = \theta^a$, $a \neq 0$ or $p(\phi) \propto \phi^{-1}$. If $\psi = \log \theta$ then $p(\psi) \propto \exp(-\psi) \exp(\psi) = 1$.
27. (a) Let $\theta_i|\mu \sim N(\mu, b)$, b known and suppose that the θ_i 's are independent given μ . Assuming that $p(\mu) \propto k$, $p(\theta, \mu) \propto \prod_{i=1}^n p(\theta_i|\mu, b)$, or $\theta|\mu \sim N(\mu \mathbf{1}_n, b \mathbf{I}_n)$.
- (b) Since $p(\mu|\mathbf{y}) \propto p(\mathbf{y}|\mu)p(\mu)$ where $Y_i|\mu \sim N[\mu, (a+b)]$, $i = 1, \dots, n$, independent, it follows from Theorem 2.1 that $\mu|\mathbf{y} \propto N[\bar{y}, (a+b)/n]$.
- (c) Note that $p(\theta_i|\mu, \mathbf{y}) \propto p(\theta_i|\mu)p(\mathbf{y}|\theta_i)$. Once again, it follows from Theorem 2.1 that $\theta_i|\mu, \mathbf{y} \sim N[\mu_i^*, b^*]$, where $\mu_i^* = (a\mu + b y_i)/(a+b)$ and $b^* = ab/(a+b)$. Hence, $E[\theta_i|\mu, \mathbf{y}] = \mu_i^*$, $i = 1, \dots, n$.
- (d) $E[\theta_i|\mathbf{y}] = E[E[\theta_i|\mathbf{y}, \mu]] = E(\mu_i^*) = b y_i/(a+b)$.

Chapter 4

4. The posterior $p(\theta|\mathbf{x}) \propto \theta^t (1-\theta)^{n-t} I_\theta(0, 1)$ where $t = \sum_{i=1}^n x_i$, that is, a beta($t+1, n-t+1$).

- (a) $E[L(\theta, d)|\mathbf{x}] \propto \int_0^1 \{(\theta-d)^2 / [\theta(1-\theta)]\} \theta^t (1-\theta)^{n-t} d\theta = \int_0^1 (\theta-d)^2 \theta^{t-1} (1-\theta)^{n-t-1} d\theta$. This integral is proportional to the expected value of the square loss with respect to the beta $(t, n-t)$ distribution. So, it is minimized when $d = t/n = \bar{x}$. The risk of the Bayes estimator is $R(\bar{x}) = B(t+1, n-t+1) \int_0^1 (\theta-\bar{x})^2 \theta^t (1-\theta)^{n-t} d\theta$. Multiplying and dividing by $B(t, n-t)$, we get $R(\bar{x}) = [B(t+1, n-t+1)/B(t, n-t)] V(\theta|\mathbf{x})$. After some simplifications, it follows that $R(\bar{x}) = 1/n^2$.

- (b) The predictive distribution is obtained via

$$p(x_{n+1}|\mathbf{x}) = \int_0^1 p(x_{n+1}|\theta) p(\theta|\mathbf{x}) d\theta.$$

Assuming that $X_{n+1} \sim \text{Ber}(\theta)$ independent of \mathbf{X} , it follows that $p(x_{n+1}|\mathbf{x}) = B(t+1, n-t+1) \int_0^1 \theta^{x_{n+1}+t} (1-\theta)^{n-x_{n+1}-t+1} d\theta$. Solving the integral gives $p(x_{n+1}|\mathbf{x}) = B(t+1, n-t+1)/B(t+$

- $x_{n+1} + 1, n - x_{n+1} - t + 2$. Finally, $p(x_{n+1}|x) = (t+1)^{x_{n+1}}(n-t+1)^{1-x_{n+1}}/(n+2)$. The mean and variance of the predictive distribution are $(t+1)/(n+2)$ and $(t+1)(n-t+1)/(n+2)^2$, respectively.
- (c) Let t_1, \dots, t_k be the counts associated with each of the k possible values of X . Then, straightforward generalizations give for each θ_i the Bayes estimator t_i/n , $i = 1, \dots, k$, with associated risks $1/n^2$. The predictive probability function generalizes to $p(x_{n+1} = i|x) = (t_i + 1)/(n+k)$, $i = 1, \dots, k$.
7. (a) $p(\theta|x) \propto \theta^{-1} I_\theta(0, 1)$. The proportionality constant is obtained making $1 = \int p(\theta|x)d\theta = \int_{x-1}^{x+1} k\theta^{-1}d\theta = k \log \theta|_{x-1}^{x+1}$. So, $k^{-1} = \log[(x+1)/(x-1)]$, $x > 1$.
- (b) Posterior mean, mode and median are easily obtained. $E(\theta|x) = \int_{x-1}^{x+1} k\theta\theta^{-1}d\theta = k[(x+1) - (x-1)] = 2k = 2 \log[(x-1)/(x+1)]$. To evaluate the mode, it is enough to observe that $p(\theta|x)$ is a strictly decreasing function of θ . So its maximum occurs at $x-1$. The median, in turn, is the solution of $\int_{x-1}^m k\theta^{-1}d\theta = 1/2$ or $k \log[m/(x-1)] = 1/2$. After some simplifications it follows that $m = [(x-1)(x+1)]^{1/2}$.
12. The density functions for each type of bulb are $p(x|\psi) = \psi \exp(-\psi x)$, $\psi^{-1} = \theta, 2\theta$ and 3θ , respectively.
- (a) Observing one bulb of each type, it follows that $l(\theta; x) \propto \exp[-(x_1 + x_2/2 + x_3/3)/\theta]/6\theta^3$. The MLE is the solution of the equation $d \log l(\theta; x)/d\theta = -3/\theta + (X_1 + X_2/2 + X_3/3)/\theta^2 = 0$. That is, $\hat{\theta} = (X_1 + X_2/2 + X_3/3)/3$.
- (b) Assuming the prior $\psi \sim G(\alpha, \beta)$, it follows that the posterior distribution is $p(\psi|x) \propto \psi^{3+\alpha-1} \exp[-\psi(x_1 + x_2/2 + x_3/3 + \beta)]$.
- (c) The Bayes estimator is the posterior mean, $E[\psi|x] = (\alpha + 3)/(\beta + x_1 + x_2/2 + x_3/3)$, if a square loss function is assumed. For the 0-1 loss function the Bayes estimator is the posterior mode, $(\alpha + 2)/(\beta + x_1 + x_2/2 + x_3/3)$.
16. It is worth remembering that the moment generating function for the Pois(θ) is $\varphi_X(t) = E[\exp(-tX)] = \exp[-\theta(1 - \exp(-t))]$.
- (a) Since $\varphi_Y(t) = \prod_{i=1}^n \varphi_{X_i}(t) = \exp[-n\theta(1 - \exp(-t))]$, then $E[\exp(-cY)] = \varphi(c) = \exp[-n\theta(1 - \exp(-c))]$.
- The estimator $\exp(-cY)$ will be unbiased for $\exp(-\theta)$ iff $n\theta(1 - \exp(-c)) = \theta$ or $c = \log[n/(n-1)]$. Therefore, $(1 - 1/n)^Y$ is an unbiased estimator of $e^{-\theta}$.
- (b) By the Cramer-Rao inequality, the variance lower bound of an unbiased estimator of $h(\theta)$ is given by $h'(\theta)^2/I(\theta)$. In the present case $h(\theta) = \exp(-\theta)$, so $h'(\theta) = -\exp(-\theta)$ and the expected information is obtained as $E[-d^2 \log l(\theta; x)/d\theta^2] = E[\sum x_i/\theta^2] = n/\theta$. Therefore $V[\exp(-cY)] \geq \theta \exp(-2\theta)/n$, for $c = \log[n/(n-1)]$.

(c) The variance of $\exp(-cY)$ can be calculated as $V[\exp(-cY)] = \varphi_Y(2c) - \varphi_Y^2(c) = \exp[-n\theta(1 - \exp(-2c))] - \exp[-2n\theta(1 - \exp(-c))]$, for $c = \log[n/(n-1)]$. After some algebra, it follows that $V[\exp(-cY)] = \exp[-\theta(2-n)] - \exp(-2\theta)$. The ratio between the Cramer-Rao lower bound and the variance of the estimator is less than 1, so the estimator is not efficient for $\exp(-\theta)$.

24. (a) $p(\theta|x) \propto p(x|\theta) \propto \theta^x(1-\theta)^{n-x}$. Assuming that $X = n$ it follows that $p(\theta|x) \propto \theta^n$, which is a monotonic increasing function of θ . So the maximum posterior density interval will be in the form $[a, 1]$, for some $a < 1$ such that $P[\theta \geq a|x] = 1 - \alpha$.
- (b) Let $\psi = \theta/(1-\theta)$. Then, $P[a/(1-a) \leq \psi|x] = P[a/(1-a) \leq \theta/(1-\theta)|x] = P(a \leq \theta \leq 1|x) = 1 - \alpha$.
- (c) It is easy to evaluate $p(\psi|x) = p_\theta(\psi/(1+\psi)|x) |d\theta/d\psi|$ or $p(\psi|x) \propto \psi^n/(1+\psi)^{n+2}$, which is in the form of an $F[2(n+1), 2]$ distribution.
- (d) The interval obtained in (b) is not an HDP because the F density is not monotonically decreasing in θ .
27. (a) a exponentially distributed lifetimes with observed mean lifetime equal to b lead to a likelihood $l(\theta) \propto \theta^a \exp(-\theta ab)$. The prior distribution for θ can then be obtained as $p(\theta) \propto \theta^a \exp(-\theta ab)$. This distribution is in the form of a $G(a+1, ab)$ distribution.
- (b) Using results of the conjugate analysis, it follows that $p(\theta|x)$ is a $G(a+t+1, n+ab)$ where $t = \sum x_i$.
- (c) Since $T \sim \text{Pois}(n\theta)$, $p(\theta|t) \propto [(n\theta)^t \exp(-n\theta)][\theta^a \exp(-ab\theta)]$ or $\theta|t \sim G(a+t+1, n+ab)$.
- (d) The posterior distributions obtained in (b) and (c) are the same, which is not surprising since T is a sufficient statistic for θ .
31. $P((\theta - \mu_1)^2 \leq 5V_1|x) = P(|\theta - \mu_1|/\sqrt{V_1} \leq \sqrt{5}|x) = P(t_{n_1}(0, 1) \leq \sqrt{5n_1/(n_1-2)})$. By trial and error, $n_1 = 11$ is the largest integer ensuring that the above probability is larger than 0.95. Therefore, $n = n_1 - n_0 = 6$.

Chapter 5

2. The log-likelihood function is given by $L(\alpha, \beta; x) = n[\alpha \log \beta - \log \Gamma(\alpha)] + (\alpha-1)T_1 - \beta T_2$, where $T_1 = \sum_i \log X_i$ and $T_2 = \sum_i X_i$. The maximum likelihood equations will be $\partial L(\alpha, \beta; X)/\partial \alpha = n[\log \beta - \Gamma'(\alpha)/\Gamma(\alpha)] + T_1$ and $\partial L(\alpha, \beta; X)/\partial \beta = n\alpha/\beta - T_2$. The MLE of β as a function of α will be $\hat{\beta}(\alpha) = \alpha/\bar{X}$, where $\bar{X} = T_2/n$. The profile log-likelihood of α is $L(\alpha, \hat{\beta}(\alpha)) = n[\alpha \log(\alpha/\bar{X}) - \log(\Gamma(\alpha))] + (\alpha-1)T_1 - n\alpha$. Differentiating the profile log-likelihood with respect to α we get $\partial L(\alpha, \hat{\beta}(\alpha); X)/\partial \alpha = n[\log \alpha - \Gamma'(\alpha)/\Gamma(\alpha)] + T_1 - n \log \bar{X}$. A numerical optimization method, such as Newton-Raphson, can then be used after a numerical approximation for the digamma function, to obtain the MLE for α . The MLE estimator for β follows from the equation $\hat{\beta} = \hat{\alpha}/\bar{X}$ and the invariance of MLE's.

9. We have $\theta = \lambda^a$ and $p(\mathbf{X}|\theta) \propto e^{-n\theta} \theta^T$ where $T = \sum_i X_i$.

- (a) The likelihood function is $l(\lambda; \mathbf{X}) = k \exp(-n\lambda^a)\lambda^{aT}$.
 - (b) The Jeffreys prior is defined as $p(\lambda) \propto I(\lambda)^{1/2}$ where $I(\lambda) = E[-L''(\lambda; \mathbf{X})]$, L is the log-likelihood and derivatives are taken with respect to λ . Then, $L(\lambda; \mathbf{X}) = -n\lambda^a + aT \log \lambda$. Its first and second derivative are respectively given by $L'(\lambda; \mathbf{X}) = -na\lambda^{a-1} + aT/\lambda$ and $L''(\lambda; \mathbf{X}) = -na(a-1)\lambda^{a-2} - aT/\lambda^2$ and the MLE is $\hat{\lambda} = (T/n)^{1/a}$. Also, $I(\lambda) = na^2\lambda^{a-2}$ and the Jeffreys prior is $p(\lambda) \propto \lambda^{a-1}$.
 - (c) The third derivative is $L'''(\lambda; \mathbf{X}) = -na(a-1)(a-2)\lambda^{a-3} + 2aT/\lambda^3$. Evaluating it at $\hat{\lambda} = \bar{x}^{1/a}$ gives $L'''(\hat{\lambda}; \mathbf{X}) = [-a(a-1)(a-2)T + 2aT]/(T/n)^{3/a}$. It is easy to verify that $L'''(\hat{\lambda}; \mathbf{X}) = 0$ iff $a = 3$ or $a = 0$ (degenerate solution).
 - (d) The transformation that makes the third derivative null improves the asymptotic approximation. So, with respect to skewness, the parametrization $\lambda = \theta^{1/3}$ should be used to improve approximations.
16. We have already obtained that $V(\hat{\theta}) = 0.126/n$. Now, $V(\tilde{\theta}) = V[w(X)]/n$ where $w(X) = (2\pi)^{-1}X^2/(1+X^2)^2$ and $V[w(X)] = E[w^2(X)] - \theta^2$, since the importance sampling estimator is unbiased. $E[w^2(X)] = \int_2^\infty (x^2/(1+x^2))^2/(2/x^2)dx = 0.021875$. Therefore, $V(\tilde{\theta}) = [0.021875 - (0.1476)^2]/n = 9 \times 10^{-5}/n$, which is substantially smaller than $V(\hat{\theta})$.
20. First assume that $\mathbf{y} \neq \mathbf{x}$. The transition kernel of the chain is given by $q^*(\mathbf{x}, \mathbf{y}) = q(\mathbf{x}, \mathbf{y})\alpha(\mathbf{x}, \mathbf{y})$. Then, suppose that $\alpha(\mathbf{x}, \mathbf{y}) < 1$ so that $\alpha(\mathbf{x}, \mathbf{y}) = [p(\mathbf{y})q(\mathbf{y}, \mathbf{x})]/[p(\mathbf{x})q(\mathbf{x}, \mathbf{y})]$. Then, $\alpha(\mathbf{y}, \mathbf{x}) = 1$ and $p(\mathbf{x})q^*(\mathbf{x}, \mathbf{y}) = p(\mathbf{x})q(\mathbf{x}, \mathbf{y})\alpha(\mathbf{x}, \mathbf{y}) = p(\mathbf{y})q(\mathbf{y}, \mathbf{x})\alpha(\mathbf{y}, \mathbf{x}) = p(\mathbf{y})q^*(\mathbf{y}, \mathbf{x})$. This equality is also trivially satisfied when $\mathbf{y} = \mathbf{x}$. Integrating both sides with respect to \mathbf{x} gives $p(\mathbf{y}) = \int p(\mathbf{x})q^*(\mathbf{x}, \mathbf{y})d\mathbf{x}$, which means that $p(\mathbf{y})$ is the stationary distribution of the chain.

Chapter 6

3. (a) The p -value is the statistic $\alpha(T) = F_T(T)$ where F_T is the distribution function of the test statistic T . By the integral probability transform (see, for example, DeGroot (1986, p. 154)), it follows that $\alpha(T) = F_T(T) \sim U(0, 1)$. Since $\alpha(T_i)$, $i = 1, \dots, k$, are function of independent statistics and all with the same distribution they constitute, a random sample of the $U(0, 1)$ distribution.
- (b) The hypothesis is rejected if $\alpha(T) > \alpha$, the significance level. So, small values of $\alpha(T)$ should lead to rejection and hence large values of $-2 \log \alpha(T)$ should lead to rejection. Therefore, large values of F must lead to rejection of the hypothesis.
- (c) It is well known that if $X \sim U(0, 1)$ then $Y = -\log X \sim \text{Exp}(1) = G(1, 1)$ or $2Y \sim G(1, 1/2)$. Since F is the sum of k iid gamma

- distributed random quantities then $F \sim G(k, 1/2)$ or χ_{2k}^2 .
10. (a) The maximum likelihood ratio is $\lambda(X) = \theta_0 \exp[-X(\theta_0 - \theta_1)]/\theta_1$ and the null hypothesis is accepted iff $\lambda(X) > c$ or $X < (\theta_0 - \theta_1)^{-1} \log[\theta_0/(c\theta_1)] = t$, where t is such that $P(X > t|\theta_0) = \alpha$. Then, $\exp(-\theta_0 t) = \alpha$ or $t = -\theta_0^{-1} \log \alpha$.
- (b) The p -value is $P[X > 3|\theta_0] = \exp(-3\theta_0) = 0.0498$.
- (c) The Bayes factor is $BF(H_0, H_1) = 2e^{-3/2} = 0.45$.
- (d) The odds ratio is $p(H_0|x)/p(H_1|x) = p(H_0)/p(H_1)BF(H_0, H_1) = BF(H_0, H_1)$, since $p(H_0) = p(H_1)$. So,
- $$p(H_0|x) = [(BF(H_0, H_1))^{-1} + 1]^{-1} = [(0.45)^{-1} + 1]^{-1} = 0.31.$$
- (e) Using the result in (b), the null hypothesis is rejected at the $\alpha = 5\%$ level because the p -value is less than α , which is somewhat conflicting with the posterior probability of 0.31 for the null hypothesis, because the latter is smaller than 0.5 but is far from recommending rejection.
13. (a) The likelihood function under H_0 is $p(x_1, x_2|\theta) = \theta^t(1-\theta)^{2-t}$, $T = X_1 + X_2$ and $\hat{\theta} = t/2$ is the MLE. Under H_1 , $p(x_1, x_2|\theta) = \theta^{x_1}(1-\theta)^{1-x_1}\theta^{x_2}(1-\theta)^{1-x_2}$ and MLE of (θ_1, θ_2) is (X_1, X_2) .
- (b) Using a uniform prior, the posterior distribution follows, from Bayes' theorem, as $p(\theta|\mathbf{x}, H_0) \propto \theta^t(1-\theta)^{2-t}$ and $p(\theta_1, \theta_2|\mathbf{x}, H_1) \propto \theta^{x_1}(1-\theta)^{1-x_1}\theta^{x_2}(1-\theta)^{1-x_2}$.
- (c) The GMLE is the mode of the posterior distribution. In this exercise, the GMLE's coincide with the MLE's obtained in (b).
- (d) The predictive distribution is given by $p(\mathbf{x}) = E_\theta[p(\mathbf{x}|\theta)]$. So $p(x_1, x_2|H_0) = \int_0^1 \theta^t(1-\theta)^{2-t} d\theta = \Gamma(t+1)\Gamma(3-t)/\Gamma(4)$ and $p(x_1, x_2|H_1) = \Gamma(x_1+1)\Gamma(2-x_1)\Gamma(1+x_2)\Gamma(2-x_2)/\Gamma(3)^2$. Numerically, we get the table of probabilities below

(x_1, x_2)	(0,0)	(0,1)	(1,0)	(1,1)
H_0	1/3	1/6	1/6	1/3
H_1	1/4	1/4	1/4	1/4

- (e) The result follows immediately from the table above. The result states that the data favours H_0 (equality of distributions) twice more when $x_1 = x_2$ than when $x_1 \neq x_2$.
15. The MLE of θ is $\hat{\theta} = \sum_i X_i / \sum_i t_i$ and the Fisher information is $I(\theta) = \sum_i t_i / \theta$. From Section 5.3.1, the asymptotic distribution of the MLE is $\hat{\theta} \sim N[\theta, I^{-1}(\theta)]$. A $100(1 - \alpha)\%$ asymptotic confidence interval is $\hat{\theta} - z_{\alpha/2}\hat{I}^{-1/2}(\theta) \leq \theta \leq \hat{\theta} + z_{\alpha/2}\hat{I}^{-1/2}(\theta)$ where \hat{I} is an estimator of I given by $\sum_i t_i / \hat{\theta}$. If $\sum_i x_i = 10$ and $\sum_i t_i = 5$ then $\hat{\theta} = 2$ and $\hat{I}(\theta) = 5/2$. A 95% asymptotic confidence interval for θ is $(0.76, 3.24)$. Since $\theta_0 = 1$ belongs to the confidence interval described above we accept the null hypothesis at

the 5% level. With a 1% level, the corresponding interval is (0.37, 3.63) and $\theta_0 = 1$ will again be accepted.

21. Let the cell frequencies be denoted by N_1, \dots, N_p , with $N_i > 0$ and $\sum_{i=1}^p N_i = n$ and parameter $\theta_1, \dots, \theta_p$, $\theta_i > 0$ with $\sum_{i=1}^p \theta_i = 1$. The log-likelihood function is $L(\theta; \mathbf{N}) = k + \sum_{i=1}^p N_i \log \theta_i$, the MLE of θ is $\hat{\theta} = \mathbf{N}/n$ and the score function is $U(\mathbf{N}; \theta) = (N_1/\theta_1, \dots, N_{p-1}/\theta_{p-1})' - \mathbf{1}N_p/\theta_p$. The Fisher information is $I(\theta) = n \text{diag}(1/\theta_1, \dots, 1/\theta_{p-1}) - n\mathbf{1}'/\theta_p$ and its inverse is

$$n^{-1}[\text{diag}(\theta_1, \dots, \theta_{p-1}) - (\theta_1, \dots, \theta_{p-1})(\theta_1, \dots, \theta_{p-1})'].$$

Then, it follows that $W_E = \sum_{i=1}^p (N_i - n\theta_{i,0})^2/N_i$ and $W_U = \sum_{i=1}^p (N_i - n\theta_{i,0})^2/n\theta_{i,0}$.

Chapter 7

1. The prior has density given by $p(\theta) = \alpha\beta^\alpha/\theta^{\alpha+1}I_\theta(\beta, \infty)$ and the likelihood is given by $l(\theta; \mathbf{X}) = I_\theta(T, \infty)/\theta^n$, for $T = \max_i X_i$. Then, $\theta|\mathbf{x} \sim Pa(\beta_1, \alpha_1)$, where $\alpha_1 = \alpha + n$ and $\beta_1 = \max\{t, \beta\}$.

(a) From Bayes' theorem it follows that $p(\mathbf{x}) = p(\mathbf{x}|\theta)p(\theta)/p(\theta|\mathbf{x}) = \alpha\beta^\alpha/\alpha_1\beta_1^{\alpha_1}$. Assuming that Y and \mathbf{X} are conditionally independent given θ , it follows that $p(y|\mathbf{x}) = \alpha_1\beta_1^{\alpha_1}/\alpha_2\beta_2^{\alpha_2}$, with $\alpha_2 = \alpha_1 + 1$ and $\beta_2 = \max\{y, \beta_1\}$.

(b) If $T > \beta$ then $\beta_1 = t$, $\beta_2 = \max\{t, y\}$ and

$$p(y|\mathbf{x}) = \alpha_1 t^{\alpha_1}/[\alpha_2(\max\{t, y\})^{\alpha_2}].$$

So, $P(Y > t|\mathbf{x}) = (\alpha_1/\alpha_2)t^{\alpha_1} \int_t^\infty y^{-\alpha_2} dy = \alpha_1/[\alpha_2(\alpha_2 - 1)]$.

- (c) If $\alpha \rightarrow 0$ then $\alpha_1 \rightarrow n$, $\alpha_2 \rightarrow n + 1$, $p(y|\mathbf{x}) \rightarrow n\beta_1^n/[(n + 1)\beta_2^{n+1}]$ and $P(Y > t|\mathbf{x}) \rightarrow n/[(n + 1)n] = 1/(n + 1)$. If $\alpha, \beta \rightarrow 0$ then $p(y|\mathbf{x}) \rightarrow nt^n/[(n + 1)(\max\{t, y\})^{n+1}]$ or $p(y|\mathbf{x}) \rightarrow n/[(n + 1)t]$, if $y < t$ and $nt^n/[(n + 1)y^{n+1}]$, if $y > T$. Also, $P(Y > t|\mathbf{x}) \rightarrow 1/(n + 1)$.

4. The easiest way to predict Y in a classical way is to replace θ by its estimator $\hat{\theta}$. In this example, $p(y|\hat{\theta}) = \binom{m}{y}(x/n)^y(1-x/n)^{m-y}$, where $\hat{\theta} = x/n$. This distribution has mean mx/n and variance $mx(n-x)/n^2$. (This can be contrasted with the Bayesian prediction under the improper prior $p(\theta) \propto \theta^{-1}(1-\theta)^{-1}$ that leads to $E(Y|\mathbf{x}) = mx/n$ and $V(Y|\mathbf{x}) = mx(n-x)(1+c)/n^2$ where $c = (m-1)/(n+1) \geq 0$. This again is similar but overdispersed with respect to the classical result.)

Chapter 8

2. Define $F(\beta) = [S(\beta) - S_c]/ps^2 \sim F(p, n-p)$. Therefore, $P[F(\beta) \leq \bar{F}_\alpha(p, n-p)] = 1 - \alpha$. Solving for β , it means that $\{\beta : (\beta - \hat{\beta})' \mathbf{X}' \mathbf{X}(\beta - \hat{\beta}) \leq ps^2 \bar{F}_\alpha(p, n-p)\}$ defines a $100(1 - \alpha)$ % confidence region for β . This region has the form of an p -dimensional ellipsoid centred around $\hat{\beta}$.
7. Developing the product and conveniently collecting the terms, $Q = (\beta - \mu_0)' \mathbf{C}_0(\beta - \mu_0) + (\beta - \hat{\beta})' \mathbf{X}' \mathbf{X}(\beta - \hat{\beta}) = \beta' [\mathbf{C}_0 + \mathbf{X}' \mathbf{X}] \beta - 2\beta' [\mathbf{C}_0 \mu_0 + \mathbf{X}' \mathbf{X} \hat{\beta}] + \mu_0' \mathbf{C}_0 \mu_0 + \hat{\beta}' \mathbf{X}' \mathbf{X} \hat{\beta}$. Using μ_1 and C_1 as defined and completing the squares, it follows that $Q = Q_1 + Q_2$ where $Q_1 = \beta' C_1^{-1} \beta - 2\beta' C_1^{-1} \mu_1 + \mu_1' C_1^{-1} \mu_1$ and $Q_2 = \mu_0' \mathbf{C}_0 \mu_0 + \hat{\beta}' \mathbf{X}' \mathbf{X} \hat{\beta} - \mu_1' C_1^{-1} \mu_1$. Now, $Q_1 = (\beta - \mu_1)' C_1 (\beta - \mu_1)$, which proves the result.
16. The group means $\bar{y}_j \sim N(\beta_j, (m\phi)^{-1})$ are sufficient and independent statistics for β_1, \dots, β_k . Therefore, $E(\beta_j|\phi, \mathbf{y}) = E(\beta_j|\phi, \bar{y}_1, \dots, \bar{y}_k) = \dots = E\{E(\beta_j|\mu, \phi, \bar{y}_j)|\bar{y}_1, \dots, \bar{y}_k\}$. From Theorem 2.1, $E(\beta_j|\mu, \phi, \bar{y}_j) = (m\bar{y}_j + c_1\mu)/(m + c_1)$. So, $E(\beta_j|\phi, \mathbf{y}) = [m\bar{y}_j + c_1 E(\mu|\phi, \bar{y}_1, \dots, \bar{y}_k)]/(m + c_1)$. From Theorem 2.1 again, $E(\mu|\phi, \bar{y}_1, \dots, \bar{y}_k) = \bar{y}$. Therefore, $E(\beta_j|\phi, \mathbf{y}) = (m\bar{y}_j + c_1\bar{y})/(m + c_1)$.
18. The required distribution has density

$$p(\theta_{t-1}|D_t) = \int p(\theta_{t-1}|\theta_t, D_t) p(\theta_t|D_t) d\theta_t.$$

By Bayes' theorem and using the fact that y_t and θ_{t-1} are conditionally independent given θ_t , it follows that the first term in the integrand is $p(\theta_{t-1}|\theta_t, D_t) = p(\theta_{t-1}|\theta_t, D_{t-1}) \propto p(\theta_{t-1}|D_{t-1}) p(\theta_t|\theta_{t-1}, D_{t-1})$. Application of Bayes' theorem for the normal observation θ_t and parameter θ_{t-1} gives $(\theta_{t-1}|\theta_t, D_t) \sim N[m_{t-1} + C_{t-1}(\theta_t - a_t)/R_t, C_{t-1} - C_{t-1}^2/R_t]$. Therefore, $(\theta_{t-1}|D_t)$ has mean

$$E[\theta_{t-1}|D_t] = E[E(\theta_{t-1}|\theta_t)|D_t] = m_{t-1} + C_{t-1}(m_t - a_t)/R_t$$

and

$$\begin{aligned} V[\theta_{t-1}|D_t] &= E[V(\theta_{t-1}|\theta_t)|D_t] + V[E(\theta_{t-1}|\theta_t)|D_t] \\ &= C_{t-1} - C_{t-1}^2(R_t - C_{t-1})/R_t^2. \end{aligned}$$

The normality follows from the linear relation of θ_t on the conditional mean of θ_{t-1} and the marginal normality of θ_t .

List of distributions

This list includes the distributions that are most often used in statistics and frequently appeared in this book. They are listed in alphabetical order and are not divided into groups such as continuous \times discrete, univariate \times multivariate. This information will be clear from the context.

1. Bernoulli

X is said to have a Bernoulli distribution with success probability θ , denoted by $X \sim \text{Ber}(\theta)$, if its probability function is given by

$$p(x|\theta) = \theta^x(1-\theta)^{1-x}, \quad x = 0, 1,$$

for $\theta \in [0, 1]$. This distribution has mean θ and variance $\theta(1-\theta)$.

2. Beta

X is said to have a beta distribution with parameters α and β , denoted by $X \sim \text{beta}(\alpha, \beta)$, if its density function is given by

$$p(x|\alpha, \beta) = kx^{\alpha-1}(1-x)^{\beta-1}, \quad x \in [0, 1],$$

for $\alpha, \beta > 0$. The constant k is given by

$$k^{-1} = B(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)} \quad \text{and} \quad \Gamma(c) = \int_0^\infty x^{c-1}e^{-x}dx, \quad c > 0.$$

The functions $B(\cdot, \cdot)$ and $\Gamma(\cdot)$ are respectively known as the beta and gamma functions. This distribution has mean $\alpha/(\alpha + \beta)$ and variance $\alpha\beta/[(\alpha + \beta)^2(\alpha + \beta + 1)]$.

3. Beta-binomial

X is said to have a beta-binomial distribution with parameters n , α and β , denoted by $X \sim BB(n, \alpha, \beta)$, if its probability function is given by

$$p(x|n, \alpha, \beta) = \binom{n}{x} \frac{B(\alpha+x, \beta+n-x)}{B(\alpha, \beta)} \quad x = 0, 1, \dots, n,$$

for $n \geq 1$, $\alpha, \beta > 0$. This distribution has mean $n\alpha/(\alpha + \beta)$ and variance $n\alpha\beta[1 + (n-1)/(\alpha + \beta - 1)]/(\alpha + \beta)^2$. The expression for the variance

248 List of distributions

can be rewritten in the form $n\bar{\theta}(1-\bar{\theta})(1+\epsilon)$, where $\bar{\theta} = \alpha/(\alpha + \beta)$ and $\epsilon = (n-1)/(\alpha + \beta - 1)$. Since $\epsilon > 0$, $\forall n > 1$, the variance of the beta-binomial distribution is bigger than the variance of the binomial, when one compares θ with $\bar{\theta}$, for $\alpha + \beta > 1$.

4. Binomial

X is said to have a binomial distribution with parameter n and success probability θ , denoted by $X \sim \text{bin}(n, \theta)$, if its probability function is given by

$$p(x|n, \theta) = \binom{n}{x} \theta^x (1-\theta)^{n-x}, \quad x = 0, 1, \dots, n,$$

for $n \geq 1, \theta \in [0, 1]$. This distribution has mean $n\theta$ and variance $n\theta(1-\theta)$. This family of distributions includes the Bernoulli as the special case $n = 1$.

5. Dirichlet

$\mathbf{X} = (X_1, \dots, X_p)'$ is said to have a Dirichlet distribution with parameter $\theta = (\theta_1, \dots, \theta_p)'$, denoted by $D(\theta)$, if its joint density function is given by

$$p(\mathbf{x}) = \frac{\Gamma(\theta_+)}{\prod_i \Gamma(\theta_i)} \prod_{i=1}^p x_i^{\theta_i-1}, \quad x_i \in [0, 1], \quad i = 1, \dots, p, \quad \sum_{i=1}^p x_i = 1,$$

for $\theta_i \in [0, 1], i = 1, \dots, p$ and $\theta_+ = \sum_i \theta_i$. This distribution has mean θ/θ_+ and variance-covariance matrix given by $\theta_+^{-2}(\theta_+ + 1)^{-1}[\theta_+\text{diag}(\theta) - \theta\theta']$, where $\text{diag}(\mathbf{c})$ denotes a diagonal matrix having the elements of the vector \mathbf{c} in the main diagonal. This family of distributions includes the beta as the special case $p = 2$.

6. Exponential

X is said to have an exponential distribution with parameter θ , denoted by $X \sim \text{Exp}(\theta)$, if its density function is given by

$$p(x|\theta) = \theta e^{-\theta x}, \quad x > 0,$$

for $\theta > 0$. This distribution has mean $1/\theta$ and variance $1/\theta^2$.

7. Gamma (or χ^2)

X is said to have a gamma distribution with parameters α and β , denoted by $X \sim G(\alpha, \beta)$, if its density function is given by

$$p(x|\alpha, \beta) = k x^{\alpha-1} e^{-\beta x}, \quad x > 0,$$

for $\alpha, \beta > 0$. The constant k is given by $k = \beta^\alpha / \Gamma(\alpha)$. This distribution has mean α/β and variance α/β^2 . This family of distributions includes the exponential as the special case $\alpha = 1$.

The χ_p^2 distribution is equivalent to the $G(p/2, 1/2)$ distribution. Therefore, the $G(\alpha, \beta)$ distribution corresponds to a $2\beta \chi_{2\alpha}^2$ distribution.

8. Multinomial

$\mathbf{X} = (X_1, \dots, X_p)'$ is said to have multinomial distribution with parameter n and probabilities $\theta = (\theta_1, \dots, \theta_p)'$, denoted by $M(n, \theta)$, if its joint

probability function is given by

$$p(\mathbf{x}|\theta) = \frac{n!}{\prod_{i=1}^p x_i!} \prod_{i=1}^p \theta_i^{x_i}, \quad x_i = 0, 1, \dots, n, \\ i = 1, \dots, p, \quad \sum_{i=1}^p x_i = n,$$

for $\theta_i \in [0, 1], i = 1, \dots, p, \sum_{i=1}^p \theta_i = 1$. This distribution has mean $n\theta$ and variance-covariance matrix given by $n[\text{diag}(\theta) - \theta\theta']$. This family of distributions includes the binomial as the special case $p = 2$.

9. Negative binomial

X is said to have a negative binomial distribution with parameters r and θ , denoted by $X \sim NB(r, \theta)$, if its probability function is given by

$$p(x|r, \theta) = \binom{r+x-1}{x} \theta^r (1-\theta)^x, \quad x = 0, 1, \dots,$$

for $r \geq 1, \theta \in [0, 1]$. This distribution has mean $r(1-\theta)/\theta$ and variance $r(1-\theta)/\theta^2$.

10. Normal

X is said to have a normal distribution with mean μ and variance σ^2 , denoted by $X \sim N(\mu, \sigma^2)$, if its density function is given by

$$p(x|\mu, \sigma^2) = (2\pi\sigma^2)^{-1/2} \exp\left\{-\frac{1}{2\sigma^2}(x-\mu)^2\right\}, \quad x \in R,$$

for $\mu \in R$ and $\sigma^2 > 0$. When $\mu = 0$ and $\sigma^2 = 1$, the distribution is referred to as standard normal.

$\mathbf{X} = (X_1, \dots, X_p)'$ is said to have a multivariate normal distribution with mean vector μ and variance-covariance matrix Σ , denoted by $N(\mu, \Sigma)$, if its density function is given by

$$(2\pi)^{-p/2} |\Sigma|^{-1/2} \exp\left\{-\frac{1}{2}(\mathbf{x}-\mu)' \Sigma^{-1} (\mathbf{x}-\mu)\right\}, \quad \mathbf{x} \in R^p,$$

for $\mu \in R^p$ and $\Sigma > 0$, where $|\mathbf{A}|$ denotes the determinant of \mathbf{A} .

11. Pareto

X is said to have a Pareto distribution with parameters a and θ , denoted by $\text{Pa}(a, \theta)$, if its density function is given by

$$p(x|\theta, a) = a\theta^a/x^{1+a}, \quad x > \theta,$$

for $a, \theta > 0$. This distribution has mean $a\theta/(a-1)$, when $a > 1$, and variance $a\theta^2/[(a-1)^2(a-2)]$, when $a > 2$.

12. Poisson

X is said to have a Poisson distribution with parameter θ , denoted by $X \sim \text{Pois}(\theta)$, if its probability function is given by

$$p(x|\theta) = e^{-\theta} \frac{\theta^x}{x!}, \quad x = 0, 1, 2, \dots,$$

for $\theta > 0$. This distribution has mean and variance given by λ .

13. Snedecor F

X is said to have Snedecor F (or simply F) distribution with v_1 and v_2 degrees of freedom, denoted $F(v_1, v_2)$, if its density function is given by

$$p(x|v_1, v_2) = \frac{\Gamma[(v_1 + v_2)/2]}{\Gamma(v_1/2)\Gamma(v_2/2)} v_1^{v_1/2} v_2^{v_2/2} x^{(v_1/2)-1} (v_2 + v_1 x)^{-(v_1+v_2)/2}, \\ x > 0$$

for $v_1, v_2 > 0$. This distribution has mean $v_2/(v_2 - 2)$, when $v_2 > 2$ and variance $[2v_2^2(v_1 + v_2 - 2)]/[v_1(v_2 - 4)(v_2 - 2)^2]$, when $v_2 > 4$.

If $X_1 \sim \chi_{v_1}^2$ and $X_2 \sim \chi_{v_2}^2$ are independent then $(X_1/v_1)/(X_2/v_2) \sim F(v_1, v_2)$. This family of distributions includes the square of the Student t as the special case $v_1 = 1$. If $v_2 \rightarrow \infty$, then $v_1 F(v_1, v_2) \xrightarrow{D} \chi_{v_1}^2$.

14. Student t

X is said to have a Student t (or simply t) distribution with mean μ , scale parameter σ^2 and v degrees of freedom, denoted by $X \sim t_v(\mu, \sigma^2)$; if its density function is given by

$$p(x|v, \mu, \sigma^2) = \frac{\Gamma[(v+1)/2]}{\Gamma(v/2)\sqrt{\pi}} \frac{v^{v/2}}{\sigma} \left[v + \frac{(x-\mu)^2}{\sigma^2} \right]^{-(v+1)/2}, \quad x \in R,$$

for $v > 0$, $\mu \in R$ and $\sigma^2 > 0$. This distribution has mean μ , when $v > 1$, and variance $v/(v-2)$, when $v > 2$. This family includes the Cauchy distribution as the special case $v = 1$, denoted by Cauchy (μ, σ^2) .

$\mathbf{X} = (X_1, \dots, X_p)'$ is said to have a multivariate Student t distribution with mean vector μ and scale matrix Σ and v degrees of freedom, denoted by $t_v(\mu, \Sigma)$, if its density function is given by

$$f(\mathbf{x}|v, \mu, \Sigma) = \frac{\Gamma[(v+p)/2]}{\Gamma(v/2)\pi^{p/2}} v^{v/2} |\Sigma|^{-1/2} \\ \times [v + (\mathbf{x}-\mu)' \Sigma^{-1} (\mathbf{x}-\mu)]^{-(v+p)/2}, \quad \mathbf{x} \in R^p,$$

for $v > 0$, $\mu \in R^p$ and $\Sigma > 0$. This distribution has mean μ , when $v > 1$, and variance $v\Sigma/(v-2)$, when $v > 2$.

15. Uniform

X is said to have a uniform distribution with parameters θ_1 and θ_2 , denoted by $U[\theta_1, \theta_2]$, if its density function is given by

$$p(x|\theta_1, \theta_2) = \frac{1}{\theta_2 - \theta_1}, \quad x \in [\theta_1, \theta_2],$$

for $\theta_1 < \theta_2$. When $\theta_1 = 0$ and $\theta_2 = 1$, the distribution is referred to as unit uniform. This distribution has mean $(\theta_1 + \theta_2)/2$ and variance $(\theta_2 - \theta_1)^2/12$.

16. Weibull

X is said to have a Weibull distribution with parameters α and β , denoted by $\text{Wei}(\alpha, \beta)$, if its density function is given by

$$p(x|\alpha, \beta) = \beta \alpha x^{\alpha-1} \exp(-\beta x^\alpha), \quad x > 0,$$

for $\alpha, \beta > 0$. This distribution is sometimes parametrized in terms of α and $\theta = 1/\beta^\alpha$ and includes the exponential as the special case $\alpha = 1$. This distribution has mean $\beta^{-1/\alpha} \Gamma(1+\alpha^{-1})$ and variance $\beta^{-2/\alpha} [\Gamma(1+2\alpha^{-1}) - \Gamma^2(1+\alpha^{-1})]$.

References

- Abramowitz, M. and Stegun, I.A., editors (1965). *Handbook of Mathematical Functions*, National Bureau of Standards Applied Mathematics Series, Number 55. Washington: U.S. Government Printing Office.
- Aitchison, J. and Dunsmore, I.R. (1975). *Statistical Prediction Analysis*. Cambridge: Cambridge University Press.
- Barnett, V. (1973). *Comparative Statistical Inference*. New York: Wiley.
- Bartlett, M.S. (1947). Multivariate analysis. *Journal of the Royal Statistical Society, 9*, B (supplement), 76–97.
- Bayes, T. (1763). An essay towards solving in the doctrine of chances. *Philosophical Transactions of the Royal Society London, 53*, 370–418.
- Berger, J. O. (1985). *Statistical Decision Theory and Bayesian Analysis*. New York: Springer-Verlag.
- Berger, J. and Delampady, M. (1987). Testing precise hypothesis. *Statistical Science, 2*, 317–352.
- Berger, J. and Sellke, T. (1987). Testing a point null hypothesis: the irreconcilability of significance levels and evidence. *Journal of the American Statistical Association, 82*, 112–122.
- Bernardo, J. M. (1979). Reference posterior distributions for Bayesian inference (with discussion). *Journal of the Royal Statistical Society, Series B, 41*, 113–147.
- Bernardo, J. M. and Smith, A. M. F. (1994). *Bayesian Theory*. Chichester: Wiley.
- Bickel, P. J. and Doksum, K. A. (1977). *Mathematical Statistics*. Englewood Cliffs: Prentice Hall.
- Box, G. E. P. and Tiao, G. C. (1992). *Bayesian Inference in Statistical Analysis*. Reading: Addison-Wesley.
- Broemeling, L. D. (1985). *Bayesian Analysis of Linear Models*. New York: Marcel Dekker.
- Buse, A. (1982). The likelihood ratio, Wald and Lagrange multiplier tests: an expository note. *The American Statistician, 36*, 3, 153–157.
- Casella, G. and George, E. I. (1992) Explaining the Gibbs sampler. *The American Statistician, 46*, 167–174.
- Chib, S. and Greenberg, E. (1995) Understanding the Metropolis–Hastings algorithm. *The American Statistician, 49*, 327–335.
- Cordeiro, G. M. (1987). On the correction to the likelihood ratio statistics. *Biometrika, 74*, 265–274.
- Cordeiro, G. M. and Ferrari, S. L. P. (1991). A modified score test having chi-squared distribution to order n^{-1} . *Biometrika, 78*, 573–582.
- Cox, D. R. and Hinkley, D. V. (1974). *Theoretical Statistics*. London: Chapman & Hall.

- Davison, A. C. and Hinkley, D. V. (1997). *Bootstrap Methods and their Application*. Cambridge: Cambridge University Press.
- de Finetti, B. (1937). La prévision: ses lois logiques, ses sources subjectives. *Annals of the Institute Poincaré*, 7, 1, 1–68.
- de Finetti, B. (1974). *Theory of Probability* (vols. 1 and 2). New York: J. Wiley.
- DeGroot, M. H. (1970). *Optimal Statistical Decisions*. New York: J. Wiley.
- DeGroot, M. H. (1986). *Probability and Statistics* (2nd edition). Reading: Addison-Wesley.
- Dempster, A. P., Laird, N. M. and Rubin, D. B. (1977). Maximum likelihood estimation from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society, Series B*, 39, 1–38 (with discussion).
- Draper, N. R. and Smith, H. (1966). *Applied Regression Analysis*. New York: Wiley.
- Efron, B. (1979). Bootstrap methods: another look at the jackknife. *Annals of Statistics*, 7, 1–26.
- Efron, B. (1982). The Jackknife, the Bootstrap and Other Resampling Plans, Monograph 38. Philadelphia: SIAM.
- Ferguson, T. S. (1967). *Mathematical Statistics: A Decision-Theoretic Approach*. New York: Academic Press.
- Fisher, R. A. (1936). The use of multiple measurements in taxonomic problems. *Annals of Eugenics*, 7, 179–188.
- Gamerman, D. (1997). *Markov Chain Monte Carlo: Stochastic Simulation for Bayesian Inference*. London: Chapman & Hall.
- Gamerman, D. and Migon, H. (1993). Dynamic hierarchical models. *Journal of the Royal Statistical Society, Series B*, 55, 629–643.
- Garthwaite, P. H., Jolliffe, I. T. and Jones, B. (1995). *Statistical Inference*. London: Prentice Hall.
- Geisser, S. (1993) *Predictive Inference: An Introduction*. London: Chapman & Hall.
- Gelfand, A. E. and Smith, A. M. F. (1990). Sampling-based approaches to calculating marginal densities. *Journal of the American Statistical Association*, 85, 398–409.
- Geman, S. and Geman, D. (1984). Stochastic relaxation, Gibbs distribution and the Bayesian restoration of images. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 6, 721–741.
- Gilks, W. R., Richardson, S. and Spiegelhalter, D. J. (editors). (1996). *Markov Chain Monte Carlo in Practice*. London: Chapman & Hall.
- Harrison, P. J. and Stevens, C. F. (1976). Bayesian forecasting. *Journal of the Royal Statistical Society, Series B*, 38, 205–247 (with discussion).
- Hastings, W. K. (1970) Monte Carlo sampling methods using Markov chains and their applications. *Biometrika*, 57, 97–109.
- Heath, D. L. and Sudderth, W. D. (1976). De Finetti's theorem for exchangeable random variables. *American Statistician*, 30, 333–345.
- Jeffreys, H. (1961). *Theory of Probability*. Oxford: Clarendon Press.
- Kalbfleisch, J. G. (1975). *Probability and Statistical Inference*, volume 2 (2nd edition). New York: Springer-Verlag.
- Lawley, D. N. (1956). A general method for approximating to the distribution of likelihood ratio criteria. *Biometrika*, 43, 295–303.
- Lee, P. M. (1997). *Bayesian Statistics: An Introduction* (2nd edition). Oxford: Oxford University Press.
- Lehmann, E. (1986). *Testing Statistical Hypothesis* (2nd edition). New York: Wiley.

- Lindley, D. V. (1957). A statistical paradox. *Biometrika*, 44, 187–192.
- Lindley, D. V. (1965) *Introduction to Probability and Statistics* (Parts 1 and 2). Cambridge: Cambridge University Press.
- Lindley, D. V. (1980). Approximate Bayesian methods. *Bayesian Statistics* (Eds. J. M. Bernardo et al.), 223–245 (with discussion).
- Lindley, D. V. (1993). The analysis of experimental data: the appreciation of tea and wine. *Teaching Statistics*, 15, 1, 22–259.
- Lindley, D. V. and Phillips, L. D. (1976). Inference for a Bernoulli Process (a Bayesian view). *The American Statistician*, 30, 112–119.
- Lindley, D. V. and Smith, A. F. M. (1972) Bayes estimates for the linear model (with discussion). *Journal of the Royal Statistical Society, Series B*, 34, 1–41.
- Little, R. J. A. and Rubin, D. B. (1988). *Statistical Analysis with Missing Data*. New York: J. Wiley.
- McCullagh, P. and Nelder, J. A. (1989). *Generalized Linear Models* (2nd edition). London: Chapman & Hall.
- Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H. and Teller, E. (1953). Equation of state calculations by fast computing machine. *Journal of Chemical Physics*, 21, 1087–91.
- Migon, H. S. and Tachibana, V. M. (1997). Bayesian approximations in randomized response model. *Computational Statistics and Data Analysis*, 24, 401–409.
- Naylor, J. C. and Smith, A. M. F. (1982). Applications of a method for efficient computation of posterior distributions. *Applied Statistics*, 31, 214–235.
- Neyman, J. and Pearson, E. S. (1928). On the use and interpretation of certain test criterion for purposes of statistical inference. *Biometrika*, 20, 175–240 and 263–294.
- O'Hagan, A. (1994). *Bayesian Inference*. Volume 2B of Kendall's *Advanced Theory of Statistics*. London: Edward Arnold.
- Quenouille, M. H. (1949). Approximate test of correlation in time series. *Journal of the Royal Statistical Society, Series B*, 34, 1–41.
- Quenouille, M. H. (1956). Notes on the bias estimation. *Biometrika*, 43, 353–360.
- Raiffa, H. and Schlaifer, R. (1961). *Applied Statistical Decision Theory*. Cambridge, MA: Harvard University Press.
- Rao, C. R. (1973). *Linear Statistical Inference* (2nd. edition). New York: Wiley.
- Ripley, B. D. (1987). *Stochastic Simulation*. New York: Wiley.
- Silvey, S. D. (1970). *Statistical Inference*. London: Chapman & Hall.
- Tanner, M. A. (1996). *Tools for Statistical Inference—Methods for Exploration of Posterior Distributions and Likelihood Functions* (3rd edition). New York: Springer Verlag.
- Thisted, R. A. (1976). *Elements of Statistical Computing*. New York: Chapman & Hall.
- Tierney, L. and Kadane, J. B. (1986). Accurate approximations for posterior moments and marginal densities. *Journal of the American Statistical Association*, 81, 82–86.
- West, M. and Harrison, P. J. (1997). *Bayesian Forecasting and Dynamic Models* (2nd edition). New York: Springer Verlag.
- Zellner, A. (1971). *An Introduction to Bayesian Inference in Econometrics*. New York: Wiley.

Author index

Abramovitz, M., 146
 Aitchison, J., 28, 200
 Bahadur, R. R., 36
 Barnett, V. D., xi, 27
 Bartlett, M. S., 43, 188, 190
 Bayes, Rev. Thomas, 27
 Berger, J., 2, 55, 84, 102, 121, 185,
 188
 Bernardo, J. M., 35, 67, 68, 71, 140
 Bickel, P. J., 19, 168
 Box, G. E. P., 32, 66
 Broemeling, L. D., 211
 Buse, A., 190
 Casella, G., 164
 Chib, S., 157
 Copas, J. B., 226
 Cordeiro, G. M., 190
 Cox, D. R., 43, 86, 96
 Darmois, G., 40
 Davison, A. C., 125, 151
 De Finetti, B., 5, 32
 DeGroot, M. H., 36, 84, 168
 Delampady, M., 188
 Dempster, A. P., 130, 131
 Doksum, K. A., 19, 168
 Draper, N. R., 211
 Dunsmore, I. R., 28, 200
 Efron, B., 147, 151, 153
 Ferguson, T. S., 84
 Ferrari, S. L. P., 190

Fisher, R. A., 36, 96, 153
 Gamerman, D., 125, 147, 156, 160,
 231
 Garthwaite, P. H., 126
 Geisser, S., 28, 200
 Gelfand, A. E., 125, 157
 Geman, D., 157
 Geman, S., 157
 George, E. I., 164
 Gilks, W. R., 160
 Goldstein, M., 208
 Greenberg, E., 157
 Harrison, P. J., 227
 Hartigan, J., 208
 Hastings, W. K., 157
 Heath, D. L., 34
 Hinkley, D. V., 43, 86, 96, 125, 151
 James, W., 226
 Jeffreys, H., 66, 68, 71, 182
 Jollife, I. T., 126
 Jones, B., 126
 Kadane, J. B., 143
 Kalbfleisch, J. G., 43
 Koopman, B. O., 40
 Laird, N. M., 130, 131
 Lawley, D. N., 190
 Lehmann, E. L., 36, 95, 168
 Lindley, D. V., 7, 31, 57, 71, 135,
 184, 187, 188, 224
 Little, R. J. A., 130

McCullagh, P., 213
 Metropolis, N., 157
 Migon, H. S., 143, 231
 Morris, C. N., 226
 Naylor, J. C., 147
 Nelder, J. A., 213
 Neyman, J., 168, 170, 173
 O'Hagan, A., 226
 Pearson, E. S., 168, 170, 173
 Phillips, L. D., 7
 Pitman, E. J. G., 40
 Quenouille, M. H., 152
 Raiffa, H., 58
 Rao, C. R., 96, 131
 Ripley, B. D., 147
 Rubin, D. B., 130, 131
 Scheffe, H., 36

Tachibana, V. M., 143
 Tanner, M. A., 125
 Thisted, R. A., 125, 130
 Tiao, G. C., 32, 66
 Tierney, L., 143
 Wald, A., 188
 West, M., 227
 Wilks, S. S., 188
 Zellner, A., 68

Subject index

Alternative hypotheses, 168
 Analytical approximations, 133
 asymptotic theory, 133–40
 by Kullback–Liebler divergence, 140–41
 Laplace approximation, 141–4
 Ancillarity, 38
 Asymptotic
 efficiency, 139
 tests, 188–91
 theory, 133–40
 Bayes factor, 182, 183–4, 185–6, 223
 Bayesian approach (subjective approach), 15
 estimation, 87, 89
 hypothesis testing, 167–8, 181–6
 confidence intervals, 186, 187–8
 inference, general problem of, 125, 126
 Laplace approximation, 142
 linear models, 218–24
 Monte Carlo with importance sampling, 151
 parameter estimation, 42, 44
 prediction, 197–201, 203–5, 208
 resampling methods, 151, 153, 155, 156, 158–9
 sufficiency, 35–8
 Bayes risk, 93

decision theory, 80, 81, 83
 Bayes rule, 80, 81
 Bayes' theorem, 26–7
 prediction, 28–30
 sequential nature, 30–32
 Bias
 classical hypothesis testing, 176
 comparison of estimators, 93
 jackknife, 152
 Bidimensional rule, 145
 Bonferroni inequality, 103, 105
 Bootstrap, 151, 153–4
 weighted, 151, 154–6
 Cartesian product rule, 145, 147
 Central limit theorem, 138
 Cholesky decomposition, 14, 147
 Classical approach (frequentist approach), 7–8, 15
 estimation, 82, 86, 92–3
 confidence interval, 100, 104–5
 hypothesis testing, 167–81
 parameter elimination, 42–4
 prediction, 201–2, 205–6
 probability, 3, 4
 resampling methods, 151
 sufficiency, 37–8
 Complete families of distributions, 95–6
 Conditional
 density, 11
 independence, 11, 12
 likelihood, 43

posterior distributions, 42
 Confidence interval
 asymptotic, 135, 138–9, 140
 bootstrap, 153, 154
 Bayesian approach, 100–101, 102–3
 classical approach, 100, 104, 105
 highest posterior density (HPD), 101–2
 highest predictive density (HPD), 201, 203, 205
 jackknife, 152, 153
 Bayesian linear models, 219, 222
 classical linear models, 214, 215, 218
 normal model: one sample case, 106, 107, 108, 109, 110
 normal model: two samples case, 114, 115, 116, 117
 optimization techniques, 129–30
 Conjugate family of distributions, 55, 56–8
 natural, 58
 Consistency of estimators, 98–9
 Correlation coefficient, 153
 Cramer–Rao inequality, 96
 Credibility interval, 100
 Critical region, 168, 174–5
 Cromwell's rule, 57
 Decision rules, 80, 81
 Decision theory, 79–86
 Bayesian prediction, 199–200
 De Finetti's representation theorem, 34, 53
 Delta method, 136, 141
 Density function, 14
 Distribution,
 Behrens–Fisher, 114, 115
 Bernoulli, 247, 248
 beta-binomial, 59, 247–8
 beta, 57, 247
 binomial, 248, 249
 bivariate normal, 111–12
 Cauchy, 250
 chi-squared distribution, 107, 113, 188, 189
 Dirichlet, 62, 248
 exponential, 248
 Gamma, 248
 multinomial, 248–9
 multivariate normal, 12, 13, 249
 multivariate Student t , 250
 negative binomial, 8, 249
 normal- χ^2 , 63
 normal, 12–13, 249
 normal-gamma, 63, 64
 Pareto, 249
 Poisson, 41, 61, 250
 Snedecor F distribution, 116, 250
 Student t distribution, 250
 uniform, 250–51
 Dynamic linear models, 227
 definition and examples, 227–9
 evolution and updating equations, 229–30
 special aspects, 231
 Efficiency of estimators, 97–8
 asymptotic, 139
 Elimination of parameters, 41–4
 EM algorithm, 130–3
 Empirical distribution function, 92
 Estimation, 1, 79
 classical point, 86–92
 comparison of estimators, 92–9
 decision theory, 79–86
 interval, 99–105
 Exchangeability, 32–5
 partial 35
 prior distribution, 53
 Exponential family, 39–40

Rao–Blackwell theorem, 94–5,
158–9

Reference priors, *see*
Non-informative priors

Risk
Bayes, 80, 81, 83, 93
comparison of estimators, 93–8
decision theory, 80, 81–3
linear, 206–7
quadratic, 94, 99

Scale
model, 20–21, 71
parameter, 21

Score function, 25–6
generalized, 90

Shrinkage estimators, 226

Simulation methods, 147
Monte Carlo methods, 147–9
Monte Carlo with importance
sampling, 149–51
Markov chain Monte Carlo,
156–60
resampling methods, 151–6

Stein, 226

Subjective approach, *see* Bayesian
approach

Sufficiency, 35–41
Fisher information, 23
partitions, 37, 39
risk, 94

Taylor expansion
dynamic linear models, 227
hypothesis testing, 188, 191
Laplace approximation, 142
optimization techniques, 127

Tchebychev's inequality, 99

Test procedures, 167
consistent, 191
goodness-of-fit, 190–91
invariant, 177
optimal, 169–70
similar, 177
t test, 177, 179
Type I errors of, 169, 170, 172
Type II errors of, 169, 172
unbiased, 175–6, 179
uniformly more powerful (UMP),
173–6

Unbiased estimators, 93

Uncertainty, 1, 2

Uniformly minimum variance
unbiased (UMVU)
estimators, 93

Variance–covariance matrices,
11–12, 13

Weibull distribution, 251

completeness, 95
 conjugacy, 59–65
 jackknife, 153
 sufficiency, 39–41

Factorization criterion, Neyman's, 37, 38, 39, 40

Fisher information, 23–5
 asymptotic theory, 137, 139
 consistency, 98
 expected, 24, 25
 observed, 24, 25

Fisher scoring, 128, 129

Frequentist approach, *see* Classical approach

Gamma function, 247

Gauss–Hermite rules, 146–7

Generalized linear models, 212–13

Generalized maximum likelihood estimator (GMLE), 87–90
 decision theory, 83–4
 EM algorithm, 130

Generalized score function, 90

Gibbs sampler algorithm, 157–160

Hierarchical
 linear models, 224–6
 priors, 71–3

Homoscedasticity, 90

Hyperparameters, 55, 58, 72

Hypothesis testing, 1, 167–8
 asymptotic tests, 188–91
 Bayesian, 181–6
 classical, 168–81
 composite, 167, 173–7

Importance sampling, 149–51

Influence diagrams, 11, 12

Information, 2
 additivity of, 24

Jackknife, 151, 152–3

Kullback–Liebler divergence, 140–41

Laplace approximation, 141–4

Least squares,
 method of, 90–92
 estimator (LSE), 90–91
 weighted estimator, 91, 92
 ordinary, 91

Likelihood
 function, 21–3
 maximum, 87, 89, 90
 principle, 23

Likelihood ratio (LR)
 classical hypothesis testing, 170, 171, 174–5, 177
 monotone, 174–5, 177

Lindley's paradox, 184

Linear algebra, basic results in, 10–11, 13–14

Linear Bayes methodology, 208

Linear growth model (LGM), 227–8

Linear models, 211–13
 Bayesian, 218–24
 classical, 213–18
 dynamic, 227–31
 hierarchical, 224–6

Linear prediction, 206–8

Location
 model, 20, 70
 parameter, 20, 21
 –scale model, 21

Loss functions, 80–84
 0–1, 83–4
 absolute, 84
 quadratic, 84, 93

Marginal
 density, 11
 likelihood, 182
 posterior distributions, 42

Markov chain Monte Carlo (MCMC), 156–60

Matrices
 linear algebra, 13
 orthogonal, 14
 positive definite, 13–4

singular, 14
 square, 13
 symmetric, 13

Maximum likelihood estimator (MLE), 86–90
 asymptotic tests, 188, 189, 190, 191
 asymptotic theory, 137–8, 139, 140
 consistency, 98
 EM algorithm, 130, 131
 generalized, *see* Generalized maximum likelihood estimator
 optimization techniques, 127–9

Mean squared error (MSE), 93

Metropolis–Hastings algorithm, 157–8, 165

Minimax estimator, 94

Moments, method of, 92

Monte Carlo methods, 147–9
 with importance sampling, 149–51

Markov chain, 156–60

Newton–Raphson algorithm, 127
 confidence intervals, determining, 129
 likelihood equation, solution of, 127, 128, 129

Neyman's factorization criterion, 37, 38, 39, 40

Non-informative priors, 65–71

Nuisance parameters, 42, 43
 non-informative priors, 68, 71

Null hypothesis, 168

Numerical integration methods, 144
 Gauss–Hermite rules, 146–7
 Newton–Cotes type methods, 144–5

One-way ANOVA, 215

Optimality criterion, linear prediction, 206

Optimization techniques, 126–7

Parameter elimination, 41–4

Posterior distribution, 27, 31–3
 asymptotic theory, 134, 135–6, 137

decision theory, 79

Laplace approximation, 143–4

Markov chain Monte Carlo methods, 159
 parameter elimination, 42

Posterior odds, 182

Prediction, 1, 197
 Bayesian, 197–201
 classical, 201–2
 linear, 206–8
 linear model, 217–18

Prior distribution, 27, 30–33, 53

Conjugacy, 59–65
 distribution function approach to specification, 54

hierarchical, 71–3

histogram approach, specification of, 53–4

Jeffreys, 66–9, 76

linear models, 218–19, 220, 221–2

non-informative, *see* Non-informative priors

subjective specification, 53–5

specification through functional forms, 55–8

Prior odds, 182

Probability, 1
 basic results in, 10–13
 concept of, 3–5
 function, 14

p-value, 172

Quadrature technique, *see* Numerical integration methods

Randomized tests, 173

Rao–Blackwellized estimator, 158, 164